

# [CSED211] Introduction to Computer Software Systems

## Lecture 12: Linking

Prof. Jisung Park



**CAOS**  
COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2023.11.13

# Lecture Agenda: Linking

---

- Linking
  - Motivation
  - What It Does
  - How It Works
  - Dynamic Linking
- Case Study: Library Interpositioning

# C Program Example

---

main.c

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main() {
    int val = sum(array, 2);
    return val;
}
```

sum.c

```
int sum(int *a, int n) {
    int i, s = 0;
    for (i = 0; i < n; i++)
        s += a[i];
    return s;
}
```

# Static Linking

---

- Programs are translated and linked using a compiler driver

```
$ gcc -Og -o prog main.c sum.c  
$ ./prog
```

`main.c`

`sum.c`

**Source Files**

# Static Linking

- Programs are translated and linked using a compiler driver

```
$ gcc -Og -o prog main.c sum.c  
$ ./prog
```

main.c

sum.c

Source Files

Translators


Translators

*Preprocessing (cpp), compile (cc1), assembly (as)*

# Static Linking

- Programs are translated and linked using a compiler driver

```
$ gcc -Og -o prog main.c sum.c  
$ ./prog
```



Source Files


*Preprocessing (cpp), compile (cc1), assembly (as)*

Separately-Compiled **Relocatable** Object Files

# Static Linking

- Programs are translated and linked using a compiler driver

```
$ gcc -Og -o prog main.c sum.c
$ ./prog
```



Source Files

*Preprocessing (cpp), compile (cc1), assembly (as)*

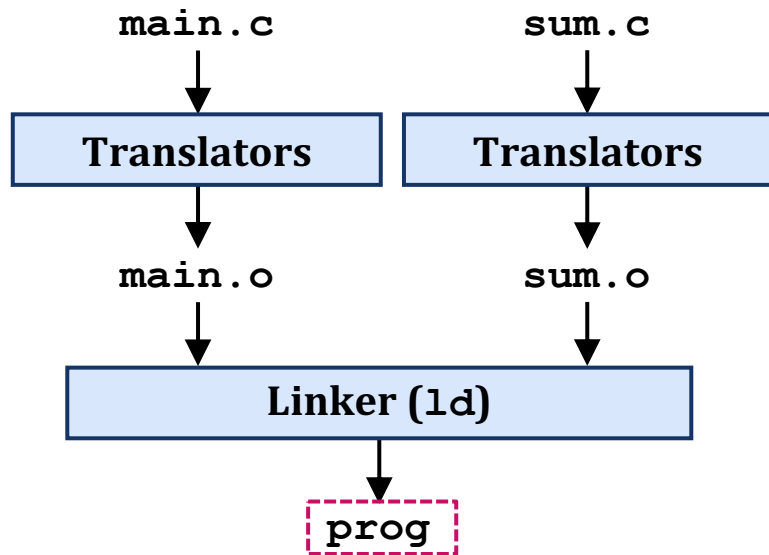
Separately-Compiled **Relocatable** Object Files

*Merges code and data for all functions defined in multiple modules (main.c and sum.c)*

# Static Linking

- Programs are translated and linked using a compiler driver

```
$ gcc -Og -o prog main.c sum.c
$ ./prog
```



Source Files

*Preprocessing (cpp), compile (cc1), assembly (as)*

Separately-Compiled **Relocatable** Object Files

*Merges code and data for all functions defined in multiple modules (main.c and sum.c)*

Fully-Linked **Executable** Object File



# Why Linkers?

---

- Reason 1: modularity
  - Program can be written as a collection of smaller source files, rather than one monolithic mass
  - Can build libraries of common functions (more on this later)
    - e.g., math library, standard C library

# Why Linkers? (Cont.)

---

- Reason 2: Efficiency

- Time: separate compilation
  - Change one source file, compile, and the relink
  - No need to recompile other source files
  - Can compile multiple files concurrently
- Space: libraries
  - Common functions can be aggregated into a single file in two ways
    - Option 1: static linking
      - Executable files and running memory images contain only the library code they actually use
    - Option 2: dynamic linking
      - Executable files contain no library code
      - During execution, a single copy of library code can be shared across all executing processes

# What Do Linkers Do?

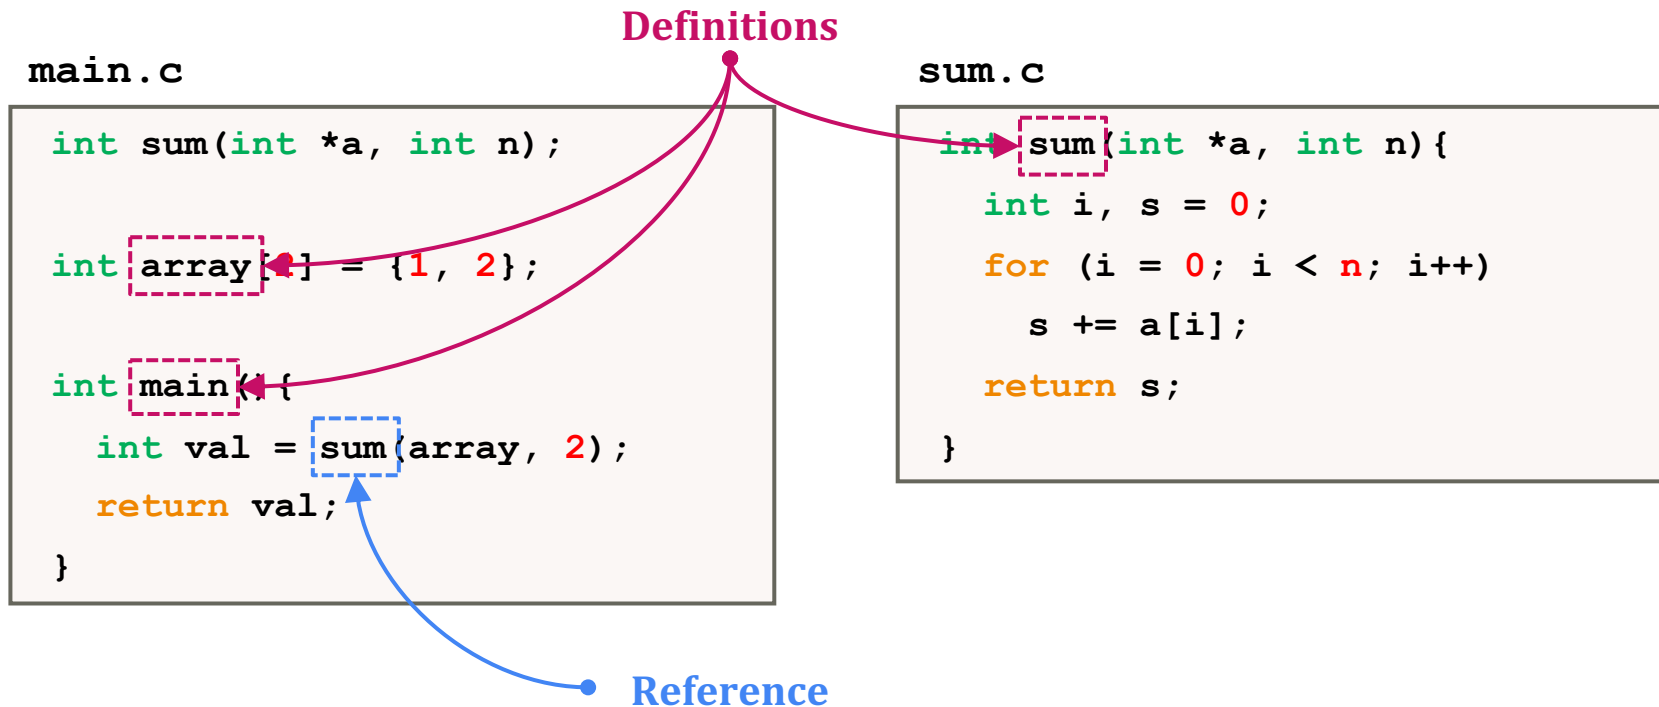
- Step 1. symbol resolution

- Programs define and reference symbols (variables and functions)

```
void swap() {...}; /* define symbol swap */  
swap();           /* reference symbol swap */  
int *xp = &x;     /* define symbol xp, reference x */
```

- Compiler stores symbol definitions in **symbol table**
  - An array of structs, each includes the name, size, and location of a symbol
- **Linker associates each symbol reference with exactly one symbol definition**

# Symbol Resolution: C Program Example



# What Do Linkers Do?

- **Step 1. symbol resolution**

- Programs define and reference symbols (variables and functions)

```
void swap() {...}; /* define symbol swap */  
swap();           /* reference symbol swap */  
int *xp = &x;     /* define symbol xp, reference x */
```

- Compiler stores symbol definitions in **symbol table**
  - An array of structs, each includes the name, size, and location of a symbol
- **Linker associates each symbol reference with exactly one symbol definition**

- **Step 2. relocation**

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable
- Updates all references to these symbols to reflect their new positions

# Three Kinds of Object Files (Modules)

---

- **Relocatable object file** (`.o` file)
  - Contains code and data in a form that can be combined with other relocatable object files to form executable object file
  - Each `.o` file is produced from exactly one source (`.c`) file
- **Executable object file** (`a.out` file)
  - Contains code and data in a form that can be copied directly into memory and then executed
- **Shared object file** (`.so` file)
  - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either loading time or run-time
  - Called **Dynamic Link Libraries (DLLs)** in Windows

# Executable and Linkable Format (ELF)

---

- Standard binary format for object files
- One unified format for
  - Relocatable object files (`.o`)
  - Executable object files (`a.out`)
  - Shared object files (`.so`)
- Generic name: ELF binaries

# ELF Object File Format

- **ELF header**: word size, byte ordering, file type, machine type, etc.
- **Segment header table**: page size, virtual addresses memory segments (sections), segment sizes
- **.text section**: code
- **.rodata section**: read-only data, e.g., jump tables, etc.
- **.data section**: initialized global variables
- **.bss section**
  - Uninitialized global variables
  - Block Started by Symbol
  - Better Save Space
  - Has section header but occupies no space

|                      |
|----------------------|
| ELF header           |
| Segment header table |
| .text section        |
| .rodata section      |
| .data section        |
| .bss section         |
| .symtab section      |
| .rel.txt section     |
| .rel.data section    |
| .debug section       |
| Section header table |



# ELF Object File Format (Cont.)

- **.symtab section**
  - Symbol table
  - Procedure and static variable names
  - Section names and locations
- **.rel.text section**
  - Relocation information for **.text** section
  - Addresses of instructions that will need to be modified in the executable
  - Instructions for modifying
- **.rel.dat section**
  - Relocation information for **.data** section
  - Addresses of pointer data that will need to be modified in the merged executable

|                      |
|----------------------|
| ELF header           |
| Segment header table |
| .text section        |
| .rodata section      |
| .data section        |
| .bss section         |
| .symtab section      |
| .rel.txt section     |
| .rel.data section    |
| .debug section       |
| Section header table |

# ELF Object File Format (Cont.)

- **.debug section**
  - Info for symbolic debugging (`gcc -g`)
- **Section header table**
  - Offsets and sizes of each section

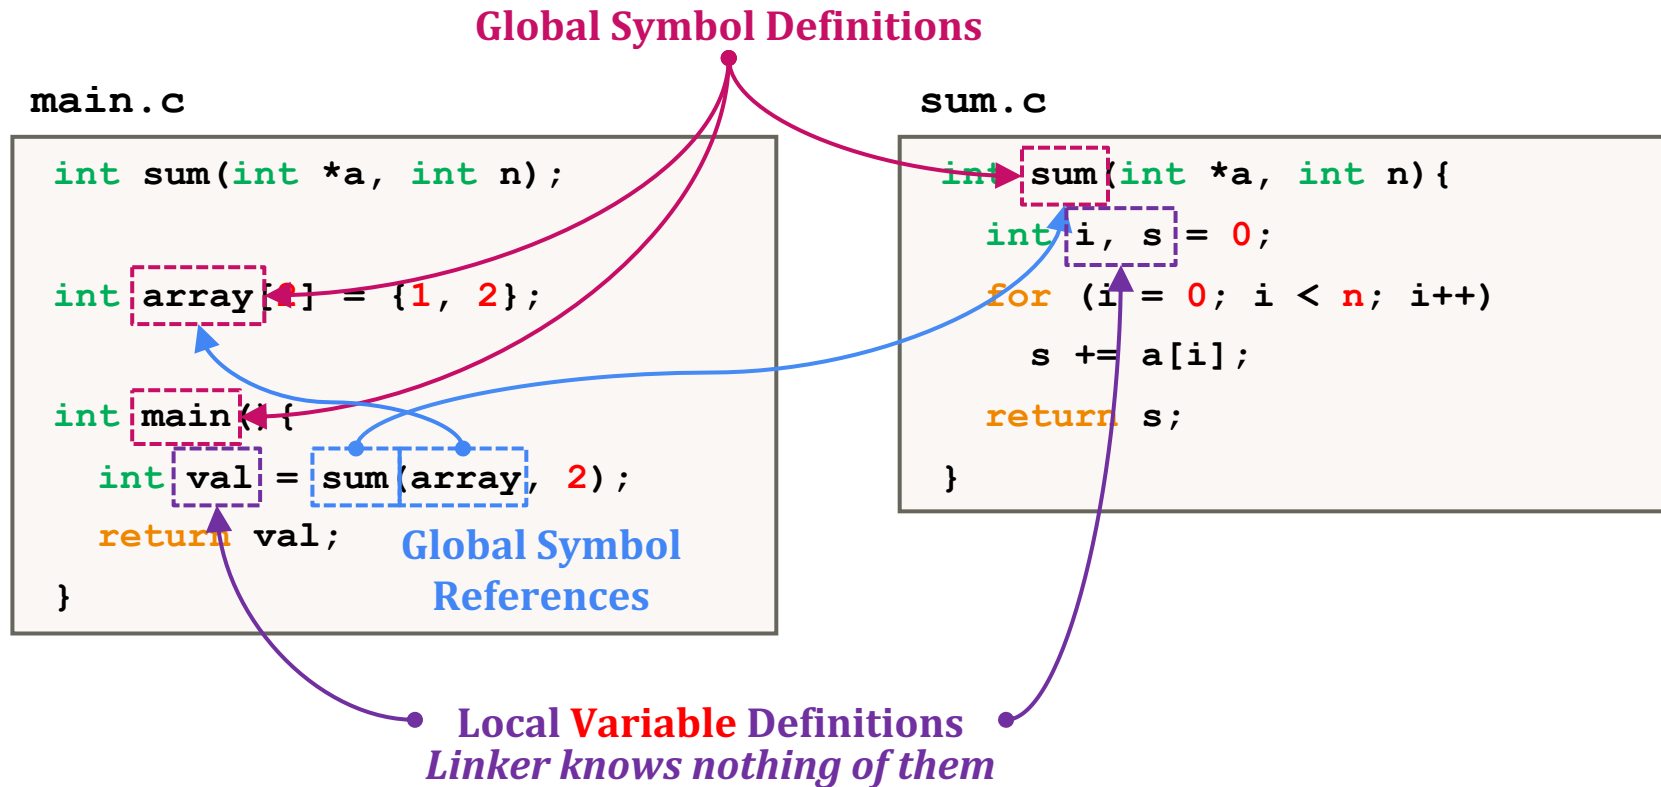
|                      |
|----------------------|
| ELF header           |
| Segment header table |
| .text section        |
| .rodata section      |
| .data section        |
| .bss section         |
| .symtab section      |
| .rel.txt section     |
| .rel.data section    |
| .debug section       |
| Section header table |

# Linker Symbols

---

- Global symbols
  - Defined by module *M* that can be referenced by other modules
  - e.g., non-static C functions and non-static global variables
- External symbols
  - Global symbols referenced by module *M* but defined by another module
- Local symbols
  - Defined and referenced exclusively by module *M*
  - e.g., C functions and global variables defined with the static attribute
  - Local linker symbols are not local program variables

# Step 1: Symbol Resolution



# Symbol Identification

- Which of the following names will be in the symbol table of `symbols.o`?
  - `time`
  - `foo`
  - `a`
  - `argc`
  - `argv`
  - `b`
  - `main`
  - `printf`
  - Any others?

`symbols.c`

```
int time;

int foo(int a){
    int b = a + 1;
    return b;
}

int main(int argc, char** argv){
    printf("%d\n", foo(5));
    return 0;
}
```

# Symbol Identification

- Which of the following names will be in the symbol table of `symbols.o`?
  - `time`
  - `foo`
  - `a`
  - `argc`
  - `argv`
  - `b`
  - `main`
  - `printf`
  - Any others?
- Can find this with `readelf`

```
$ readelf -s symbols.o
```

`symbols.c`

```
int time;

int foo(int a){
    int b = a + 1;
    return b;
}

int main(int argc, char** argv){
    printf("%d\n", foo(5));
    return 0;
}
```

# Local Symbols

- Local non-static C variables vs. Local static C variables
  - Local non-static: stored on the stack
  - Local static: stored in either `.bss` or `.data`

*Compiler allocates space in `.data`  
for **each definition of `x`***

*Creates local static symbols in the symbol table  
with unique names, e.g., `x`, `x.1721`, and `x.1724`*

```
static int x = 15;



int f() {
    static int x = 17;
    return x++;
}

int g() {
    static int x = 19;
    return x += 14;
}

int h() {
    return x += 27;
}
```

# How Linker Resolves Duplicate Symbol Definitions

- Program symbols are either **strong** or **weak**
  - **Strong**: procedures and initialized globals
  - **Weak**: uninitialized globals





# Linker's Symbol Rules

---

- **Rule 1: multiple strong symbols are not allowed**
  - Each item can be defined only once
  - Otherwise, linker error
- **Rule 2: a strong symbol and multiple weak symbols → the strong symbol**
  - References to the weak symbol resolve to the strong symbol
- **Rule 3: multiple weak symbols → pick an arbitrary one**
  - Can override this with `gcc -fno-common`

# Linker Puzzles

p1.c

```
int x;  
p1() {}
```

```
int x;  
p1() {}
```

```
int x, y;  
p1() {}
```

```
int x = 7, y = 5;  
p1() {}
```

```
int x = 7;  
p1() {}
```

p2.c

```
p1() {}
```

```
int x;  
p2() {}
```

```
double x;  
p2() {}
```

```
double x;  
p2() {}
```

```
int x;  
p2() {}
```

**Error:** two strong symbols (p1)

References to `x` will refer to the same uninitialized `int`; is this what you really want?

Writes to `x` in p2.c **might overwrite y**!

Writes to `x` in p2.c **will overwrite y**!

References to `x` will refer to the initialized variable (in p1.c)

**Nightmare scenario: two identical weak structs compiled with different alignment rules.**

# Type Mismatch Example

- Compiles without any errors or warnings, but what gets printed?

mismatch-main.c

```
long int x; /* Weak symbol */

int main(int argc, char** argv){
    printf("%ld\n", x);
    return 0;
}
```

mismatch-variable.c

```
/* Global strong symbol */
double x = 3.14;
```

# Global Variables

---

- Avoid if you can
- Otherwise
  - Use **static** if you can
  - Initialize if you define a global variable
  - Use **extern** if you reference an external global variable
    - Treated as **weak symbol**
    - But also causes linker error if not defined in some file

# Use of extern in .h Files

c1.c

```
#include "global.h"

int f(){
    return g + 1;
}
```

global.h

```
extern int g;
int f();
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main(int argc, char** argv){
    if(init)
        // do something, e.g., g=31;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

# Use of extern in .h Files

c1.c

```
#include "global.h"

int f(){
    return g + 1;
}
```

global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
extern int g;
static int init = 0;
#endif
```

c2.c

```
#define INITIALIZE
#include <stdio.h>
#include "global.h"

int main(int argc, char** argv){
    if(init)
        // do something, e.g., g=31;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

# Use of extern in .h Files

c1.c

```
int g;  
static int init = 0;  
int f(){  
    return g + 1;  
}
```

global.h


```
#ifndef INITIALIZE  
    int g = 23;  
    static int init = 1;  
#else  
    extern int g;  
    static int init = 0;  
#endif
```

c2.c


```
#define INITIALIZE  
#include <stdio.h>  
  
int g = 23;  
static int init = 1;  
  
int main(int argc, char** argv){  
    if(init)  
        // do something, e.g., g=31;  
    int t = f();  
    printf("Calling f yields %d\n", t);  
    return 0;  
}
```

# Step 2: Relocation


## Relocatable Object Files



main.o



sum.o



.text

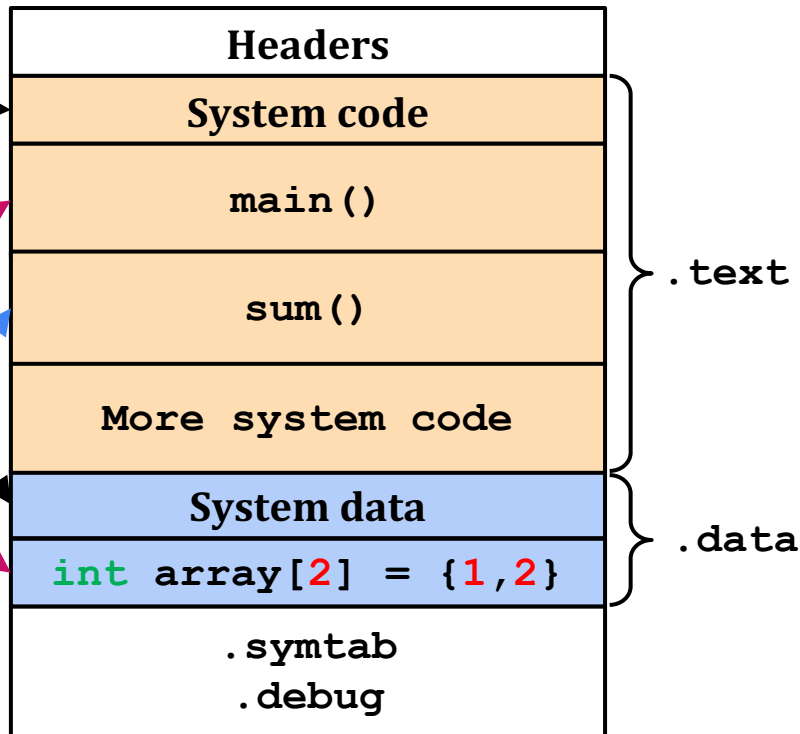
.data

.text

.data

.text

## Executable Object File





# Relocation Entries

```
int array[2] = {1, 2};    main.c
```

```
int main(){
    int val = sum(array, 2);
    return val;
}
```

```
0000000000000000 <main>:                                main.o
0:  48 83 ec 08      sub    $0x8,%rsp
4:  be 02 00 00 00   mov    $0x2,%esi
9:  bf 00 00 00 00   mov    $0x0,%edi    # %edi = &array
                        a: R_X86_64_32 array                # Relocation entry

e:  e8 00 00 00 00   callq 13 <main+0x13> # sum()
                        f: R_X86_64_PC32 sum-0x4            # Relocation entry
13: 48 83 c4 08      add    $0x8,%rsp
17: c3              retq
```

# Relocated .text Section

00000000004004d0 <main>:

```
4004d0: 48 83 ec 08      sub    $0x8,%rsp
4004d4: be 02 00 00 00    mov    $0x2,%esi
4004d9: bf 18 10 60 00    mov    $0x601018,%edi    # %edi = &array
4004de: e8 05 00 00 00    callq 4004e8 <sum>      # sum()
4004e3: 48 83 c4 08      add    $0x8,%rsp
4004e7: c3               retq
```

00000000004004e8 <sum>:

```
4004e8: b8 00 00 00 00    mov    $0x0,%eax
4004ed: ba 00 00 00 00    mov    $0x0,%edx
4004f2: eb 09             jmp     4004fd <sum+0x15>
4004f4: 48 63 ca         movslq %edx,%rcx
4004f7: 03 04 8f         add    (%rdi,%rcx,4),%eax
4004fa: 83 c2 01         add    $0x1,%edx
4004fd: 39 f2            cmp    %esi,%edx
4004ff: 7c f3            jl     4004f4 <sum+0xc>
400501: f3 c3            repz retq
```

Using PC-relative addressing for sum() :

$0x4004e8 = 0x4004e3 + 0x05$

Source: objdump -dx prog

# Loading Executable Object Files

## Executable Object File

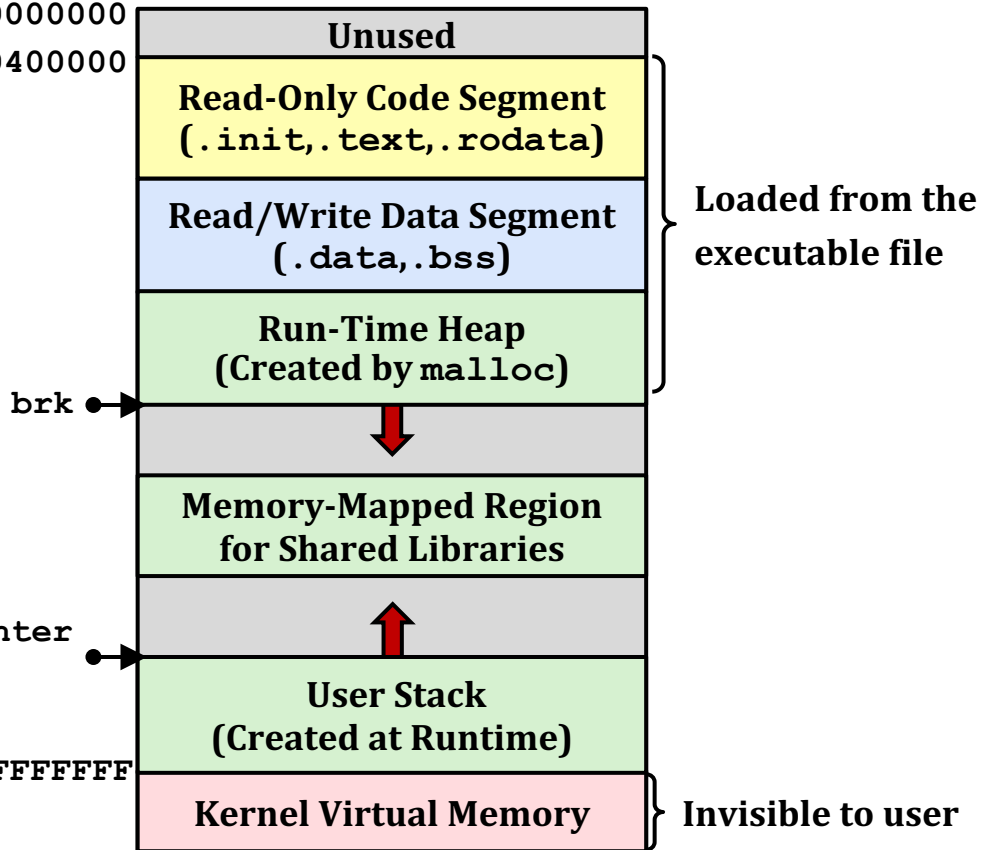
|   |
|---|
| ELF header  |
| Program header table<br>(required for executables)  |
| .init section                                       |
| .text section                                       |
| .rodata section                                     |
| .data section                                       |
| .bss section  |
| .symtab section                                     |
| .debug section                                      |
| .line section                                       |
| .strtab section                                     |
| Section header table<br>(required for relocatables) |

0x0000000000000000

0x000000000000400000

Stack pointer  
(%rsp)

0x00007FFFFFFFFFFFFF



# Packaging Commonly Used Functions

---

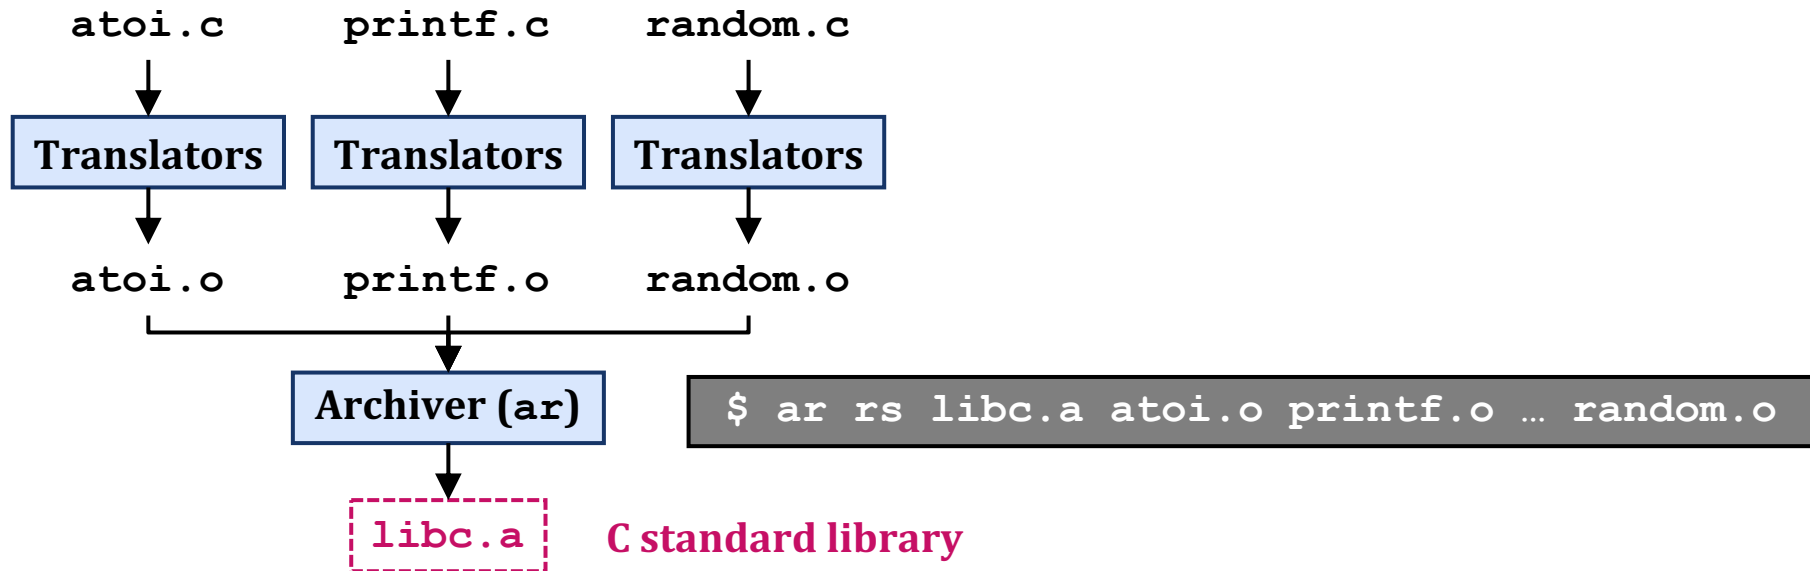
- How to package functions commonly used by programmers?
  - e.g., math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far
  - **Option 1:** put all functions into a single source file
    - Programmers link big object file into their programs
    - Space and time inefficient
  - **Option 2:** put each function in a separate source file
    - Programmers explicitly link appropriate binaries into their programs
    - More efficient, but burdensome on the programmer

# Old-fashioned Solution: Static Libraries

---

- **Static libraries** (.a archive files)
  - Concatenate related relocatable object files into a single file with an index
    - Which is called an **archive**
  - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives
  - If an archive member file resolves reference, link it into the executable

# Creating Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace `.o` file in archive

# Commonly Used Libraries

- **libc.a** (the C standard library)
  - 4.6 MB archive of 1,496 object files
  - I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math, etc.
- **libm.a** (the C math library)
  - 2 MB archive of 444 object files
  - Floating point math (sin, cos, tan, log, exp, sqrt, etc.)

```
$ ar -t libc.a | sort
```

```
...
```

```
fork.o
```

```
...
```

```
fprintf.o
```

```
fpu_control.o
```

```
fputc.o
```

```
freopen.o
```

```
fscanf.o
```

```
fseek.o
```

```
fstab.o
```

```
...
```

```
$ ar -t libm.a | sort
```

```
...
```

```
e_acos.o
```

```
e_acosf.o
```

```
e_acosh.o
```

```
e_acoshf.o
```

```
e_acoshl.o
```

```
e_acosl.o
```

```
e_asin.o
```

```
e_asinf.o
```

```
e_asinl.o
```

```
...
```

# Linking with Static Libraries

```
main2.c
#include <stdio.h>
#include "vector.h"

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main() {
    addvec(x, y, z, 2);
    printf("z = [%d %d]\n", z[0], z[1]);
    return 0;
}
```

```
void addvec(int* x, int* y, int* z, int n) {
    int i;
    for (i = 0; i < n; i++)
        z[i] = x[i] + y[i];
}
```

addvec.c

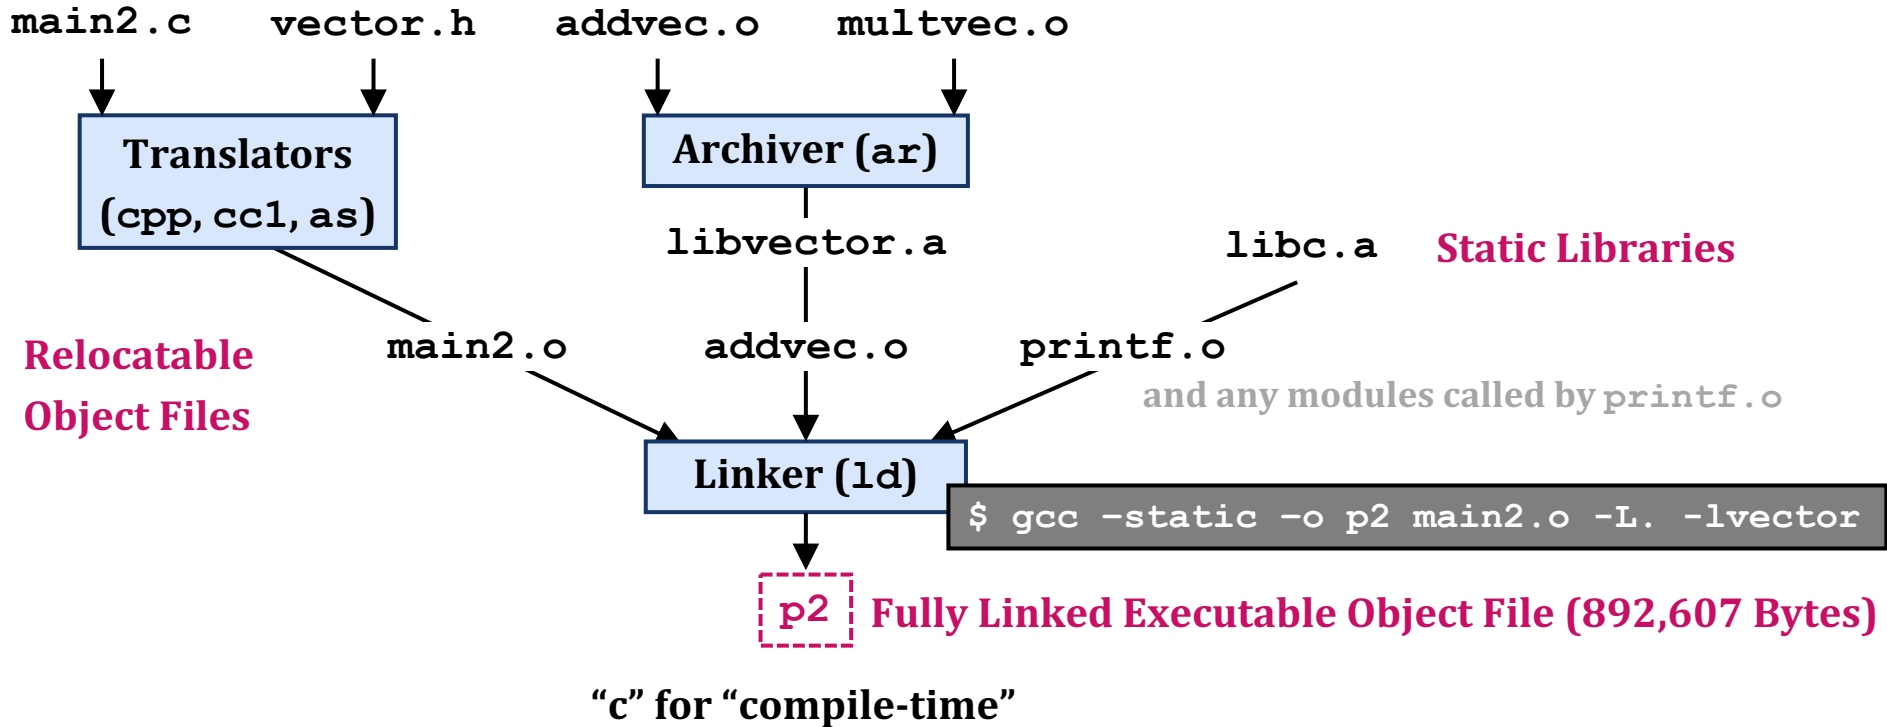
```
void multvec(int *x, int *y, int *z, int n) {
    int i;
    for (i = 0; i < n; i++)
        z[i] = x[i] * y[i];
}
```

multvec.c

libvector.a



# Linking with Static Libraries



# Using Static Libraries

- Linker's algorithm for resolving external references
  - Scan `.o` files and `.a` files in the **command line order**
  - During the scan, keep a list of the current unresolved references
  - As each new `.o` or `.a` file, `obj`, is encountered, try to resolve each unresolved reference in the list against the symbols defined in `obj`
  - If any entries in the unresolved list at end of scan, then error
- Problem
  - **Command line order** matters
  - Moral: put **libraries at the end of the command line**

```
$ gcc -L. libtest.o -lmime
$ gcc -L. -lmime libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

# Modern Solution: Shared Libraries

---

- Static libraries have the following disadvantages:
  - Duplication in the stored executables (every function need std libc)
  - Duplication in the running executables
  - Minor bug fixes of system libraries require each application to explicitly relink
    - Rebuild everything with glibc?
    - <https://security.googleblog.com/2016/02/cve-2015-7547-glibc-getaddrinfo-stack.html>
- Modern solution: **shared libraries**
  - Object files that contain code and data that are loaded and linked into an application **dynamically**, at either **load-time** or **run-time**
  - Also called: dynamic link libraries (DLLs) in Windows and **.so** files in Linux

# Shared Libraries (Cont.)

---

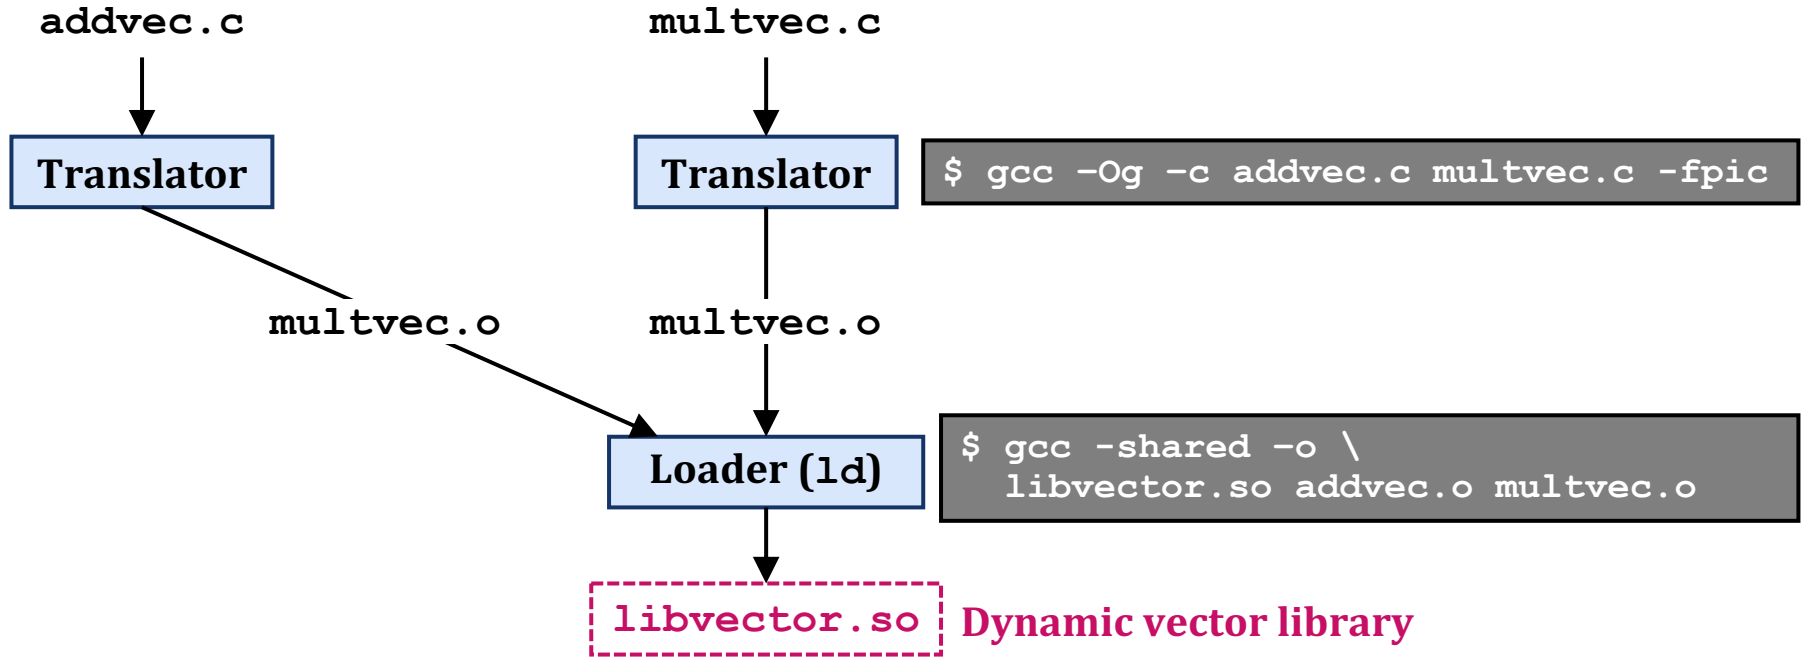
- Dynamic linking can occur when executable is **first loaded and run**
  - **Load-time linking**
  - Common case for Linux, handled automatically by the dynamic linker
    - i.e., `ld-linux.so`
  - Standard C library (`libc.so`) usually dynamically linked
- Dynamic linking can also occur **after program has begun**
  - **Run-time linking**
  - In Linux, this is done by calls to the `dlopen()` interface
    - Distributing software
    - High-performance web servers
    - Runtime library interpositioning
- Shared library routines can be shared by multiple processes
  - More on this when we learn about virtual memory

# What Dynamic Libraries Are Required?

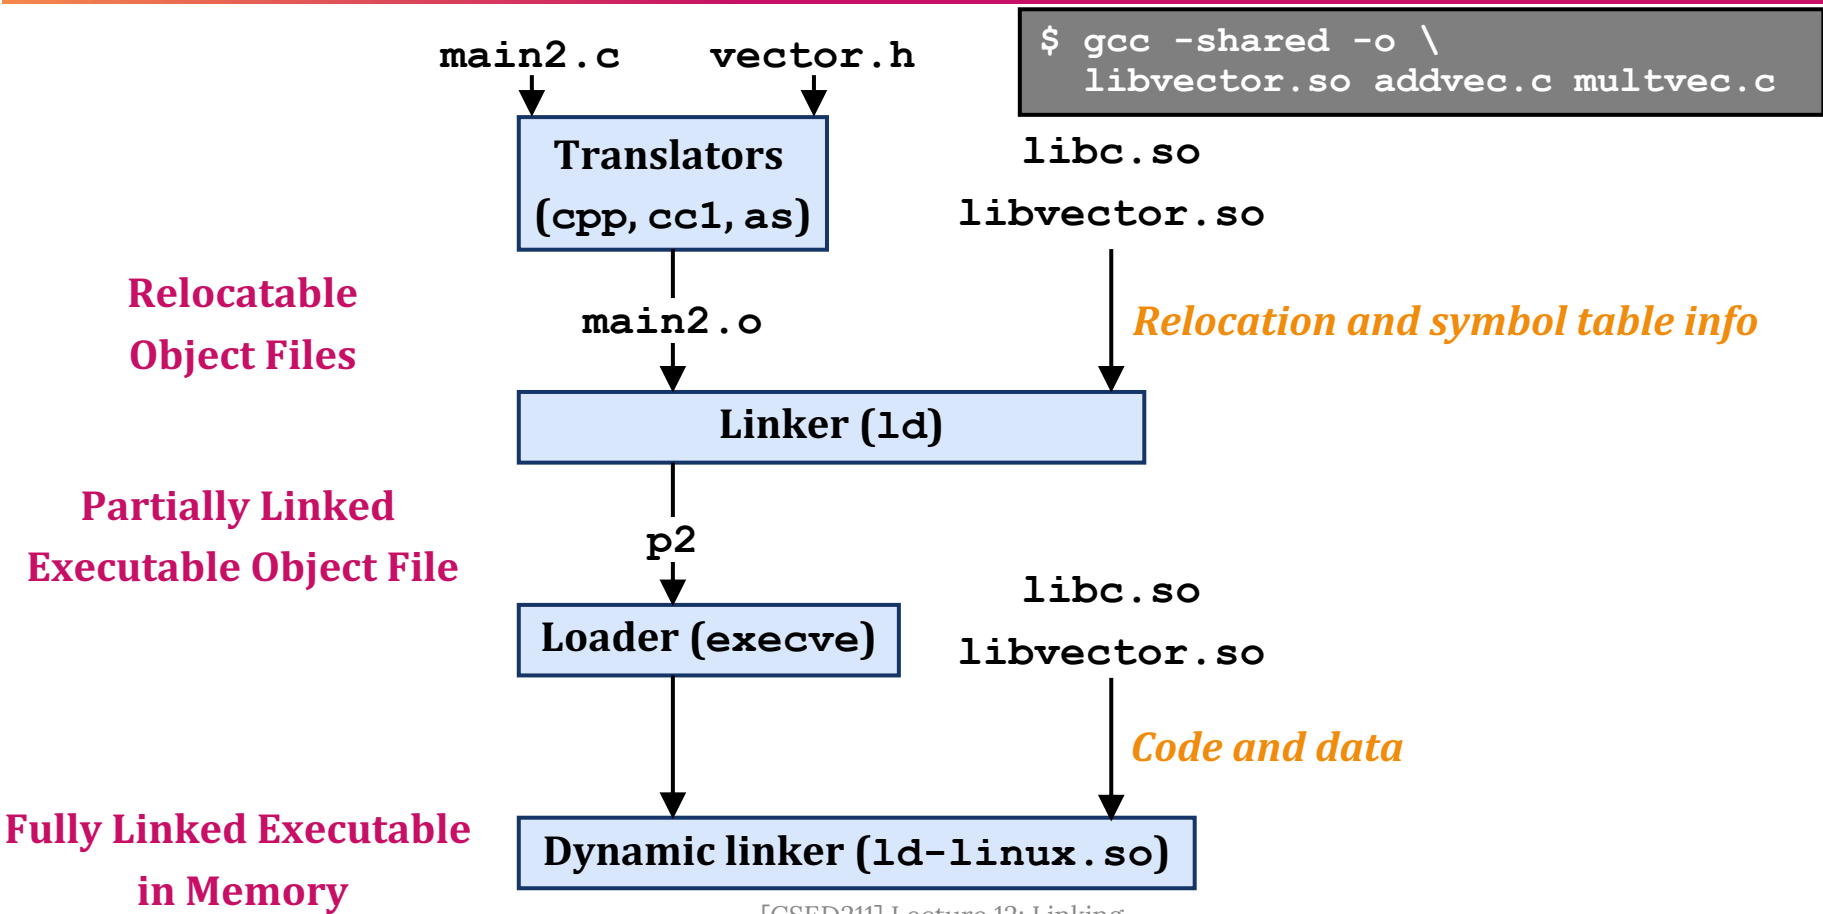
- **.interp** section
  - Specifies the dynamic linker to use (i.e., `ld-linux.so`)
- **.dynamic** section
  - Specifies the names, etc. of the dynamic libraries to use
  - Follow an example of `prog`  
(NEEDED)      Shared library: [libm.so.6]
- Where are the libraries found?
  - Use “`ldd`” to find out

```
$ ldd prog
linux-vdso.so.1 => (0x00007ffcf2998000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f99ad927000)
/lib64/ld-linux-x86-64.so.2 (0x00007f99adcef000)
```

# Dynamic Linking at Load-time



# Dynamic Linking at Load-Time



# Dynamic Linking at Run-Time

```
dll.c
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main(){
    void* handle;
    void (*addvec)(int *, int *, int *, int);
    char* error;

    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if(!handle){
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```

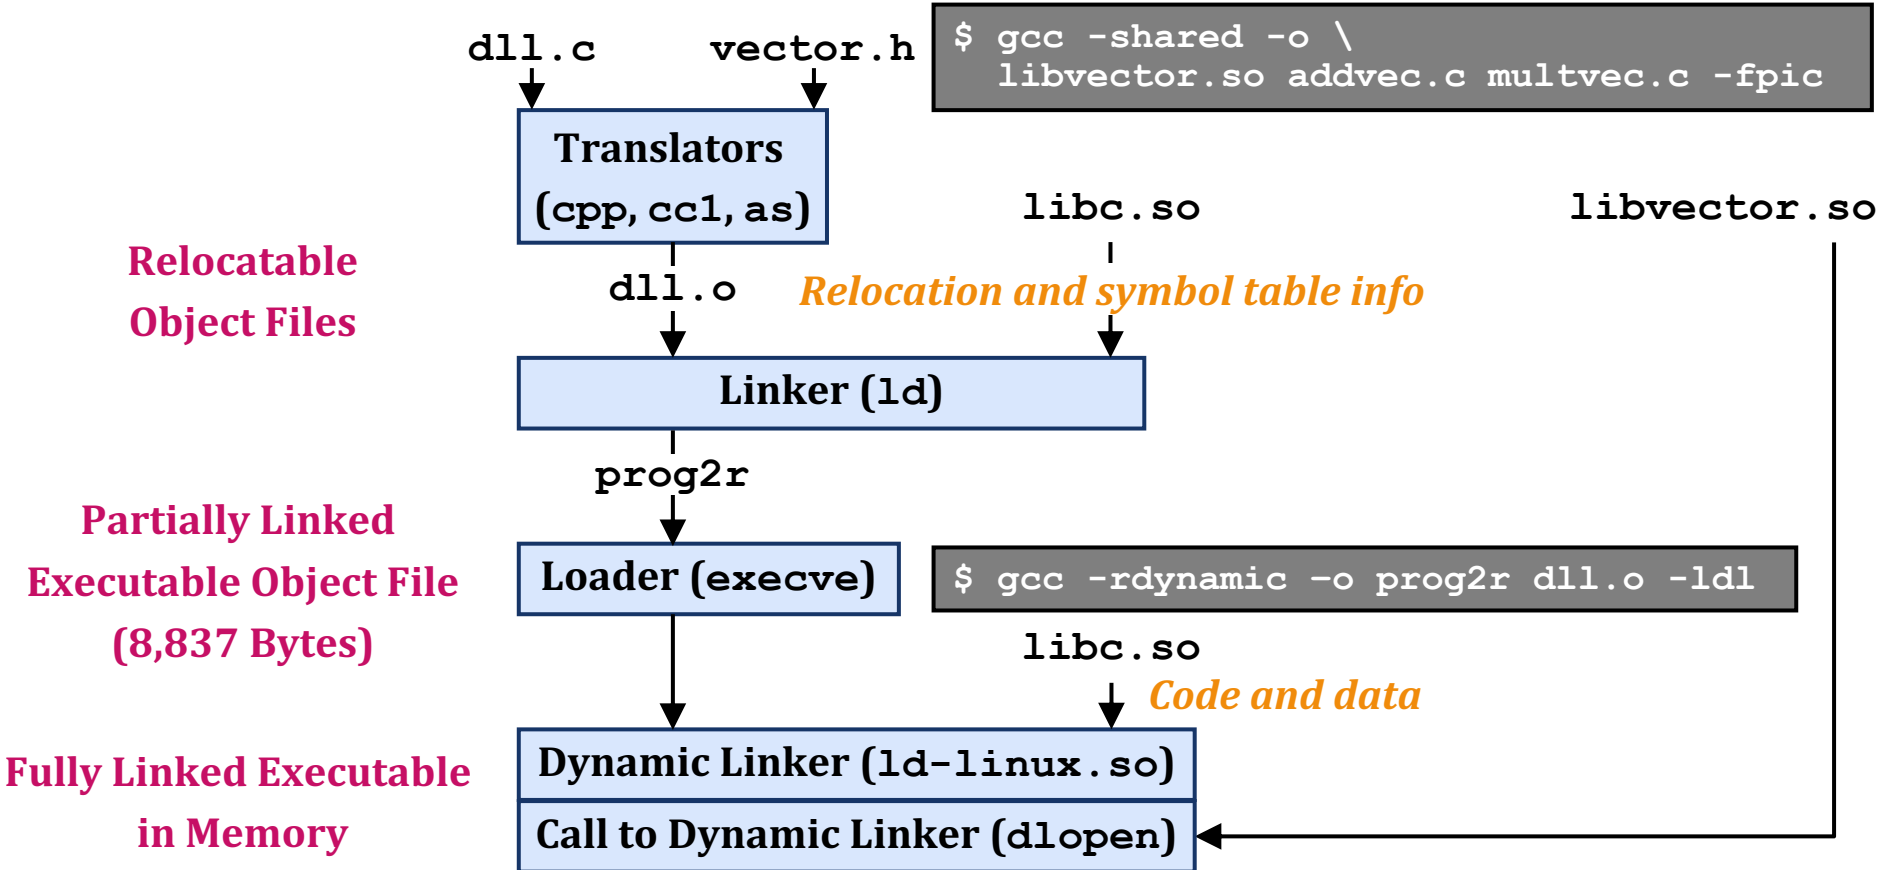


# Dynamic Linking at Run-Time

dll.c (Cont.)

```
...
/* get a pointer to the addvec() function we just loaded */
addvec = dlsym(handle, "addvec");
if((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}
/* Now we can call addvec() just like any other function */
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);
/* unload the shared library */
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}
return 0;
}
```

# Dynamic Linking at Run-Time



# Linking Summary

---

- Linking allows programs to be constructed from multiple object files
- Linking can happen at different times in a program's lifetime
  - Compile time (when a program is compiled)
  - Load time (when a program is loaded into memory)
  - Run time (while a program is executing)
- Understanding linking can help you avoid nasty errors and make you a better programmer

# Lecture Agenda: Linking

---

- Linking
  - Motivation
  - What It Does
  - How It Works
  - Dynamic Linking
- Case Study: Library Interpositioning

# Case Study: Library Interpositioning

---

- Powerful linking technique that allows programmers to intercept calls to arbitrary functions
- Interpositioning can occur at:
  - **Compile time**: when the source code is compiled
  - **Link time**: when the relocatable object files are statically linked to form an executable object file
  - **Load/run time**: when an executable object file is loaded into memory, dynamically linked, and then executed

# Some Interpositioning Applications

---

- Security
  - Confinement (sandboxing)
    - Interpose calls to `libc` functions
  - Behind the scenes encryption
    - Automatically encrypt otherwise unencrypted network connections
- Debugging
  - In 2014, two Facebook engineers debugged a treacherous 1-year old bug in their iPhone app using interpositioning
  - Code in the SPDY networking stack was writing to the wrong location
  - Solved by intercepting calls to Posix write functions (`write`, `writew`, `pwrite`)

Source: Facebook engineering blog post at

<https://code.facebook.com/posts/313033472212144/debugging-file-corruption-on-ios/>

# Some Interpositioning Applications (Cont.)

---

- Monitoring and profiling
  - Count number of calls to functions
  - Characterize call sites and arguments to functions
  - `malloc` tracing
    - Detecting memory leaks
    - Generating address traces
- Error checking
  - C Programming Lab used customized versions of `malloc/free` to do careful error checking

# Example program

- **Goal:** trace the addresses and sizes of the allocated and freed blocks
  - w/o breaking the program
  - w/o modifying the source code
- **Three solutions:** interpose on the library **malloc** and **free** functions at
  - Compile time
  - Link time
  - Load/run time

int.c

```
#include <stdio.h>
#include <malloc.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int i;
    for (i = 1; i < argc; i++) {
        void* p = malloc(atoi(argv[i]));
        free(p);
    }
    return 0;
}
```



# Compile-Time Interpositioning

```
#ifdef COMPILETIME
#include <stdio.h>
#include <malloc.h>

/* malloc wrapper function */
void* mymalloc(size_t size){
    void* ptr = malloc(size);
    printf("malloc(%d)=%p\n", (int)size, ptr);
    return ptr;
}

/* free wrapper function */
void myfree(void* ptr){
    free(ptr);
    printf("free(%p)\n", ptr);
}
#endif
```

mymalloc.c

# Compile-Time Interpositioning

```
#define malloc(size) mymalloc(size)
#define free(ptr) myfree(ptr)
```

mymalloc.h

```
void* mymalloc(size_t size);
void myfree(void* ptr);
```

```
$ make intc
gcc -Wall -DCOMPILETIME -c mymalloc.c
gcc -Wall -I. -o intc int.c mymalloc.o
```

```
$ make runc
```

```
./intc 10 100 1000
```

```
malloc(10)=0x1ba7010
```

```
free(0x1ba7010)
```

```
malloc(100)=0x1ba7030
```

```
free(0x1ba7030)
```

```
malloc(1000)=0x1ba70a0
```

```
free(0x1ba70a0)
```

Search for <malloc.h> leads to /usr/include/malloc.h

Search for <malloc.h> leads to

# Link-Time Interpositioning

```
mymalloc.c

#ifdef LINKTIME
#include <stdio.h>

void* __real_malloc(size_t size);
void __real_free(void* ptr);

/* malloc wrapper function */
void* __wrap_malloc(size_t size){
    void* ptr = __real_malloc(size); /* Call libc malloc */
    printf("malloc(%d) = %p\n", (int) size, ptr);
    return ptr;
}

/* free wrapper function */
void __wrap_free(void* ptr){
    __real_free(ptr); /* Call libc free */
    printf("free(%p)\n", ptr);
}
#endif
```

# Link-Time Interposition

```
$ make intl
gcc -Wall -DLINKTIME -c mymalloc.c
gcc -Wall -c int.c
gcc -Wall -Wl,--wrap,malloc -Wl,--wrap,free -o intl int.o mymalloc.o
$ make runl
./intl 10 100 1000
malloc(10) = 0x91a010
free(0x91a010)
...
```

Search for `<malloc.h>` leads to  
`/usr/include/malloc.h`

- The `-Wl` flag passes argument to linker, replacing each comma with a space
- The `--wrap,malloc` argument instructs linker to resolve references in a special way:
  - References to `malloc` should be resolved as `__wrap_malloc`
  - References to `__real_malloc` should be resolved as `malloc`

# Load/Run-Time Interpositioning

mymalloc.c

```
#ifdef RUNTIME
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h> } Observe no '#include <malloc.h>'
#include <dlfcn.h>

/* malloc wrapper function */
void* malloc(size_t size){
    void* (*mallocp)(size_t size);
    char* error;
    mallocp = dlsym(RTLD_NEXT, "malloc"); /* Get addr of libc malloc */
    if((error = dlerror()) != NULL){
        fputs(error, stderr);
        exit(1);
    }
    char* ptr = mallocp(size); /* Call libc malloc */
    printf("malloc(%d) = %p\n", (int) size, ptr);
    return ptr;
}
```

# Load/Run-Time Interpositioning

```
/* free wrapper function */
```

mymalloc.c (Cont.)

```
void free(void* ptr)
```

```
{
```

```
    void (*freep)(void*) = NULL;
```

```
    char* error;
```

```
    if(!ptr)
```

```
        return;
```

```
    freep = dlsym(RTLD_NEXT, "free"); /* Get address of libc free */
```

```
    if((error = dlerror()) != NULL) {
```

```
        fputs(error, stderr);
```

```
        exit(1);
```

```
    }
```

```
    freep(ptr); /* Call libc free */
```

```
    printf("free(%p)\n", ptr);
```

```
}
```

```
#endif
```

# Load/Run-Time Interpositioning

```
$ make intr
gcc -Wall -DRUNTIME -shared -fpic -o mymalloc.so mymalloc.c -ldl
gcc -Wall -o intr int.c
$ make runr
(LD_PRELOAD="./mymalloc.so" ./intr 10 100 1000)
malloc(10) = 0x91a010
free(0x91a010)
...
```

Search for `<malloc.h>` leads to `/usr/include/malloc.h`

- The `LD_PRELOAD` environment variable tells the dynamic linker to resolve unresolved refs (e.g., to `malloc`) by looking in `mymalloc.so` first
- Type into (some) shells as:  
(`setenv LD_PRELOAD "./mymalloc.so"; ./intr 10 100 1000`)

# Interpositioning Summary

---

- **Compile time**

- Apparent calls to `malloc/free` get macro-expanded into calls to `mymalloc/myfree`
- Simple approach that must have access to source & recompile

- **Link time**

- Use linker trick to have special name resolutions
  - `malloc` → `__wrap_malloc`
  - `__real_malloc` → `malloc`

- **Load/Run time**

- Implement custom version of `malloc/free` that use dynamic linking to load library `malloc/free` under different names
- Can use with **any** dynamically linked binary  
(`setenv LD_PRELOAD "./mymalloc.so"; gcc -c int.c`)



# Linking Summary

---

- Usually: just happens, no big deal
- Sometimes: strange errors
  - Bad symbol resolution
  - Ordering dependence of linked `.o`, `.a`, and `.so` files
- For power users: interpositioning to trace programs w/ or w/o source

# [CSED211] Introduction to Computer Software Systems

## Lecture 12: Linking

Prof. Jisung Park



**CAOS**  
COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2023.11.13