

[CSED211] Introduction to Computer Software Systems

Lecture 13: Exceptional Control Flow – Exceptions and Processes

Prof. Jisung Park



CAOS
COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.11.22


Lecture Agenda

- Exceptional Control Flow
- Exceptions
- Processes
- Process Control

Control Flow

- Processors do **only one thing**:
 - From startup to shutdown, a CPU simply **reads and executes (interprets) a sequence of instructions**, one at a time
 - This sequence is the CPU's **control flow** (or **flow of control**)

Physical control flow



Altering the Control Flow

- **Up to now:** two mechanisms for changing control flow
 - Jumps and branches
 - Call and returnBoth react to changes in **program state**
- Insufficient for a useful system

Difficult to react to changes in **system state**

 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires
 - ...
- System needs mechanisms for **exceptional control flow**

Exceptional Control Flow

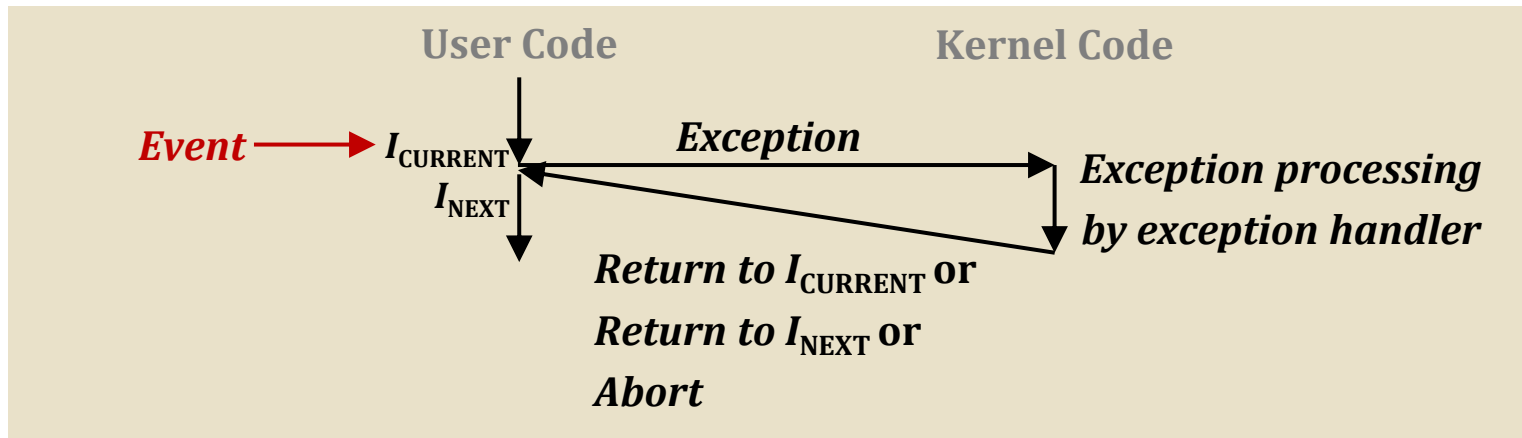
- Exists at all levels of a computer system
- Low level mechanisms
 - 1. **Exceptions**
 - Control-flow change in response to a system event (i.e., system-state change)
 - Implemented using combination of hardware and OS software
- Higher level mechanisms
 - 2. **Process context switch**
 - Implemented by OS software and hardware timer
 - 3. **Signals**
 - Implemented by OS software
 - 4. **Nonlocal jumps**: `setjmp()` and `longjmp()`
 - Implemented by C runtime library

Lecture Agenda

- Exceptional Control Flow
- Exceptions
- Processes
- Process Control


Exceptions

- An **exception** requires a transfer of control to the **OS kernel** in response to a **specific event** (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: divide by zero, arithmetic overflow, page fault, I/O request completions, typing Ctrl-C, etc.



Exception Table

- Data structure to point the handling routine for each type of event
 - Each type of event (i.e., exception) has a unique number k as its ID (set by either CPU or OS)
 - k = index into the exception table (a.k.a. **interrupt vector**)
 - Handler k is called each time exception k occurs

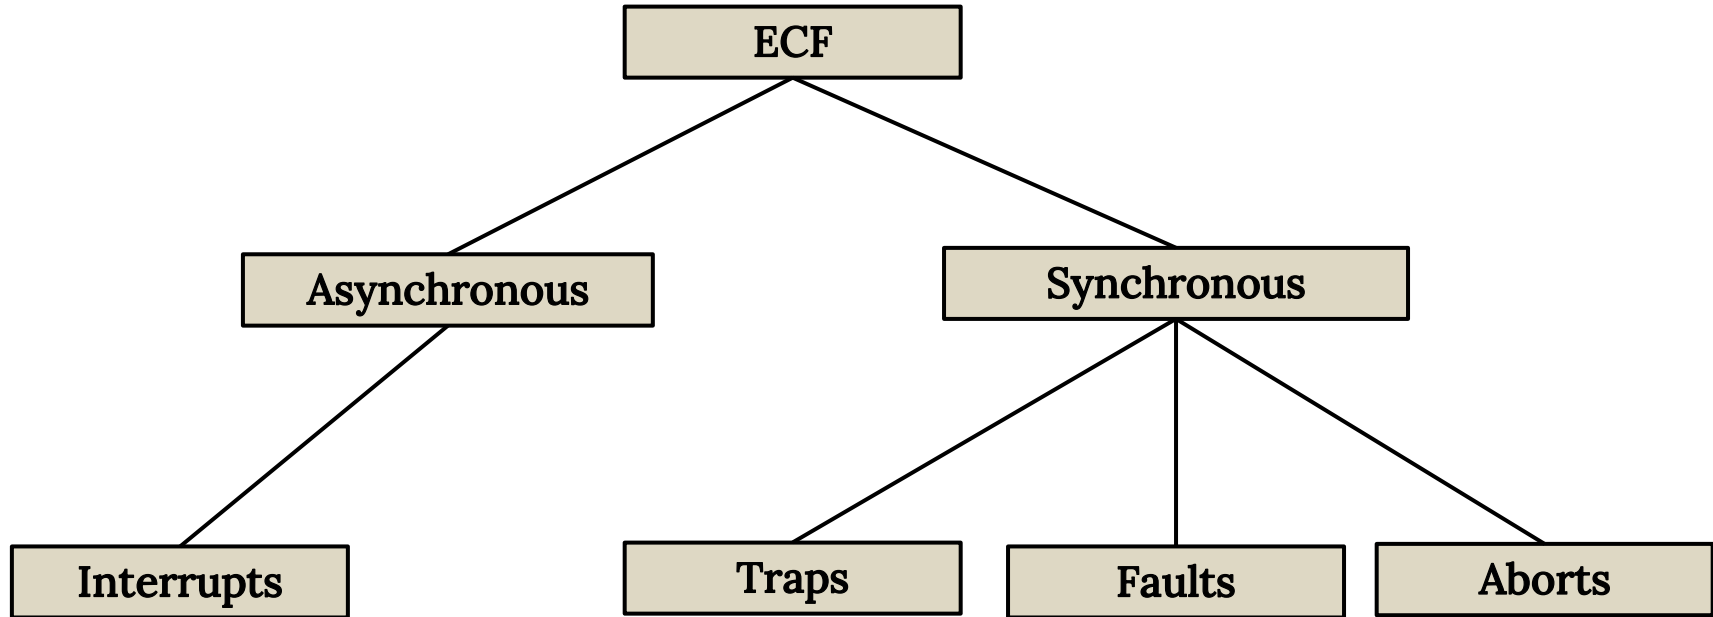


Exception Table: IA32 (Excerpt)

- Check Intel manual for more information

Exception Number	Description	Exception Class
0	Divide error	Fault
13	General protection fault	Fault
14	Page fault	Fault
18	Machine check	Abort
32–127	OS-defined	Interrupt or trap
128 (0x80)	System call	Trap
129–255	OS-defined	Interrupt or trap

(Partial) Taxonomy



Asynchronous Exceptions (a.k.a. Interrupts)

- Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - Handler returns to the **next instruction**
- Examples
 - **Timer interrupt**
 - Every few milliseconds, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - **I/O interrupt from external devices**
 - Hitting Ctrl-C at the keyboard
 - Arrival of a packet from a network
 - Arrival of data from a disk

Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction
 - **Traps**
 - Intentional
 - Examples: **system calls**, breakpoint traps, special instructions
 - Returns control to **the next instruction**
 - **Faults**
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating-point exceptions
 - Either re-executes faulting (i.e., **current**) instruction or aborts
 - **Aborts**
 - Unintentional and unrecoverable
 - Examples: parity error, machine check
 - Aborts current program

System Calls

- Each x86-64 system call has a unique ID number

Number	Name	Description
0	read	Read a file
1	write	Write a file
2	open	Open a file
3	close	Close a file
4	stat	Get info about a file
57	fork	Create a process
59	execve	Execute a program
60	_exit	Terminate a process
62	kill	Send signal to a process

System Call Example: File Open

- User calls: `open(filename, options)`
 - Calls `__open` function, which invokes system call instruction `syscall`

```
00000000000e5d70 <__open>:
```

```
...
```


```
e5d79:  b8 02 00 00 00    sub    $0x2,%eax                # open is syscall #2
```

```
e5d7e:  0f 05              syscall                          # Return value in %rax
```

```
e5d80:  48 3d 01 f0 ff     cmp    $0xffffffffffff001,%rax
```

```
...
```

```
e5dfa:  c3                retq
```




- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, and `%r9`
- Return value in `%rax`
- Negative return value indicates an error corresponding to negative `errno`

System Call Example: File Open

- User calls: `open(filename, ...)`
 - Calls `__open` function, which...

```
00000000000e5d70 <__open>:
```

```
...  
e5d79:  b8 02 00 00 00    sub    $0x2, %rax  
e5d7e:  0f 05             syscal  
e5d80:  48 3d 01 f0 ff    cmp    $0xffffffff, %rax  
...  
e5dfa:  c3               retq
```



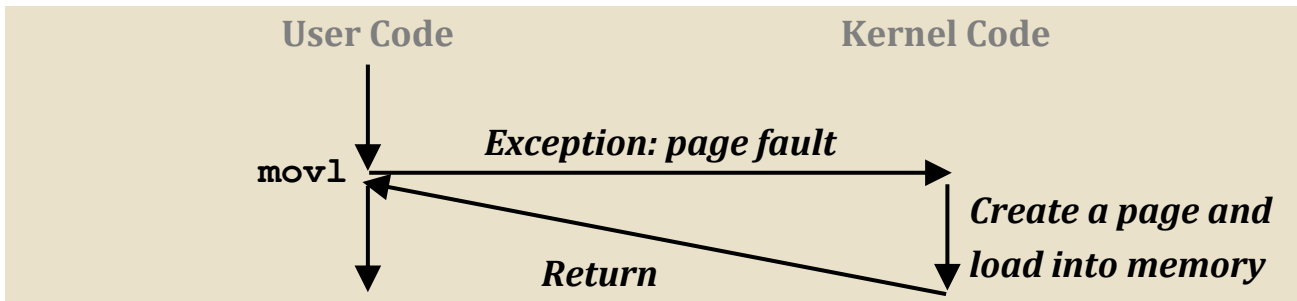
- Almost like a function call
 - Transfer of control
 - On return, executes the next instruction
 - Passes arguments using the same calling convention
 - Gets result in `%rax`
- One important exception
 - Executed by kernel
 - Different set of privileges
- And other differences
 - **Function address** is in `%rax`
 - Uses `errno`
 - etc.

Fault Example: Page Fault

- User writes to a memory location
 - When portion (page) of user's memory is currently on disk. Not in memory.

```
int a[1000];  
main() {  
    a[500]=13;  
}
```

80483b7: c7 05 10 9d 04 08 0d movl \$0xd,0x8049d10




- Page handler must load the page into physical memory
- Returns to **the faulting instruction**
- Successful on second try

Fault Example: Invalid Memory Reference

- User writes to an **invalid** memory location

```
int a[1000];  
main() {  
    a[5000]=13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360




- Page handler detects invalid address
- Sends **SIGSEGV signal** to user process
- User process exits with **segmentation fault**

Lecture Agenda

- Exceptional Control Flow
- Exceptions
- **Processes**
- Process Control

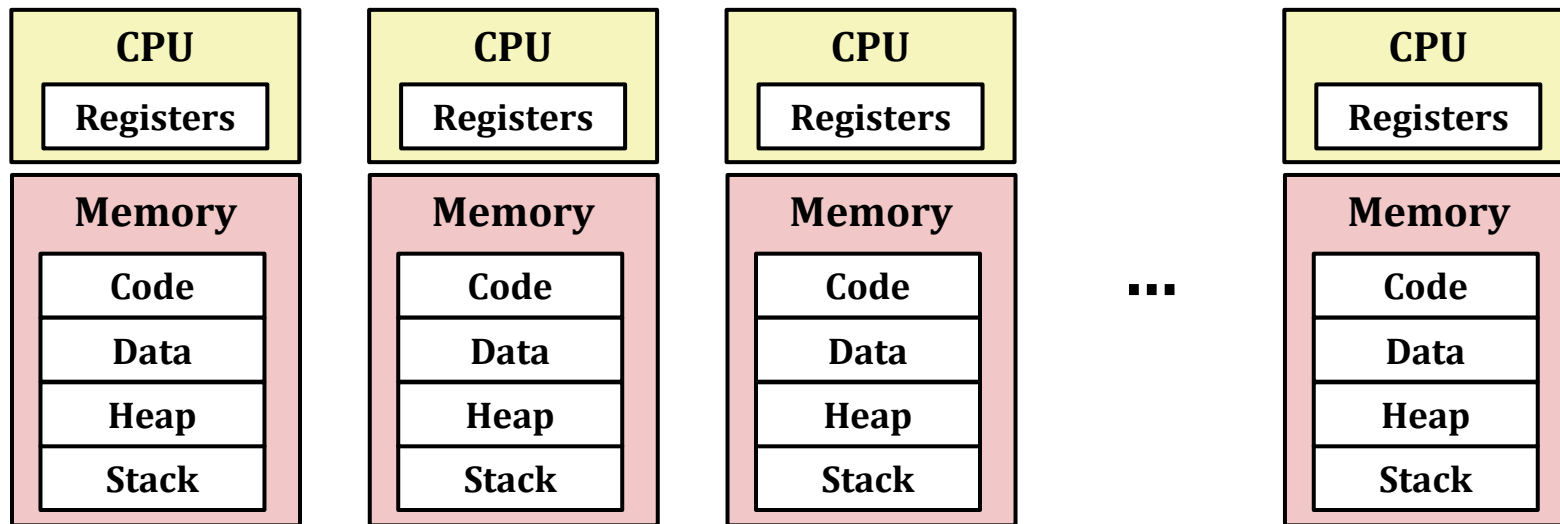
Processes

- **An instance of running program**
 - One of the most profound ideas in computer science
 - **Not** the same as **program** or **processor**
- Process provides each program with **two key abstractions**
 - **Logical control flow**
 - Each program seems to have **exclusive use of the CPU**
 - Provided by kernel mechanism called **context switching**
 - **Private address space**
 - Each program seems to have **exclusive use of main memory**
 - Provided by kernel mechanism called **virtual memory**



Multiprocessing: The Illusion

- Computer runs many processes simultaneously
 - Applications for one or more users: e.g., web browsers, email clients, editors, etc.
 - Background tasks: e.g., monitoring network & I/O devices



Multiprocessing Example

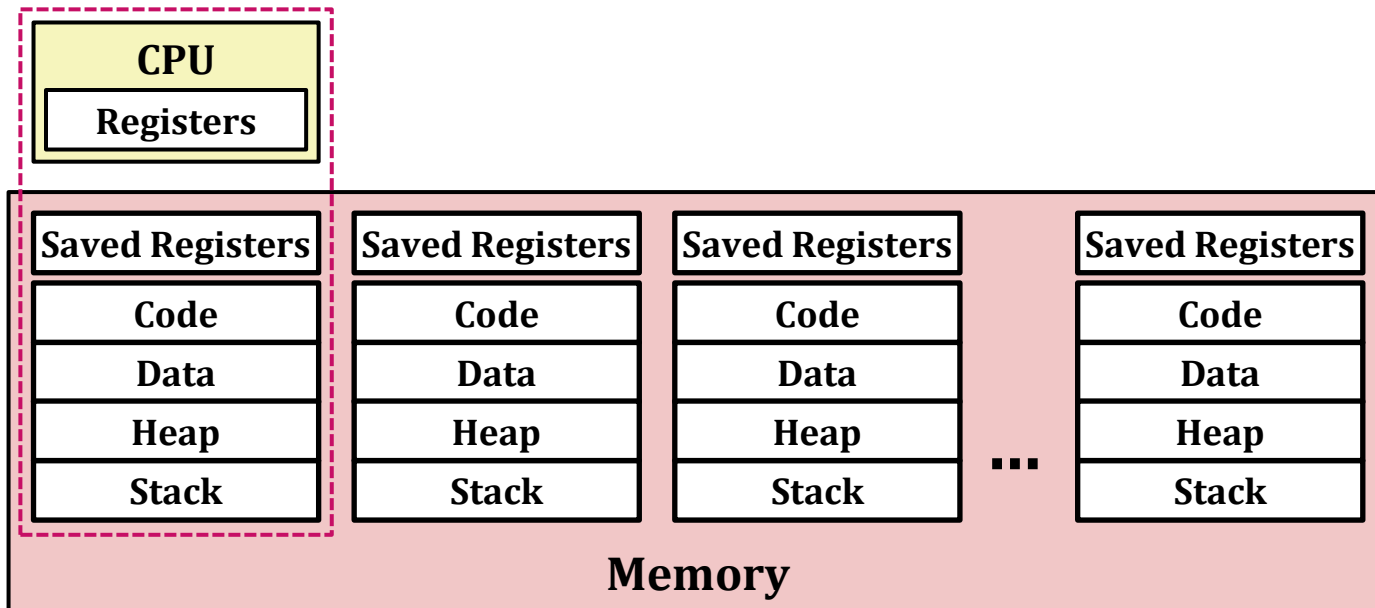
- Running program **top** on Mac
 - System has 810 processes, five of which are active
 - Identified by Process ID (PID)

```
jisung — top — 111x33
Processes: 810 total, 5 running, 805 sleeping, 4116 threads
Load Avg: 2.29, 2.18, 2.38  CPU usage: 16.46% user, 2.62% sys, 80.90% idle
SharedLibs: 797M resident, 152M data, 156M linkedit.
MemRegions: 200354 total, 5896M resident, 500M private, 5721M shared.
PhysMem: 23G used (2968M wired, 2498M compressor), 604M unused.
VM: 310T vsize, 4729M framework vsize, 0(0) swapins, 0(0) swapouts.
Networks: packets: 3653323/3213M in, 2249336/882M out. Disks: 2466133/39G read, 1197430/23G written.

PID    COMMAND      %CPU    TIME    #TH    #WQ    #PORT  MEM    PURG    CMPRS    PGRP    PPID    STATE    BOOSTS
1629   HwpMac2014   99.2    12:19:51 3/1    1      210    27M    0B      21M     1629    1      running *1534[3]
561    WindowServer 7.3     20:03:43 17     6      4081-  2085M+ 324M-  70M     561    1      sleeping *0[1]
1975   Discord Help 5.6     36:02:48 44     2      715    269M+  0B      75M     1625   1625   sleeping *1[4]
53296  top          4.3     00:01:27 1/1    0      29     13M    0B      0B      53296  53280  running  *0[1]
0      kernel_task  3.5     28:37:85 569/8  0      0       16M    0B      0B      0      0      running  0[0]
1161   WindowManage 2.8     01:18:28 5      2      385    38M    0B      3776K   1161    1      sleeping *0[5185+]
1      launchd      2.4     21:44:08 3/1    2/1    4187   28M-   0B      4720K   1      0      running  0[0]
53315  screencaptur 2.2     00:00:14 5      3      170-   13M+   0B      0B      53315  1      sleeping *0[8+]
940    distnoted    1.5     10:17:13 2      1      650-   5025K  0B      832K    940    1      sleeping *0[1]
2221   nosmain      1.3     09:57:39 5      1      207    34M    0B      19M     2175   1      sleeping *0[1]
53318  screencaptur 1.2     00:00:05 7      6      74-    7346K+ 16K     0B      1633   1633   sleeping *0[105+]
554    bluetoothd   1.1     03:53:77 11     5      412    16M    240K    6912K   554    1      sleeping *0[1]
1620   Microsoft Po 0.8     13:33:01 33     8      1896   1260M  103M    200M    1620    1      sleeping *0[2730]
1619   AdobeAcrobat 0.8     06:45:49 28     4      388    326M   0B      159M    1619    1      sleeping 0[2335]
535    launchservic 0.7     01:37:05 5      4      887    9777K  0B      816K    535    1      sleeping *2550[318]
545    distnoted    0.6     04:03:04 2      1      236    1969K  0B      528K    545    1      sleeping *0[1]
493    logd          0.6     05:56:01 4/1    3      2594   14M    0B      13M     493    1      running  *0[1]
1792   Discord Help 0.6     04:55:87 13     2      213    337M   13M     35M     1625   1625   sleeping *1[5]
1753   Microsoft Te 0.5     04:38:59 12     1      239    226M   6160K   22M     1622   1622   sleeping *1[5]
753    BTLEServer   0.5     00:44:78 3      2      69     4290K  0B      1248K   753    1      sleeping *0[1]
1939   LogiTune     0.5     01:38:37 32     1      381    62M    0B      32M     1939    1      sleeping *0[1441]
1927   logioptionsp 0.4     02:12:43 260    1      832    95M    0B      29M     1927    1      sleeping *0[1]
1941   Dropbox      0.4     01:42:15 134    1      741    336M   80K     88M     1941    1      sleeping *0[2438+]
1616   Google Chrom 0.4     04:32:51 44     1      933    207M   0B      53M     1616    1      sleeping *0[1863]
```

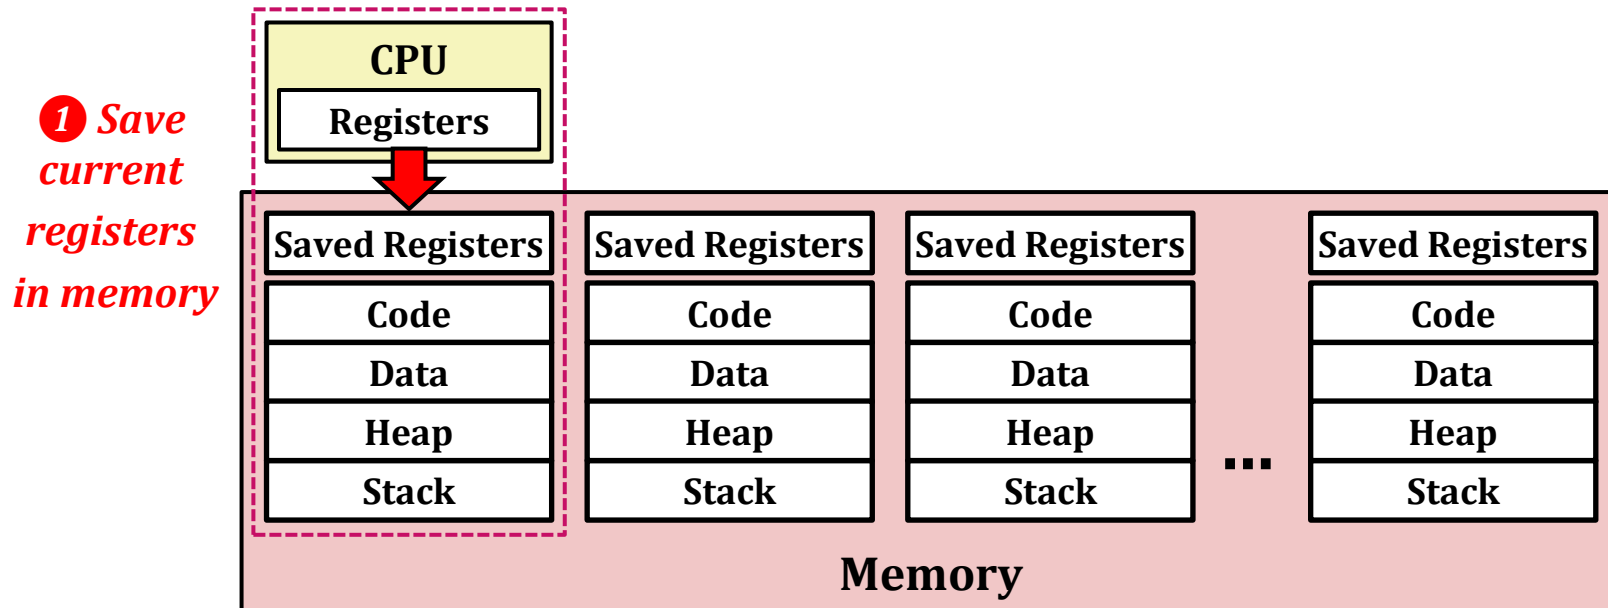
Multiprocessing: The (Traditional) Reality

- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory



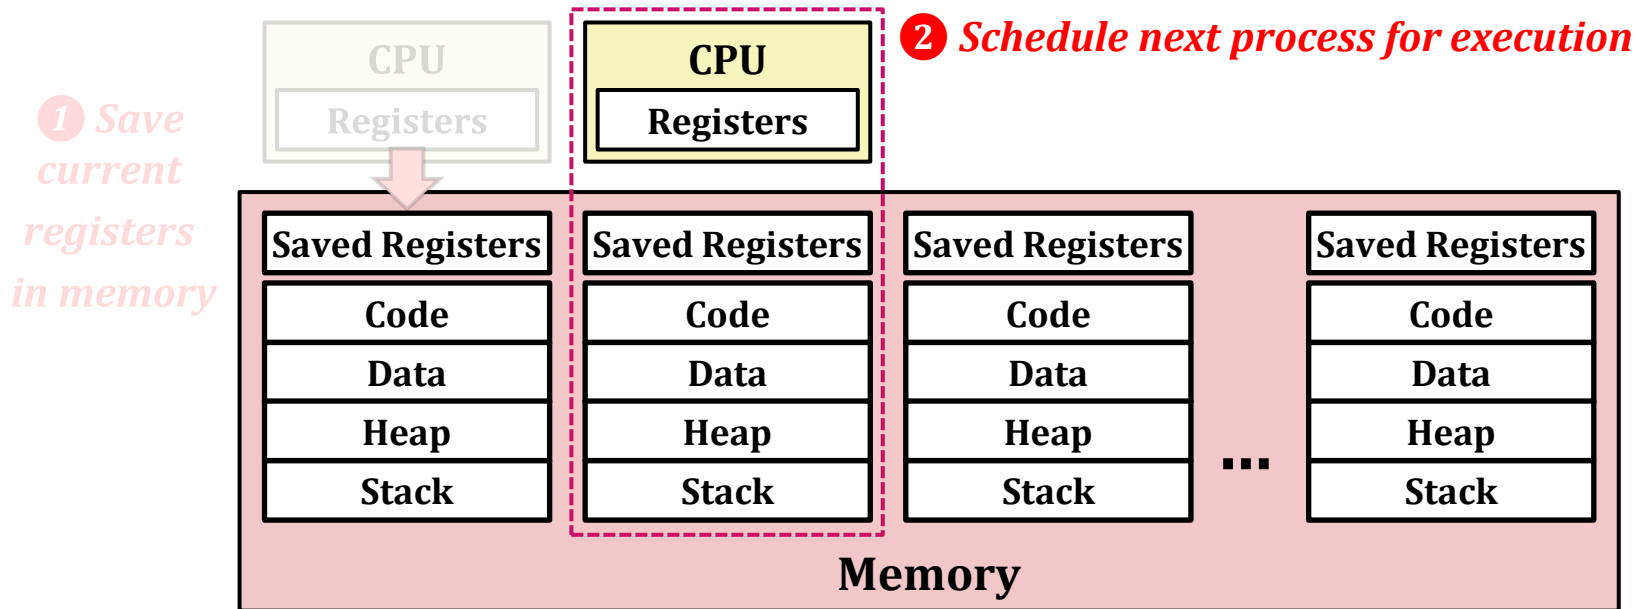
Multiprocessing: The (Traditional) Reality

- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory



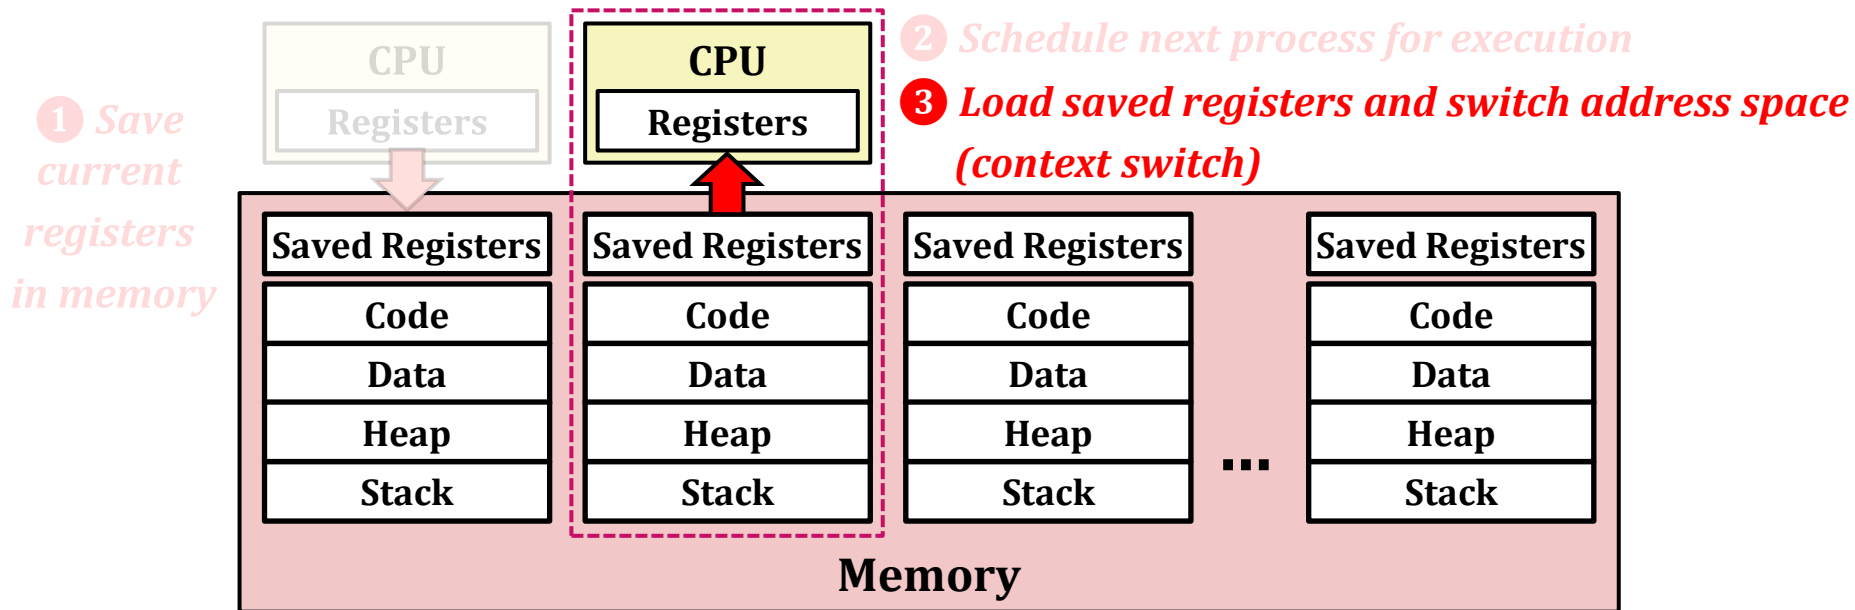
Multiprocessing: The (Traditional) Reality

- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory



Multiprocessing: The (Traditional) Reality

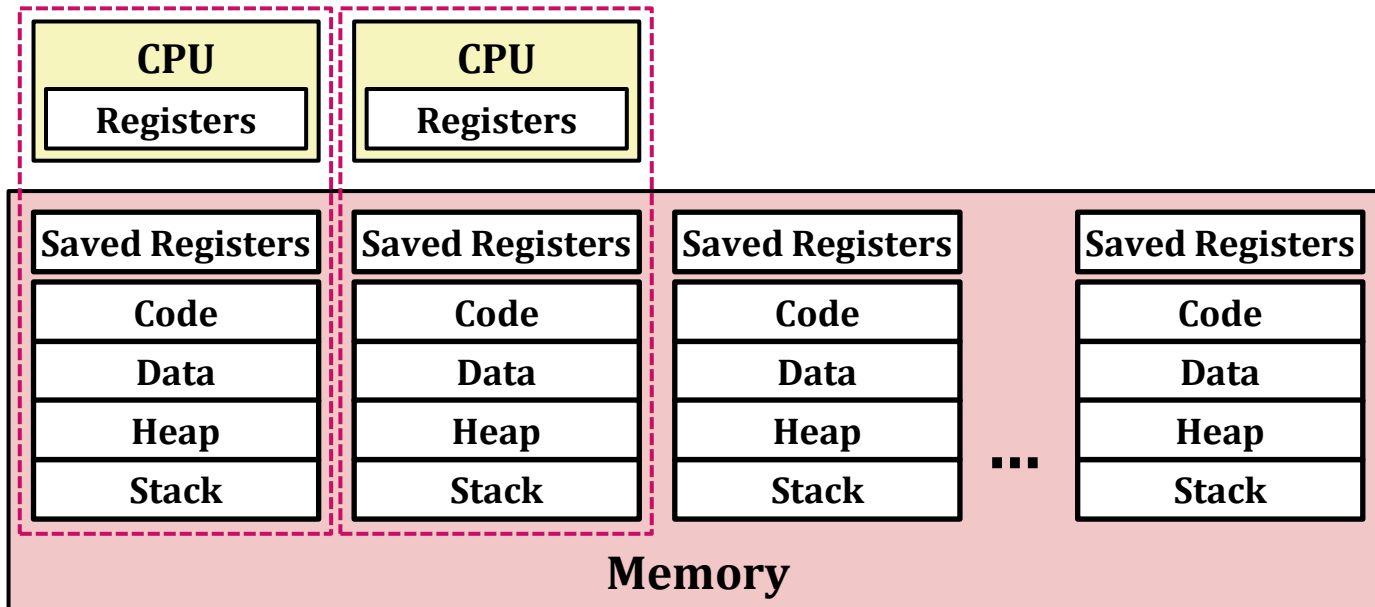
- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory



Multiprocessing: The (Traditional) Reality


- **Multicore processors**

- Multiple CPUs on single chip
- Share main memory (and some of the caches)
- Each can execute a separate process: kernel schedules processors onto cores




Concurrent Processes

- Each process is a logical control flow
- Two processes **run concurrently** if their flows overlap in time
 - Otherwise, they are **sequential**
- Examples (running on single core)
 - Concurrent: A & B, A & C
 - Sequential: B & C



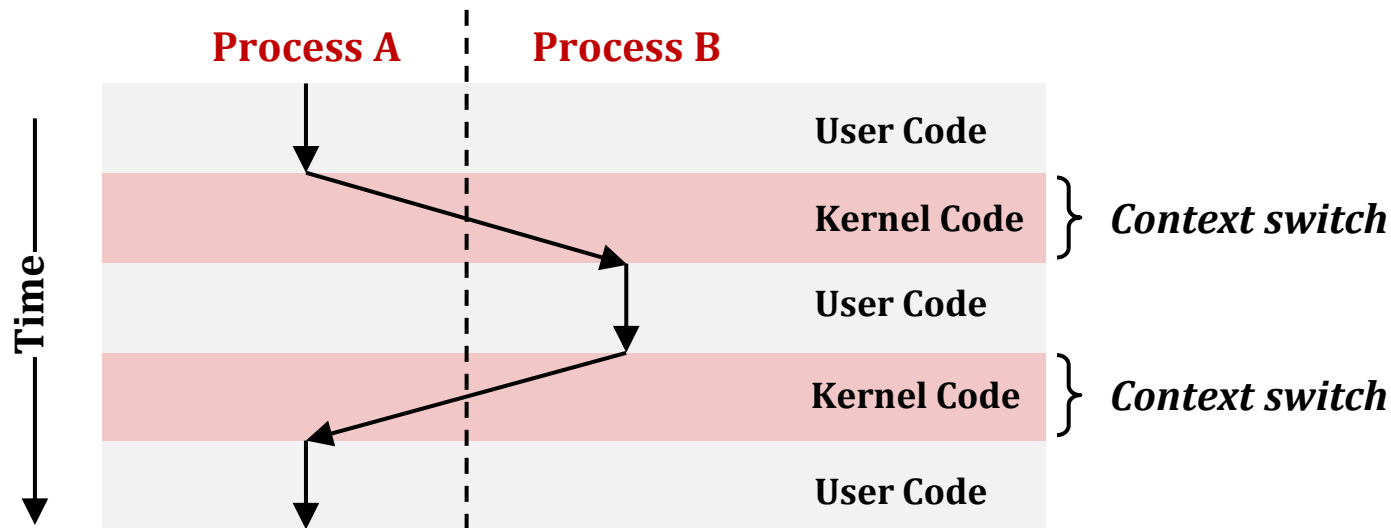
Concurrent Processes: User View

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other



Context Switching

- Processes are managed by a shared chunk of OS code called kernel
 - The kernel is not a separate process, but rather runs as part of some user process
- Control flow passes from one process to another via context switch



Lecture Agenda

- Exceptional Control Flow
- Exceptions
- Processes
- Process Control

System Call Error Handling

- On error, Linux system-level functions typically **return -1** and set global variable **errno** to indicate cause.
- Hard and fast rule
 - Must check the return status of every system-level function
 - Only exception is the handful of functions that return **void**
- Example

```
if((pid = fork()) < 0){  
    fprintf(stderr, "fork error: %s\n", strerror(errno));  
    exit(0);  
}
```

Error-Reporting Functions

- Can simplify somewhat using an **error-reporting function**

```
/* Unix-style error */  
void unix_error(char *msg) {  
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));  
    exit(-1);  
}
```

```
if((pid = fork()) < 0)  
    unix_error("fork error");
```

- But must think about application; not always appropriate to exit when something goes wrong.

Error-Handling Wrappers

- Simplify the code even further using Stevens-style error-handling wrappers

```
pid_t Fork(void) {  
    pid_t pid;  
  
    if((pid = fork()) < 0)  
        unix_error("fork error");  
    return pid;  
}
```

```
pid = Fork();
```

- **NOT** what you generally want to do in a real application

Obtaining Process IDs

- `pid_t getpid(void)`
 - Returns PID of current process
- `pid_t getppid(void)`
 - Returns PID of parent process

Creating and Terminating Processes

- From a programmer's perspective, a process is in **one of three states**
- **Running**: the process is either executing or waiting to be executed (and will eventually be scheduled (i.e., chosen to execute) by the kernel)
- **Stopped**: the process execution is suspended and will not be scheduled until further notice (next lecture when we study signals)
- **Terminated**: the process is stopped permanently

Terminating Processes

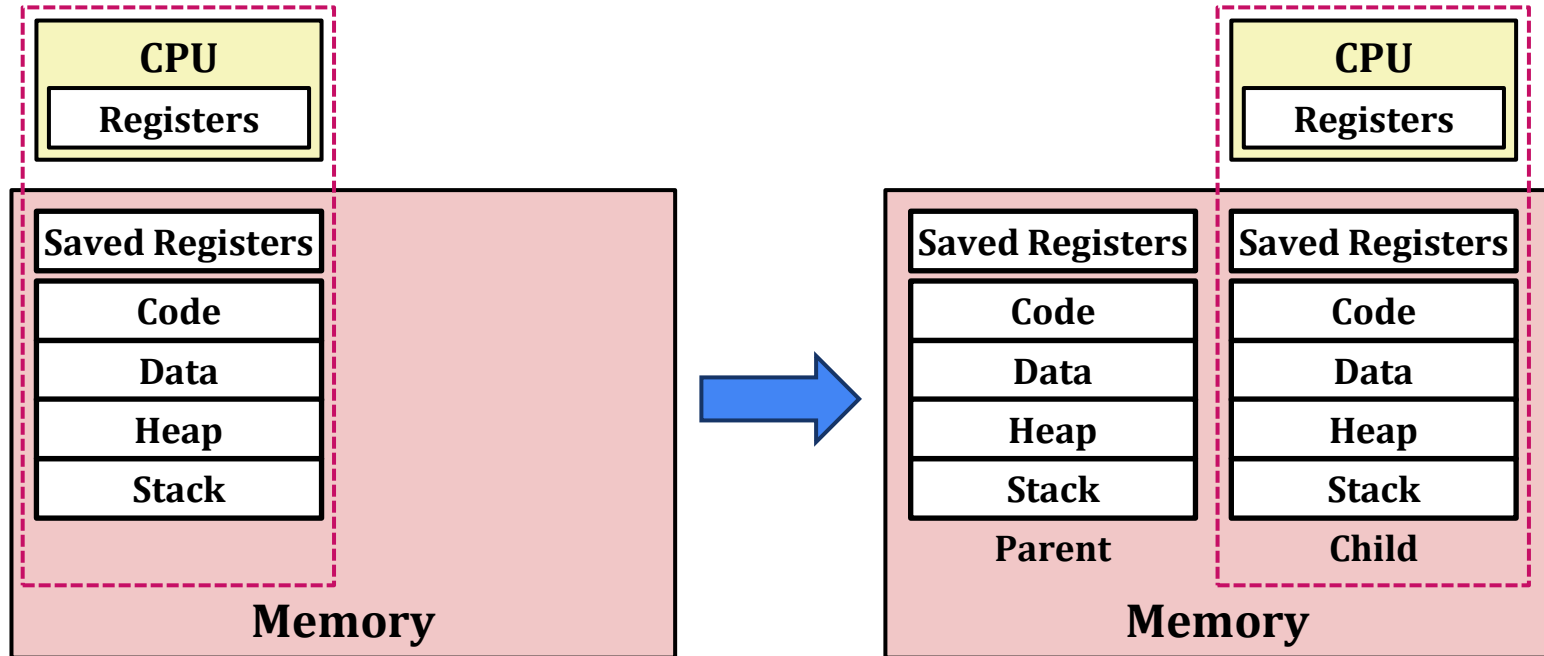
- Process becomes terminated for one of three reasons:
 - Receiving a signal whose default action is to terminate (next lecture)
 - Returning from the `main` routine
 - Calling the `exit` function
- `void exit(int status)`
 - Terminates with an `exit status` of `status`
 - Convention: normal return status is 0, nonzero on error
 - Another way to explicitly set the exit status is to return an integer value from the main routine
- `exit` is called **once** but **never** returns

Creating Processes

- **Parent process** creates a new running **child process** by calling **fork**
- **int fork(void)**
 - Returns **0 to the child process** and **child's PID to parent process**
 - Child is **almost identical** to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- **fork** is interesting (and often confusing): it is **called once** but **returns twice**

Conceptual View of fork

- Make complete copy of execution state
 - Designate one as parent and one as child
 - Resume execution of parent or child



fork Example

```
int main(int argc, char** argv){
    pid_t pid;
    int x = 1;

    pid = Fork();
    if(pid == 0) {
        /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }
    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
fork.c
```

- Call once, return twice
- Concurrent execution: **cannot** predict execution order of parent and child

```
$ ./fork
parent: x=0
child : x=2
```

```
$ ./fork
child : x=2
parent: x=0
```

```
$ ./fork
parent: x=0
child : x=2
```

```
$ ./fork
parent: x=0
child : x=2
```

```
$ ./fork
child : x=2
parent: x=0
```

...

Making fork More Nondeterministic

- Problem

- Linux scheduler does not create much run-to-run variance
- Hides potential race conditions in nondeterministic programs
 - e.g., does fork return to child first or to parent?

- Solution

- Create custom version of library routine that inserts **random delays** along different branches
 - e.g., for parent and child in fork
- Use **runtime interpositioning** to have program use special version of library code

Variable Delay fork

```
/* fork wrapper function */
pid_t fork(void) {
    initialize();
    int parent_delay = choose_delay();
    int child_delay = choose_delay();
    pid_t parent_pid = getpid();
    pid_t child_pid_or_zero = real_fork();
    if(child_pid_or_zero > 0) { /* Parent */
        if(verbose) {
            printf("Fork. Child pid=%d, delay=%dms. Parent pid=%d, delay=%dms\n",
                child_pid_or_zero, child_delay, parent_pid, parent_delay);
            fflush(stdout);
        }
        ms_sleep(parent_delay);
    } else { /* Child */
        ms_sleep(child_delay);
    }
    return child_pid_or_zero;
}
```

myfork.c

forkx2 Example

```
int main(int argc, char** argv){
    pid_t pid;
    int x = 1;

    pid = Fork();
    if(pid == 0){
        /* Child */
        printf("child : x=%d\n", ++x);
        printf("child : x=%d\n", ++x);
        return 0;
    }
    /* Parent */
    printf("parent: x=%d\n", --x);
    printf("parent: x=%d\n", --x);
    return 0;
}
```

```
$ ./fork
```

```
parent: x=0
```

```
parent: x=-1
```

```
child : x=2
```

```
child : x=3
```

```
forkx2.c
```

- Call once, return twice
- Concurrent execution: **cannot** predict execution order of parent and child
- Duplicate but separate address space
 - **x** has a value of 1 when fork returns in parent and child
 - Subsequent changes to **x** are independent
- Shared open files
 - **stdout** is the same in both parent and child

Making fork with Process Graphs


- **Process graph**: a useful tool to capture the partial ordering of statements in a concurrent program
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means **a happens before b**
 - Edges can be labeled with current value of variables
 - `printf` vertices can be labeled with output
 - Each graph begins with a vertex with no inward edge
- Any **topological sort** of the graph corresponds to a feasible total ordering
 - Total ordering of vertices where all edges point from left to right

Making fork with Process Graphs

```
int main(int argc, char** argv){
    pid_t pid;
    int x = 1;


    pid = Fork();
    if(pid == 0) {
        /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }
    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

fork.c




Interpreting Process Graphs

- Original graph



- Relabeled graph



Feasible total ordering:



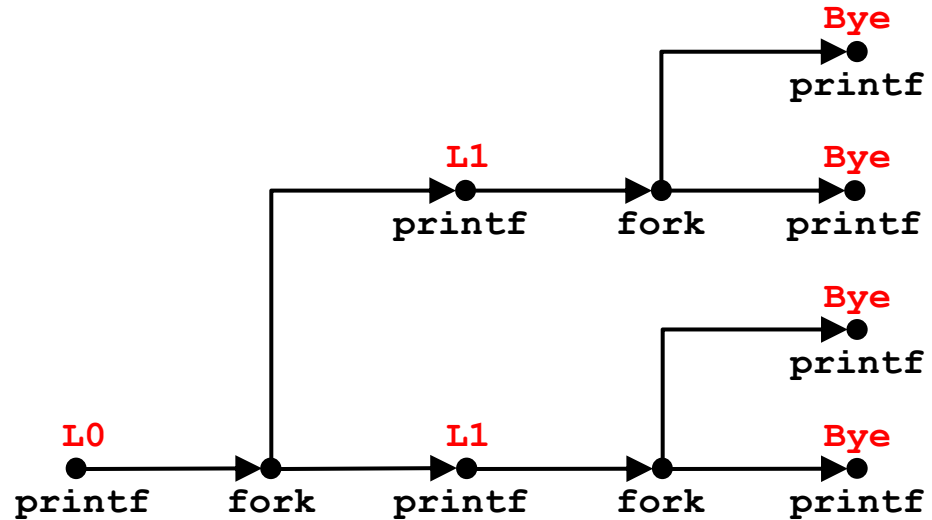
Infeasible total ordering:



fork Example: Two Consecutive forks

```
void fork2() {  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```

forks.c




- **Feasible output:** L0 → L1 → Bye → Bye → L1 → Bye
- **Infeasible output:** L0 → Bye → L1 → Bye → Bye → L1

fork Example: Nested forks in Parent

```
void fork4() {  
    printf("L0\n");  
    if(fork() != 0){  
        printf("L1\n");  
        if(fork() != 0){  
            printf("L2\n");  
        }  
    }  
    printf("Bye\n");  
}
```

forks.c

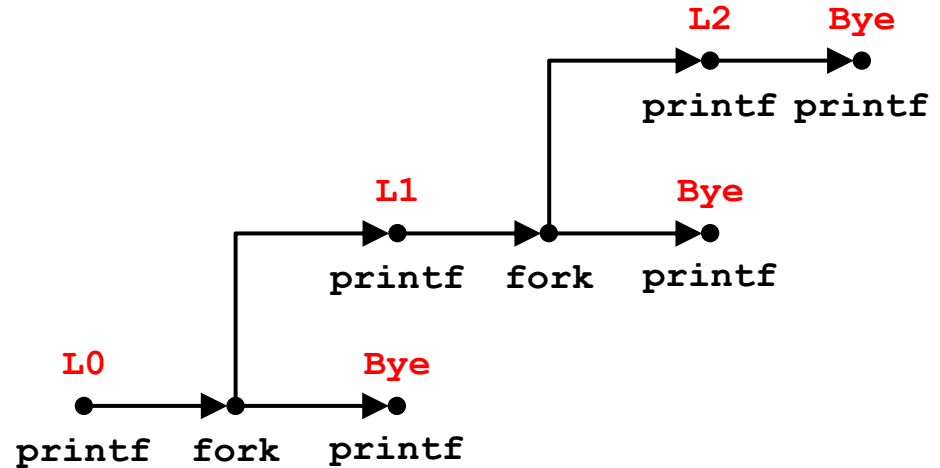


- **Feasible output:** L0 → L1 → Bye → Bye → L2 → Bye
- **Infeasible output:** L0 → Bye → L1 → Bye → Bye → L2

fork Example: Nested forks in Children

```
void fork4() {  
    printf("L0\n");  
    if(fork() == 0){  
        printf("L1\n");  
        if(fork() == 0){  
            printf("L2\n");  
        }  
    }  
    printf("Bye\n");  
}
```

forks.c



- **Feasible output:** L0 → Bye → L1 → L2 → Bye → Bye
- **Infeasible output:** L0 → Bye → L1 → Bye → Bye → L2

Reaping Child Processes

- Idea
 - When process terminates, it still consumes system resources
 - e.g., exit status, various OS tables
 - Called a **zombie**
 - Living corpse, half alive and half dead
- Reaping
 - Performed by parent on terminated child (using `wait` or `waitpid`)
 - Parent is given exit status information
 - Kernel then deletes zombie child process

Reaping Child Processes

- Idea: when process terminates, it still consumes system resources
 - Examples: exit status, various OS tables
 - Called a **zombie**: living corpse, half alive and half dead
- **Reaping**
 - Performed by parent on terminated child (using `wait` or `waitpid`)
 - Parent is given exit status information
 - Kernel then deletes zombie child process
- What if parent does not reap?
 - If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
 - So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
void fork7() {                                forks.c
    if(fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
               getpid());
        exit(0);
    } else {
        /* Parent */
        printf("Running Parent, PID = %d\n",
               getpid());
        while(1); /* Infinite loop */
    }
}
```

```
$ ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
$ ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6639 ttyp9        00:00:03 forks
 6640 ttyp9        00:00:00 forks <defunct>
 6641 ttyp9        00:00:00 ps
$ kill 6639
[1] Terminated
$ ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6642 ttyp9        00:00:00 ps
```

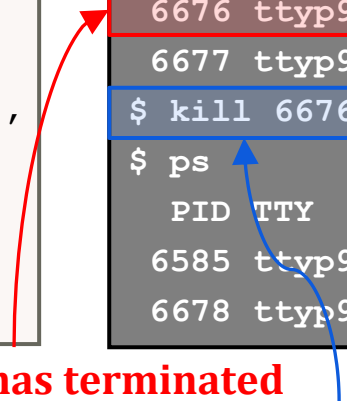
ps shows child process as “defunct” (i.e., a zombie)

Killing parent allows child to be reaped by init

Non-Terminating Child Example

```
void fork8() {                                forks.c
    if(fork() == 0){
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while(1); /* Infinite loop */
    } else {
        /* Parent */
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

```
$ ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
$ ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6676 ttyp9        00:00:06 forks
 6677 ttyp9        00:00:00 ps
$ kill 6676
$ ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6678 ttyp9        00:00:00 ps
```

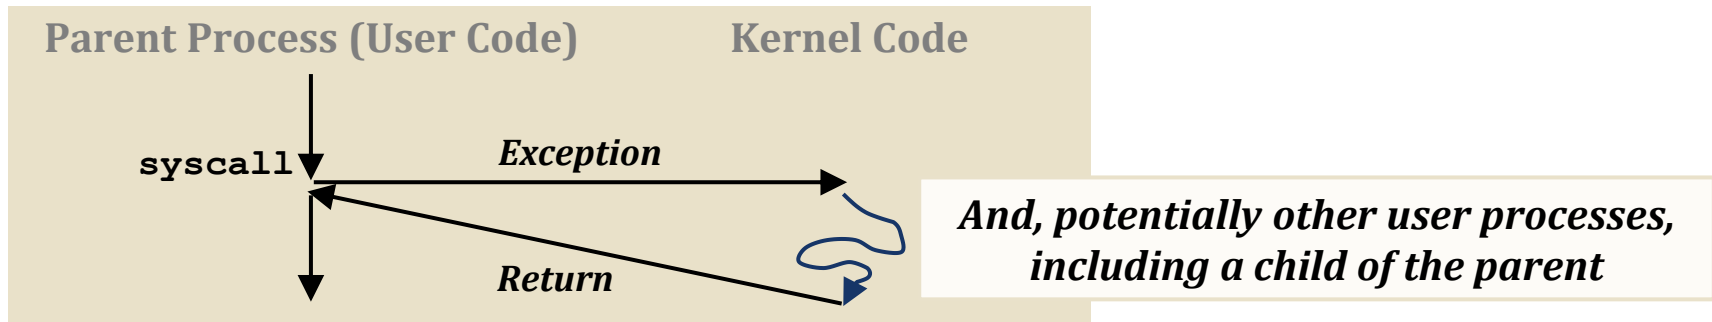


Child process still active even though parent has terminated

Must kill child explicitly; otherwise it will keep running indefinitely

wait: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Implemented as syscall




`wait`: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Implemented as syscall
 - Return value is the PID of the terminated child process
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`
 - See textbook for more details

wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if(fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

forks.c



- **Feasible output:** HC → HP → CT → Bye
- **Infeasible output:** HP → CT → Bye → HC

Another wait Example

- If multiple children completed, will take in arbitrary order
- **WIFEXITED** and **WEXITSTATUS** macros to get information about exit status

```
void fork10() {
    pid_t pid[N];
    int i, child_status;
    for(i = 0; i < N; i++)
        if((pid[i] = fork()) == 0) /* Child */
            exit(100+i);
    for(i = 0; i < N; i++){ /* Parent */
        pid_t wpid = wait(&child_status);
        if(WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n", wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```

forks.c

waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Various options (see textbook)

```
void fork10() {
    pid_t pid[N];
    int i, child_status;
    for(i = 0; i < N; i++)
        if((pid[i] = fork()) == 0) /* Child */
            exit(100+i);
    for(i = N-1; i >= 0; i--){ /* Parent */
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if(WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n", wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```

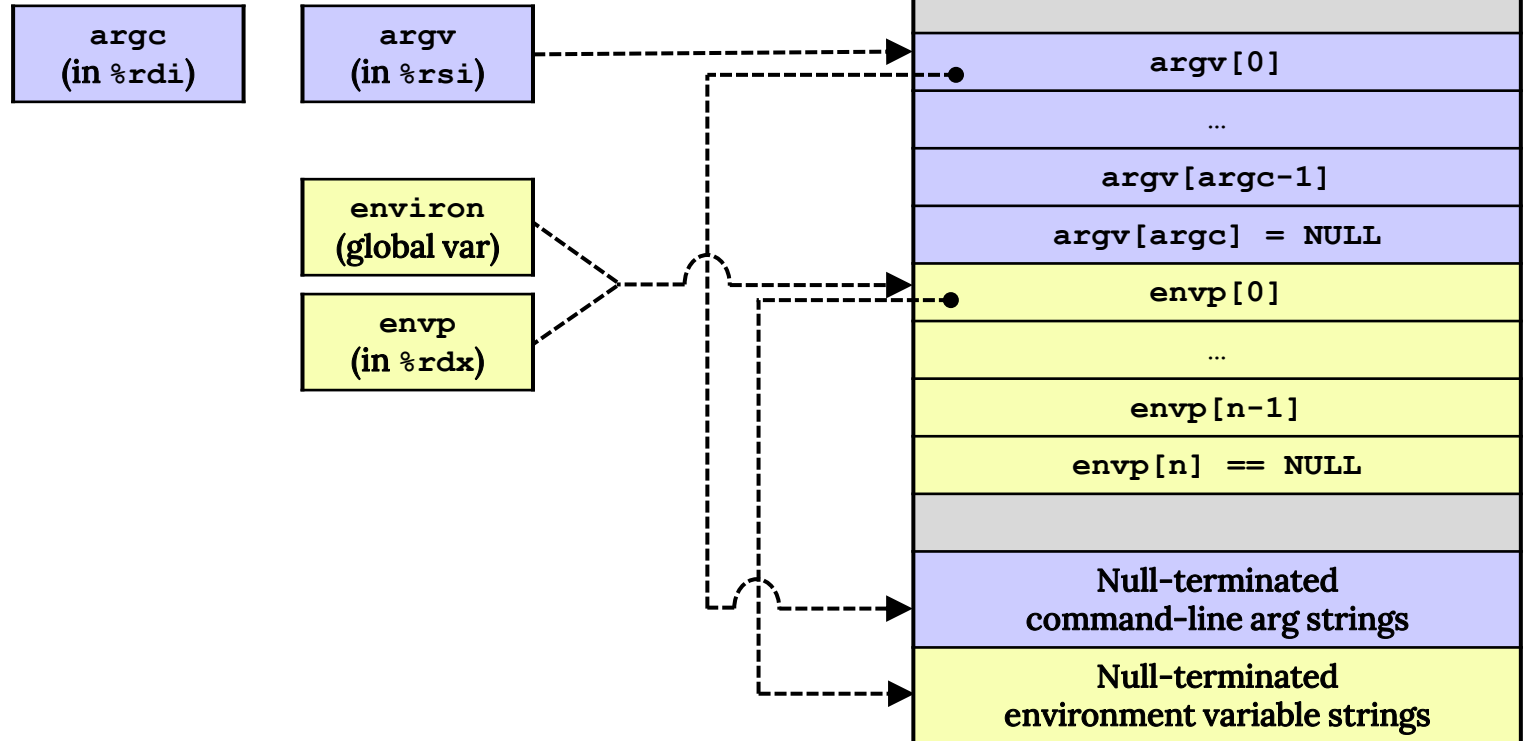
forks.c

execve: Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- Loads and runs in the current process
 - Executable file `filename`
 - Object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
 - Argument list `argv`
 - By convention `argv[0]==filename`
 - Environment variable list `envp`
 - `name=value` strings (e.g., `USER=jspark`)
 - `getenv`, `putenv`, `printenv`
- Overwrites code, data, and stack
 - Retains PID, open files and signal context
- Called once and never returns
 - Except if there is an error

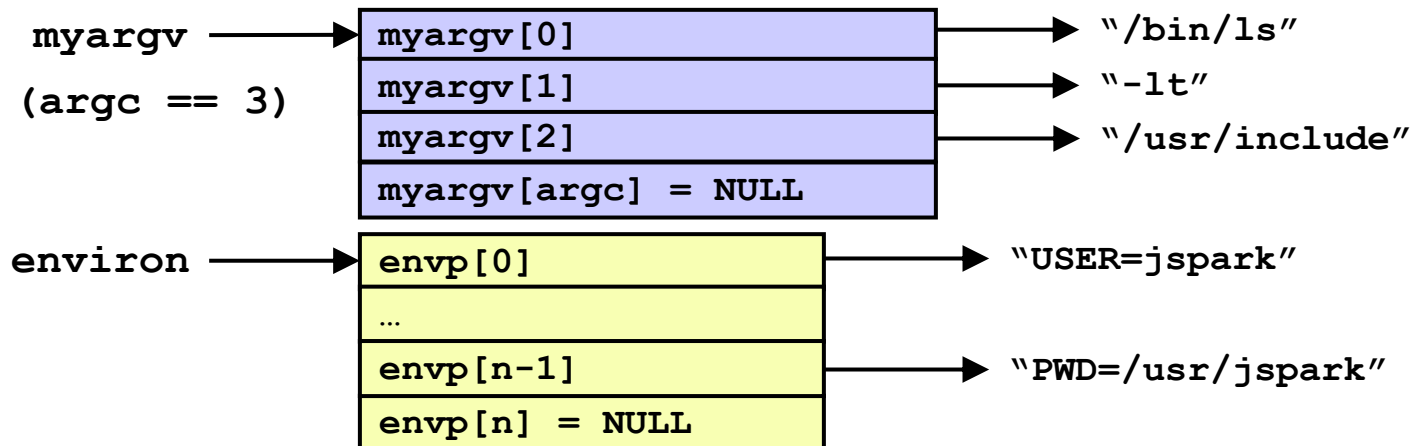
Structure of the Stack

When a New Program Starts



execve Example

- Executes `/bin/ls -lt /usr/include` in child process using current environment



```
if((pid = Fork()) == 0){    /* Child runs program */
    if(execve(myargv[0], myargv, environ) < 0){
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```

Summary

- Exceptions
 - Events that require nonstandard control flow
 - Generated externally (interrupts) or internally (traps and faults)
- Processes
 - At any given time, system has multiple active processes
 - Only one can execute at a time on a single core, though
 - Each process appears to have total control of processor + private memory space

Summary

- Spawning processes
 - Call **fork**
 - One call, two returns
- Process completion
 - Call **exit**
 - One call, no return
- Reaping and waiting for processes
 - Call **wait** or **waitpid**
- Loading and running programs
 - Call **execve** (or variant)
 - One call, (normally) no return

[CSED211] Introduction to Computer Software Systems

Lecture 13: Exceptional Control Flow – Exceptions and Processes

Prof. Jisung Park



CAOS
COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.11.22