

# [CSED211] Introduction to Computer Software Systems

## Lecture 14: Exceptional Control Flow – Signals

Prof. Jisung Park



**CAOS**

COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

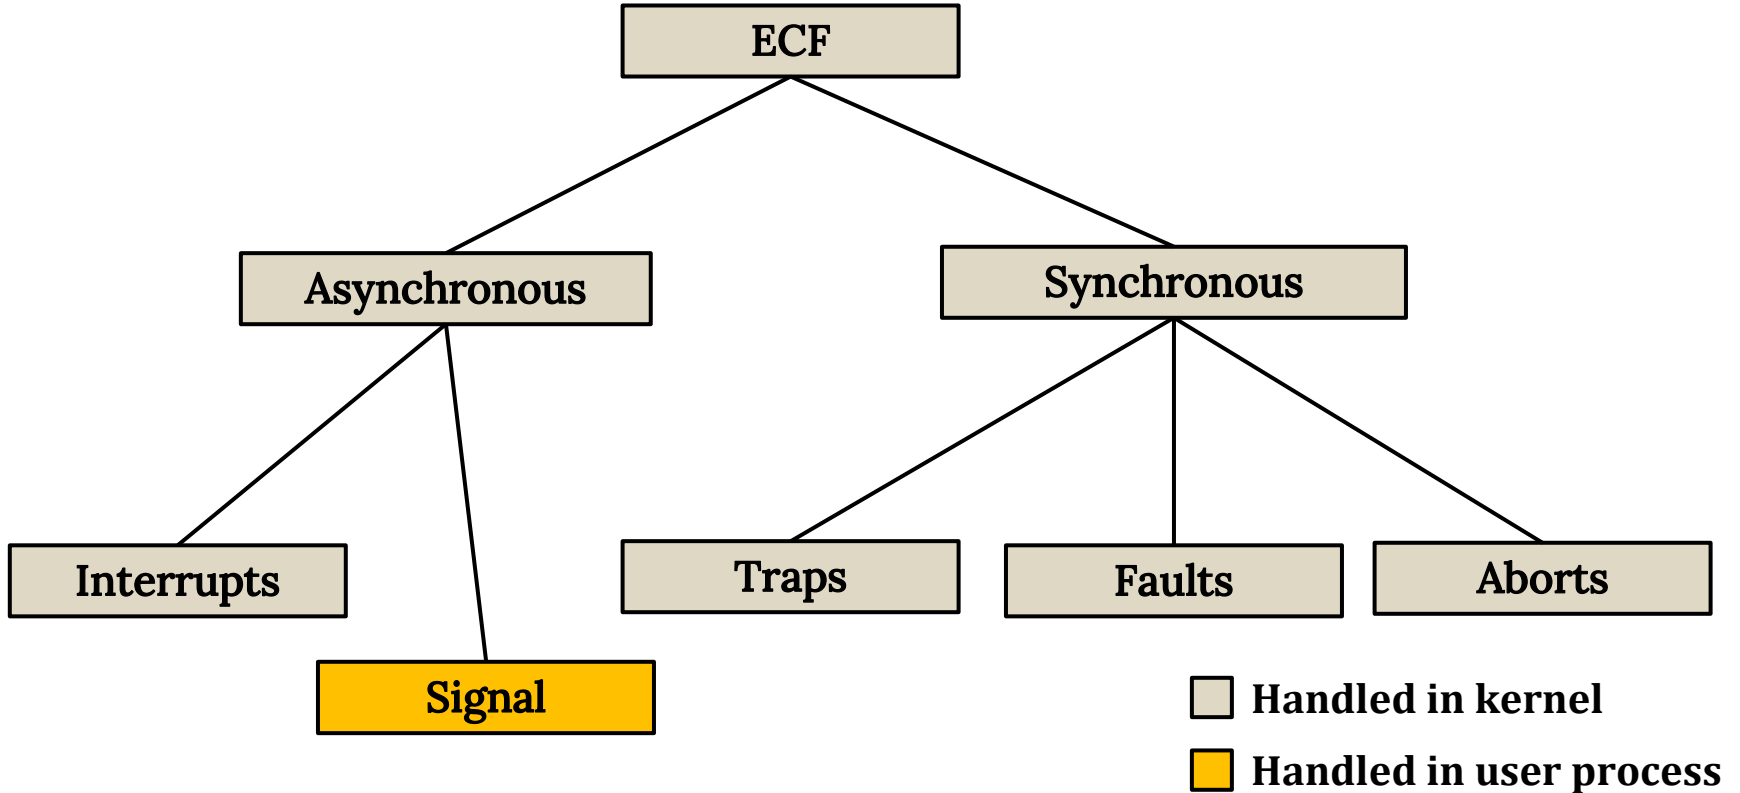
2023.11.29

# ECF Exists at All Levels of a System

---

- Exceptions
    - Hardware and operating system kernel software
  - Process context switch
    - Hardware timer and kernel software
  - Signals
    - Kernel software
  - Nonlocal jumps
    - Application code
- 
- Previous Lecture
- This Lecture
- Textbook

# Taxonomy

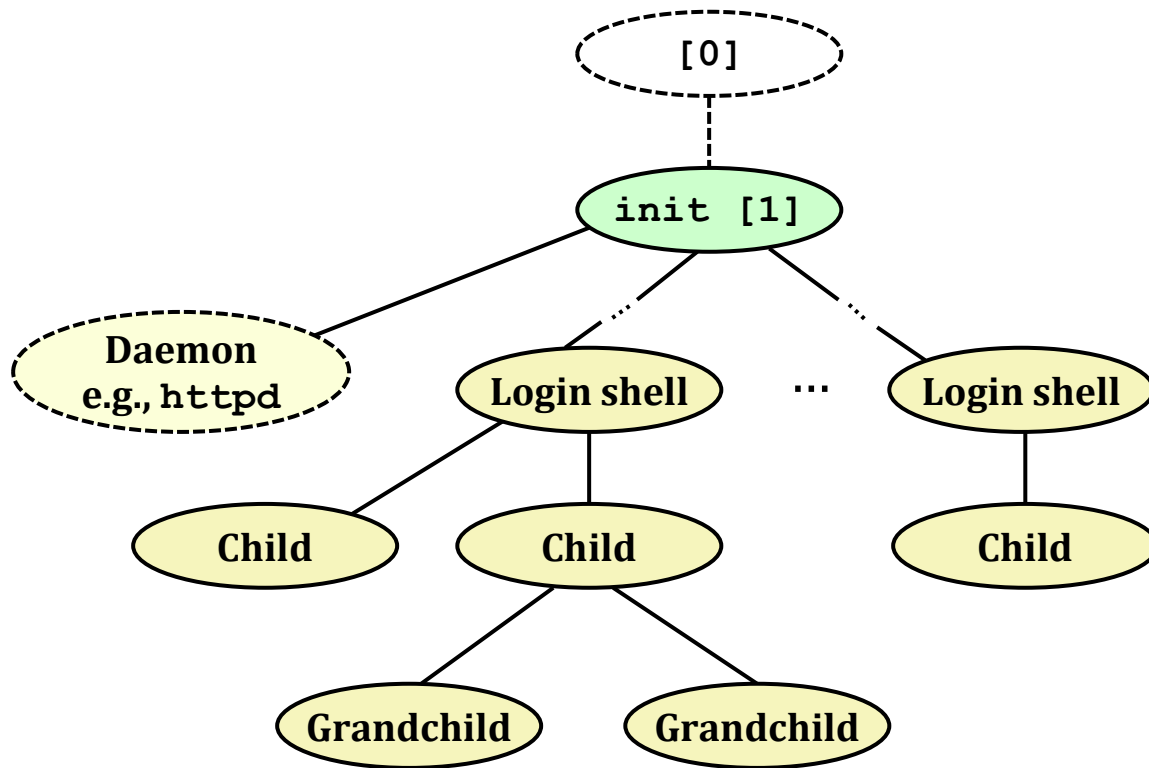


# Lecture Agenda

---

- Shells
- Signals

# Linux Process Hierarchy



Note: you can view the hierarchy using the Linux `ps tree` command

# Shell Programs

---

- **Shell**: an application program that runs programs on behalf of the user
  - **sh**: original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
  - **csh/tcsh**: BSD Unix C shell
  - **bash**: Bourne-Again shell (default Linux shell)
- Simple shell
  - Described in the textbook (pp. 753~)
  - Implementation of a very elementary shell
  - Purpose
    - Understand what happens when you type commands
    - Understand use and operation of process control operations

# Simple Shell Example

```
$ ./shellex
> /bin/ls -l csapp.c Must give full pathnames for programs
-rw-r--r-- 1 jspark users 23053 Jun 15  2023 csapp.c
> /bin/ps
  PID TTY          TIME CMD
 31542 pts/2        00:00:01 tcsh
 32017 pts/2        00:00:00 shellex
 32019 pts/2        00:00:00 ps
> /bin/sleep 10 & Run program in background
32031 /bin/sleep 10 &
> /bin/ps
  PID TTY          TIME CMD
 31542 pts/2        00:00:01 tcsh
 32024 pts/2        00:00:00 emacs
 32030 pts/2        00:00:00 shellex
 32031 pts/2        00:00:00 sleep
 32033 pts/2        00:00:00 ps
> quit
```

# Simple Shell Implementation

- Basic loop
  - Read line from command line
  - Execute the requested operation
    - Built-in command  
(only one implemented is `quit`)
    - Load and execute program from a file
  - Execution is a sequence of read & evaluation

```
int main(int argc, char** argv) {  
    /* command line */  
    char cmdline[MAXLINE];  
  
    while(1) {  
        /* read */  
        printf("> ");  
        fgets(cmdline, MAXLINE, stdin);  
        if (feof(stdin))  
            exit(0);  
  
        /* evaluate */  
        eval(cmdline);  
    }  
    ...  
}
```

shellex.c



# Simple Shell Implementation

```
void eval(char *cmdline){
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if(argv[0] == NULL)
        return; /* Ignore empty lines */

    if(!builtin_command(argv)){
        if((pid = Fork()) == 0){ /* Child runs user job */
            if(execve(argv[0], argv, environ) < 0){
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }
    }
}
```

*parseline parses 'buf' into 'argv' and returns whether or not input line ended in '&'*

shellex.c

# Simple Shell Implementation

```
void eval(char *cmdline){
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if(argv[0] == NULL)
        return; /* Ignore empty lines */

    if(!builtin_command(argv)){
        if((pid = Fork()) == 0){ /* Child runs user job */
            if(execve(argv[0], argv, environ) < 0){
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }
    }
}
```

*If the target program is a built-in command, then handle it here.*

*Otherwise, fork/exec the program*

shellex.c

# Simple Shell Implementation

```
strcpy(buf, cmdline);  
bg = parseline(buf, argv);  
if(argv[0] == NULL)  
    return;    /* Ignore empty lines */
```

*If the target program is a built-in command, then handle it here.  
Otherwise, fork/exec the program*

```
if(!builtin_command(argv)){  
    if((pid = Fork()) == 0){    /* Child runs user job */
```

*Create a child  
via fork*

```
    if(execve(argv[0], argv, environ) < 0){  
        printf("%s: Command not found.\n", argv[0]);  
        exit(0);  
    }  
}
```

*Execute the target program (specified by argv[0])  
Note: execve returns only on error*

```
/* Parent waits for foreground job to terminate */  
if(!bg){  
    int status;  
    if(waitpid(pid, &status, 0) < 0)
```

shellex.c

# Simple Shell Implementation

```
        printf("%s: Command not found.\n", argv[0]);
        exit(0);
    }
}

/* Parent waits for foreground job to terminate */
if(!bg) {
    int status;
    if(waitpid(pid, &status, 0) < 0)
        unix_error("waitfg: waitpid error");
}
else
    printf("%d %s", pid, cmdline);
}
return;
}
```

If the program runs in foreground,  
the shell must wait for its completion

If the program runs in background,  
print its PID and continue doing other stuff

Is everything okay?

shellex.c

# What Is a 'Background Job'?

- Users generally run one command at a time
  - Type command, read output, type another command
- Some programs run for a long time
  - e.g., delete this after two hours

```
$ sleep 7200; rm /tmp/junk
```

*Shell stuck for 2 hours*

- A background job is a process we do not want to wait for

```
$ (sleep 7200; rm /tmp/junk) &
```

```
[1] 907
```

```
$
```

*Ready for the next command*

# Problem with the Simple Shell Example

---

- The shell program is designed to **run indefinitely**
  - i.e., until the user input is 'quit'
  - Should **not** accumulate resources unnecessarily
    - Memory
    - Child processes
    - File descriptors
- Our example shell **correctly waits for and reaps foreground jobs**
- **What about background jobs?**
  - Will become zombies when they terminate
  - Will never be reaped because shell (typically) will not terminate
  - Will create a memory leak that could run the kernel out of memory

# ECF to the Rescue

---

- Problem
  - The shell does not know when a background job will finish
    - Which could happen at any time by nature
  - The shell's regular control flow cannot reap exited background processes in a timely fashion
    - Regular control flow: wait until running job completes, then reap it
- Solution: exceptional control flow
  - The kernel interrupts regular processing to alert the shell (parent) when a background process (child) completes
  - Such an alert mechanism is called a **signal** in UNIX

# Lecture Agenda

---

- Shells
- Signals



# Signal

- A small message that notifies a process of a system event
  - Akin to exceptions and interrupts
  - Sent from the kernel to a process (sometimes at the request from another process)
  - Signal type is identified by small integer ID's (1~30)
  - The only information in a signal is its ID (and its arrival for sure)

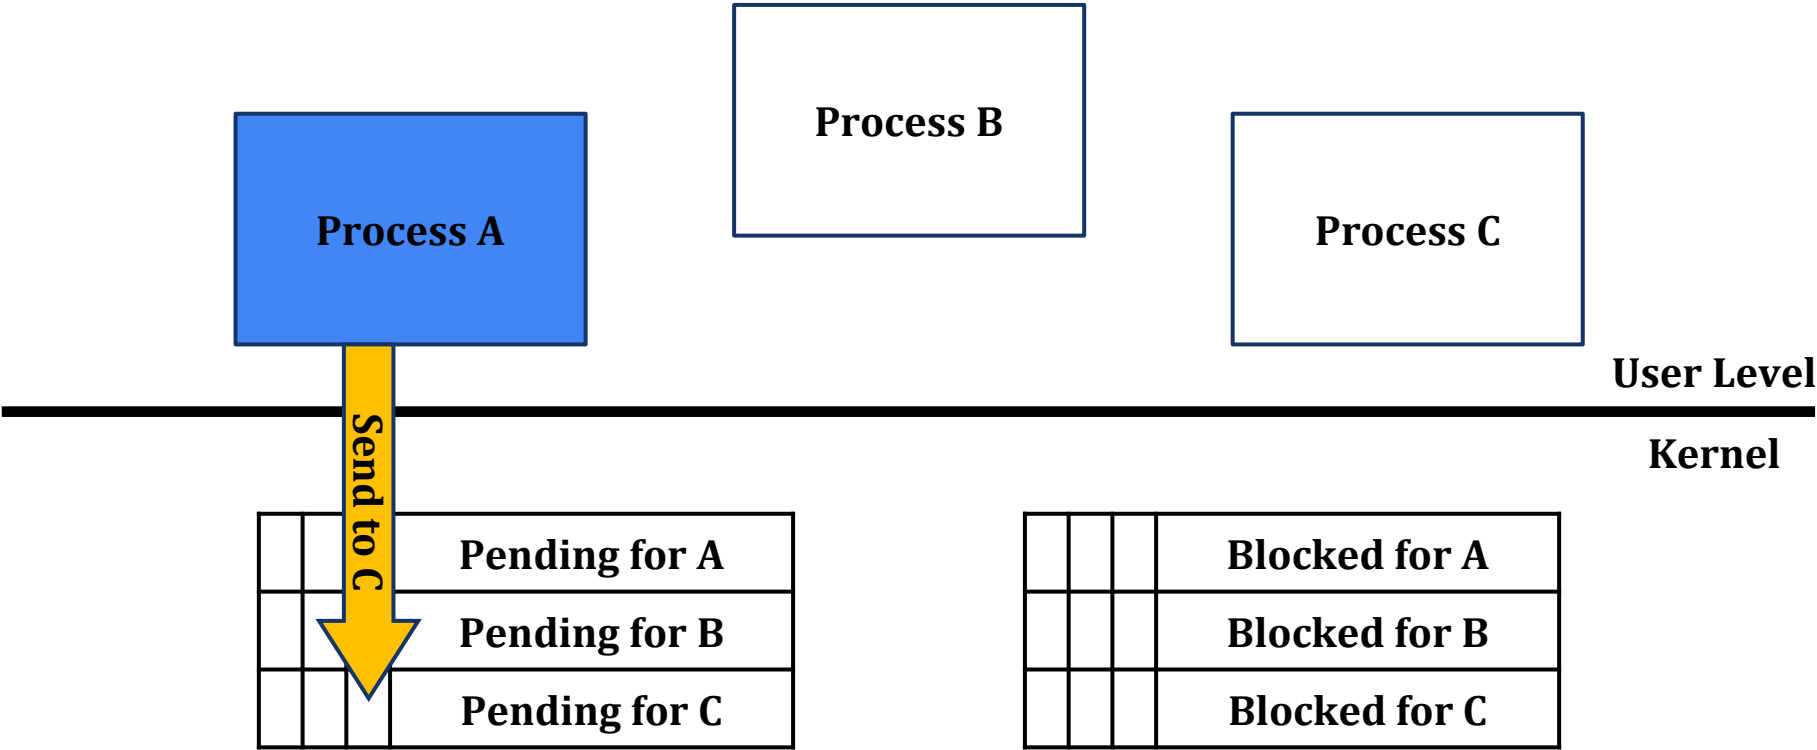
ID	Name	Default Action	Corresponding Event
2	<b>SIGINT</b>	Terminate	Interrupt (e.g., CTRL-C from keyboard)
9	<b>SIGKILL</b>	Terminate	Kill program (cannot override or ignore)
11	<b>SIGSEGV</b>	Terminate & Dump	Segmentation violation
14	<b>SIGALRM</b>	Terminate	Timer signal
17	<b>SIGCHLD</b>	Ignore	Child stopped or terminated

# Signal Concepts: Sending a Signal

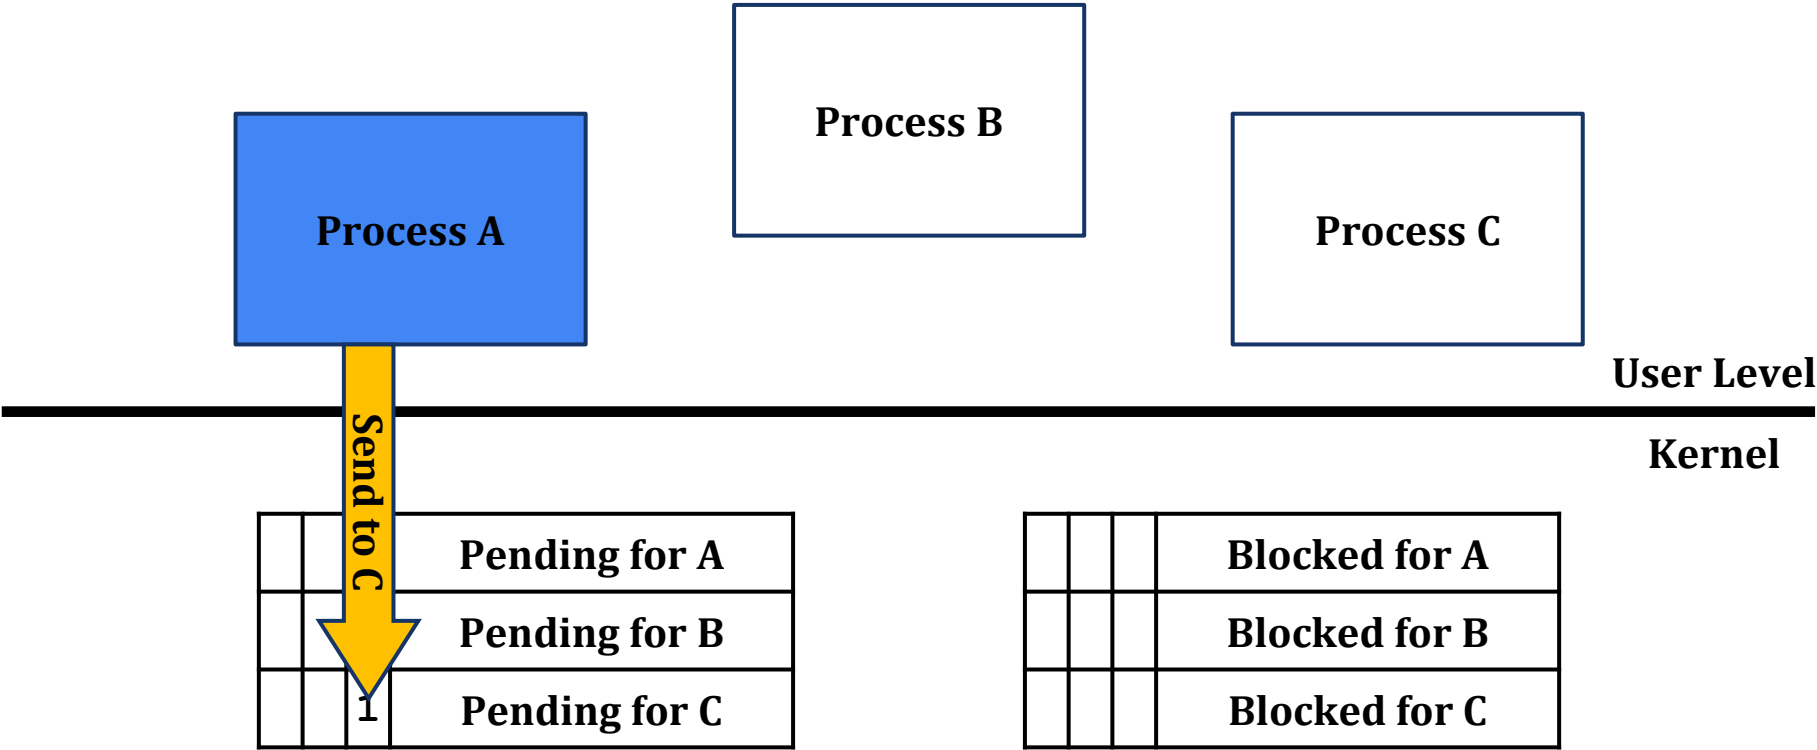
---

- Kernel **sends** (**delivers**) a signal to a destination process by updating some state in the context of the destination process
- Kernel sends a signal for one of the following reasons
  - Kernel has detected a system event such as divide-by-zero (**SIGFPE**) or the termination of a child process (**SIGCHLD**)
  - Another process has invoked the **kill** system call to explicitly request the kernel to send a signal to the destination process

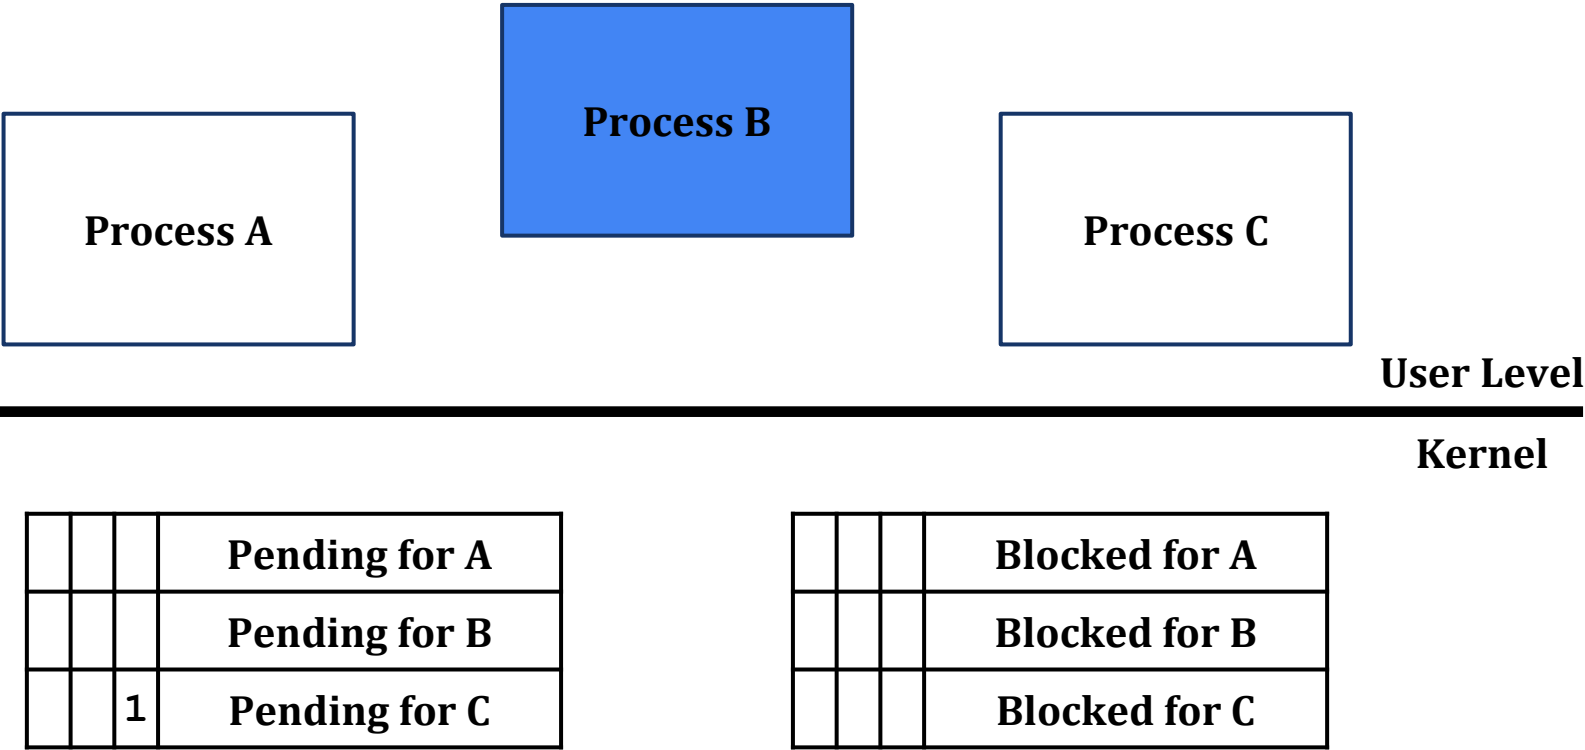
# Signal Concepts: Sending a Signal (Cont.)



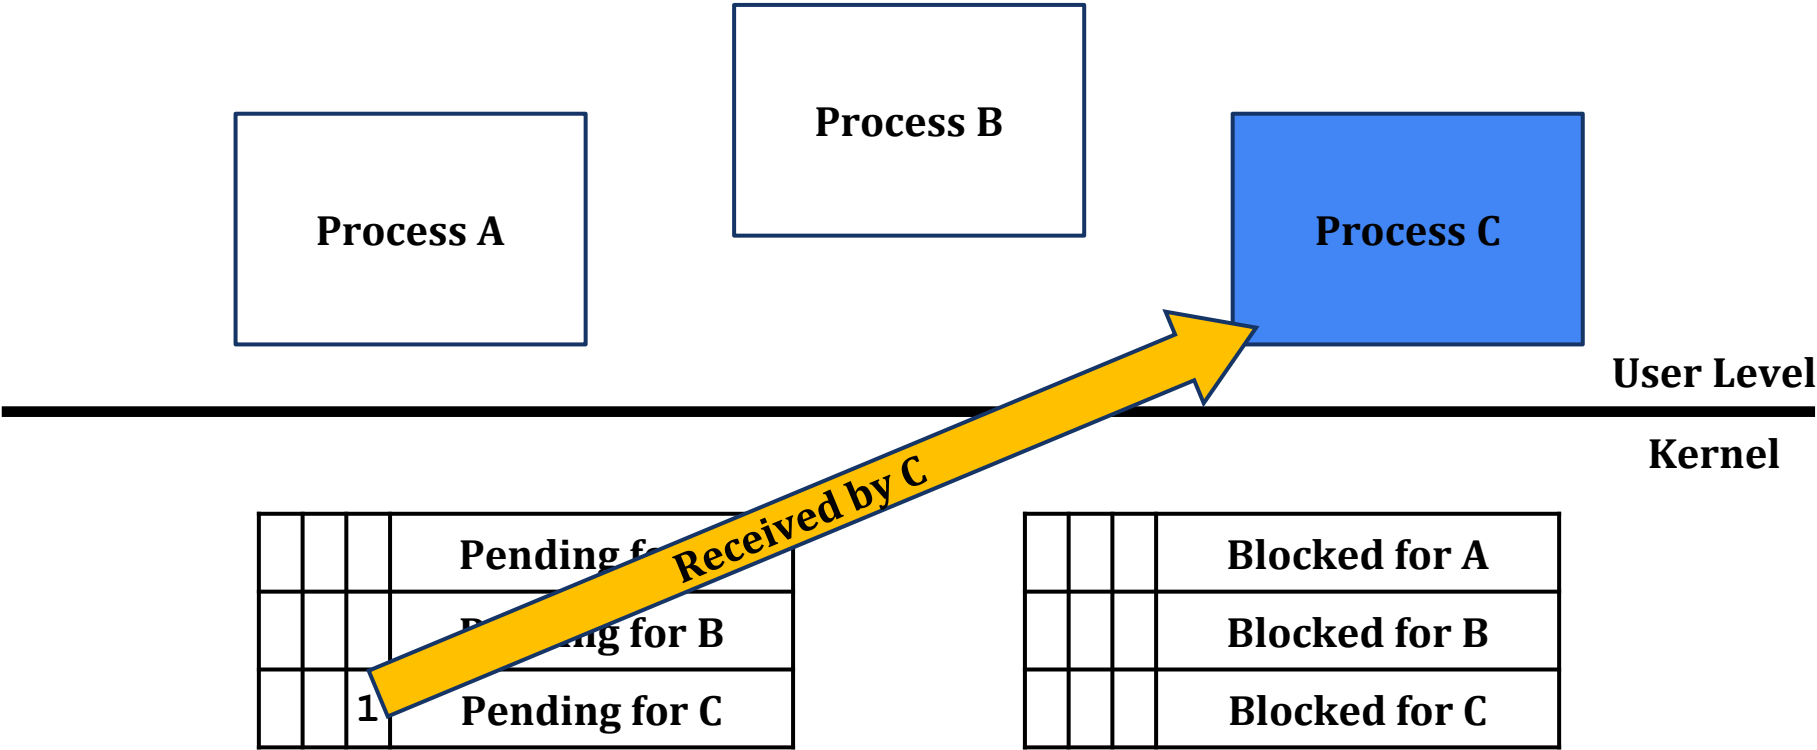
# Signal Concepts: Sending a Signal (Cont.)



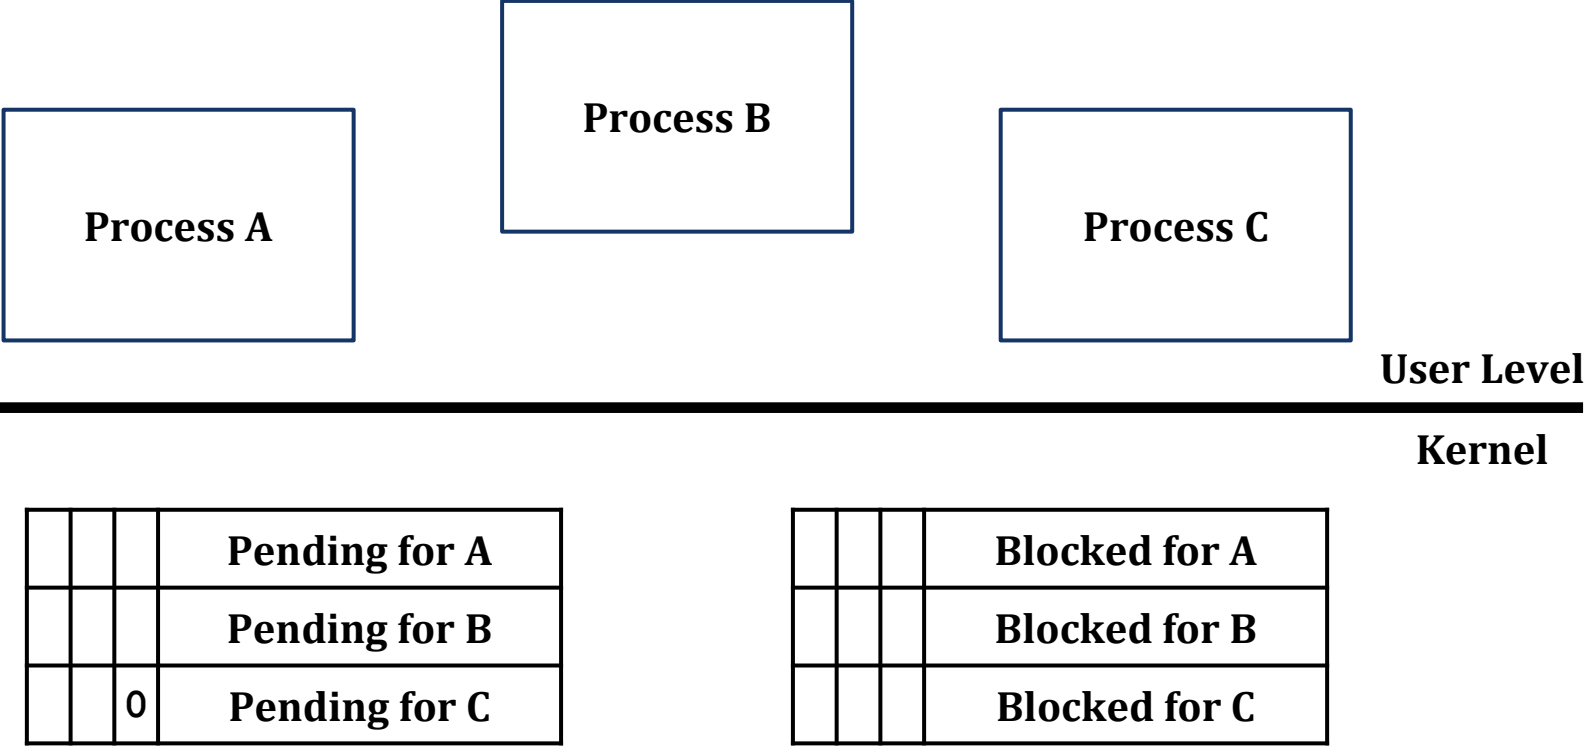
# Signal Concepts: Sending a Signal (Cont.)



# Signal Concepts: Sending a Signal (Cont.)

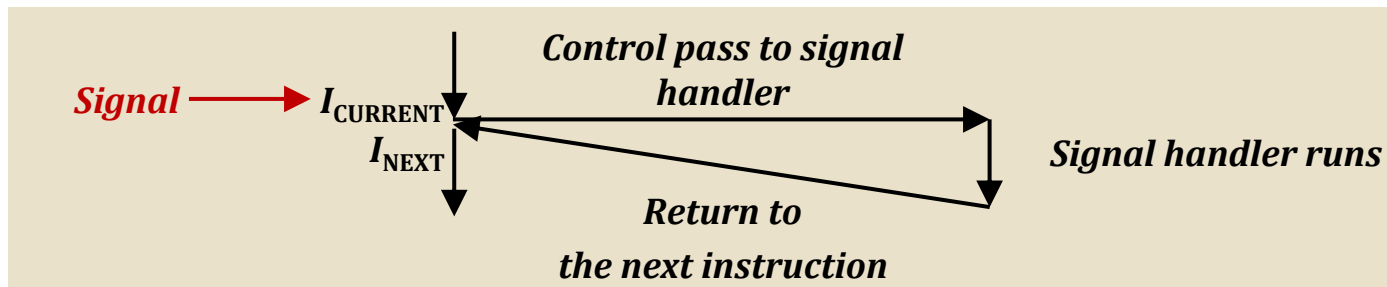


# Signal Concepts: Sending a Signal (Cont.)



# Signal Concepts: Receiving a Signal

- A destination process **receives** a signal when it is **forced by the kernel to react** in some way to the delivery of the signal
- Three possible ways to react
  - Ignore the signal (do nothing)
  - Terminate the process (with optional core dump)
  - Catch the signal by executing a user-level function called signal handler
    - Akin to a hardware exception handler being called in response to an asynchronous interrupt





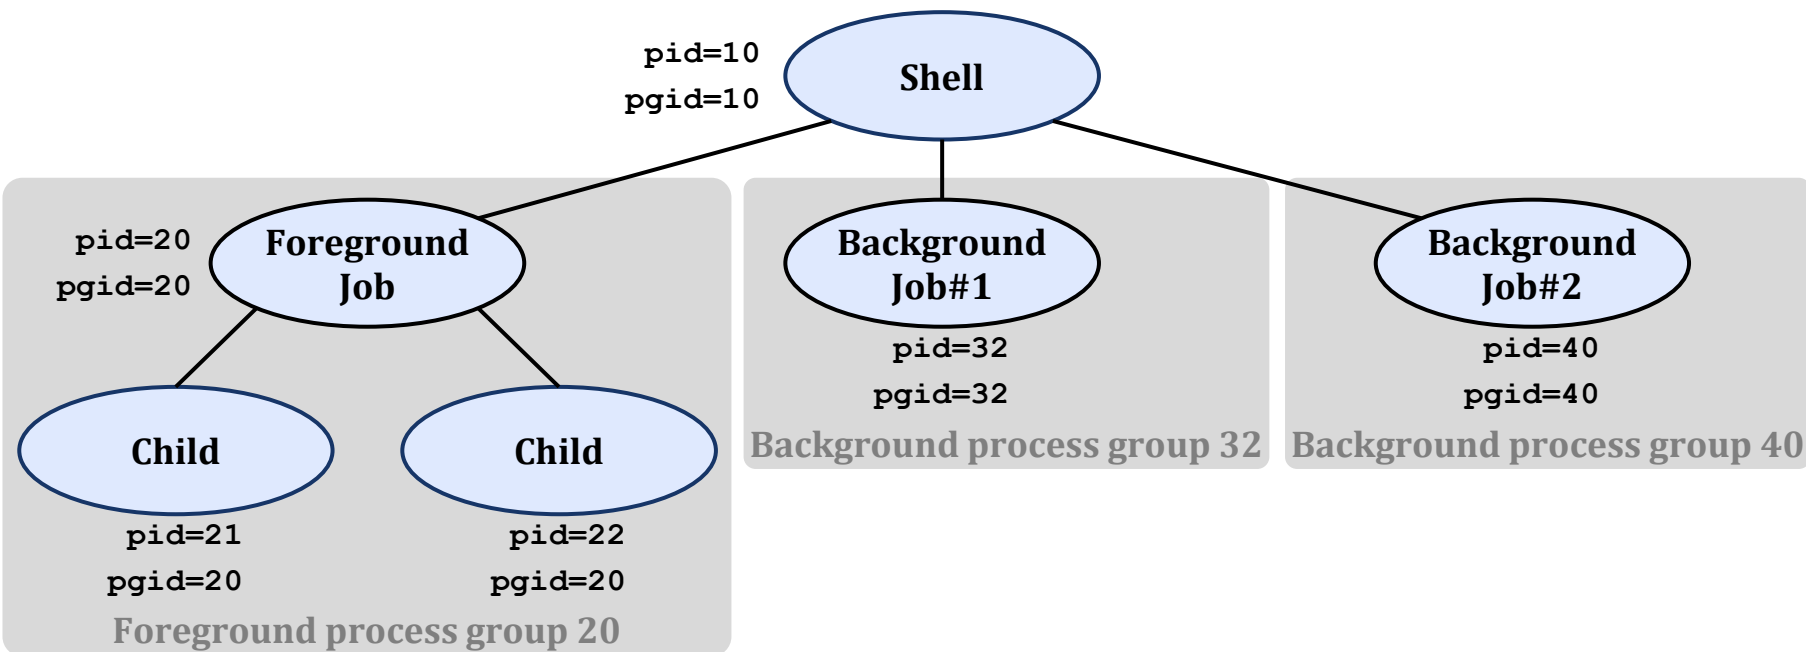
# Signal Concepts: Pending and Blocked Signals

---

- A signal is **pending** if sent but not yet received
  - There can be at most one pending signal of any particular type
  - **Important:** signals are **not queued**; if a process has a pending signal of type  $k$ , then subsequent type- $k$  signals will be discarded
  - A pending signal is received **at most** once
- A process can **block** the receipt of certain signals
  - A blocked signal can be delivered, but will not be received until unblocked
- Kernel keeps bit vectors **pending**/**blocked** in the context of each process
  - Represent the sets of pending/blocked signals
  - Kernel sets (clears) bit  $k$  in **pending** when a signal of type  $k$  is delivered (received)
  - A process can (un)block a signal of type  $k$  by setting (clearing) bit  $k$  in **blocked** (also referred to as the **signal mask**) using the **sigprocmask** function

# Sending Signals: Process Groups

- Every process belongs to exactly one process group
  - `setpgid()`: changes the process group of a process
  - `getpgrp()`: returns the current process group



# Sending Signals with `/bin/kill` Program

- Sends arbitrary signal to a process or process group
- Example

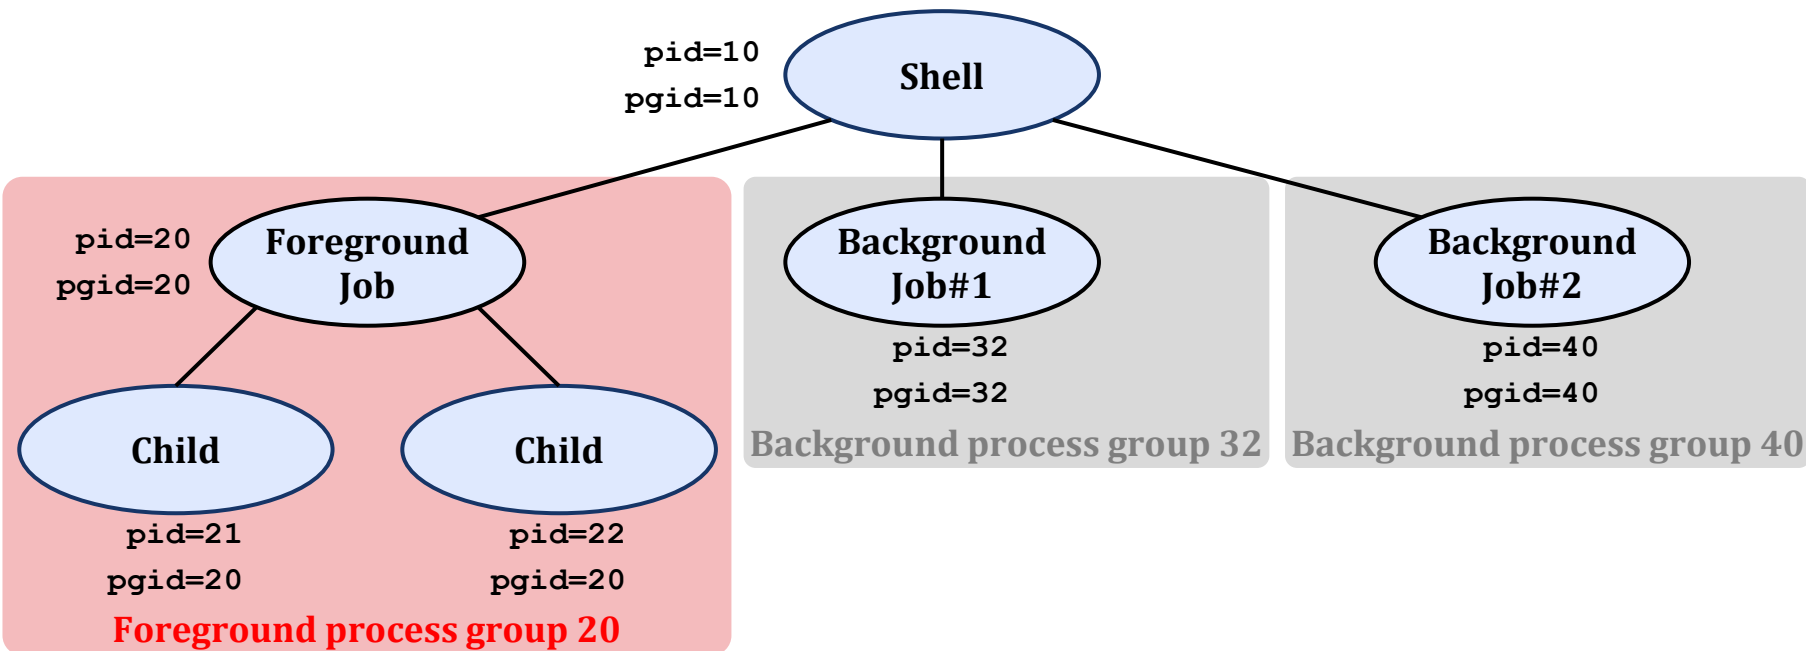
```
$ ./forks 16
Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817
$ ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24818 pts/2        00:00:02 forks
 24819 pts/2        00:00:02 forks
 24820 pts/2        00:00:00 ps
$ /bin/kill -9 -24817
$ ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24823 pts/2        00:00:00 ps
$
```

**`/bin/kill -9 24818`: sends SIGKILL to process 24818**

**`/bin/kill -9 24817`: sends SIGKILL to every process  
in process group 24817**

# Sending Signals with Keyboard

- **CTRL-C** (**CTRL-Z**) sends a **SIGINT** (**SIGTSTP**) to every job in the foreground process group; default action is to terminate (suspend) each process



# Sending Signals with Keyboard: Example

```
$ ./forks 17
Child: pid=28108 pgrp=28107
Parent: pid=28107 pgrp=28107
Suspended
```

*Types CTRL-Z*

```
$ ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28107	pts/8	T	0:01	./forks 17
28108	pts/8	T	0:01	./forks 17
28109	pts/8	R+	0:00	ps w

```
$ fg
./forks 17
```

*Types CTRL-C*

```
$ ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28110	pts/8	R+	0:00	ps w

- **STAT (process state) legend**
  - **First letter**
    - **S**: sleeping
    - **T**: stopped
    - **R**: running
  - **Second letter**
    - **s**: session leader
    - **+**: foreground proc. group
  - See 'man ps' for more details

# Sending Signals with kill Function

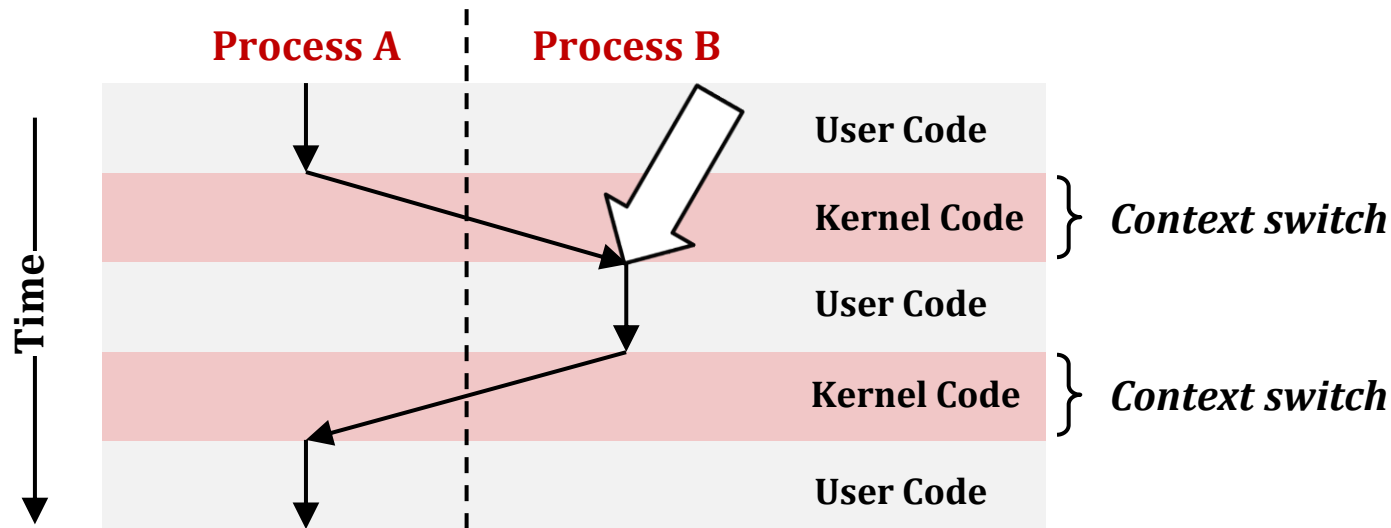
```
void fork12(){
    pid_t pid[N];
    int i, child_status;
    for(i = 0; i < N; i++)
        if((pid[i] = fork()) == 0) while(1); /* Child infinite loop */

    /* Parent terminates the child processes */
    for(i = 0; i < N; i++){
        printf("Killing process %d\n", pid[i]);
        kill(pid[i], SIGINT);
    }

    /* Parent reaps terminated children */
    for(i = 0; i < N; i++){
        pid_t wpid = wait(&child_status);
        if(WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n", wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

# Receiving Signals

- Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$



# Receiving Signals

---

- Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$
- Kernel computes  $\mathbf{pnb} = \mathbf{pending} \ \& \ \sim\mathbf{blocked}$ 
  - i.e., the set of pending non-blocked signals for process  $p$
- If  $(\mathbf{pnb} == 0)$ : pass control to the next instruction in the logical flow for  $p$
- Otherwise
  - Choose the nonzero LSB  $k$  in  $\mathbf{pnb}$  and force process  $p$  to receive signal  $k$
  - The receipt of the signal triggers the **corresponding action** by  $p$
  - Repeat for all nonzero bits in  $\mathbf{pnb}$
  - Pass control to the next instruction in logical flow for  $p$



# Default Actions

---

- Each signal type has a **predefined default action** that is one of
  - The process terminates
  - The process terminates and dumps core
  - The process stops until restarted by a **SIGCONT** signal
  - The process ignores the signal

# Installing Signal Handlers

- The **signal** function modifies the default action for signal **signum**
  - `handler_t *signal(int signum, handler_t *handler)`
- Different values for **handler**
  - **SIG\_IGN**: ignore signals of type **signum**
  - **SIG\_DFL**: revert to the default action on receipt of signals of type **signum**
  - Otherwise, **handler** is the address of a **signal handler**
    - Called when process receives signal of type **signum**
    - Referred to as **installing the handler**
    - Executing the handler is called **catching** or **handling the signal**
    - When the handler executes its return statement, control passes back to the instruction in the control flow of the process that was interrupted by receipt of the signal

# Signal Handling Example

```
/* SIGINT handler */
void sigint_handler(int sig){
    printf("So you think you can stop the bomb with ctrl-c, do you?\n");
    sleep(2);
    printf("Well...\n");
    fflush(stdout);
    sleep(1);
    printf("OK. :-)\n");
    exit(0);
}


int main(){
    /* Install the SIGINT handler */
    if(signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

    /* Wait for the receipt of a signal */
    pause();
    return 0;
}
```

sigint.c

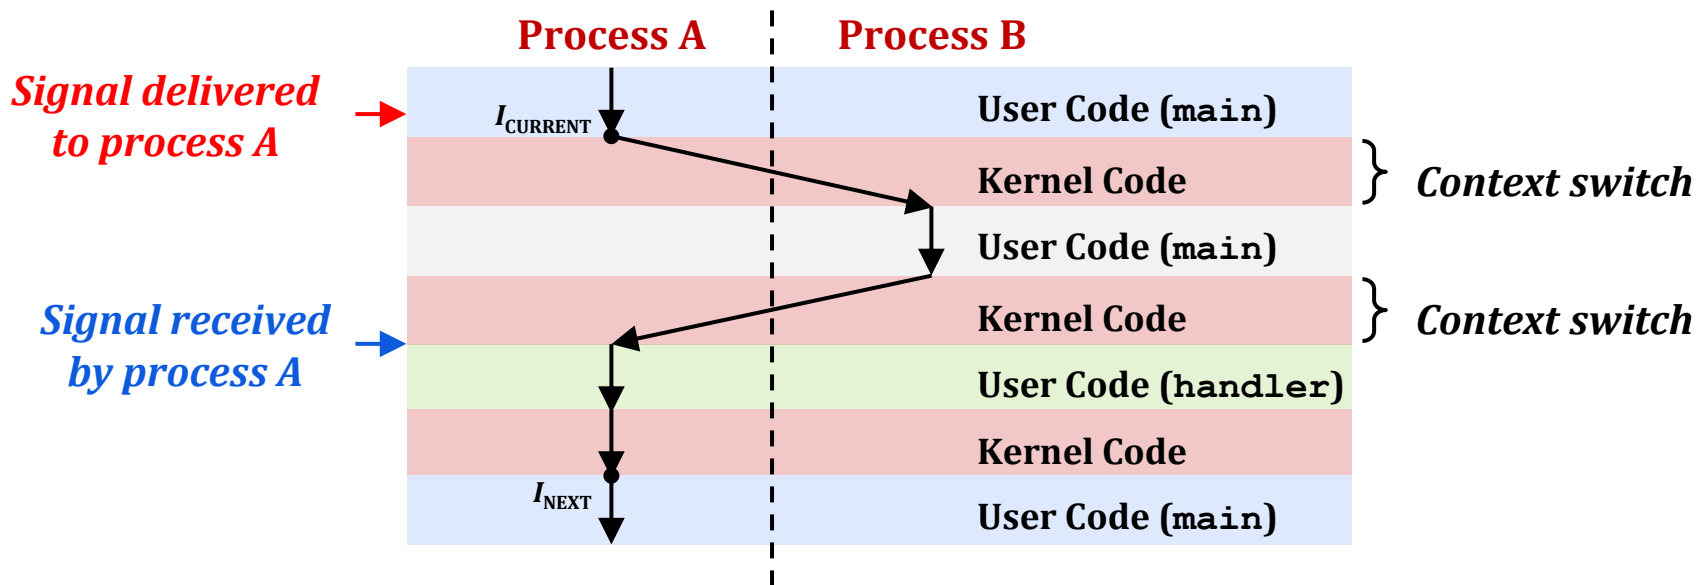
# Signal Handlers as Concurrent Flows

- A signal handler is a separate logical flow (not process)
  - Runs concurrently with the main program
  - Exists only until returns to the main program



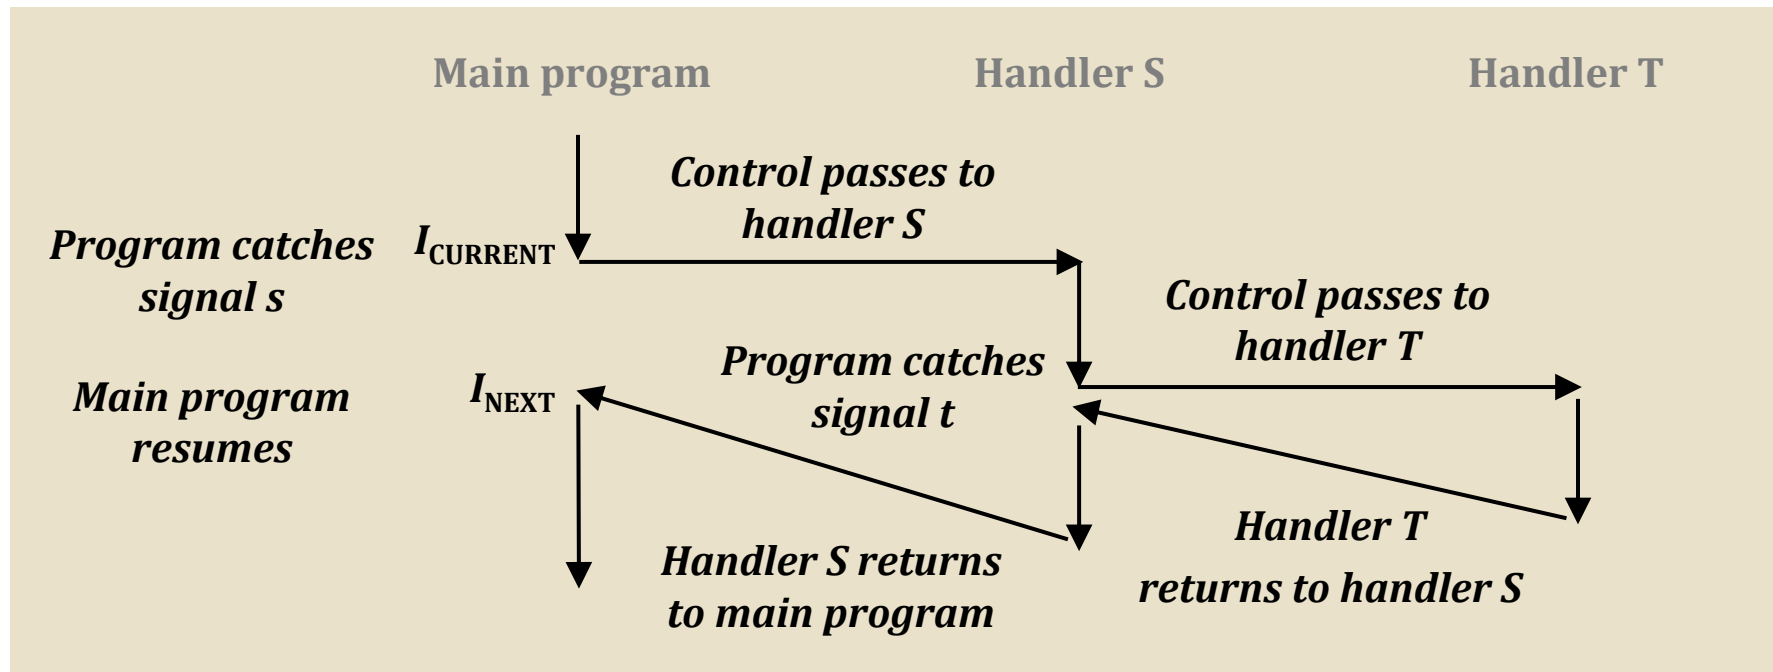
# Signal Handlers as Concurrent Flows: Another View

- A signal handler is a separate logical flow (not process)
  - Runs concurrently with the main program
  - Exists only until returns to the main program



# Nested Signal Handlers

- Handlers can be interrupted by other handlers



# Blocking and Unblocking Signals

---

- **Implicit blocking** mechanism
  - Kernel blocks **any pending signals of the type currently being handled**  
e.g., a **SIGINT** handler cannot be interrupted by another **SIGINT**
- **Explicit blocking and unblocking** mechanism: **sigprocmask** function
- Supporting functions
  - **sigemptyset**: create an empty set
  - **sigfillset**: add every signal number to set
  - **sigaddset**: add signal number to set
  - **sigdelset**: delete signal number from set

# Temporarily Blocking Signals

---

```
sigset_t mask, prev_mask;

sigemptyset(&mask);
sigaddset(&mask, SIGINT);

/* Block SIGINT and save previous blocked set */
sigprocmask(SIG_BLOCK, &mask, &prev_mask);

/* Code region that will not be interrupted by SIGINT */

/* Restore previous blocked set, unblocking SIGINT */
sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```



# Safe Signal Handling

---

- Handlers are tricky because they are concurrent with main program and share the same global data structures
  - Shared data structures can become corrupted
- We will (hopefully) discuss concurrency issues later in this course; for now, here are some guidelines to help you avoid trouble

# Guildlines for Writing Safe Handlers

---

- **G0:** Keep your handlers **as simple as possible**
  - e.g., set a global flag and return
- **G1:** Call only **async-signal-safe function in your handlers**
  - `printf`, `sprintf`, `malloc`, and `exit` are **not** safe
- **G2:** **Save and restore `errno`** on entry and exit
  - So that other handlers do not overwrite your value of `errno`
- **G3:** Protect shared data structures by **temporarily blocking all signals**
- **G4:** Declare **global variables** as **`volatile`**
  - To prevent compiler from storing them in a register
- **G5:** Declare **global flags** (only read and written) as **`volatile sig_atomic_t`**
  - So that they do not need to be protected like other globals

# Async-Signal-Safety

---

- A function is **async-signal-safe** if it meets either of the following conditions
  - **Reentrant**, i.e., all variables stored on stack frame (§12.7.2 in the textbook)
  - **Non-interruptible** by signals
- Posix guarantees 117 functions to be async-signal-safe
  - Source: `man 7 signal`
  - Popular functions on the list: `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`
  - Popular functions **not** on the list
    - `printf`, `sprintf`, `malloc`, `exit`
    - Unfortunate fact: **`write` is the only async-signal-safe output function**

# Safe Formatted Output Option#1

- Use the reentrant SIO (safe I/O library) from `csapp.c` in your handlers
  - `ssize_t sio_puts(char s[]) /* Put string */`
  - `ssize_t sio_putl(long v) /* Put long */`
  - `void sio_error(char s[]) /* Put msg & exit */`

```
/* Safe SIGINT handler */
void sigint_handler(int sig){
    sio_puts("So you think you can stop the bomb with ctrl-c, do you?\n");
    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

# Safe Formatted Output Option#2

- Use the new & improved reentrant `sio_printf`
  - Handles restricted class of `printf` format strings
    - Recognizes `%c`, `%s`, `%d`, `%u`, `%x`, and `%%`
    - Size designators `l` and `z`

```
/* Safe SIGINT handler */
void sigint_handler(int sig){
    sio_printf("So you think you can stop the bomb (process %d)
               with ctrl-%c, do you?\n", (int) getpid(), 'c');

    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

# Correct Signal Handling

- Pending signals are **not** queued
  - Only one bit in the pending bit vector for each signal type  
→ At most one pending signal of any particular type
  - You cannot use signals to count events, such as children terminating

```
volatile int ccount = 0;
void child_handler(int sig){
    int olderrno = errno;
    pid_t pid;
    if((pid = wait(NULL)) < 0)
        sio_error("wait error");
    ccount--;
    sio_puts("Handler reaped child ");
    sio_putl((long) pid);
    sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}
```

```
void fork14(){
    pid_t pid[N];
    int i;
    ccount = N;
    signal(SIGCHLD, child_handler);

    for(i = 0; i < N; i++){
        if((pid[i] = Fork()) == 0){
            sleep(1);
            exit(0); /* Child exits */
        }
    }
    while(ccount > 0); /* Parent spins */
}
```

# Correct Signal Handling

- Pending signals are **not** queued
  - Only one bit in the pending bit vector for each signal type  
→ At most one pending signal of any particular type
  - You cannot use signals to count events, such as children terminating

```
volatile int ccount = 0;
void child_handler(int sig){
    int olderrno = errno;
    pid_t pid;
    if((pid = wait(NULL)) < 0)
        sio_error("wait error");
    ccount--;
    sio_puts("Handler reaped child ");
    sio_putl((long) pid);
    sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}
```

```
void fork14(){
    pid_t pid[N];
    int i;
    ccount = N;
    signal(SIGCHLD, child_handler);

    for(i = 0; i < N; i++){
        if((pid[i] = Fork()) == 0){
            sleep(1);
            exit(0); /* Child exits */
        }
    }
    while(ccount > 0); /* Parent spins */
}
```

```
$ ./forks 14
Handler reaped child 23240
Handler reaped child 23241
Program hangs
```

# Correct Signal Handling (Cont.)

- Must call wait for all terminated child processes
  - Put `wait` in a loop to reap all terminated children

```
$ ./forks 15
Handler reaped child 23246
Handler reaped child 23247
Handler reaped child 23248
Handler reaped child 23249
Handler reaped child 23250
$
```

```
volatile int ccount = 0;
void child_handler2(int sig){
    int olderrno = errno;
    pid_t pid;
    while((pid = wait(NULL)) > 0){
        ccount--;
        sio_puts("Handler reaped child ");
        sio_putl((long) pid);
        sio_puts(" \n");
    }
    if(errno != ECHILD)
        sio_error("wait error");
    errno = olderrno;
}
```

```
void fork15(){
    pid_t pid[N];
    int i;
    ccount = N;
    signal(SIGCHLD, child_handler);

    for(i = 0; i < N; i++){
        if((pid[i] = Fork()) == 0){
            sleep(1);
            exit(0); /* Child exits */
        }
    }
    while(ccount > 0); /* Parent spins */
}
```



# Synchronizing Flows to Avoid Races

- SIGCHLD handler for a simple shell
  - Blocks all signals while running critical code

```
void handler(int sig){
    int olderrno = errno;
    sigset_t mask_all, prev_all;
    pid_t pid;

    sigfillset(&mask_all);
    while((pid = waitpid(-1, NULL, 0)) > 0){
        /* Reap child */
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        deletejob(pid); /* Delete the child from the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    if(errno != ECHILD)
        sio_error("waitpid error");
    errno = olderrno;
}
```

procmask1.c

# Synchronizing Flows to Avoid Races (Cont.)

- Simple shell with a subtle synchronization error due to its wrong assumption that parent always runs before child

```
int main(int argc, char **argv){
    int pid;
    sigset_t mask_all, prev_all;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */
    while(n--){
        if((pid = Fork()) == 0) /* Child */
            execve("/bin/date", argv, NULL);
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* Parent */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    exit(0);
}
```

procmask1.c

# Corrected Shell Program Without Race

```
int main(int argc, char **argv){
    int pid;
    sigset_t mask_all, mask_one, prev_one;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    sigemptyset(&mask_one);
    sigaddset(&mask_one, SIGCHLD);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */
    while(n--){
        sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
        if((pid = Fork()) == 0){ /* Child */
            sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
            execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* Parent */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_one, NULL);
    }
    exit(0);
}
```

procmask2.c

# Explicitly Waiting for Signals

- Handlers for program explicitly waiting for SIGCHLD to arrive

```
volatile sig_atomic_t pid;

void sigchld_handler(int s){
    int olderrno = errno;
    pid = waitpid(-1, NULL, 0); /* Main is waiting for nonzero pid */
    errno = olderrno;
}

void sigint_handler(int s){
}
```

waitforsignal.c

# Explicitly Waiting for Signals (Cont.)

```
int main(int argc, char **argv){
    sigset_t mask, prev;
    int n = N;  /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);

    while(n--){
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if(Fork() == 0) /* Child */
            exit(0);
        pid = 0; /* Parent */
        sigprocmask(SIG_SETMASK, &prev, NULL); /* Unblock SIGCHLD */
        while(!pid); /* Wait for SIGCHLD to be received (wasteful!) */
        printf("."); /* Do something after receiving SIGCHLD */
    }
    printf("\n");
    exit(0);
}
```

Similar to a shell waiting for a foreground job to terminate

waitforsignal.c

# Explicitly Waiting for Signals (Cont.)

---

- `while(!pid)`
  - Program is correct, but very wasteful
  - In a busy-wait loop
- `while(!pid) pause();`
  - Possible race condition
  - It might receive signal b/w checking `pid` and calling `pause()`
- `while(!pid) sleep();`
  - Safe, but slow
  - It will take up to one second to respond

# Waiting for Signals with `sigsuspend`

- `int sigsuspend(const sigset_t *mask)`
- Equivalent to atomic (uninterruptable) version of

```
sigprocmask(SIG_BLOCK, &mask, &prev);  
pause();  
sigprocmask(SIG_SETMASK, &prev, NULL);
```

# Waiting for Signals with sigsuspend (Cont.)

```
int main(int argc, char **argv){
    sigset_t mask, prev;
    int n = N; /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);

    while(n--){
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if(Fork() == 0) /* Child */
            exit(0);
        pid = 0; /* Parent */
        while(!pid) sigsuspend(&prev); /* Wait for SIGCHLD to be received */
        sigprocmask(SIG_SETMASK, &prev, NULL); /* Optionally unblock SIGCHLD */
        printf("."); /* Do something after receiving SIGCHLD */
    }
    printf("\n");
    exit(0);
}
```

sigsuspend.c



# Lecture Agenda

---

- Shells
  - Signals
  - Portable Signal Handling
  - Nonlocal Jumps
- } Consult your textbook!

# Summary

---

- Signals provide process-level exception handling
  - Can generate from user programs
  - Can define effect by declaring signal handler
  - Be very careful when writing signal handlers

# [CSED211] Introduction to Computer Software Systems

## Lecture 14: Exceptional Control Flow – Signals

Prof. Jisung Park



**CAOS**

COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2023.11.29