

[CSED211] Introduction to Computer Software Systems

Lecture 15: Virtual Memory – Concepts

Prof. Jisung Park



CAOS

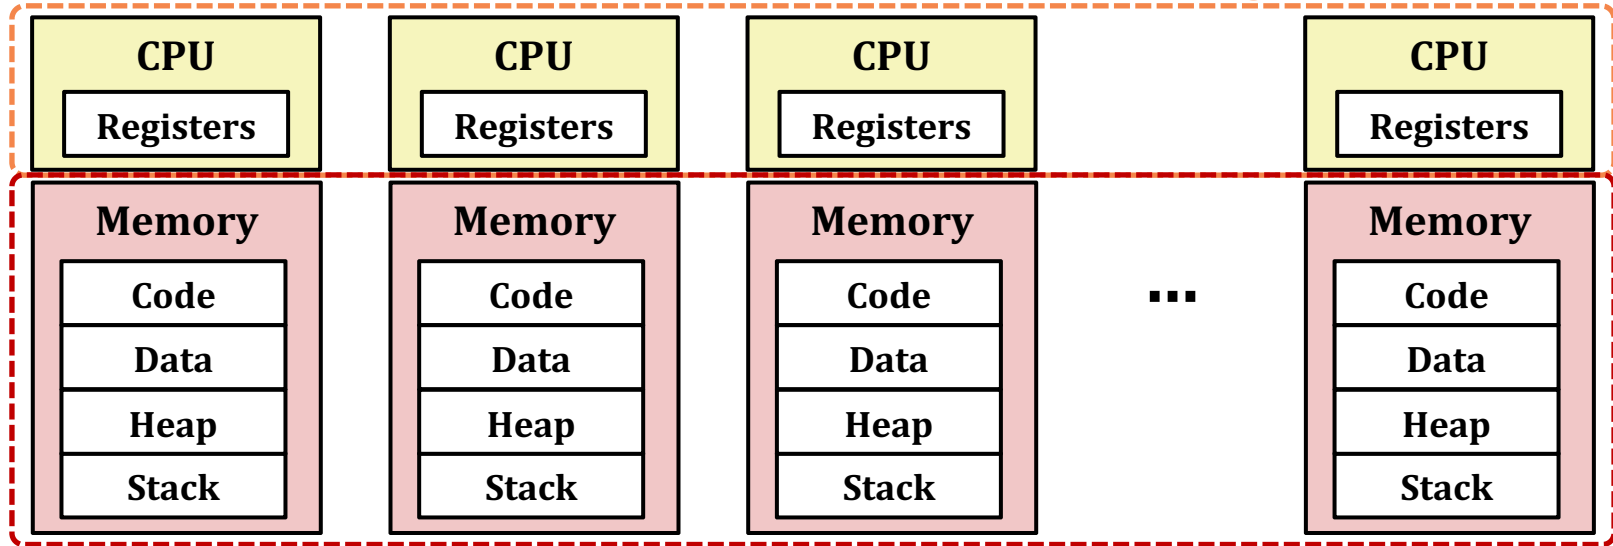
COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.12.06

Recall - Multiprocessing: The Illusion

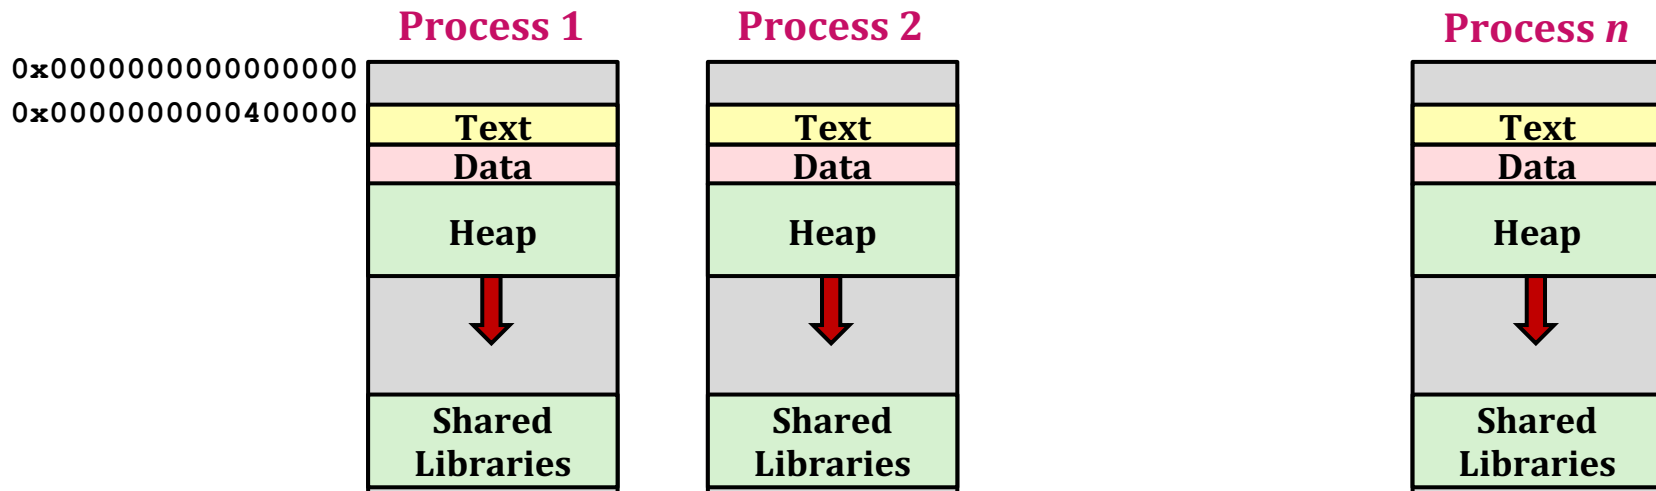
- Computer runs many processes simultaneously
 - Applications for one or more users: e.g., web browsers, email clients, editors, etc.
 - Background tasks: e.g., monitoring network & I/O devices

Process interleaving & context switch



Virtual memory (today's topic)

Virtual Memory: Concept



Exclusive (per-process), extremely large ($> 2^{47}$ bytes) memory space

How does this work?

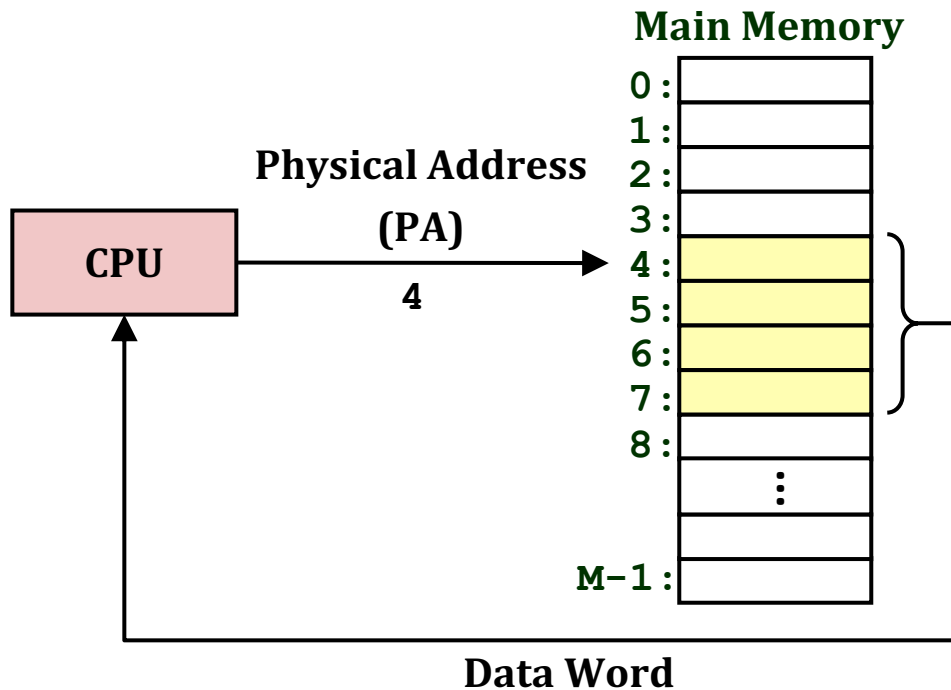


Lecture Agenda

- Address Spaces
- VM as a Tool for Caching
- VM as a Tool for Memory Management
- VM as a Tool for Memory Protection
- Address Translation

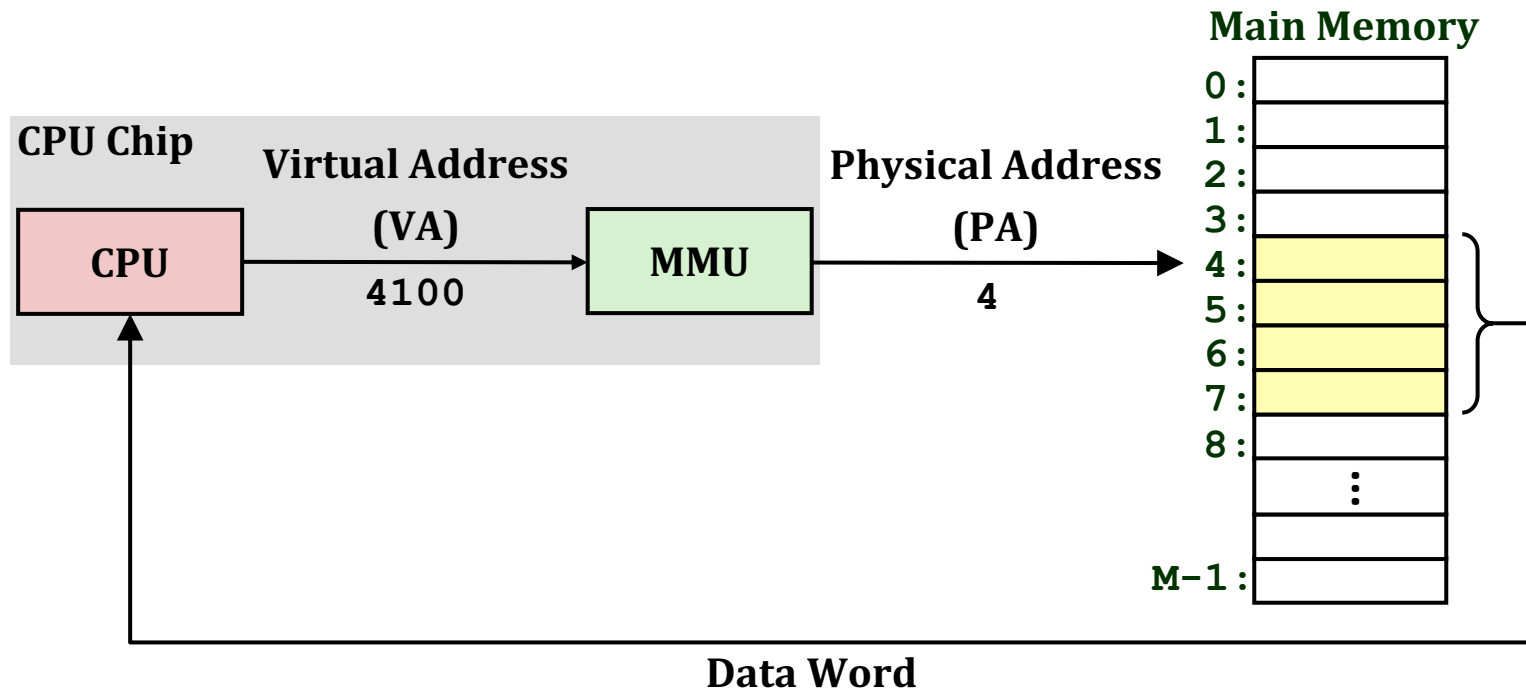
A System Using Physical Addressing

- Used in **simple systems**, e.g., cars, elevators, and digital picture frames



A System Using Physical Addressing

- One of the great ideas in computer science used in all modern servers, desktops, and laptops



Address Spaces

- **Linear address space**: an ordered set of contiguous non-negative integer addresses, i.e., $\{0, 1, 2, 3, \dots\}$
- **Virtual address space**: a set of $N = 2^n$ virtual addresses, i.e., $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space**: a set of $M = 2^m$ physical addresses, i.e., $\{0, 1, 2, 3, \dots, M-1\}$

Why Virtual Memory (VM)?

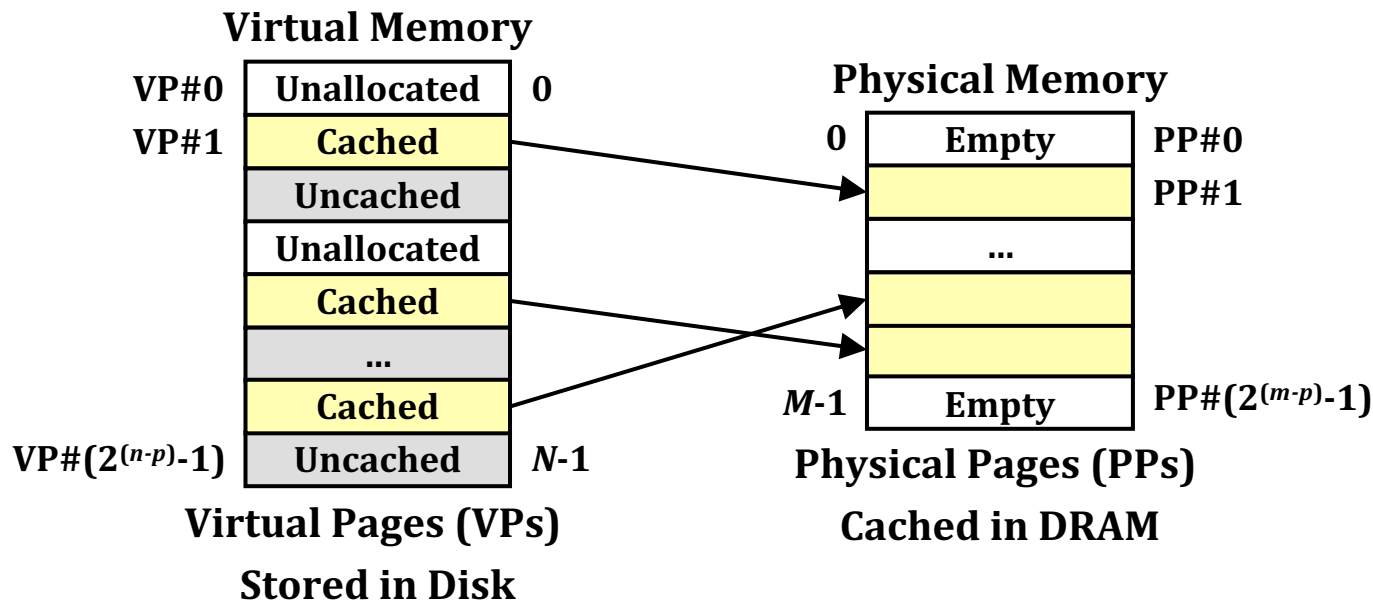
- **Efficiently uses** main memory
 - Uses DRAM as a cache for the parts of a virtual address space
- **Simplifies** memory management
 - Each process gets the same uniform linear address space
- **Isolates** address spaces
 - One process cannot interfere with another's memory
 - User programs cannot access privileged kernel information

Lecture Agenda

- Address Spaces
- VM as a Tool for Caching
- VM as a Tool for Memory Management
- VM as a Tool for Memory Protection
- Address Translation

VM as a Tool for Caching

- Conceptually, **virtual memory** is an N-byte array stored in a disk
 - Part of the array is cached in (M-byte) **physical memory** (e.g., in **DRAM**)
 - These cache blocks are called **pages** (size is $P = 2^p$ bytes)

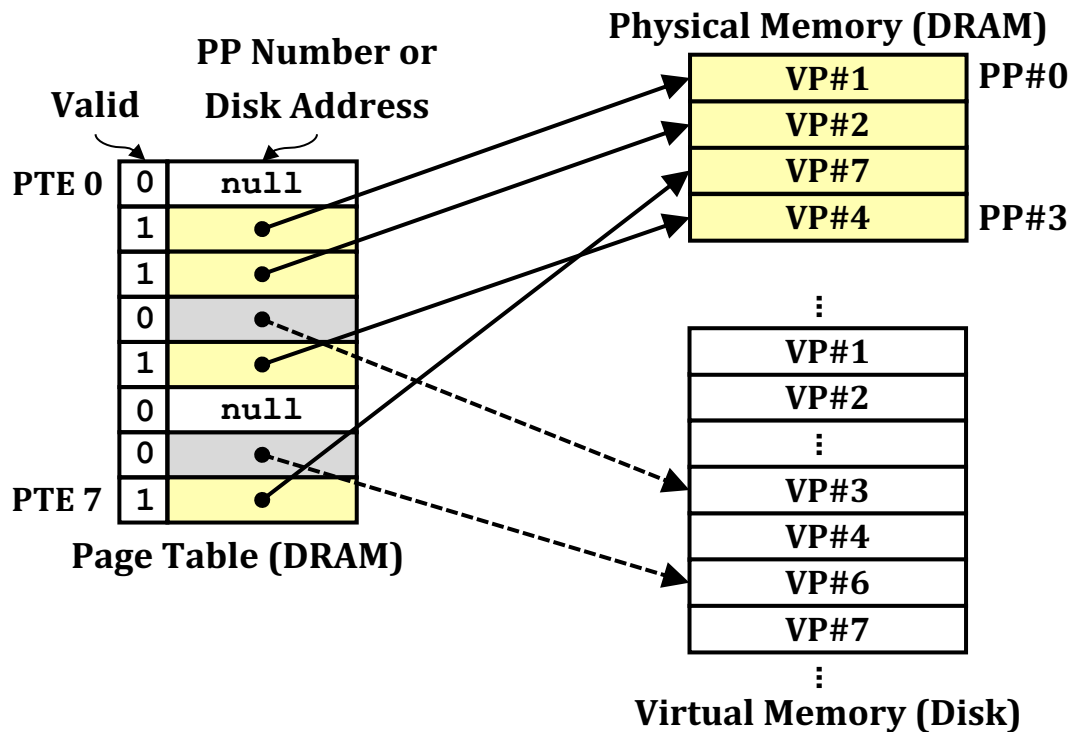


DRAM Cache Organization

- Driven by the **enormous miss penalty**
 - DRAM is about **10×** slower than SRAM
 - Disk is about **10,000×** slower than DRAM
- Consequences
 - **Large page (i.e., cache block) size**: typically 4-8 KiB, sometimes 4 MiB
 - **Fully associative**
 - Any VP can be placed in any PP
 - Requires a **large mapping function** – different from CPU caches
 - **Highly sophisticated, expensive replacement algorithms**
 - Too complicated and open-ended to be implemented in hardware
 - **Write-back** rather than write-through

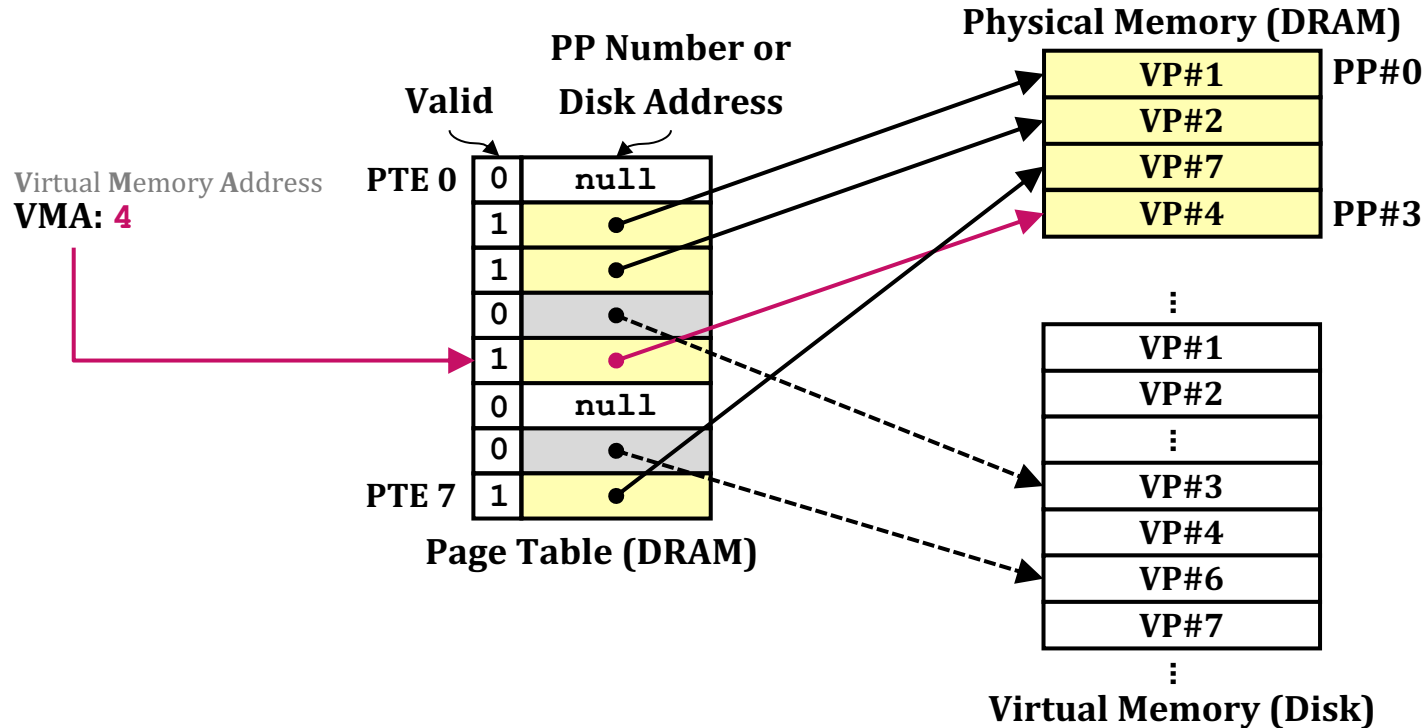
Enabling Data Structure: Page Table

- An array of **page table entries (PTEs)** storing **virtual-to-physical mappings**
 - **Per-process kernel data structure** in DRAM



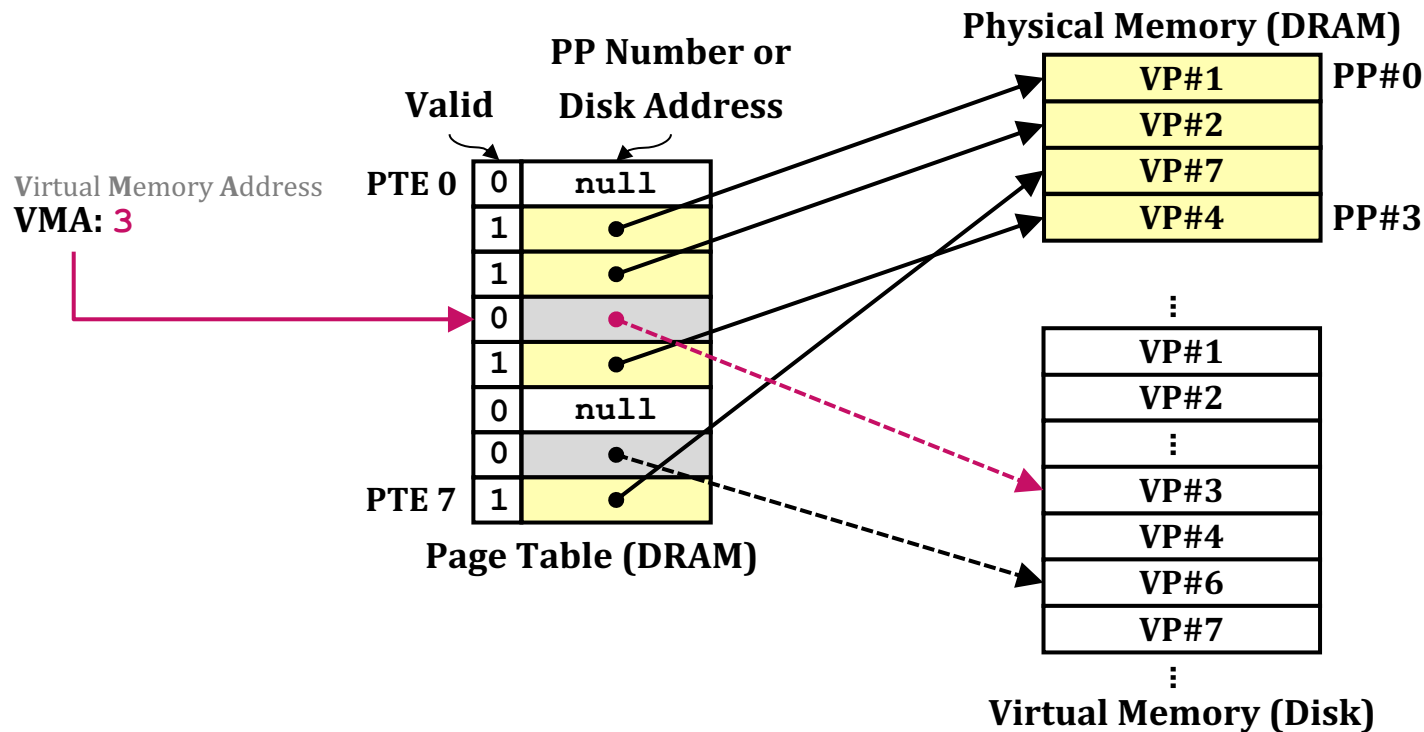
Page Hit

- Reference to a VM word **in physical memory**, i.e., DRAM cache hit




Page Fault

- Reference to a VM word **not in physical memory**, i.e., DRAM cache miss



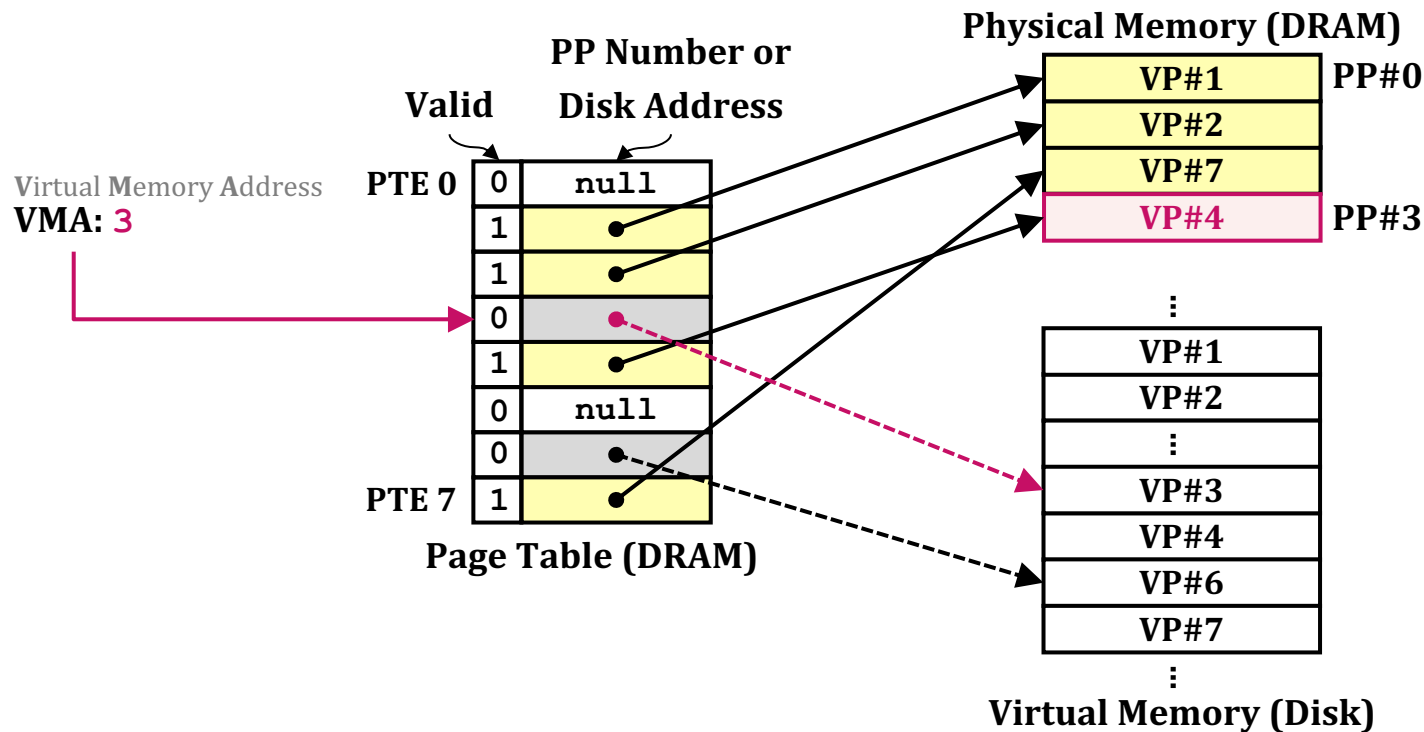
Page Fault Handling

- A page fault is treated as **an exception** (caused by a page miss)



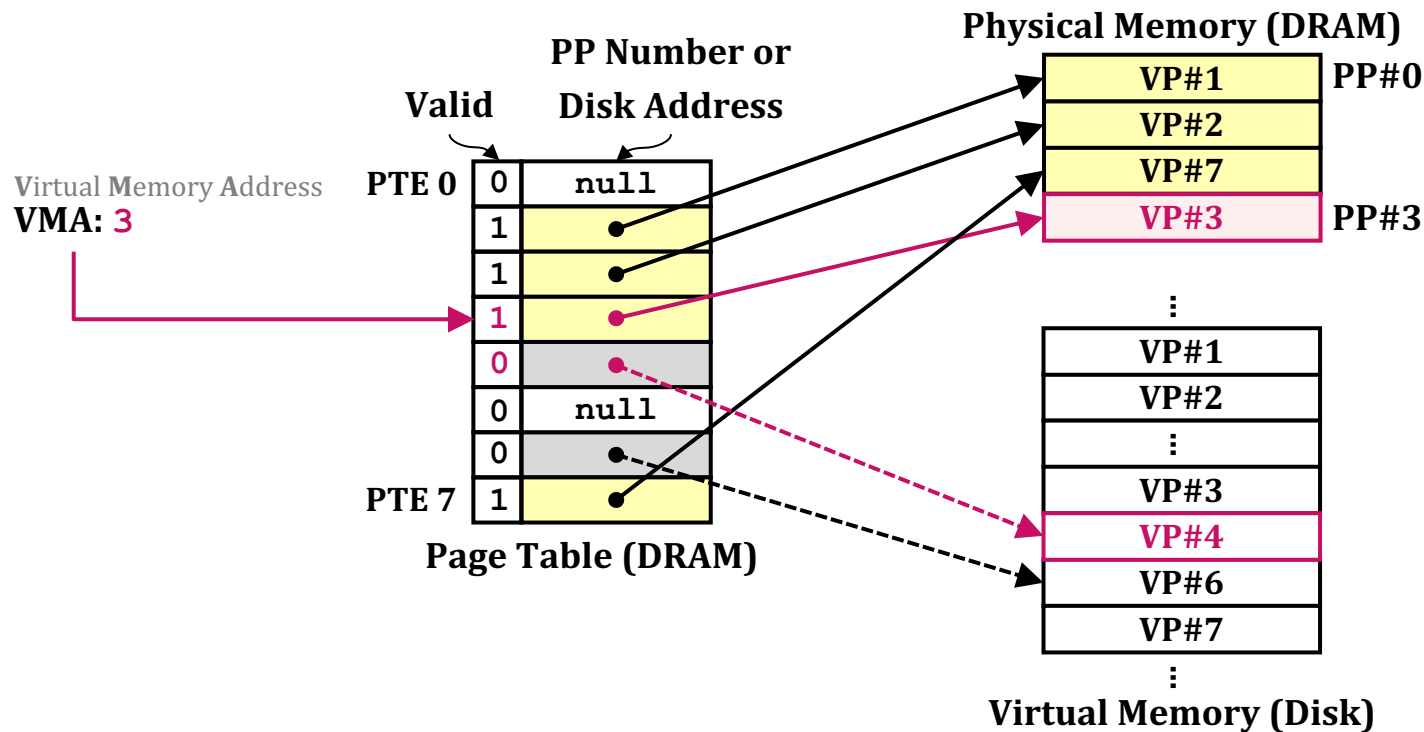
Page Fault Handling

- A page fault is treated as **an exception** (caused by a page miss)
 - Page fault handler **evicts a victim** (e.g., VP#4)



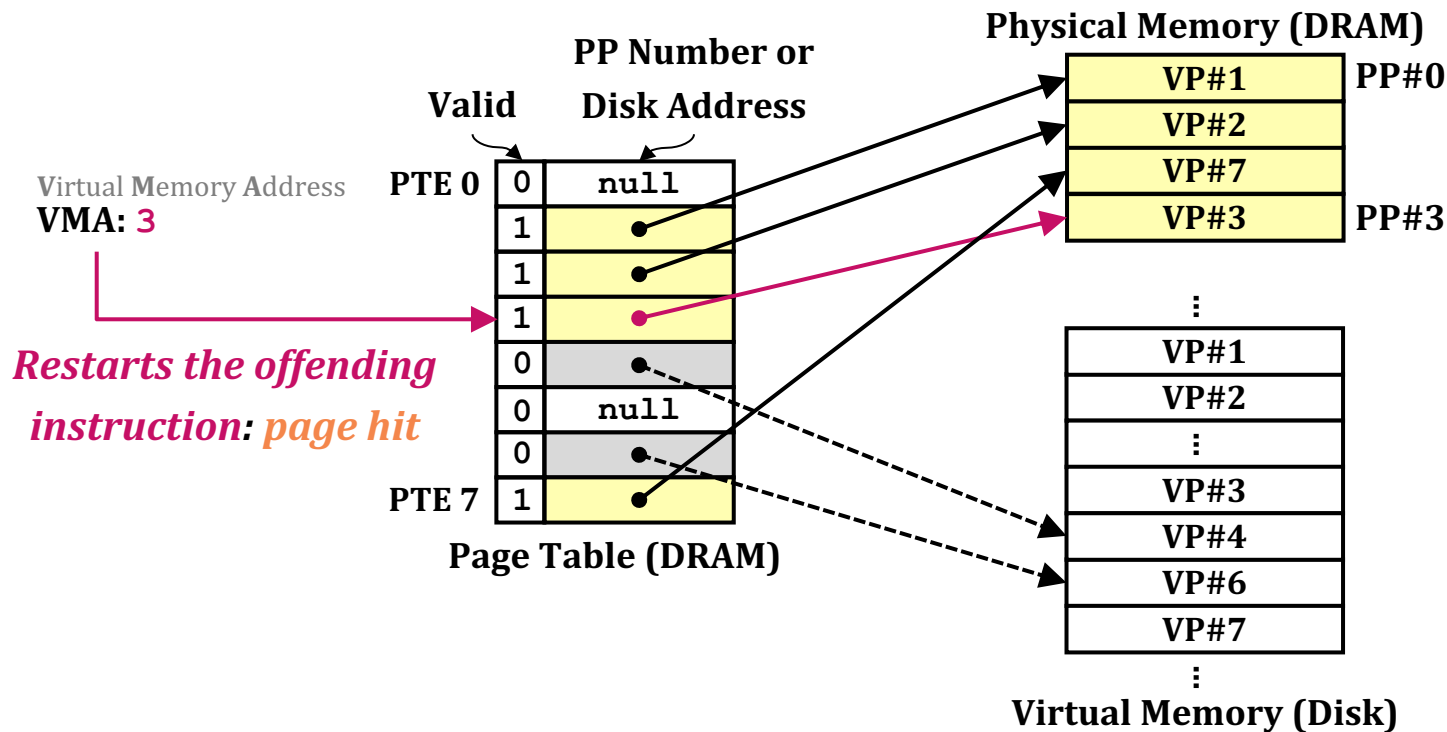
Page Fault Handling

- A page fault is treated as **an exception** (caused by a page miss)
 - Page fault handler **evicts a victim** (e.g., VP#4) and **fetch the target** (e.g., VP#3)



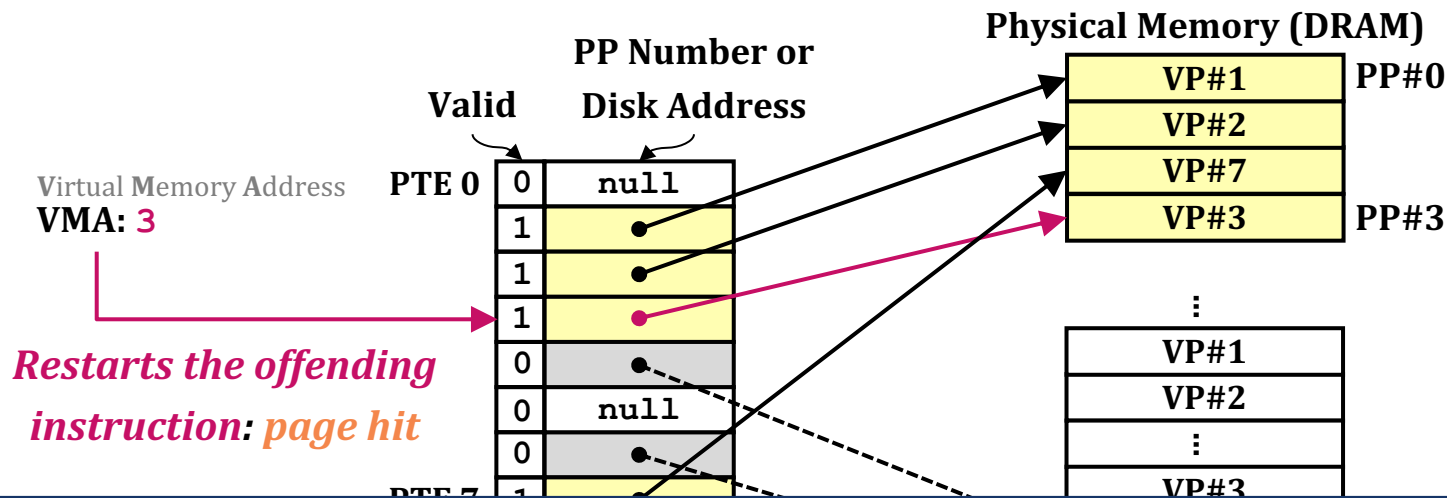
Page Fault Handling

- A page fault is treated as **an exception** (caused by a page miss)
 - Page fault handler **evicts a victim** (e.g., VP#4) and **fetch the target** (e.g., VP#3)



Page Fault Handling

- A page fault is treated as **an exception** (caused by a page miss)
 - Page fault handler **evicts a victim** (e.g., VP#4) and **fetch the target** (e.g., VP#3)

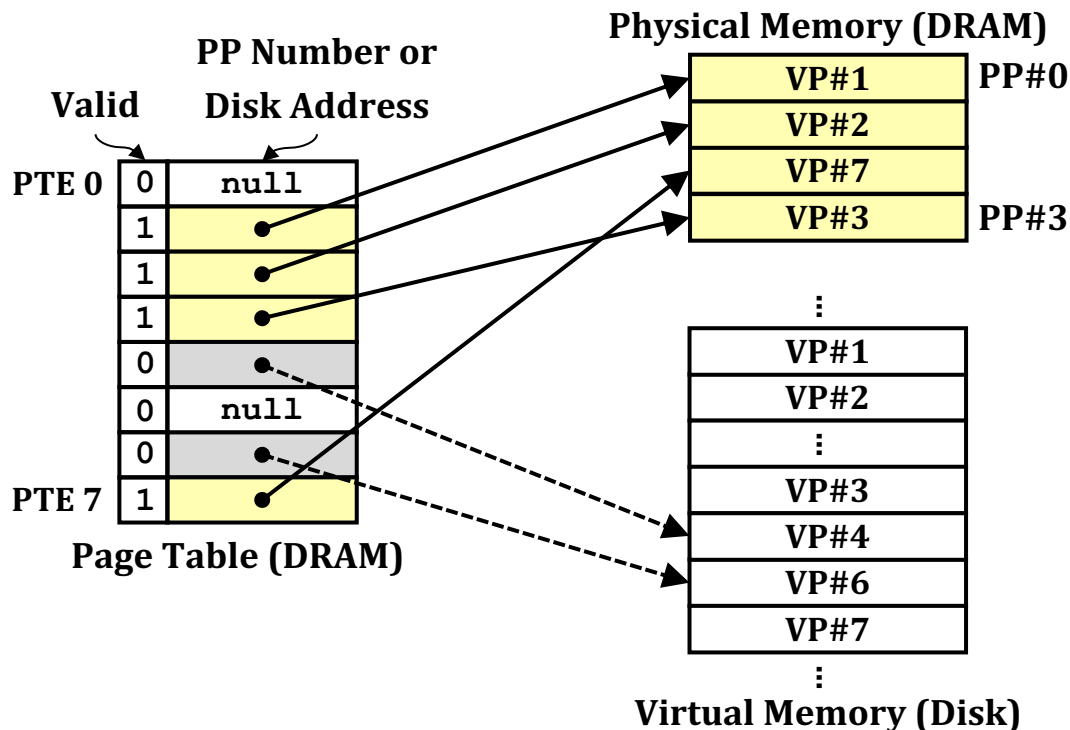


Demand paging: waiting until the miss to copy a page to DRAM

Virtual Memory (Disk)

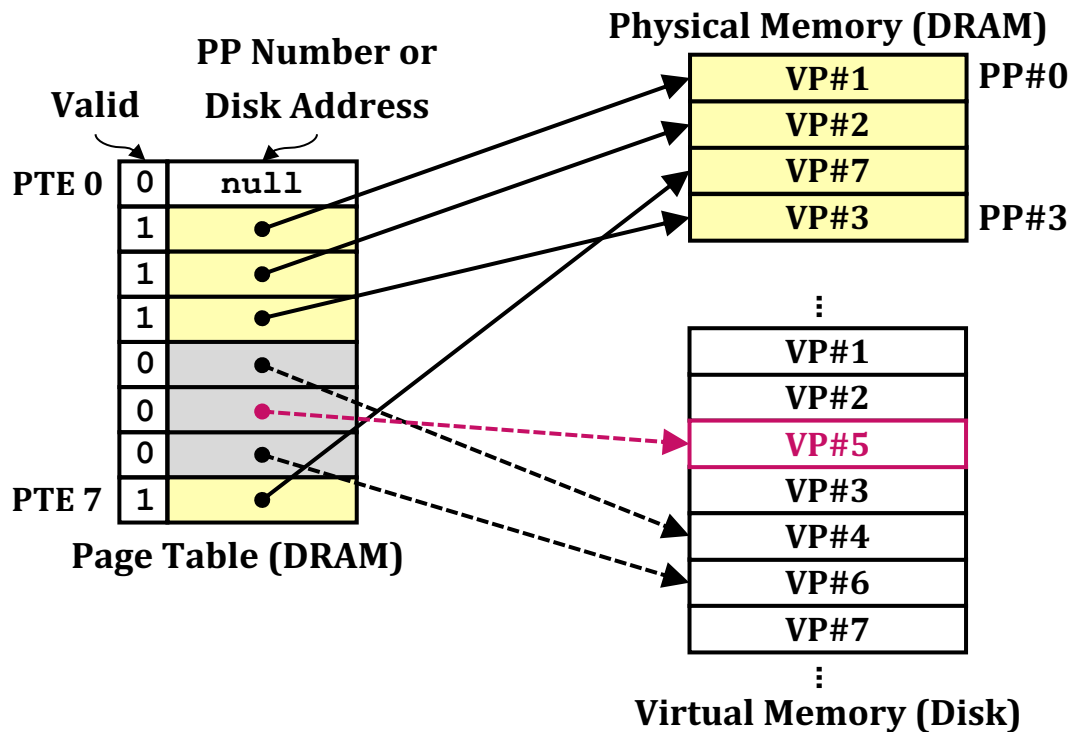
Page Allocation

- Only sets the corresponding PTE (e.g., for VP#5)
 - Subsequent page fault will bring the page into memory



Page Allocation

- Only sets the corresponding PTE (e.g., for VP#5)
 - Subsequent page fault will bring the page into memory



Locality to the Rescue Again

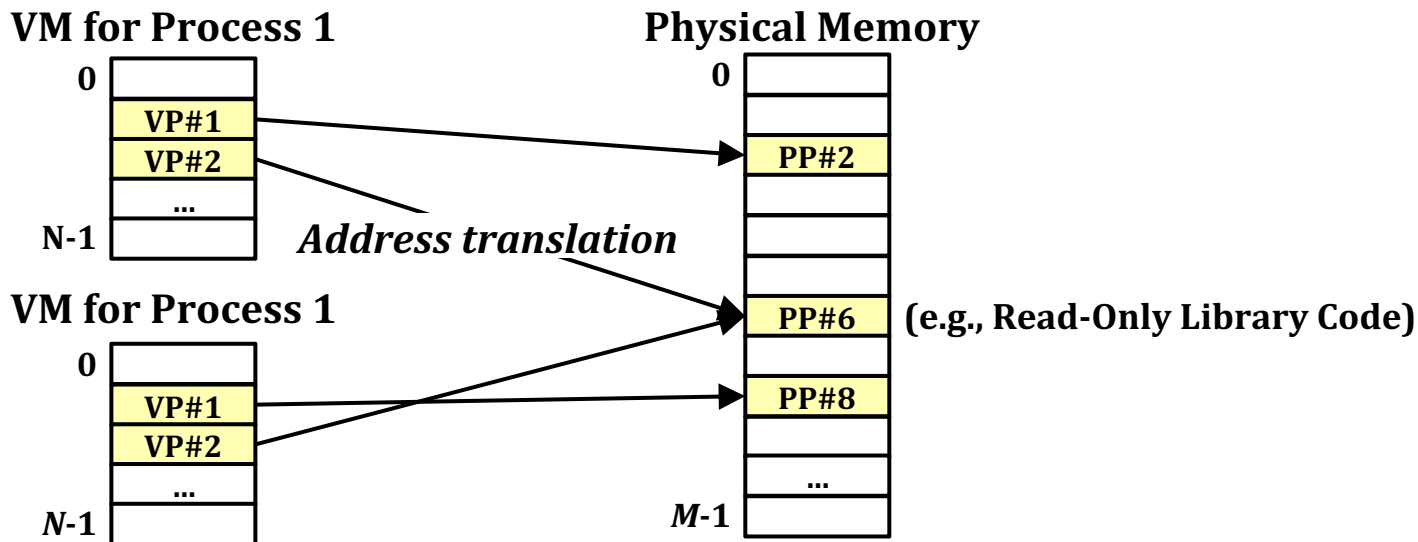
- Virtual memory **may seem terribly** inefficient but **works because of locality**
- At any point in time, programs tend to access **a set of active virtual pages**
 - Which is called the **working set**
 - Programs with better temporal locality will have smaller working sets
- If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- If (total working set size > main memory size)
 - **Cache thrashing**: performance meltdown where pages are swapped (copied) in and out continuously

Lecture Agenda

- Address Spaces
- VM as a Tool for Caching
- VM as a Tool for Memory Management
- VM as a Tool for Memory Protection
- Address Translation

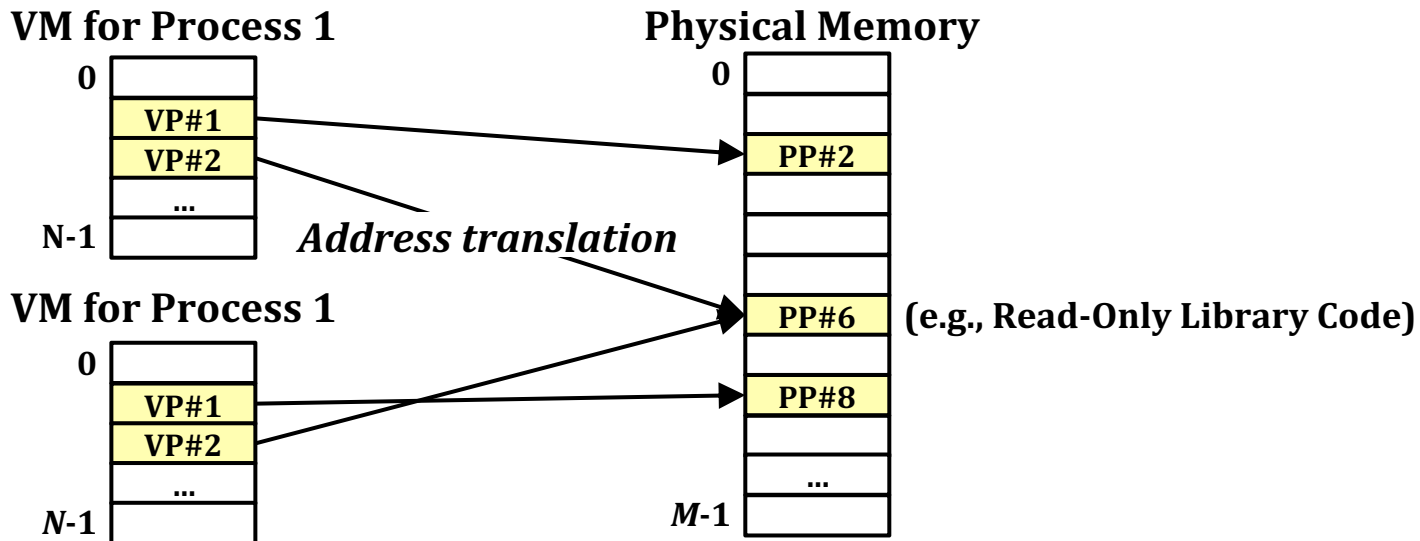
VM as a Tool for Memory Management

- Each process has its **own virtual address space**
 - Allows each process to view **memory as a simple linear array**
 - Mapping function scatters addresses through physical memory
 - Well chosen mappings simplify memory allocation and management



VM as a Tool for Memory Management (Cont.)

- Simplifying memory allocation
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- Sharing code and data among processes
 - Map virtual pages to the same physical page (e.g., PP#6)



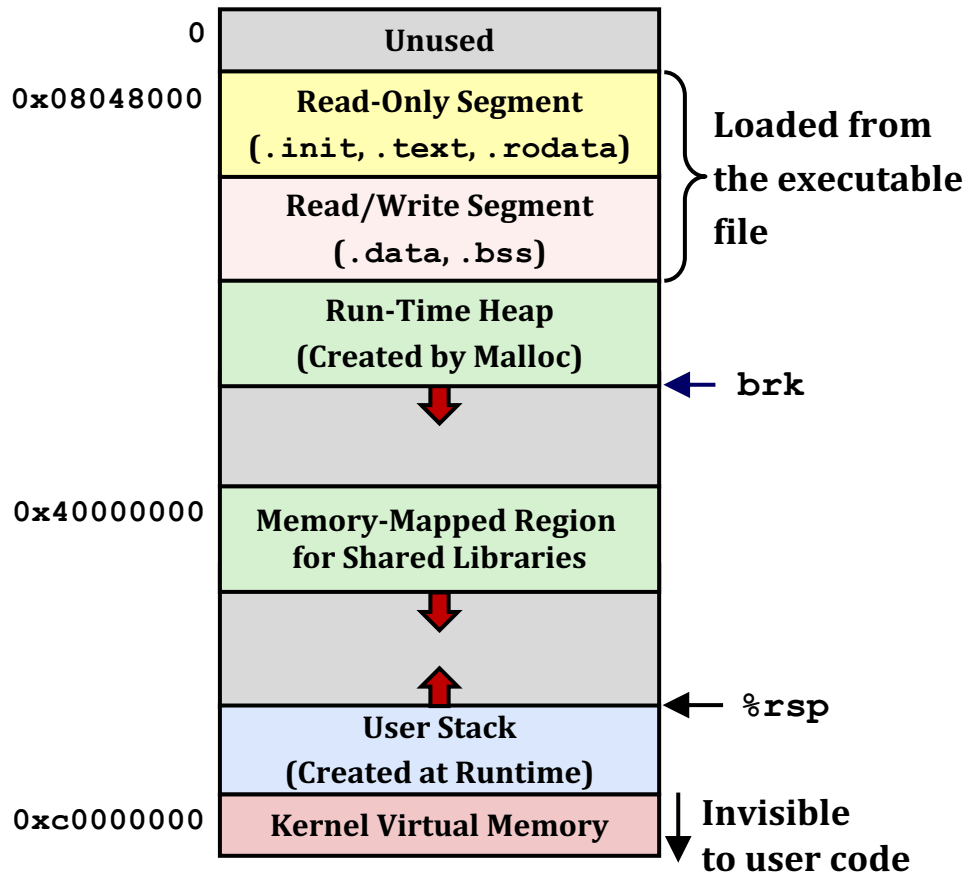
Simplifying Linking and Loading

- Linking

- Each program has similar virtual address space
- Code, stack, and shared libraries always at the same address

- Loading

- `execve()` allocates virtual pages for `.text` and `.data` sections: creates PTEs marked as invalid
- The `.text` and `.data` sections are copied, paged by page, on demand by the virtual memory system

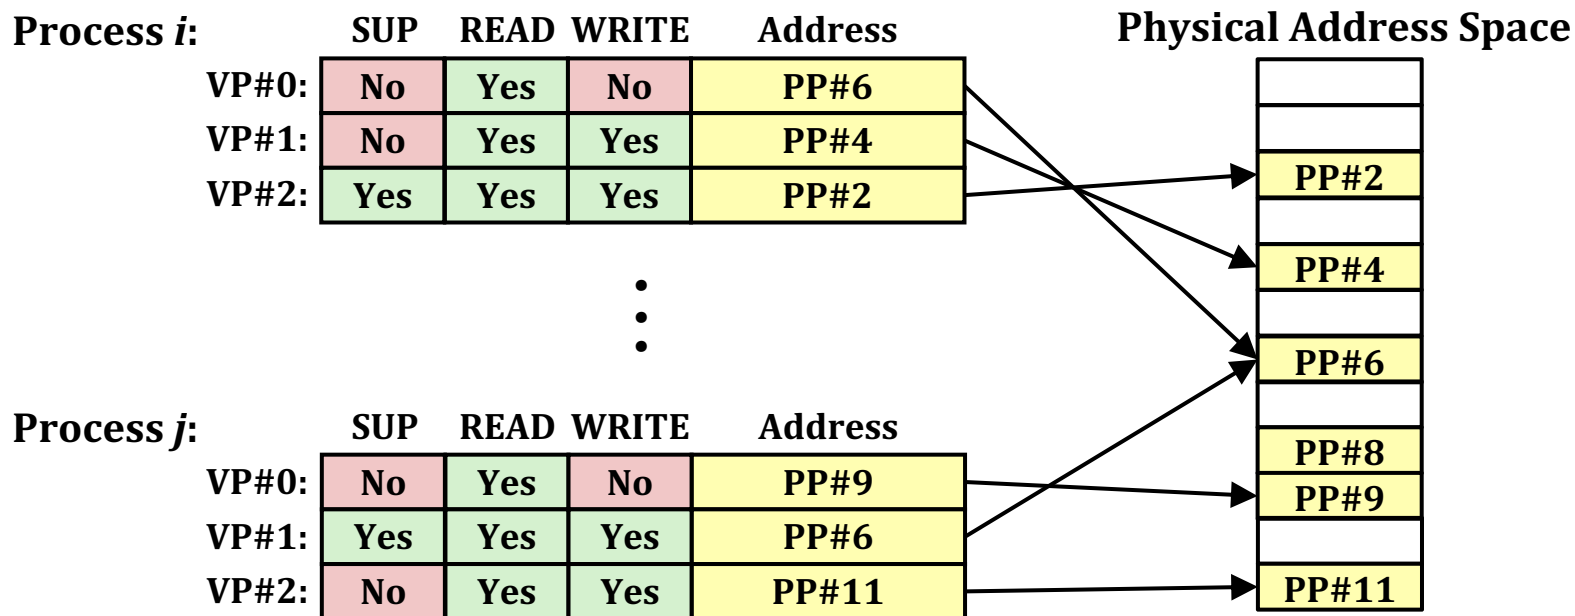


Lecture Agenda

- Address Spaces
- VM as a Tool for Caching
- VM as a Tool for Memory Management
- **VM as a Tool for Memory Protection**
- Address Translation

VM as a Tool for Memory Protection

- Extend PTEs with permission bits
 - Page fault handler checks these before remapping
 - If violated, send process **SIGSEGV** (segmentation fault)



Lecture Agenda

- Address Spaces
- VM as a Tool for Caching
- VM as a Tool for Memory Management
- VM as a Tool for Memory Protection
- Address Translation

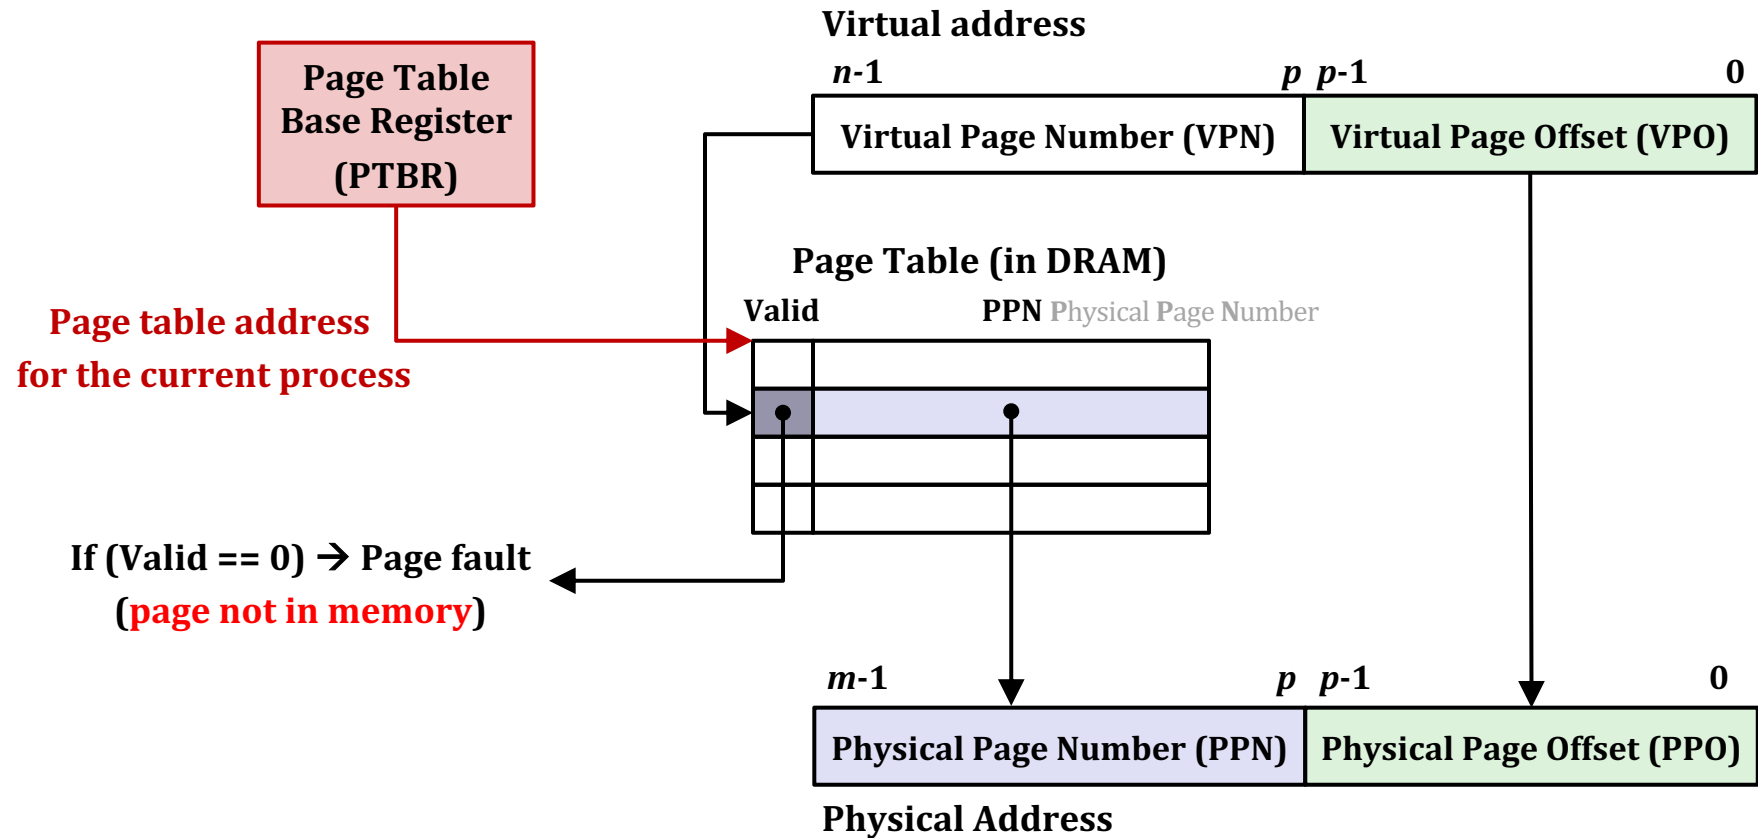
VM Address Translation

- **Virtual address space:** $V = \{0, 1, \dots, N-1\}$
- **Physical address space:** $P = \{0, 1, \dots, M-1\}$
- **Address translation**
 - $\text{MAP}: V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - $\text{MAP}(a) = a'$ if data at virtual address a is at physical address $a' \in P$
 - $\text{MAP}(a) = \emptyset$ if data at virtual address a is not in physical memory
 - Either invalid or stored on disk

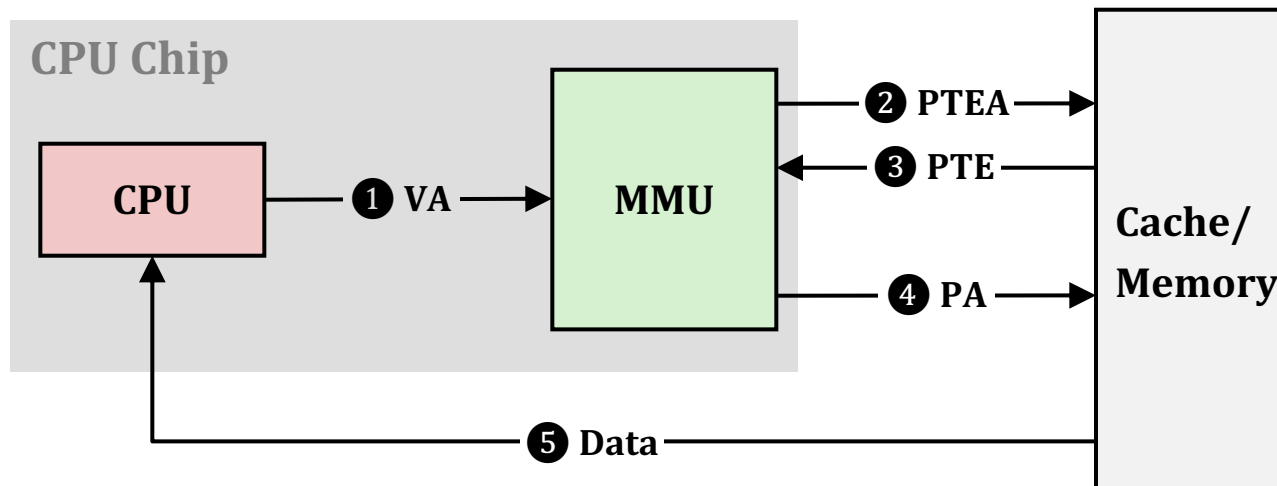
Summary of Address Translation Symbols

- Basic Parameters
 - $N = 2^n$: the number of addresses in virtual address space
 - $M = 2^m$: the number of addresses in physical address space
 - $P = 2^p$: the page size (bytes)
- Components of the virtual address (VA)
 - VPO: virtual page offset
 - VPN: virtual page number
- Components of the physical address (PA)
 - PPO: physical page offset (same as VPO)
 - PPN: physical page number
 - CO: byte offset within cache line
 - CI: cache index
 - CT: cache tag

Address Translation with a Page Table

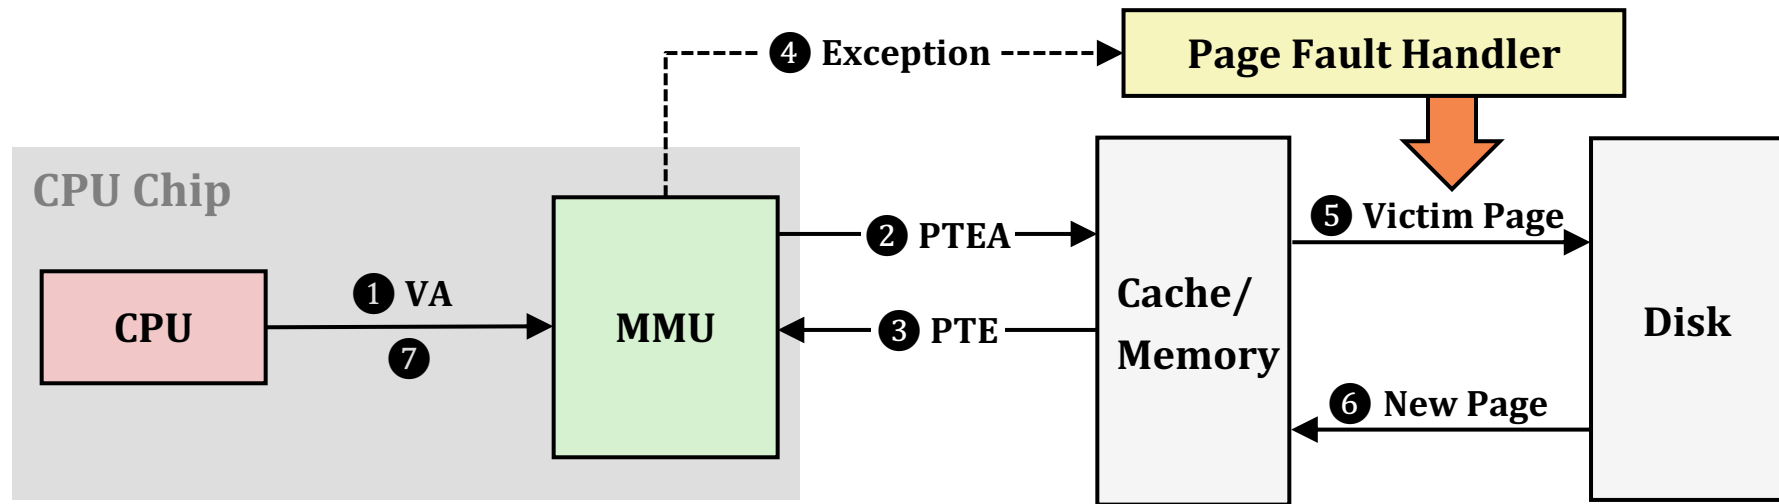


Address Translation: Page Hit



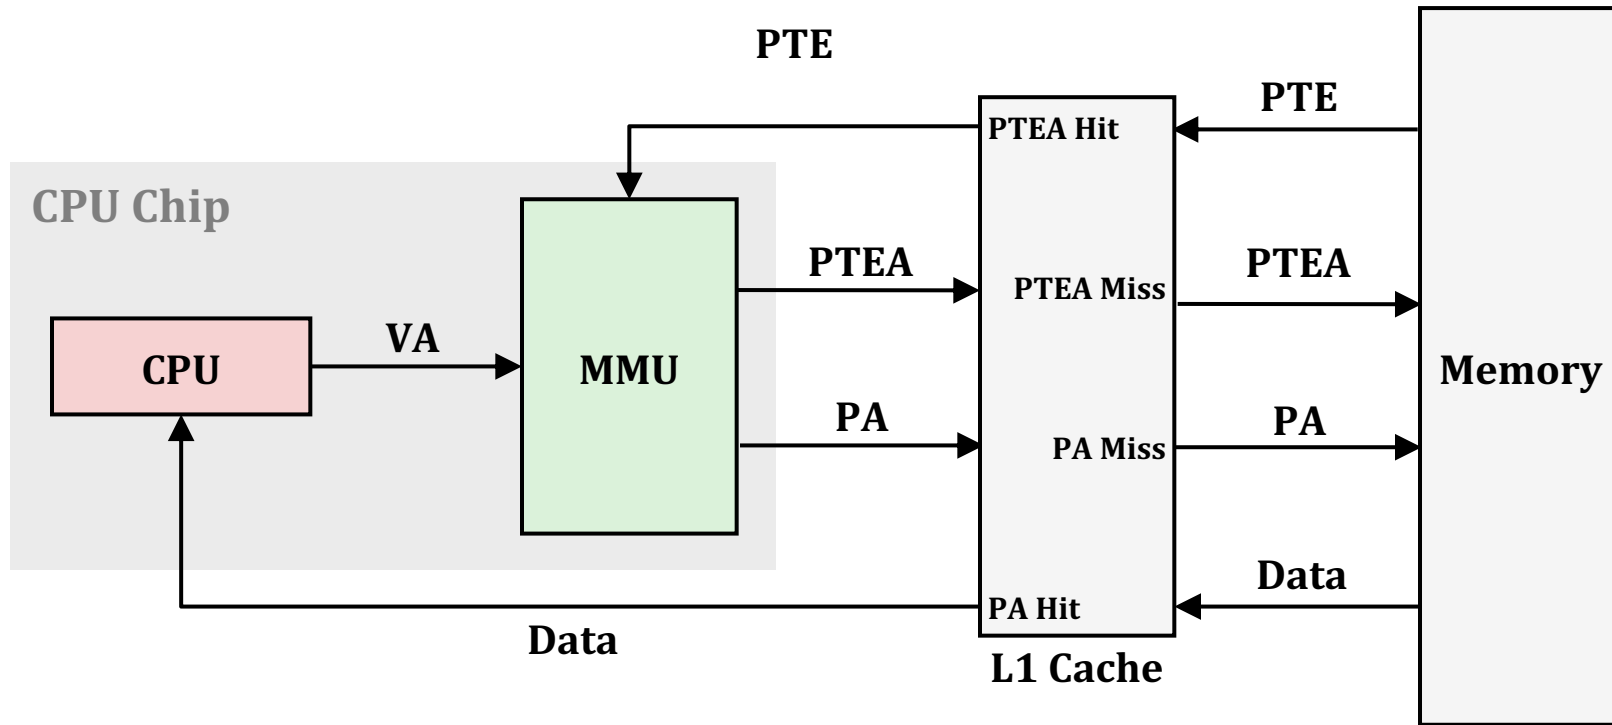
- ❶ Processor sends virtual address to MMU
- ❷, ❸ MMU fetches PTE from page table in memory
- ❹ MMU sends physical address to cache/memory
- ❺ Cache/memory sends data word to processor

Address Translation: Page Fault



- ➊ Processor sends virtual address to MMU
- ➋, ➌ MMU fetches PTE from page table in memory
- ➍ Valid bit is zero, so MMU triggers page fault exception
- ➎ Handler identifies victim (and, if dirty, pages it out to disk)
- ➏ Handler pages in new page and updates PTE in memory
- ➐ Handler returns to original process, restarting faulting instruction

Integrating VM and Cache



VA: Virtual Address, **PA:** Physical Address, **PTE:** Page Table Entry, **PTEA:** PTE Address

Speeding Up Translation with a TLB

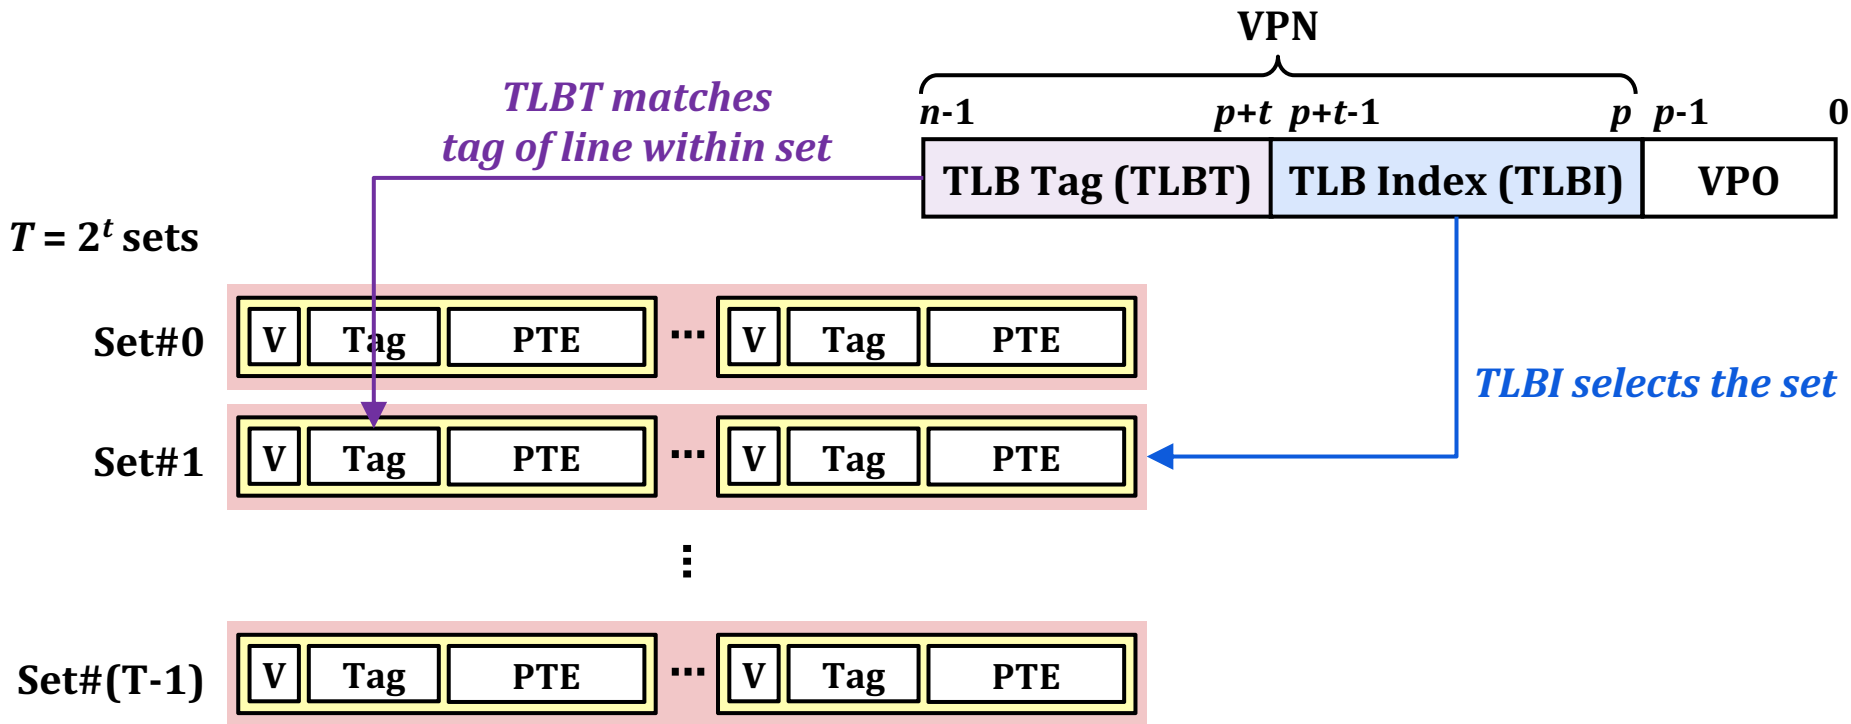
- **Problem:** PTEs are cached in L1 like any other memory word
 - PTEs **may be evicted** by other data references
 - PTE hit still requires **a small L1 delay**
- **Solution:** **translation lookaside buffer** (TLB)
 - Small hardware cache in MMU
 - Stores virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

Summary of Address Translation Symbols

- Basic Parameters
 - $N = 2^n$: the number of addresses in virtual address space
 - $M = 2^m$: the number of addresses in physical address space
 - $P = 2^p$: the page size (bytes)
- Components of the virtual address (VA)
 - **TLBI**: TLB index
 - **TLBT**: TLB tag
 - **VPO**: virtual page offset
 - **VPN**: virtual page number
- Components of the physical address (PA)
 - **PPO**: physical page offset (same as VPO)
 - **PPN**: physical page number

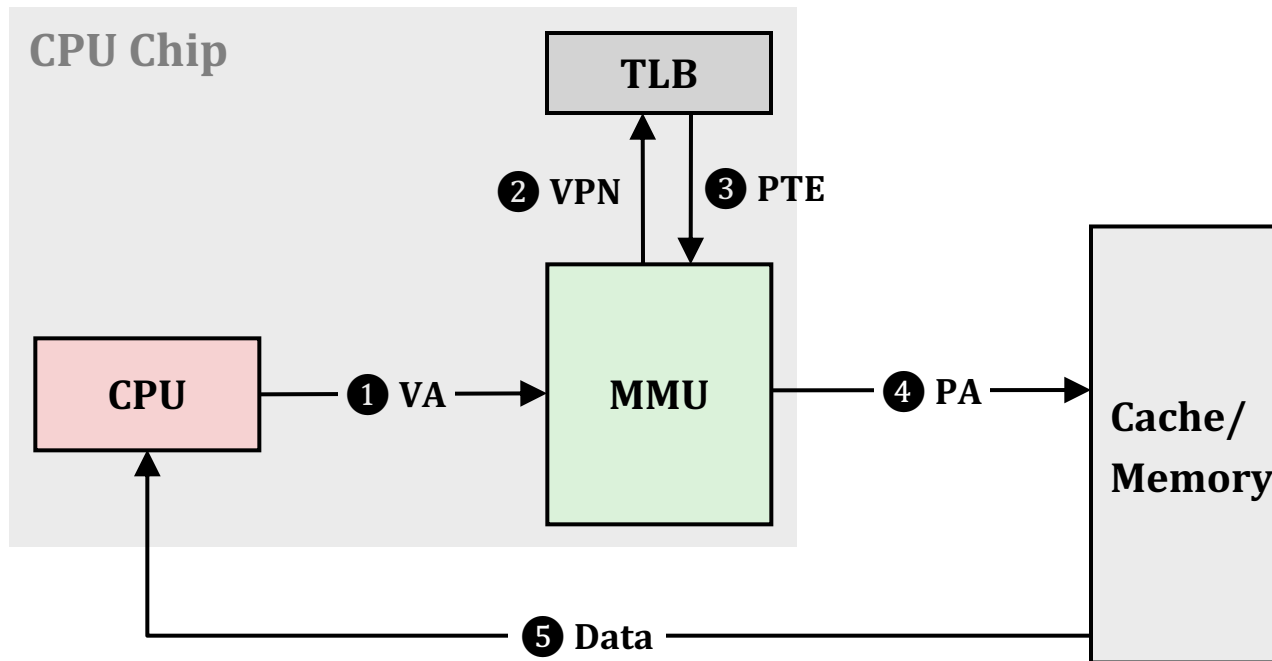
Accessing the TLB

- MMU uses the VPN portion of the virtual address to access the TLB:



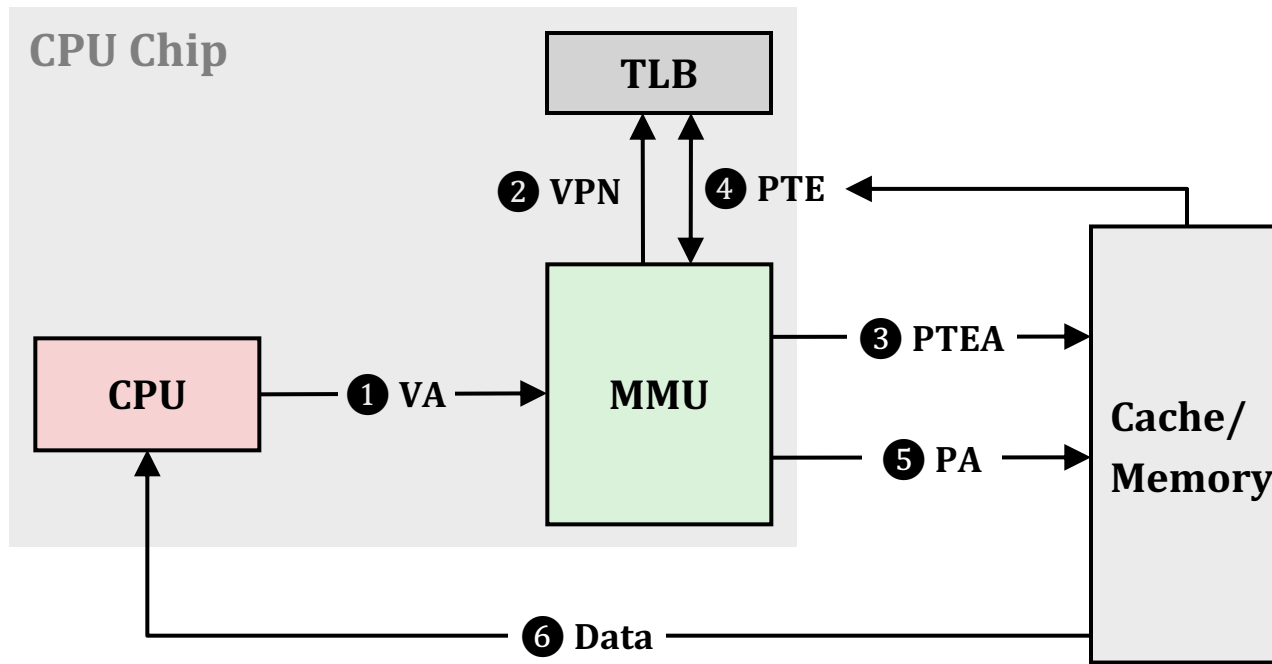
TLB Hit

- A TLB hit eliminates a memory access




TLB Miss

- A TLB miss incurs an additional memory access (the PTE)
 - Fortunately, TLB misses are rare. Why?

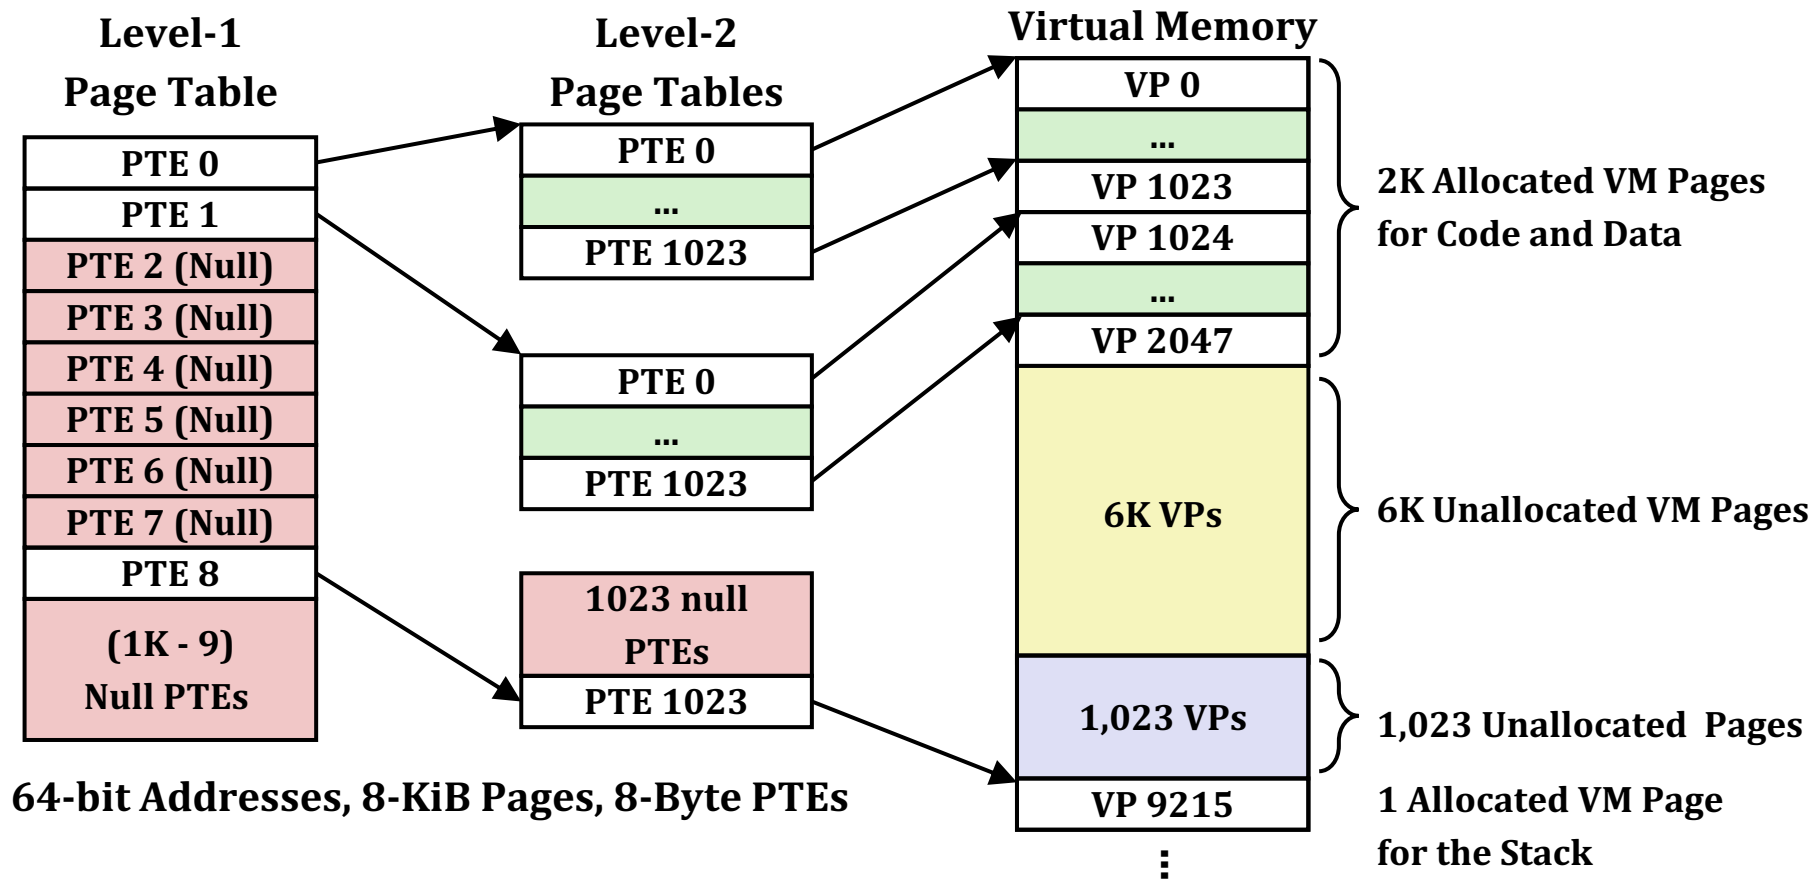


Multi-Level Page Tables

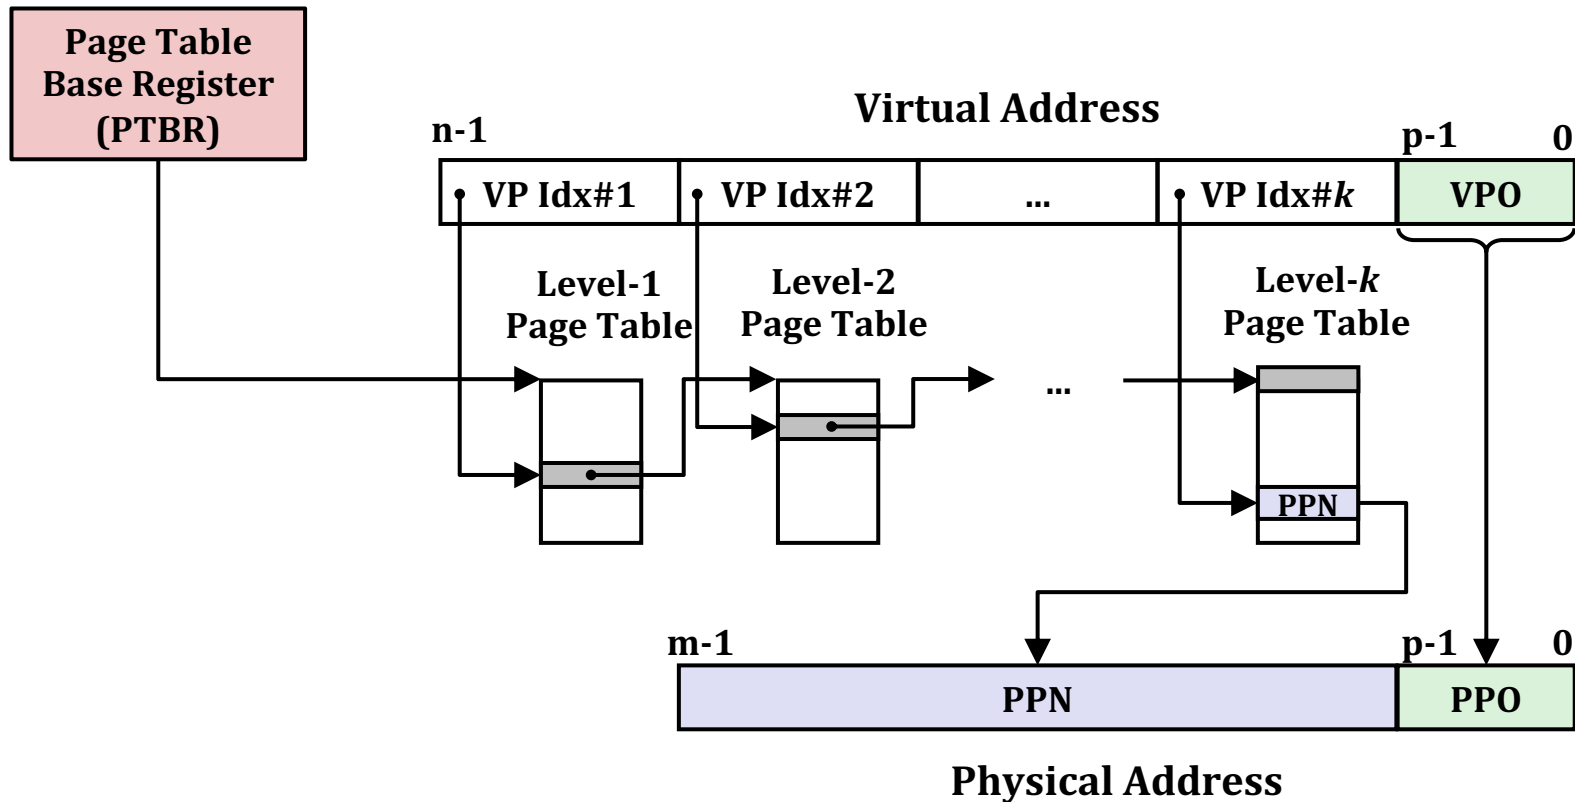
- Assumptions: 4-KiB (2^{12}) page size, 48-bit address space, 8-byte PTE
- **Problem:** Would need a 512 GB page table!
 - $2^{48} \times 2^{-12} \times 2^3 = 2^{39}$ bytes
- **Common solution:** Multi-level page tables
- Example: 2-level page table
 - Level-1 table: each PTE points to a page table, i.e., a page that stores part of Level-2 page table (always memory resident)
 - Level-2 table: each PTE points to a page (paged in and out like any other data)



A Two-Level Page Table Hierarchy



Translating with a k -Level Page Table



Summary

- Programmer's view of virtual memory
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- System view of virtual memory
 - Uses memory efficiently by caching virtual memory pages
 - Efficient only because of locality
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient inter-positioning point to check permissions

[CSED211] Introduction to Computer Software Systems

Lecture 15: Virtual Memory – Concepts

Prof. Jisung Park



CAOS

COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.12.06