

[CSED211] Introduction to Computer Software Systems

Lecture 2: Bits, Bytes, Integers

Prof. Jisung Park



CAOS

COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

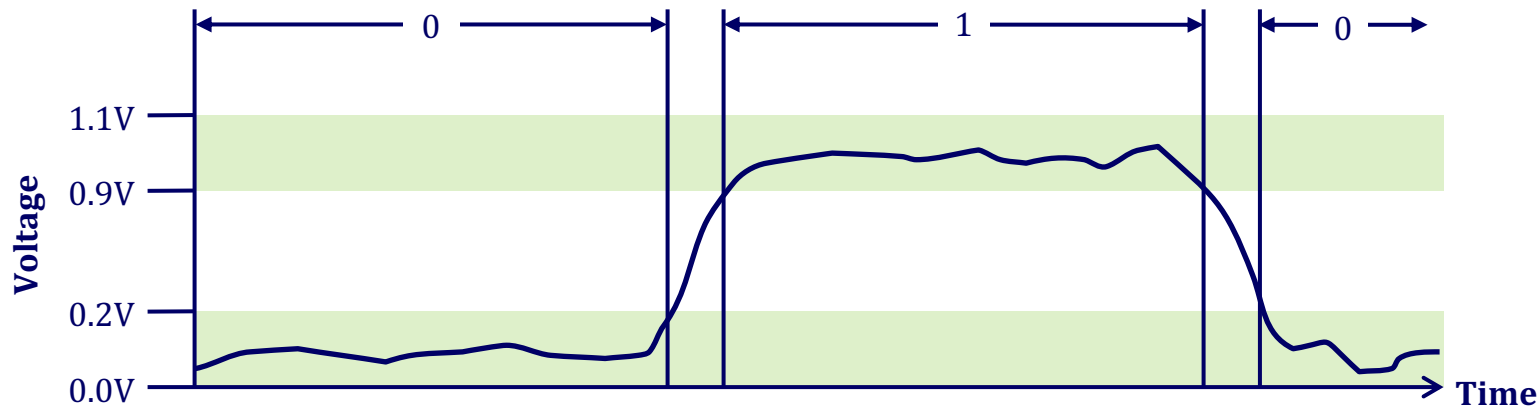
2023.09.06

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Everything is bits

- Each bit is 0 or 1
- By encoding and interpreting sets of bits in various ways, a computer
 - Determines what to do (i.e., instructions)
 - Represents and manipulates numbers, sets, strings, etc.
- Why bits? Electronic implementation
 - Easy to represent data with bi-stable elements
 - Reliably transmitted on noisy and inaccurate wires



Example: Representation of Numerical Values

- Base 2 number representation
 - Represent $15213_{(10)}$ as $11101101101101_{(2)}$
 - Represent $1.20_{(10)}$ as $1.0011001100110011[0011...]_{(2)}$
 - Represent $1.5213_{(10)} \times 10^4$ as $1.1101101101101_{(2)} \times 2^{13}$

Encoding Byte Values

- Byte = 8 bits
 - Binary: $00000000_{(2)}$ to $11111111_{(2)}$
 - Decimal: $000_{(10)}$ to $255_{(10)}$
 - Hexadecimal: $00_{(16)}$ to $\text{FF}_{(16)}$
 - Base 16 number representation
 - Use symbols '0' to '9' and 'A' to 'F'
 - Write $\text{FA1D37B}_{(16)}$ in C as
 - **0x**FA1D37B
 - **0x**fa1d37b
 - **0x**Fa1D37b

Decimal	Binary	HEX
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
10	1010	A
11	1011	B
12	1100	C
13	1101	D
14	1110	E
15	1111	F

Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	8	8
float	4	4	4
double	8	8	8
long double	–	–	10/16
Pointer (word)	4	8	8

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Boolean Algebra

- Developed by George Boole in 19th century
 - Algebraic representation of logic
 - Encode “True” as 1 and “False” as 0

And

- $A \& B = 1$ when both $A=1$ and $B=1$

$\&$	0	1
0	0	0
1	0	1

Or

- $A \mid B = 1$ when either $A=1$ or $B=1$

\mid	0	1
0	0	1
1	1	1

Not

- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0

Exclusive-Or (Xor)

- $A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

\wedge	0	1
0	0	1
1	1	0

General Boolean Algebras

- Operate on bit vectors
 - Operations applied **bitwise**

01101001	01101001	01101001	
<u>& 01010101</u>	<u> 01010101</u>	<u>^ 01010101</u>	<u>~ 01010101</u>
01000001	01111101	00111100	10101010

- All of the properties of boolean algebra hold (e.g., De Morgan's Laws)

Example: Representing & Manipulating Sets

- Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$
- $a_j = 1$ if $j \in A$
 - $a = 01101001$ for $A = \{0, 3, 5, 6\}$
 - $b = 01010101$ for $B = \{0, 2, 4, 6\}$

- Operations

- $\&$ Intersection $a \& b = 01000001$ $A \cap B = \{0, 6\}$
- $|$ Union $a | b = 01111101$ $A \cup B = \{0, 2, 3, 4, 5, 6\}$
- \wedge Symmetric difference $a \wedge b = 00111100$ $(A \cup B) - (A \cap B) = \{2, 3, 4, 5\}$
- \sim Complement $\sim a = 10010110$ $A^C = \{1, 2, 4, 7\}$

Bit-Level Operations in C

- Operations `&`, `|`, `~`, `^` are available in C
 - Apply to any “integral” data type
 - `long`, `int`, `short`, `char`, `unsigned`
 - View arguments as bit vectors
 - Arguments applied bitwise
- Examples (`char` data type)
 - `~0x41 → 0xBE`
 - `~01000001 → 10111110`
 - `~0x00 → 0xFF`
 - `~00000000 → 11111111`
 - `0x69 & 0x55 → 0x41`
 - `01101001 & 01010101 → 01000001`
 - `0x69 | 0x55 → 0x7D`
 - `01101001 | 01010101 → 01111101`

Comparison: Logic Operations in C

- Logical operators

- `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination

- Examples (`char` data type)

- `!0x41 → 0x00`
- `!0x00 → 0x01`
- `!!0x41 → 0x01`
- `0x69 && 0x55 → 0x01`
- `0x69 || 0x55 → 0x01`
- `p && *p` (avoids null pointer access)

Watch out for `&&` vs. `&` (and `||` vs. `|`, ...), one of the common oopsies in C programming

Shift Operations

- **Left shift:** $x \ll y$
 - Shift bit-vector x to left by y positions
 - Throw away extra bits on left
 - Fill with 0's on right
- **Right shift:** $x \gg y$
 - Shift bit-vector x to right by y positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- **Undefined behavior (machine-specific)**
 - When shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Encoding Integers

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

- `short` in C: 2-byte long

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011

Sign Bit



- Sign bit
 - For 2's complement, most significant bit indicates sign ($\because 2^{w-1} > \sum_{i=0}^{w-2} 2^i$)
 - 0 for nonnegative
 - 1 for negative

Encoding Example (Cont.)

x =	15213: 00111011 01101101
y =	-15213: 11000100 10010011

Weight	15213		-15213	
1	1	1	1	1
2	0	0	1	2
4	1	4	0	0
8	1	8	0	0
16	0	0	1	16
32	1	32	0	0
64	1	64	0	0
128	0	0	1	128
256	1	256	0	0
512	1	512	0	0
1024	0	0	1	1024
2048	1	2048	0	0
4096	1	4096	0	0
8192	1	8192	0	0
16384	0	0	1	16384
-32768	0	0	1	-32768
Sum		15213	Sum	-15213

Numeric Ranges

- Unsigned values

- $UMin = 0$

000...0

- $UMax = 2^w - 1$

111...1

- Two's complement values

- $TMin = -2^{w-1}$

100...0

- $TMax = 2^{w-1} - 1$

011...1

Values for $w = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for Different Word Sizes

	w			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

- Observations

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$

- C programming

- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values are platform-specific

Unsigned & Signed Numeric Values

X	$B2U(X)$	$B2T(X)$
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

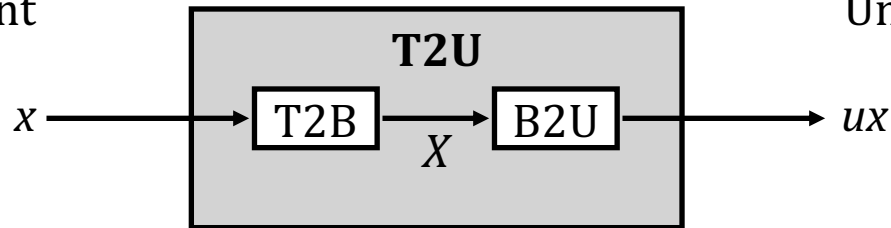
- Equivalence
 - Same encodings for positive values
- Uniqueness
 - Every bit pattern represents a unique integer value
 - Each representable integer has a unique bit encoding
- \Rightarrow Invert mappings
 - $U2B(x) = B2U^{-1}(x)$
 - Bit pattern for unsigned integer
 - $T2B(x) = B2T^{-1}(x)$
 - Bit pattern for 2's complement integer

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - **Conversion, casting**
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Mapping Between Signed & Unsigned

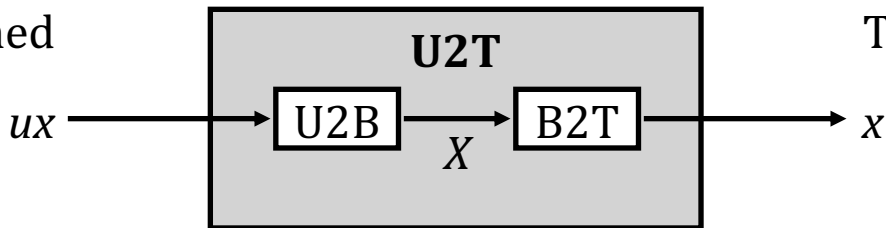
Two's Complement



Unsigned

Keep the same bit pattern

Unsigned



Two's Complement

Keep the same bit pattern

- Mappings b/w unsigned and two's complement numbers: **keep the same bit representations but reinterpret**

Mapping Signed \leftrightarrow Unsigned

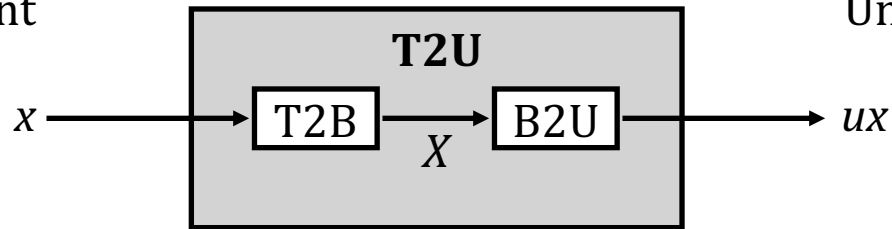
Bits	Signed		Unsigned
0000	0		0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8	T2U →	8
1001	-7	← U2T	9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15

Mapping Signed \leftrightarrow Unsigned

Bits	Signed		Unsigned
0000	0	\longleftrightarrow =	0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8	\longleftrightarrow +/- 16	8
1001	-7		9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15

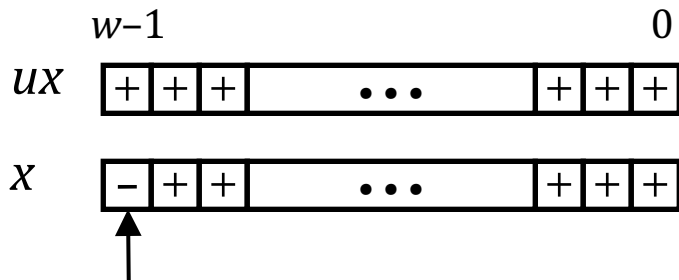
Relation between Signed & Unsigned

Two's Complement



Unsigned

Keep the same bit pattern



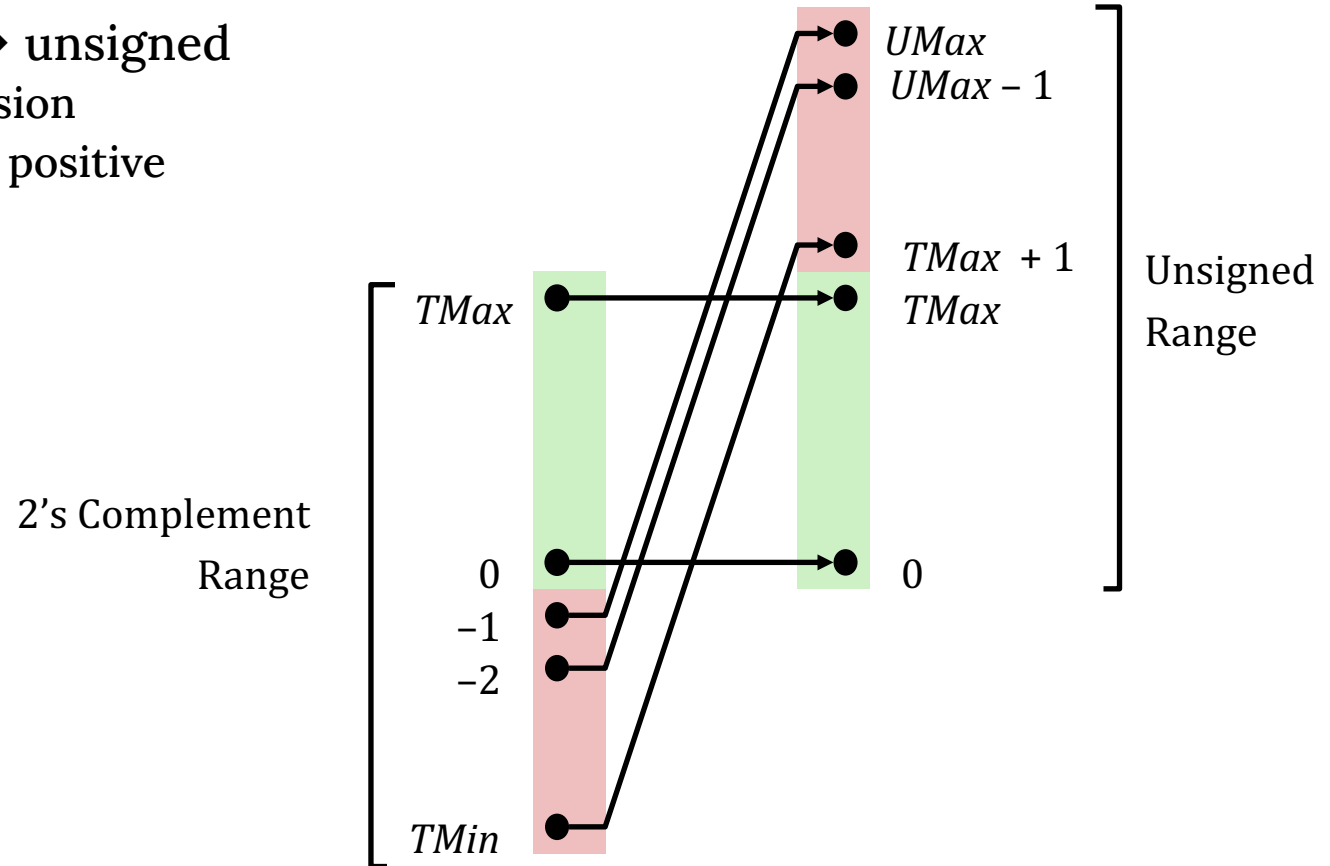
*Large negative weight
becomes*

Large positive weight

$$ux = \begin{cases} x & x \geq 0 \\ x + 2^w & x < 0 \end{cases}$$

Conversion Visualized

- 2's complement \rightarrow unsigned
 - Ordering inversion
 - Negative \rightarrow **big** positive



Signed vs. Unsigned in C

- Constants

- Considered to be signed integers by default
- Unsigned if have “U” as a suffix
`0U, 4294967259U`

- Casting

- Explicit casting between signed & unsigned: the same as U2T and T2U
`int tx, ty;`
`unsigned ux, uy;`
`tx = (int) ux;`
`uy = (unsigned) ty;`
- Implicit casting also occurs via assignments and procedure calls
`tx = ux;`
`uy = ty;`

Casting Surprises

- Expression Evaluation

- If there is a mix of unsigned and signed in single expression, **signed values implicitly cast to unsigned**
- Including comparison operations `<`, `>`, `==`, `<=`, `>=`
- Examples for $w = 32$: **$TMin = -2,147,483,648$** , **$TMax = 2,147,483,647$**

Constant#1	Constant#2	Relation	Evaluation
0	0U	==	Unsigned
-1	0	<	Signed
-1	0U	>	Unsigned
2147483647	-2147483648	>	Signed
2147483647U	-2147483648	<	Unsigned
-1	-2	>	Signed
-1	(unsigned) -2	>	Unsigned
2147483647	2147483648U	<	Unsigned
2147483647	(int) 2147483648U	>	Signed

Summary: Basic Signed-Unsigned Casting Rules

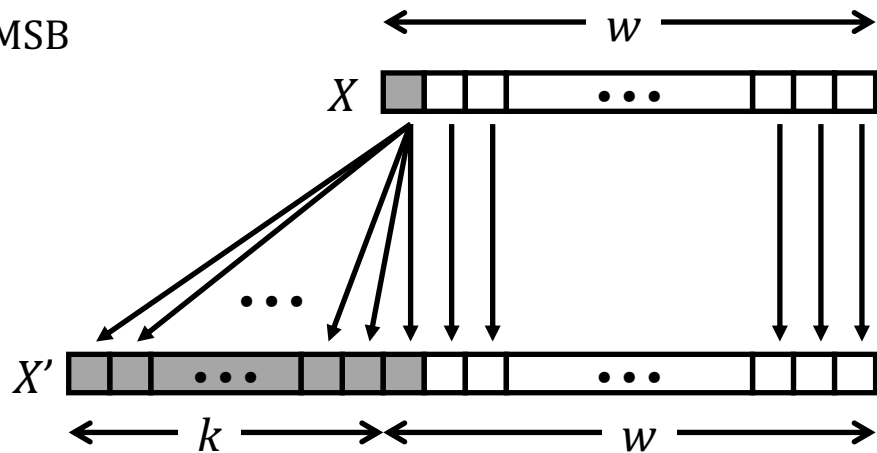
- Keep the same bit pattern
- But reinterpret
- Can have unexpected effects: adding or subtracting 2^w
- Expression containing signed and unsigned: cast to unsigned

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - **Expanding, truncating**
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Sign Extension

- Task:
 - Given w -bit signed integer x
 - Convert it to $(w+k)$ -bit integer with the same value
- Rule:
 - Make k copies of sign bit:
 - $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$



Sign Extension Example

```
short int x = 15213;  
int      ix = (int) x;  
short int y = -15213;  
int      iy = (int) y;
```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
ix	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
y	-15213	C4 93	11000100 10010011
iy	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

- Converting from smaller to larger integer data type
- C automatically performs sign extension

Summary: Basic Expansion & Truncation Rules

- Expanding (e.g., `short` to `int`)
 - Unsigned: zeros added
 - Signed: sign extension
 - Both yield the expected result
- Truncating (e.g., `unsigned int` to `unsigned short`)
 - Unsigned/signed: bits are truncated
 - Result reinterpreted
 - Unsigned: mod (%) operation
 - Signed: **similar** to mod
 - Expected behaviour **only for small numbers** yield

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Negation: Inversion & Increment

- Claim: the following holds for 2's complement

$$\sim \mathbf{x} + 1 = -\mathbf{x}$$

- Complement

- Observation: $\sim \mathbf{x} + \mathbf{x} = 1111\dots 111 = -1$

	x	1	0	0	1	1	1	0	1
		1	0	0	1	1	1	0	1
+	~x	0	1	1	0	0	0	1	0
<hr/>									
	-1	1	1	1	1	1	1	1	1

Inversion & Increment Examples

x = 15213

	Decimal	HEX	Binary
x	15213	3B 6D	00111011 01101101
~x	-15214	C4 92	11000100 10010010
~x+1	-15213	C4 93	11000100 10010011

x = 0

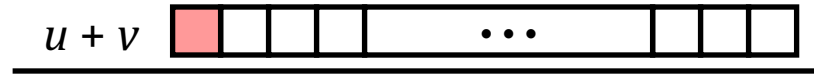
	Decimal	HEX	Binary
x	0	3B 6D	00000000 00000000
~x	-1	FF FF	11111111 11111111
~x+1	0	00 00	00000000 00000000

Unsigned Addition

Operands: w bits



True Sum: $w+1$ bits



Discard Carry: w bits

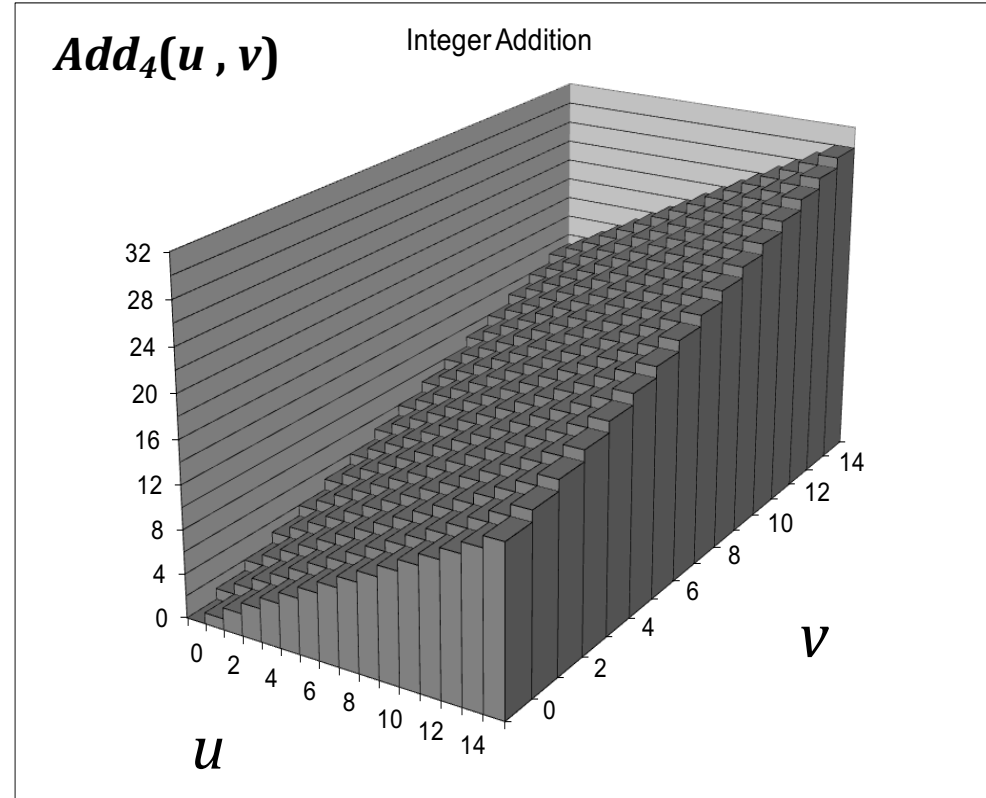


- Standard addition function
 - Ignores carry output
- Implements **modular arithmetic**
 $s = UAdd_w(u, v) = (u + v) \bmod 2^w$

$$UAdd_w(u, v) = \begin{cases} u + v & u + v < 2^w \\ u + v - 2^w & u + v \geq 2^w \end{cases}$$

Visualizing (Mathematical) Integer Addition

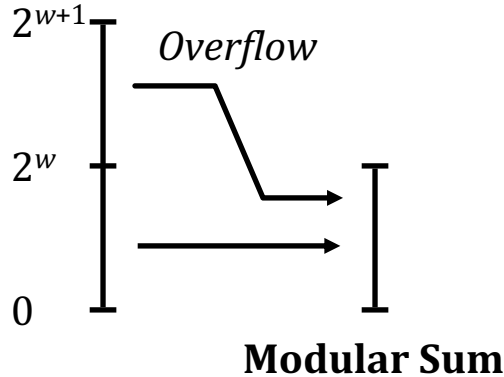
- Integer addition
 - 4-bit integers u, v
 - Compute true sum $Add_4(u, v)$
 - Values increase linearly with u and v
 - Forms planar surface



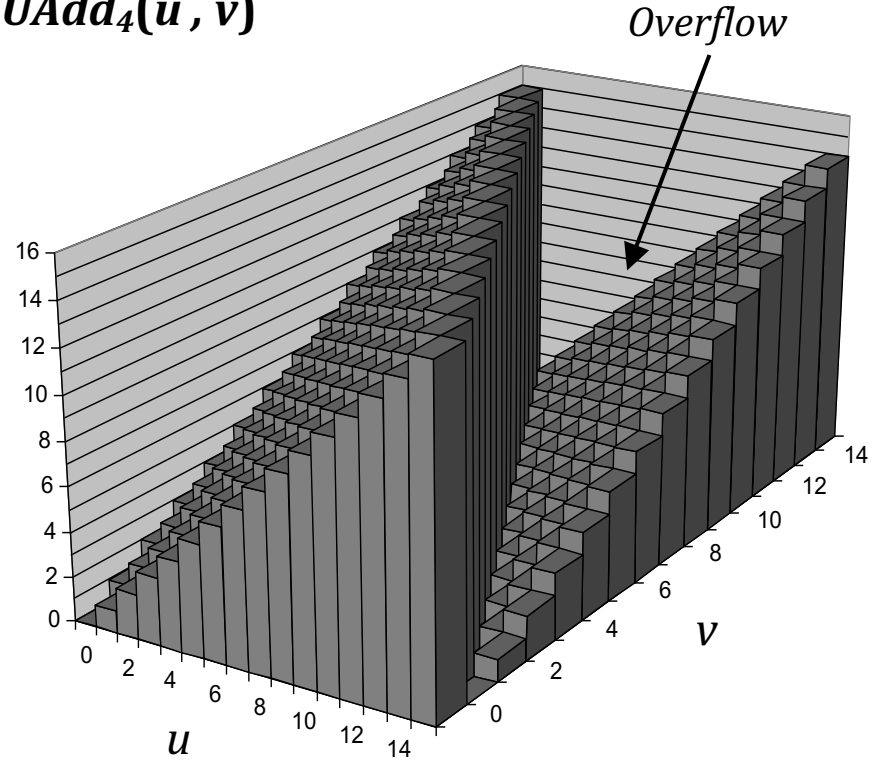
Visualizing Unsigned Addition

- Wraps around
 - If true sum $\geq 2^w$
 - At most once

True Sum



$UAdd_4(u, v)$



Two's Complement Addition

Operands: w bits



$+ v$



True Sum: $w+1$ bits

$u + v$



Discard Carry: w bits

$TAdd_w(u, v)$

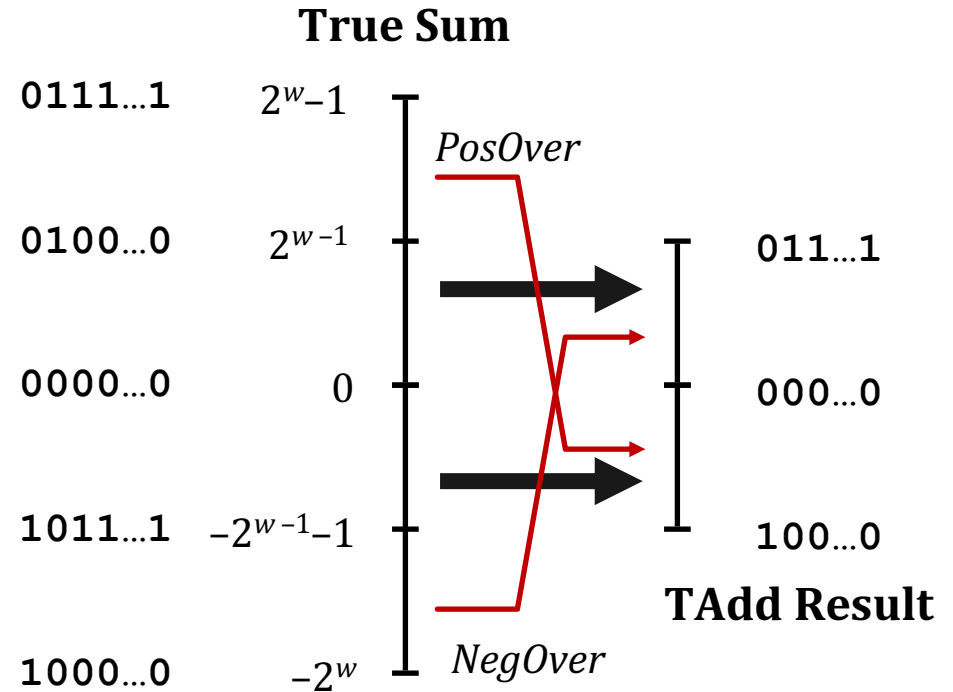


- $TAdd$ and $UAdd$ have identical bit-level behaviour
 - Signed vs. unsigned addition in C

```
int s, t, u, v;  
s = (int) ((unsigned) u + (unsigned) v);  
t = u + v
```
 - Will give $s == t$

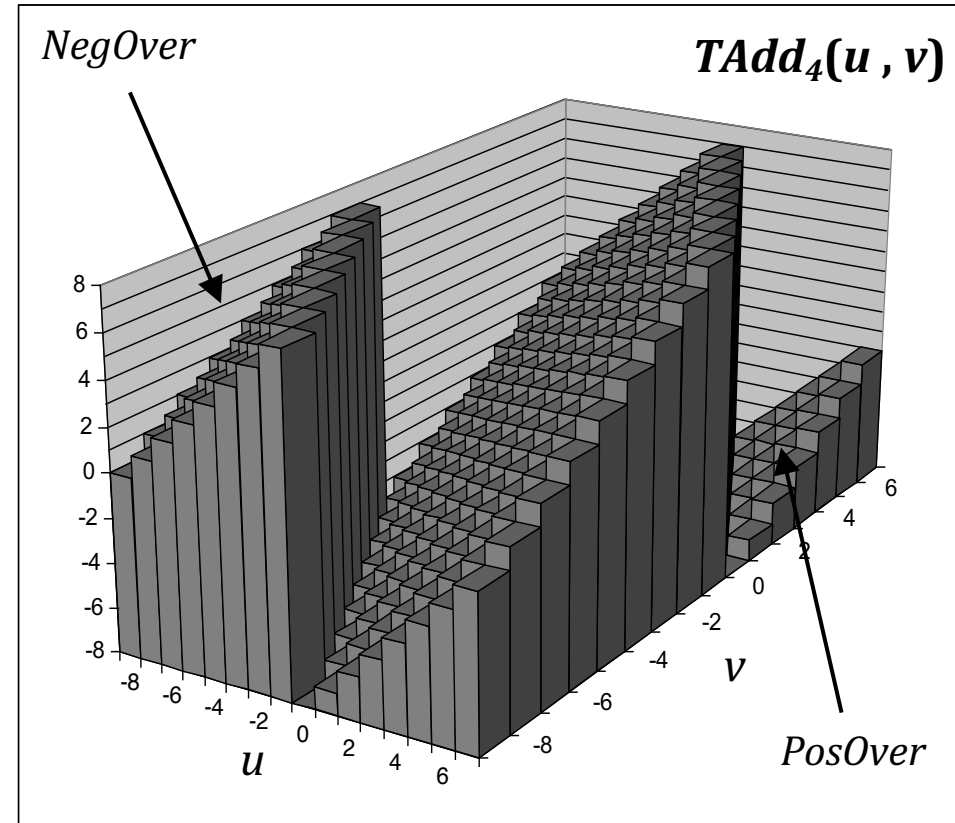
TAdd Overflow

- Functionality
 - True sum requires $w+1$ bits
 - Drop off MSB
 - Treat remaining bits as 2's complement integer



Visualizing 2's Complement Addition

- Values
 - 4-bit 2's complement
 - Range from -8 to +7
- Wraps around
 - If $\text{sum} > 2^{w-1}$
 - Becomes negative
 - At most once
 - If $\text{sum} < -2^{w-1}$
 - Becomes positive
 - At most once



Multiplication

- Computing exact product of w -bit numbers x, y
 - Either signed or unsigned
- Problem: exact results may require more than w bits
 - Unsigned: **up to $2w$ bits**
 - Result range: $0 \leq x \times y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
 - Two's complement min: **up to $(2w-1)$ bits**
 - Result range: $x \times y \geq (-2^{w-1}) \times (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
 - Two's complement max: **up to $2w$ bits, but only for $(TMin_w)^2$**
 - Result range: $x \times y \leq (-2^{w-1})^2 = 2^{2w-2}$
- Maintaining exact results
 - Would need to keep expanding word size with each product computed
 - Done in software by **arbitrary-precision (or bignum)** arithmetic packages

Unsigned Multiplication in C

Operands: w bits



True Product: $2w$ bits



Discard w MSBs: w bits



- Standard Multiplication Function
 - Ignores high order w bits
- Implements Modular Arithmetic
$$UMult_w(u, v) = (u \times v) \bmod 2^w$$

Signed Multiplication in C

Operands: w bits



True Product: $2w$ bits



Discard w MSBs: w bits



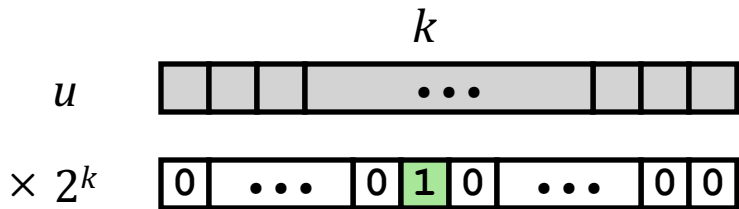
- Standard Multiplication Function
 - Ignores high order w bits
 - Some of which are different for signed vs. unsigned multiplication
 - Lower bits are the same

Power-of-2 Multiply with Shift

- Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned

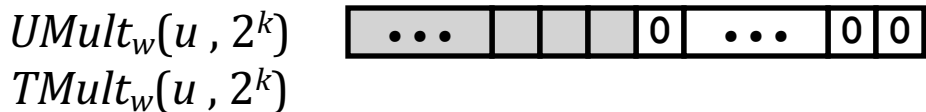
Operands: w bits



True Product: $w+k$ bits



Discard k bits: w bits



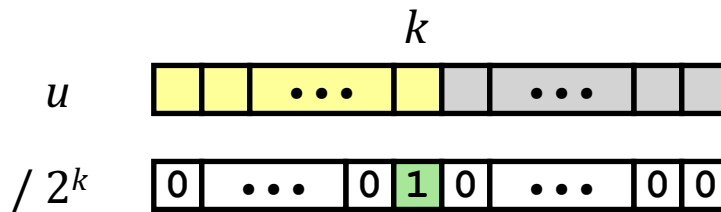
- Examples

- $u \ll 3 == u * 8$
- $u \ll 5 - u \ll 3 == u * 24$
- Most machines shift and add faster than multiply: compiler may optimize this

Unsigned Power-of-2 Divide with Shift

- Quotient of unsigned by power of 2
 - $u \gg k$ gives $\lfloor u / 2^k \rfloor$
 - Uses **logical shift**

Operands:



Division:



Result:



	Division	Computed	HEX	Binary
x	15213	15213	3B 6D	00111011 01101101
x >> 1	7606.5	7606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

Arithmetic: Basic Rules

- Addition

- Unsigned/signed: normal addition followed by truncate, the same operation on bit level
- Unsigned: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
- Signed: modified addition mod 2^w (result in a proper range)
 - Mathematical addition + possible addition **or** subtraction of 2^w

- Multiplication

- Unsigned/signed: normal multiplication followed by truncate, the same operation on bit level
- Unsigned: multiplication mod 2^w
- Signed: modified multiplication mod 2^w (result in a proper range)

Why to Use Unsigned?

- Do **NOT** use signed w/o understanding implications

- Easy to make a mistake

```
unsigned i;  
for (i = CNT-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- Can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)
```

...

- Do use unsigned to perform modular arithmetic
 - Multi-precision arithmetic
- Do use unsigned to represent sets
 - Logical right shift, no sign extension

Counting Down with Unsigned

- Proper way to use unsigned as loop index

```
unsigned i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- See Robert Seacord, *Secure Coding in C and C++*
 - C Standard guarantees that unsigned addition will behave like modular arithmetic
 - $0 - 1 \rightarrow UMax$
- Even better

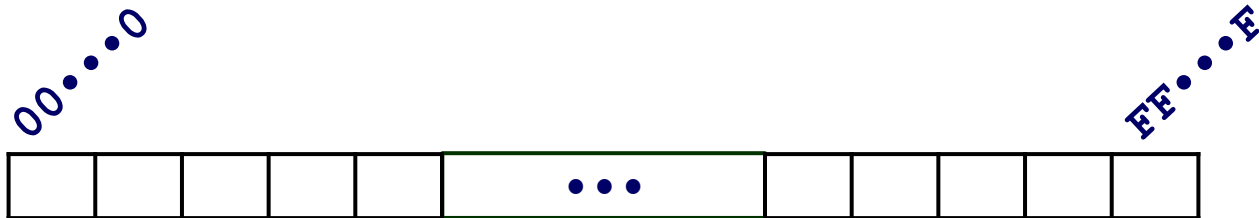
```
size_t i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- Data type `size_t` defined as unsigned value with length = word size
- Code will work even if `cnt = UMax`
- What if `cnt` is signed and `< 0`?

Lecture Agenda: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representation in memory, pointers, strings

Byte-Oriented Memory Organization



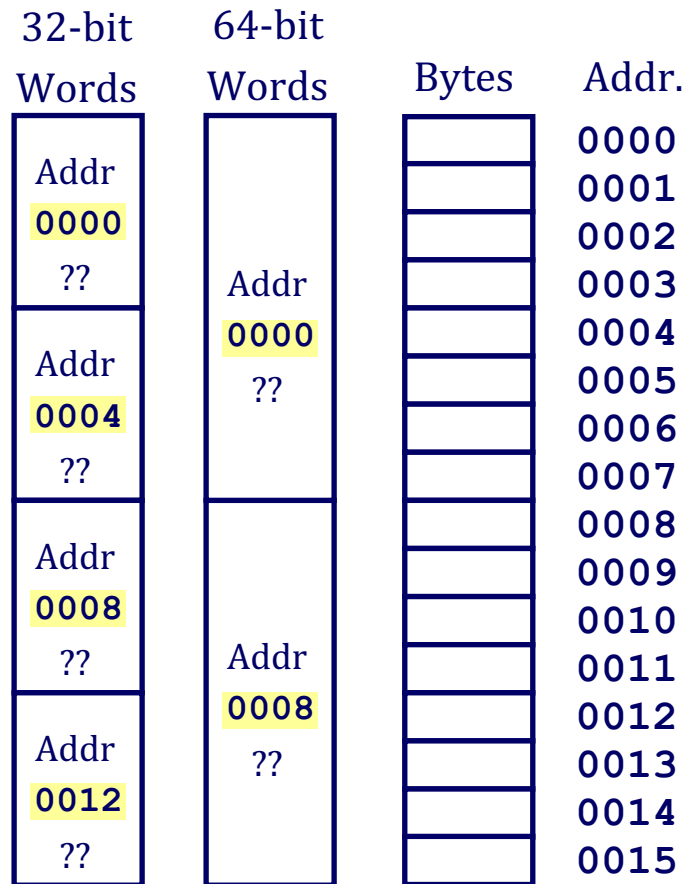
- Programs refer to **virtual addresses**
 - Conceptually, consider main memory as a very large array of bytes
 - In reality, it is not true, but a program can think of it that way
 - An address is like an index into that array
 - A pointer variable stores an address
- Note: a system can provide private address spaces to each **process**
 - Think of a process as a program being executed
 - So, a program can clobber its own data, but not that of others

Machine Words

- Any given computer has a **word size**
 - The nominal size of integer-valued data
 - More precisely: the size of addresses
 - Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4 GiB (2^{32} bytes)
 - Increasingly, machines have 64-bit word size
 - Potentially, could have 18 EB (exabytes) of addressable memory
 - Around 18.4×10^{18}
 - Machines still support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Word-Oriented Memory Organization

- Addresses specify byte locations
 - Address of the **first** byte in word
 - Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	8	8
float	4	4	4
double	8	8	8
long double	–	–	10/16
Pointer (word)	4	8	8

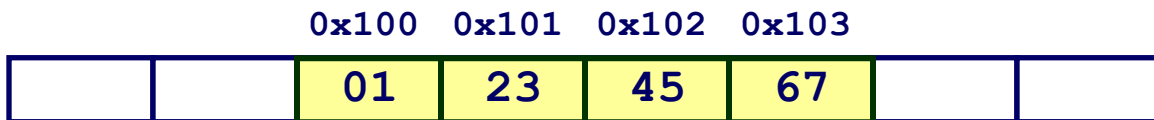
Byte Ordering

- How should bytes within a multi-byte word be ordered in memory?
- Conventions
 - **Big endian**: Sun, PowerPC Mac, internet
 - The most significant byte comes first
 - **Little endian**: x86
 - The least significant byte comes first

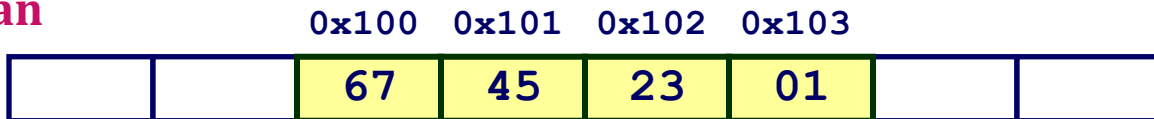
Byte Ordering Example

- Big endian
 - The most significant byte has the highest (**smallest**) address
- Little endian
 - The least significant byte has the highest (**smallest**) address
- Example
 - Variable **x** has 4-byte representation **0x01234567**
 - Address given by **&x** is **0x100**

Big Endian



Little Endian



Reading Byte-Reversed Listings

- Disassembly
 - Text representation of binary machine code
 - Generated by program that reads the machine code
- Example fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

- Deciphering numbers

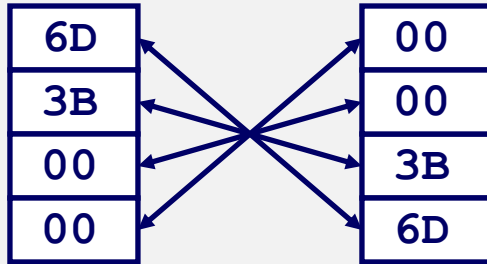
- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse:

0x12ab
0x000012ab
00 00 12 ab
ab 12 00 00

Representing Integers

```
int a = 15213;
```

IA32, x86-64 Sun



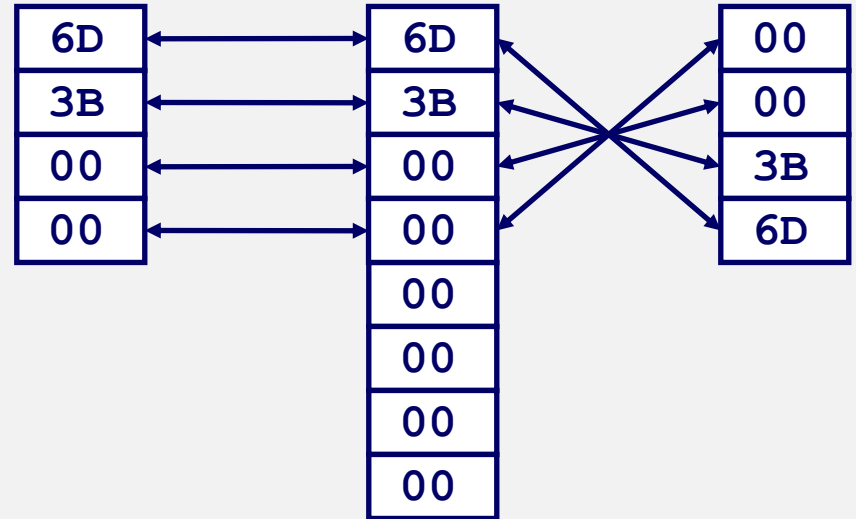
Decimal: 15213

Binary: 0011 1011 0110 1101

Hex: 3 B 6 D

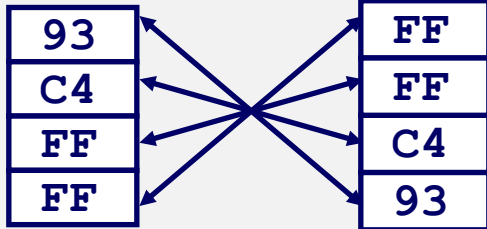
```
long int c = 15213;
```

IA32 x86-64 Sun



```
int b = -15213;
```

IA32, x86-64 Sun



Examining Data Representations

- Code to print byte representation of data
 - Casting pointer to `unsigned char *` creates a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, int len){
    int i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}
```

printf directives:

%p: Print pointer

%x: Print hexadecimal

show_bytes Execution Example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux):

```
int a = 15213;  
0x11ffffcb8      0x6d  
0x11ffffcb9      0x3b  
0x11ffffcba      0x00  
0x11ffffcbb      0x00
```

Representing Pointers

```
int b = -15213;  
int *p = &b;
```

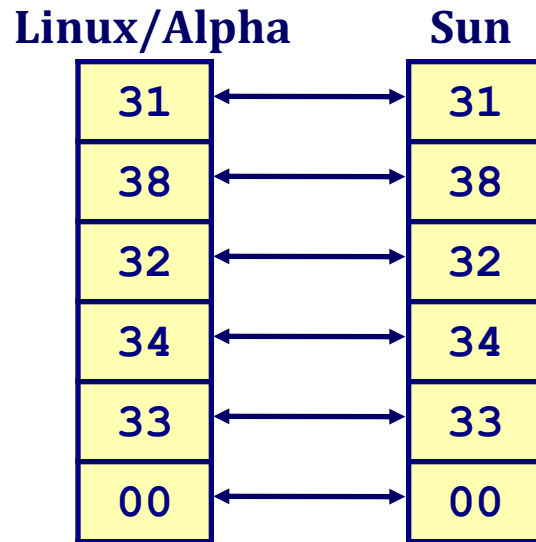
Sun	IA32	x86-64
EF	D4	0C
FF	F8	89
FB	FF	EC
2C	BF	FF
		FF
		7F
		00
		00

- Different compilers & machines assign different locations to objects
- Even get different results each time run program

Representing Strings

- Strings in C
 - Represented by an array of characters
 - Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character “0” has code 0×30
 - Digit “i” has code $0 \times 30 + i$
 - String should be null-terminated
 - Final character = “0”
- Compatibility
 - Byte ordering not an issue

```
char s[6] = "18243";
```



Integer C Puzzles

- Assume 32-bit word size, two's complement integers
- For each of the following C expressions: true or false? Why?

Initialization

```
int x = foo();  
int y = bar();  
unsigned ux = x;  
unsigned uy = y;
```

✗ $x < 0 \rightarrow ((x*2) < 0)$

✓ $ux \geq 0$

✓ $x \& 7 == 7 \rightarrow (x \ll 30) < 0$

✗ $ux > -1$

✗ $x > y \rightarrow -x < -y$

✗ $x * x \geq 0$

✗ $x > 0 \&\& y > 0 \rightarrow x + y > 0$

✓ $x \geq 0 \rightarrow -x \leq 0$

✗ $x \leq 0 \rightarrow -x \geq 0$

✗ $(x | -x) \gg 31 == -1$

✓ $ux \gg 3 == ux/8$

✗ $x \gg 3 == x/8$

✗ $x \& (x-1) != 0$

[CSED211] Introduction to Computer Software Systems

Lecture 2: Bits, Bytes, Integers

Prof. Jisung Park



CAOS

COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.09.06