

# [CSED211] Introduction to Computer Software Systems

## Lecture 7: Data

Prof. Jisung Park



**CAOS**  
COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2023.10.11

# Lecture Agenda

---

- Arrays
  - One-Dimensional
  - Multi-Dimensional (Nested)
  - Multi-Level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point

# Basic Data Types

- Integral

- Stored and operated on in **general (integer) registers**
- Signed or unsigned depending on instructions used

Intel	ASM	Bytes	C
Byte	b	1	[unsigned] char
Word	w	2	[unsigned] short
Double word	d	4	[unsigned] int
Quad word	q	8	[unsigned] long int (x86-64)

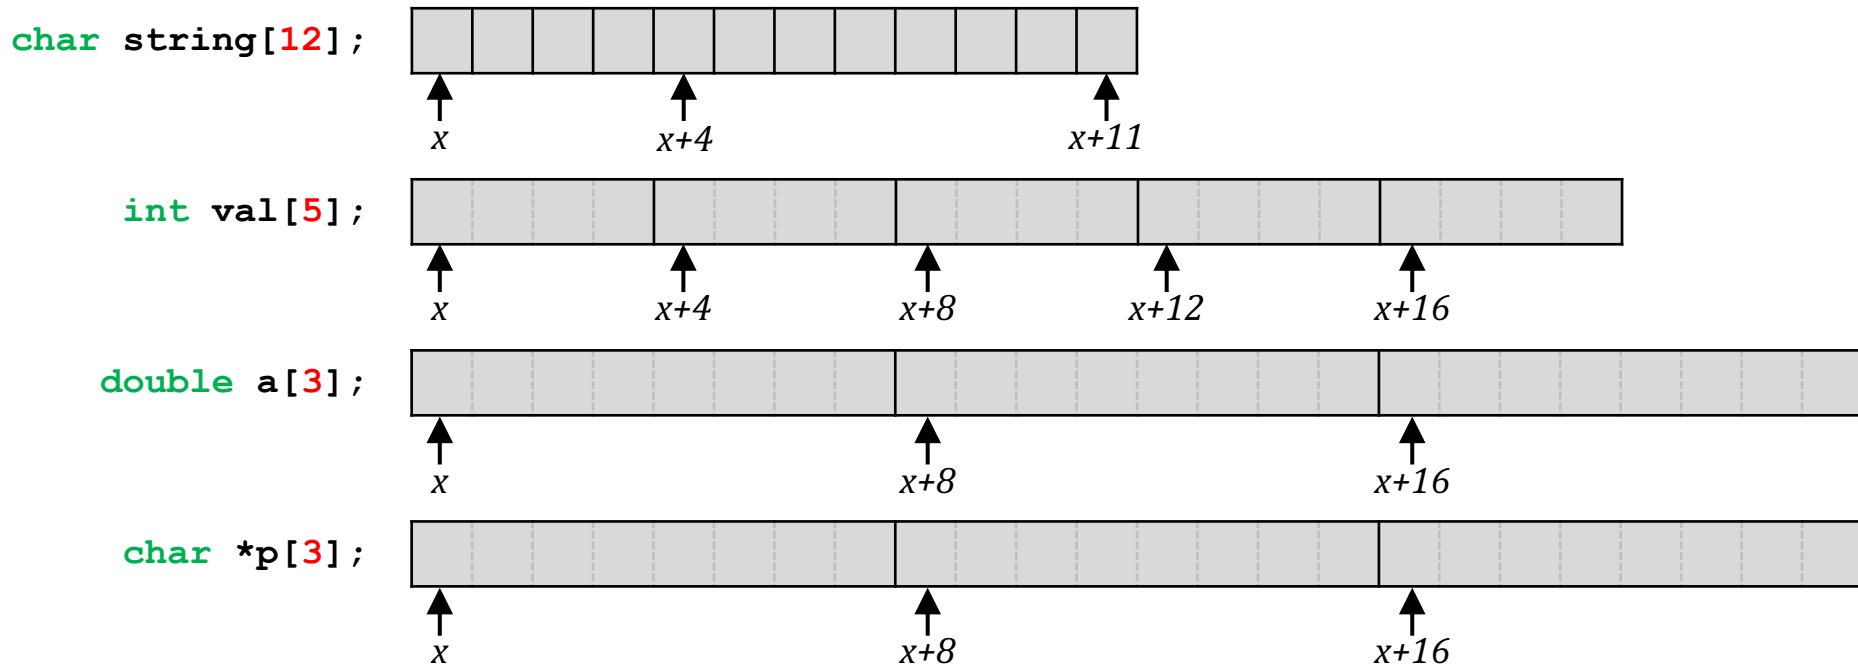
- Floating point

- Stored and operated on in **floating point registers**

Intel	ASM	Bytes	C
Single	s	4	float
Double	d	8	double
Extended	t	10/12/16	long double

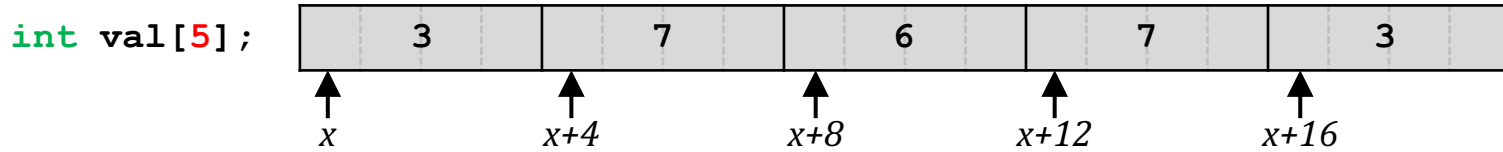
# Array Allocation

- Basic Principle: **T N[L];**
  - Array of data type **T** and length **L** named as **N**
  - Contiguously allocated region of  $(L * \text{sizeof}(T))$  bytes



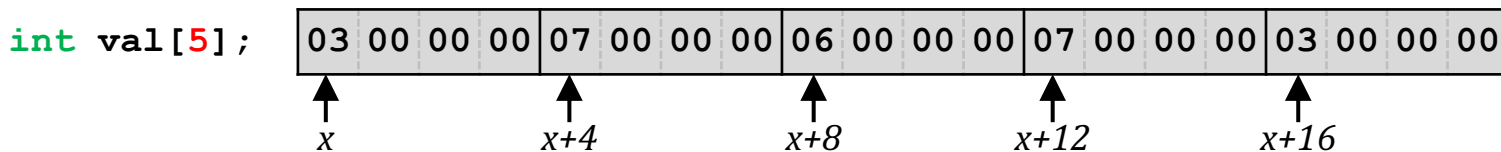
# Array Access

- Basic Principle: **T N[L] ;**
  - Array of data type **T** and length **L** named as **N**
  - Contiguously allocated region of  $(L * \text{sizeof}(T))$  bytes



# Array Access

- Basic Principle: **T N[L];**
  - Array of data type **T** and length **L** named as **N**
  - Contiguously allocated region of  $(L * \text{sizeof}(T))$  bytes



## Reference

`val[4]`

`val[5]`

`val`

`val+1`

`&val[2]`

`*(val+1)`

`val+i`

## Type

`int`

`int`

`int*`

`int*`

`int*`

`int`

`int*`

## Value

`3`

`??`

`x`

`x + 4`

`x + 8`

`7`

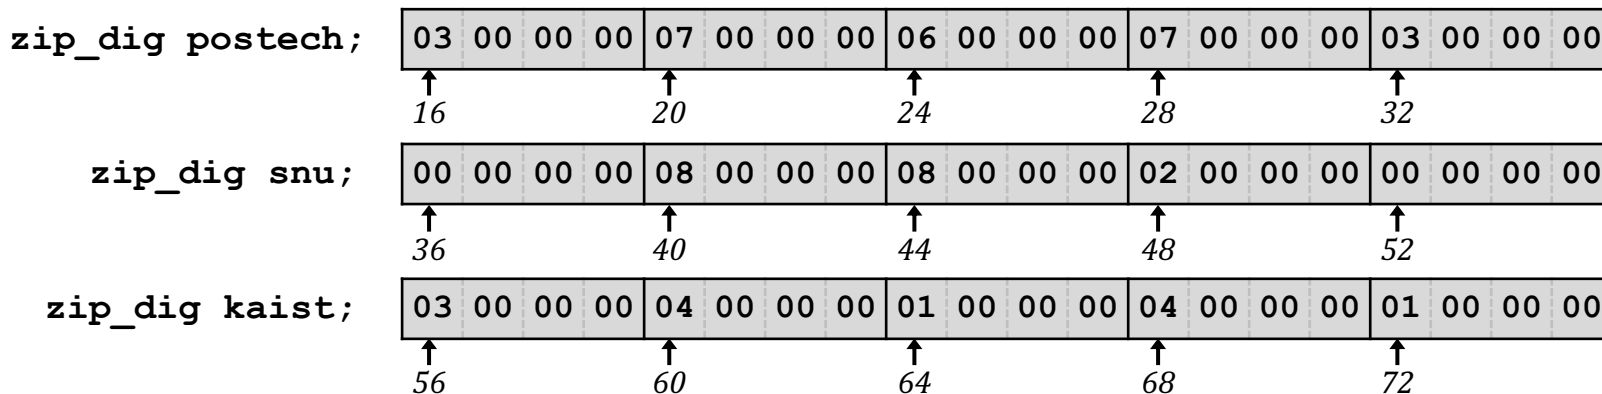
`x + 4*i`

# Array Example

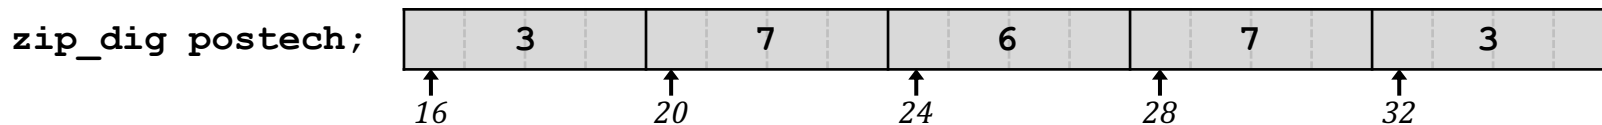
```
#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig postech = {3, 7, 6, 7, 3};
zip_dig snu      = {0, 8, 8, 2, 0};
zip_dig kaist    = {3, 4, 1, 4, 1};
```

- Declaration 'zip\_dig postech' equivalent to 'int postech[5]'
- Example arrays are allocated in successive 20-byte blocks
  - Not guaranteed to happen in general



# Array Accessing Example



```
int get_digit(zip_dig z, int dig){  
    return z[dig];  
}
```

```
# %rdi = z, %rsi = dig  
movl (%rdi,%rsi,4),%eax # z[dig]  
ret
```

- Register `%rdi` contains the target array's starting address
- Register `%rsi` contains the target array index
- Desired digit at `%rdi+4×%rsi`, i.e., `(%rdi,%rsi,4)`

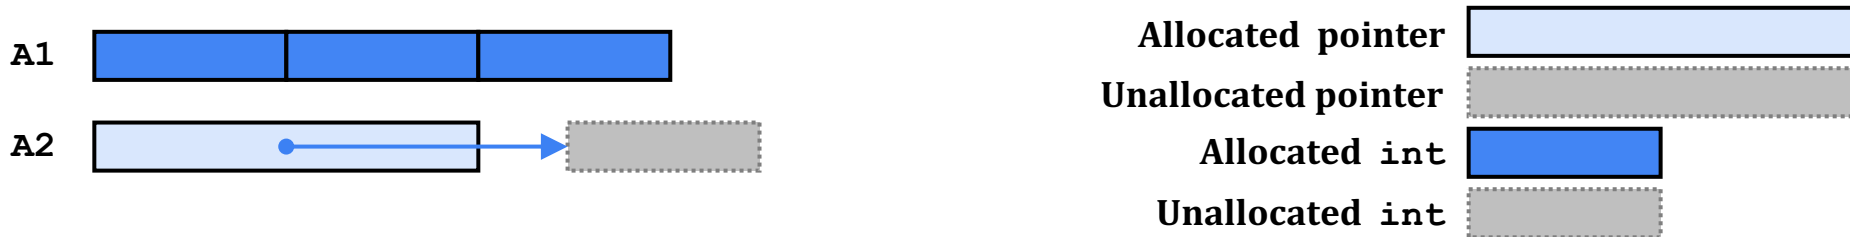


# Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# rdi = z  
movl    $0, %eax           # %eax = i  
jmp     .L3  
.L4:                          # loop:  
addl    $1, (%rdi,%rax,4)   # z[i]++  
addq    $1, %rax           # i++  
.L3:                          # middle  
cmpq    $4, %rax           # compare i and 4  
jbe     .L4                # if <=, goto loop  
rep; ret
```

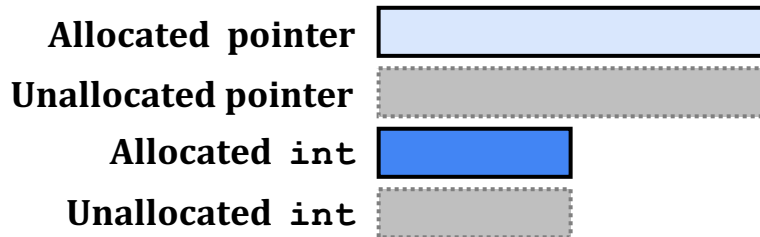
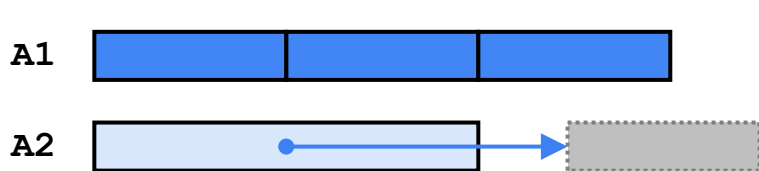
# Understanding Pointers & Arrays#1



Declaration	A1, A2			*A1, *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						

- **Comp**: can be compiled (Y/N)
- **Bad**: possible bad pointer reference (Y/N)
- **Size**: value returned by `sizeof`

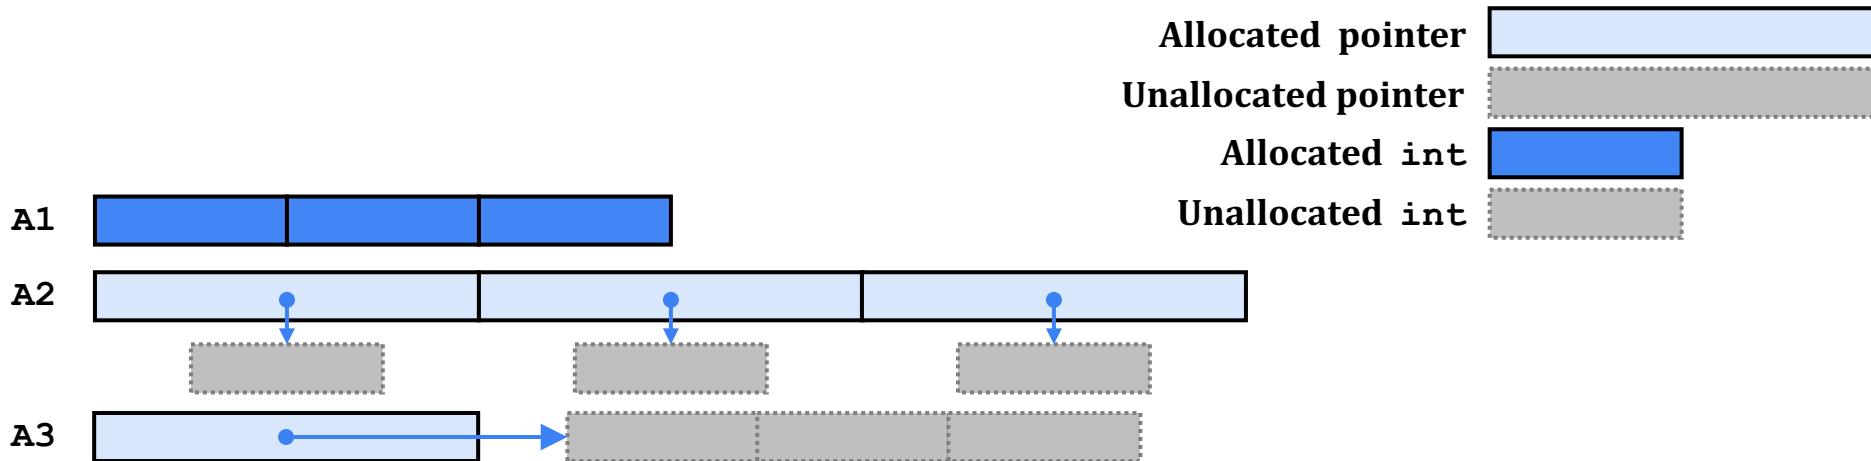
# Understanding Pointers & Arrays#1



Declaration	A1, A2			*A1, *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4
<code>int *A2</code>	Y	N	8	Y	Y	4

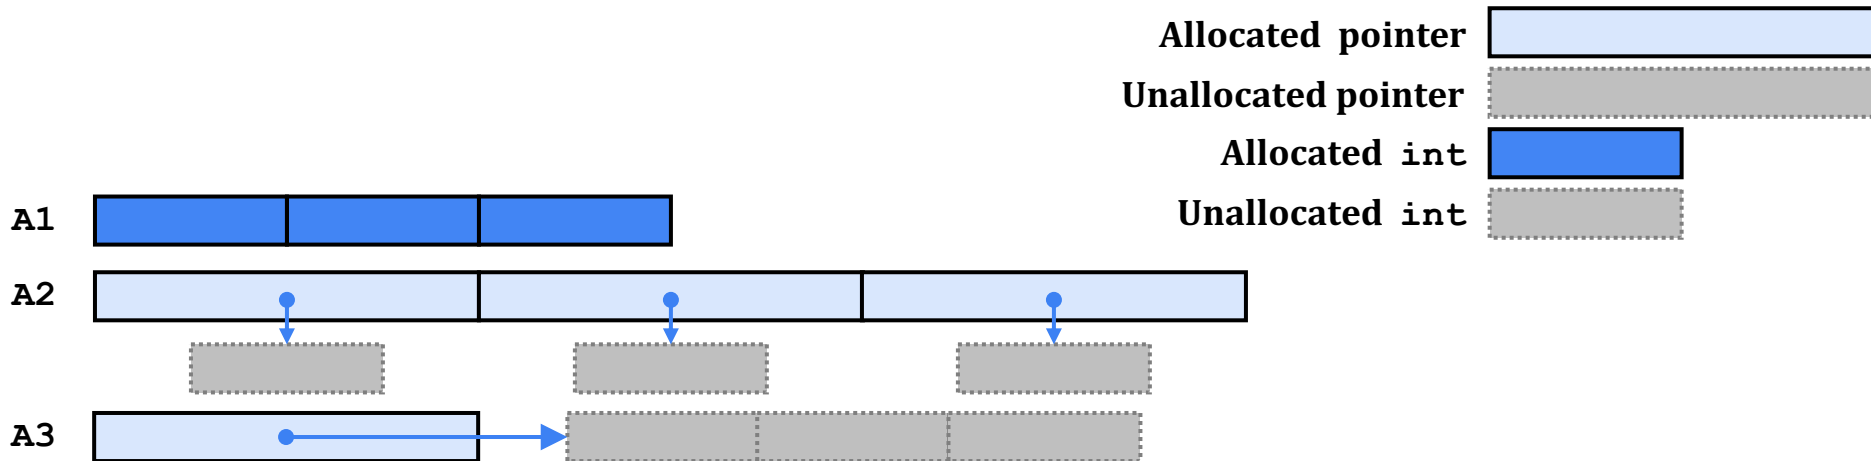
- **Comp**: can be compiled (Y/N)
- **Bad**: possible bad pointer reference (Y/N)
- **Size**: value returned by `sizeof`

# Understanding Pointers & Arrays#2



Declaration	<i>A<sub>n</sub></i>			<i>*A<sub>n</sub></i>			<i>**A<sub>n</sub></i>		
	Comp	Bad	Size	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>									
<code>int *A2[3]</code>									
<code>int (*A3)[3]</code>									

# Understanding Pointers & Arrays#2

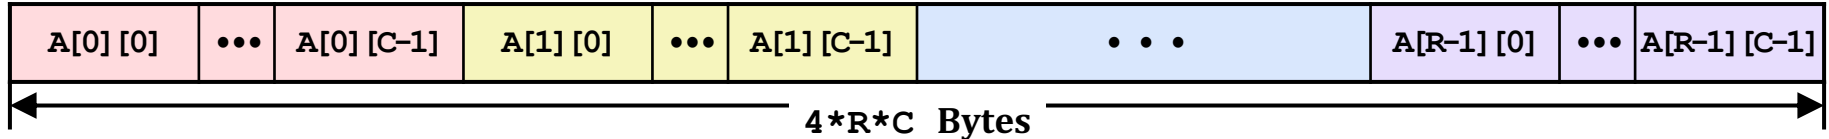
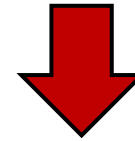
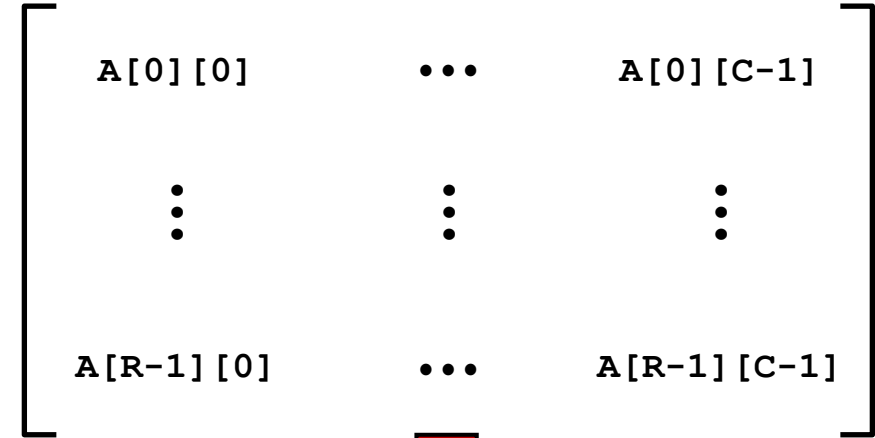


Declaration	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Comp	Bad	Size	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12	Y	Y	4

# Multidimensional (Nested) Arrays

- Declaration: **T N[R][C];**
  - A 2D array of data type **T**
  - **R** rows and **C** columns
- Array Size
  - **R\*C\*k** bytes
  - Where **sizeof(T)=k**
- Arrangement: **Row-major** ordering

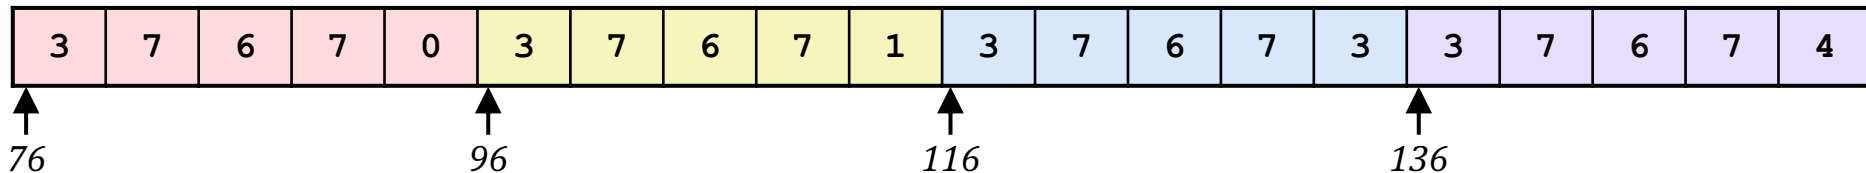
```
int A[R][C];
```



# Nested Array Example

```
#define PCOUNT 4  
  
zip_dig pohang[PCOUNT] = {{3, 7, 6, 7, 0},  
                           {3, 7, 6, 7, 1},  
                           {3, 7, 6, 7, 3},  
                           {3, 7, 6, 7, 4}};
```

```
zip_dig pohang[4];
```

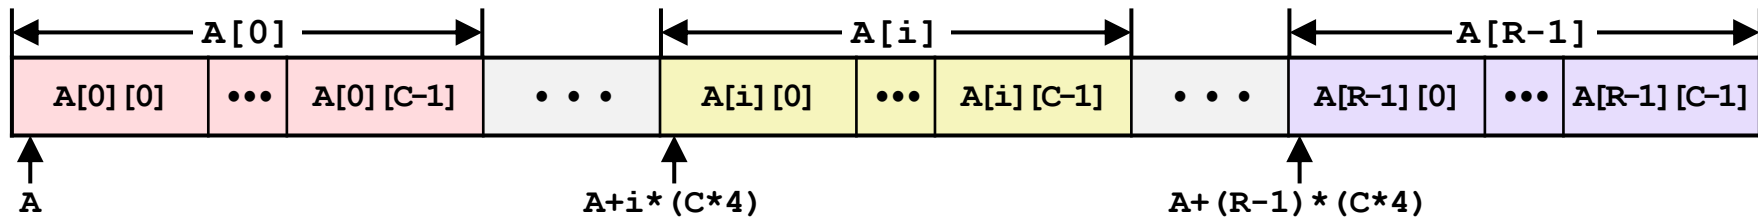


- ‘zip\_dig pohang[4]’ equivalent to ‘int pohang[4][5]’
  - Variable `pohang`: an array of 4 elements allocated contiguously
  - Each element is an array of 5 `int`’s also allocated contiguously
- **Row-major ordering** of all elements guaranteed

# Nested Array Row Access

- Row vectors **T** **N**[**R**][**C**]
  - **N**[**i**] is an array of **C** elements
  - Each element of type **T** requires **k** bytes
  - Starting address  $\mathbf{N} + \mathbf{i} * (\mathbf{C} * \mathbf{k})$

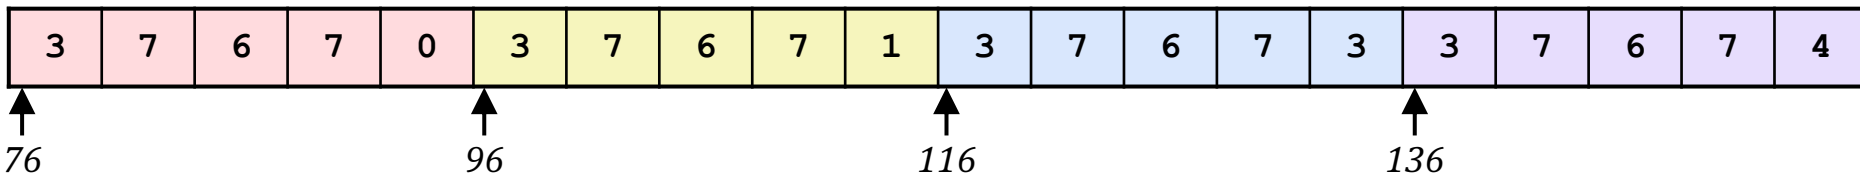
```
int A[R][C];
```





# Nested Array Row Access Code

```
zip_dig pohang[4];
```



```
int *get_pohang_zip(int index) {  
    return pohang[index];  
}
```

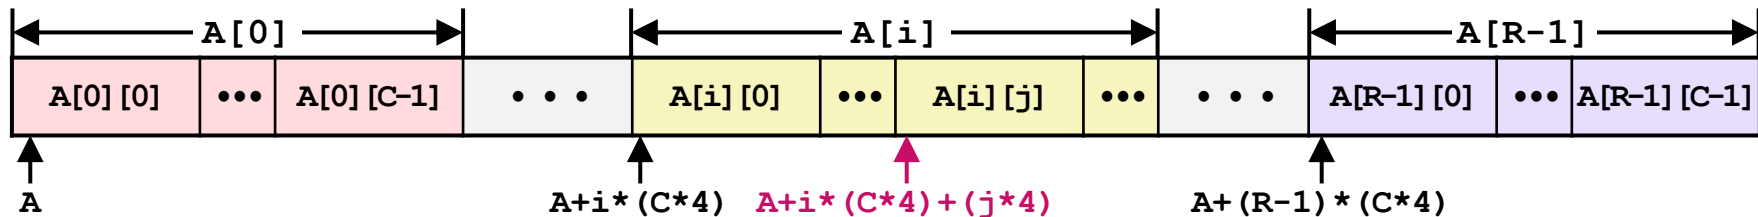
```
# %rdi = index  
leaq (%rdi,%rdi,4),%rax    # 5 * index  
leaq pohang(,%rax,4),%rax  # pohang+(20*index)  
ret
```

- Row vector **pohang**
  - **pohang[index]** is an array of 5 **int**'s
  - Starting address **pgh+20\*index**
- Machine Code
  - Computes and returns address
  - Computes as **pgh + 4\*(index+4\*index)**

# Nested Array Element Access

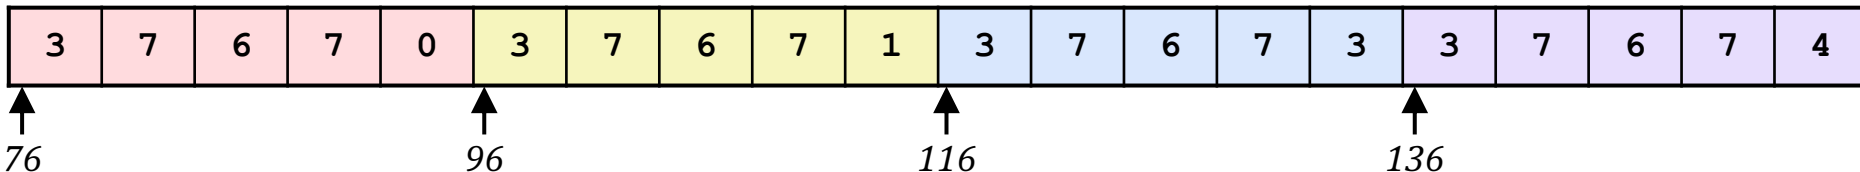
- Array elements
  - $N[i][j]$  is an element of type  $T$  that requires  $k$  bytes
  - Address  $N + i * (C * K) + j * k = N + (i * C + j) * K$

```
int A[R][C];
```



# Nested Array Element Access Code

```
zip_dig pohang[4];
```



```
int *get_pohang_dig(int index, int dig){  
    return pohang[index][dig];  
}
```

```
leaq    (%rdi,%rdi,4), %rax    # 5*index  
addl    %rax, %rsi             # 5*index+dig  
movl    pohang(,%rsi,4), %eax  # M[pohang+4*(5*index+dig)]  
ret
```

- Array Elements

- `pohang[index][dig]` is `int`
- Address: `pohang+20*index+4*dig = pohang+4*(5*index+dig)`

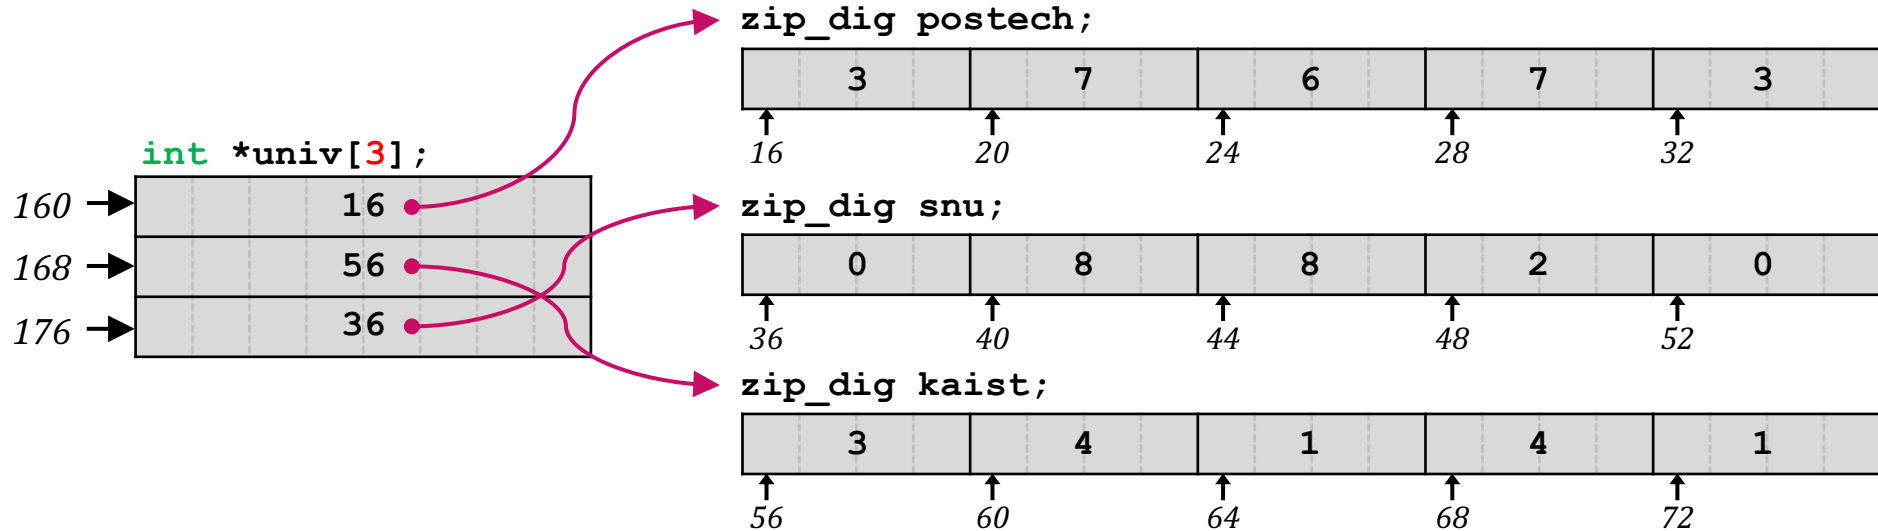
# Multi-Level Array Example

```
zip_dig postech = {3, 7, 6, 7, 3};  
zip_dig snu     = {0, 8, 8, 2, 0};  
zip_dig kaist   = {3, 4, 1, 4, 1};
```

```
#define UCOUNT 3
```

```
int *univ[UCOUNT] = {postech, kaist, snu};
```

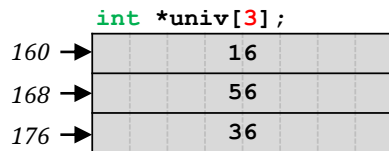
- Variable `univ` is an array of 3 pointer elements (8 bytes)
- Each points to an array of 5 `int`'s



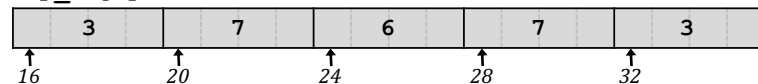
# Element Access in Multi-Level Array

```
int get_univ_digit(size_t index, size_t dig){  
    return univ[index][dig];  
}
```

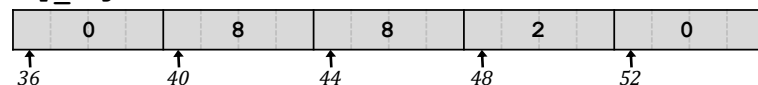
```
salq    $2, %rsi          # 4*dig  
addq    univ(,%rdi,8), %rsi # p=univ[index]+4*dig  
movl    (%rsi), %eax       # return *p  
ret
```



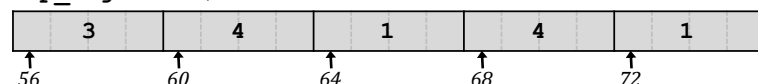
zip\_dig postech;



zip\_dig snu;



zip\_dig kaist;



- Computation

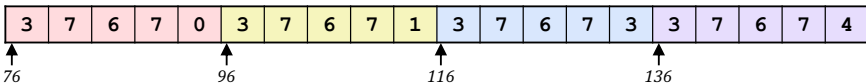
- Access `Mem[Mem[univ+8*index]+4*dig]`
- Must do two memory reads
  - First to get the pointer to the target row array
  - Second to access the target element within the array

# Array Element Accesses

## Nested Array

```
int *get_pohang_dig(int index, int dig){  
    return pohang[index][dig];  
}
```

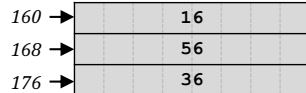
zip\_dig pohang[4];



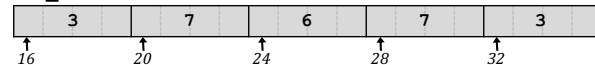
## Multi-Level Array

```
int get_univ_digit(size_t index, size_t dig){  
    return univ[index][dig];  
}
```

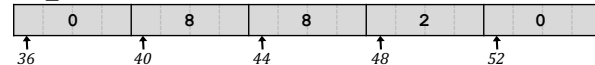
int \*univ[3];



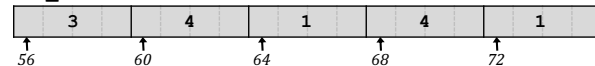
zip\_dig postech;



zip\_dig snu;



zip\_dig kaist;



Accesses looks similar in C, but addresses computations very different:

$\text{Mem}[\text{pohang} + 20 * \text{index} + 4 * \text{dig}]$

$\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{dig}]$

# N×N Matrix Code

- Fixed dimensions
  - A known value of N at compile time
- Variable dimensions, explicit indexing
  - Traditional way to implement dynamic arrays
- Variable dimensions, implicit indexing
  - Now supported by gcc

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele(fix_matrix a, int i, int j){
    return a[i][j];
}
```

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele(size_t n, int *a, int i, int j){
    return a[IDX(n,i,j)];
}
```

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

# 16×16 Matrix Access

- Array elements
  - Address  $N+i*(C*K)+j*K$
  - $C = 16, k = 4$

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele(fix_matrix a, int i, int j){
    return a[i][j];
}
```

```
# a in %rdi, i in %rsi, j in %rdx
salq    $6, %rsi          # 64*i
addq    %rsi, %rdi         # a + 64*i
movl    (%rdi,%rdx,4), %eax # M[a + 64*i + 4*j]
ret
```



# $n \times n$ Matrix Access

- Array elements
  - Address  $N + i * (C * K) + j * K$
  - $C = n, k = 4$
  - Must perform integer multiplication

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

```
# n in %rdi, a in %rsi, i in %rdx, j in %rcx
imulq    %rdx, %rdi          # n*i
leaq     (%rsi,%rdi,4), %rax  # a + 4*n*i
movl     (%rax,%rcx,4), %eax  # a + 4*n*i + 4*j
ret
```

# Example: Array Access

```
#include <stdio.h>
#define ZLEN 5
#define PCOUNT 4
typedef int zip_dig[ZLEN];

int main(int argc, char** argv) {
    zip_dig pohang[PCOUNT] = {{3, 7, 6, 7, 0},
                               {3, 7, 6, 7, 1},
                               {3, 7, 6, 7, 3},
                               {3, 7, 6, 7, 4}};

    int *linear_zip = (int *) pohang;
    int *zip2 = (int *) pohang[2];
    int result = pohang[0][0] +
                 linear_zip[7] +
                 *(linear_zip + 8) +
                 zip2[1];
    printf("result: %d\n", result);
    return 0;
}
```

```
$ ./array
result: 23
```

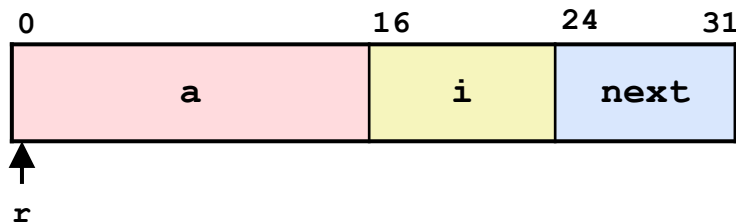
# Lecture Agenda

---

- Arrays
  - One-Dimensional
  - Multi-Dimensional (Nested)
  - Multi-Level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point

# Structure Representation

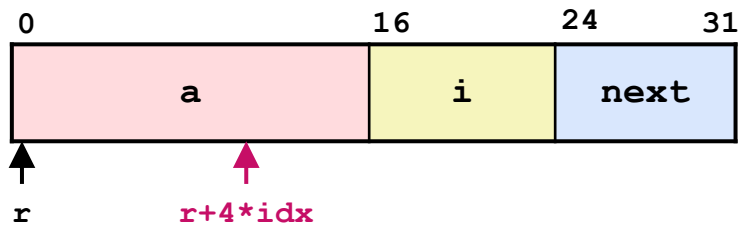
```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- Structure represented as a memory block
  - Big enough to hold all the fields
- Fields ordered **according to declaration**
  - Even if another ordering could yield a more compact representation
- Compiler determines the overall size and positions of fields
  - Machine-level program has **no understanding** of the structures in the source code

# Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



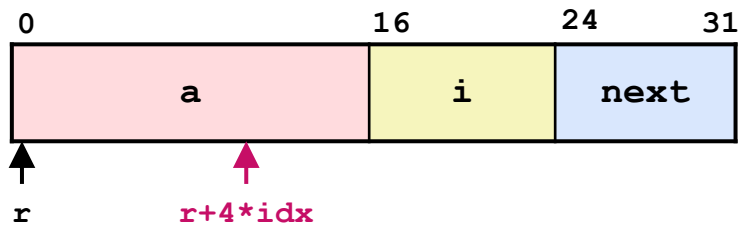
- Generating pointer to array element
  - Offset of each structure member determined at compile time
  - Compute as  $r+4*idx$

```
int *get_ap(struct rec *r, size_t idx){  
    return &r->a[idx];  
}
```

```
# r in %rdi, idx in %rsi  
leaq  (%rdi,%rsi,4), %rax  
ret
```

# Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



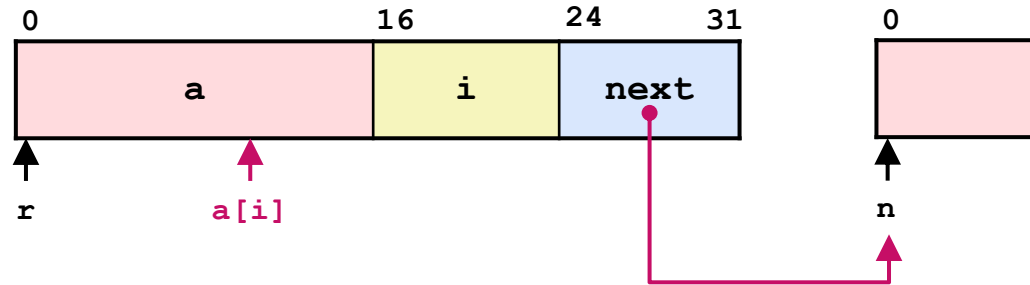
```
int *get_ap(struct rec *r, size_t idx){  
    return &r->a[idx];  
}
```

```
# r in %rdi, idx in %rsi  
leaq  (%rdi,%rsi,4), %rax  
ret
```

- Generating pointer to array element
  - Offset of each structure member determined at compile time
  - Compute as `r+4*idx`

# Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



```
void set_val(struct rec *r, int val){  
    while (r) {  
        int i = r->i;  
        r->a[i] = val;  
        r = r->next;  
    }  
}
```

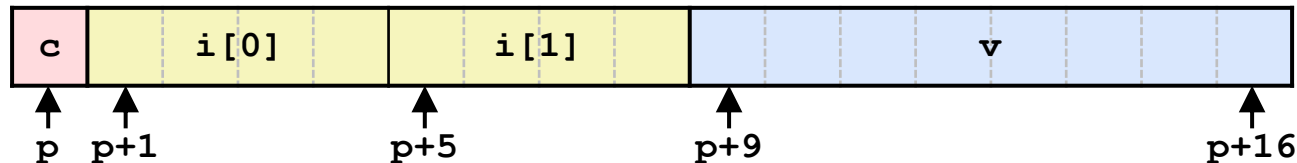
```
.L11:                                # loop:  
    movslq    16(%rdi), %rax        # i = Mem[r+16]  
    movl      %esi, (%rdi,%rax,4)   # Mem[r+4*i] = val  
    movq      24(%rdi), %rdi        # r = Mem[r+24]  
    testq     %rdi, %rdi            # Test r  
    jne       .L11                  # if !=0 goto loop
```

Register	Use(s)
%rdi	r
%rsi	val

# Structures & Alignment

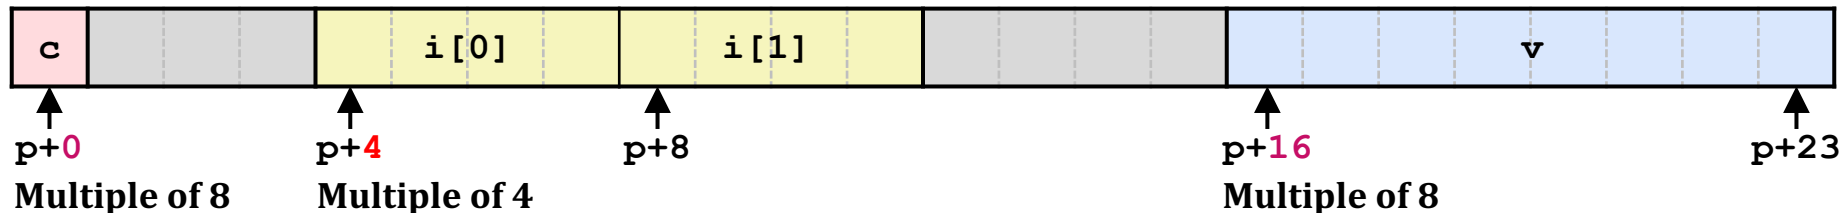
```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- Unaligned Data



- Aligned data

- Address must be **multiple of B** if primitive data type requires **B** bytes





# Alignment Principles

---

- Aligned data
  - Address must be **multiple of B** if primitive data type requires **B** bytes
  - Required on some machines; advised on x86-64
- Motivation for aligning Data
  - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
    - Inefficient to load or store data that spans quad-word boundaries
    - Virtual memory trickier when data spans 2 pages
- Compiler
  - Inserts gaps in structure to ensure correct alignment of fields

# Specific Cases of Alignment (x86-64)

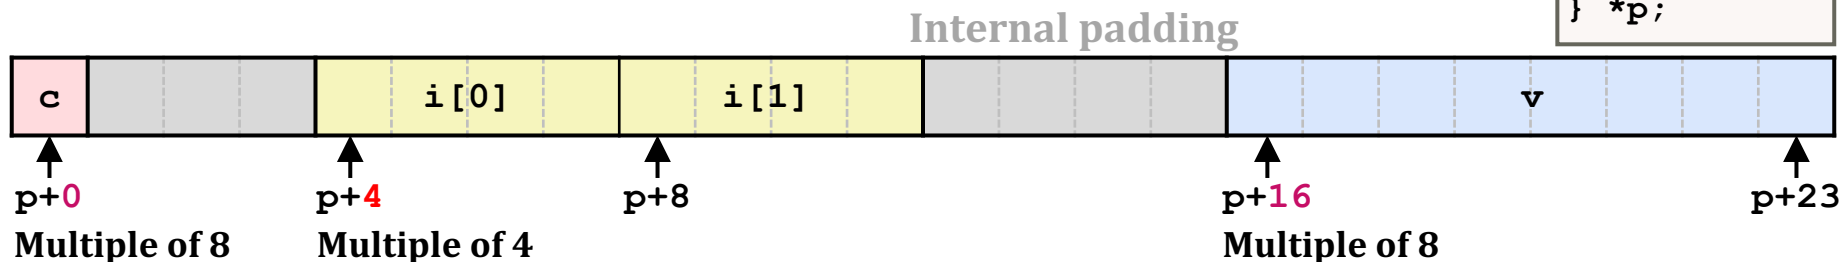
---

- 1 byte: `char`
  - No restrictions on address
- 2 bytes: `short`
  - Lowest bit of address must be  $0_2$
- 4 bytes: `int`, `float`, ...
  - Lowest 2 bits of address must be  $00_2$
- 8 bytes: `double`, `long`, `char *`, ...
  - Lowest 3 bits of address must be  $000_2$
- 16 bytes: `long double` (gcc on Linux)
  - Lowest 4 bits of address must be  $0000_2$

# Satisfying Alignment with Structures

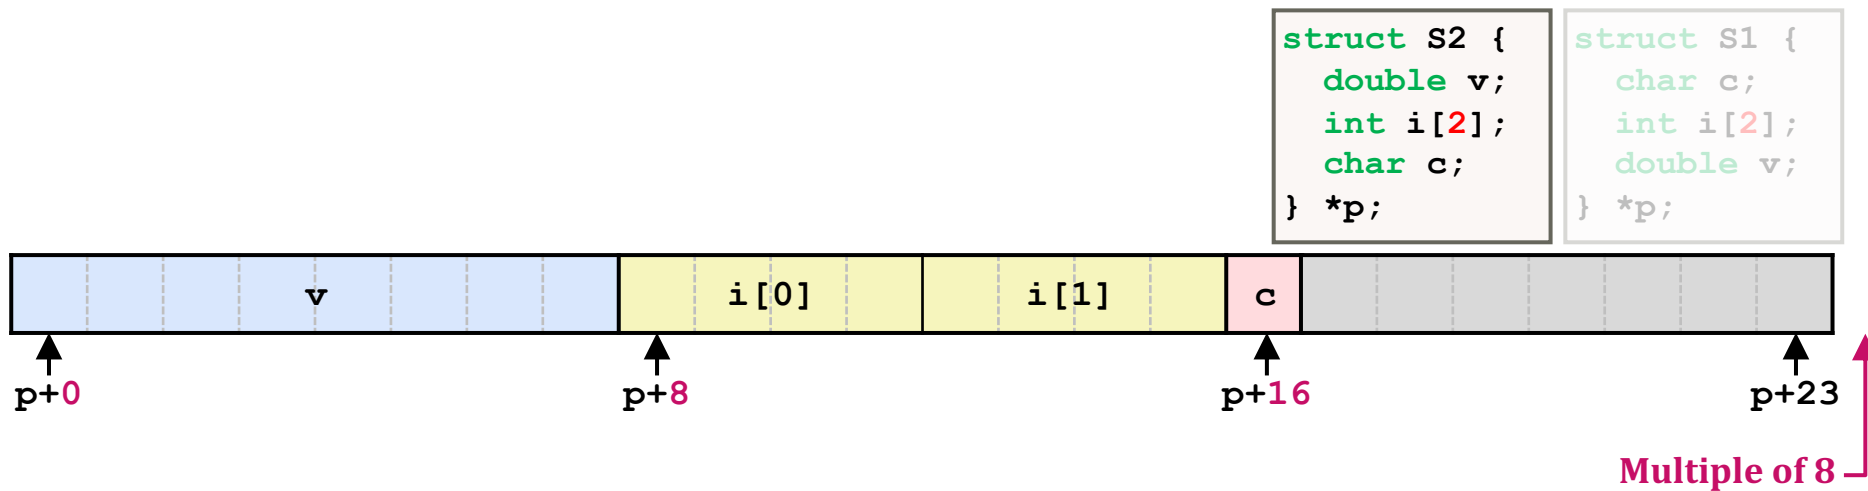
- Within structure: must satisfy each element's alignment requirement
- Overall structure placement
  - Each structure has **alignment requirement K**
    - K = Largest alignment of any element
  - Initial address & structure length must be **multiples of K**
- Example: K = 8 due to **double** element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



# Meeting Overall Alignment Requirement

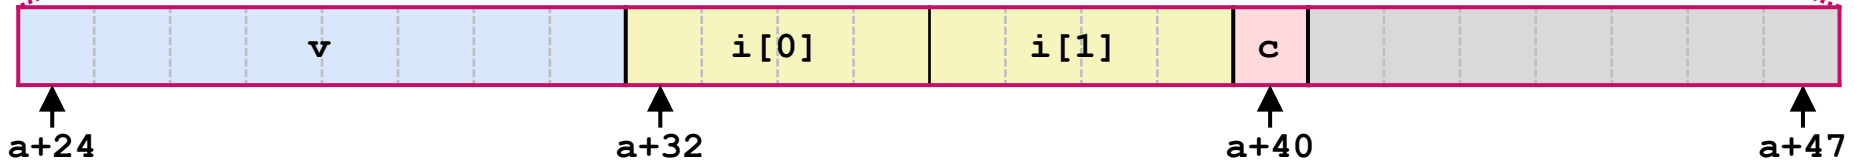
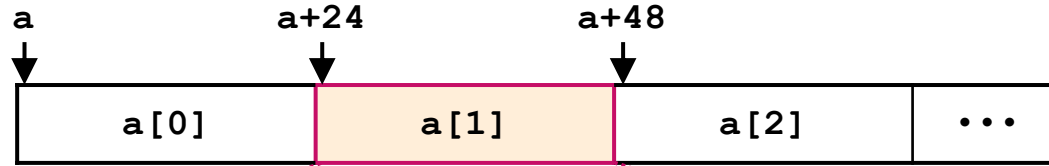
- For **largest alignment requirement K**
- Overall structure must be **multiple of K**



# Arrays of Structures

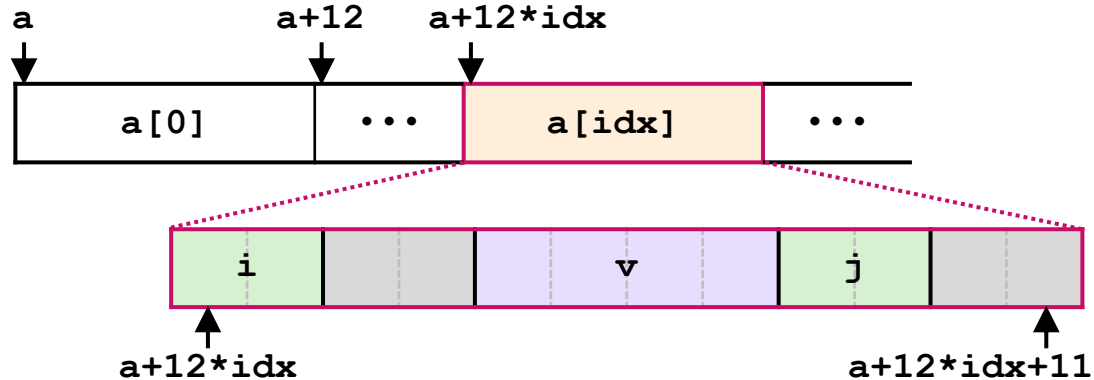
- Overall structure length **multiple of K**
- Satisfy alignment requirement for every element

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Accessing Array Elements

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx){  
    return a[idx].j;  
}
```

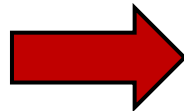
```
# %rdi = idx  
leaq    (%rdi,%rdi,2),%rax # 3*idx  
movzwl  a+8(,%rax,4),%eax  
ret
```

- Compute array offset  $12 \cdot \text{idx}$ 
  - `sizeof(S3)=12`, including alignment spacers
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`, which is resolved during linking

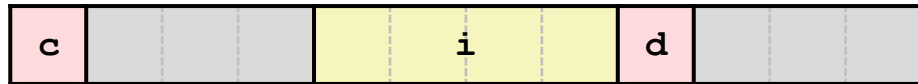
# Saving Space

- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```



# Example Struct Exam Question

## Problem 5. (8 points):

*Struct alignment.* Consider the following C struct declaration:

```
typedef struct{  
    char a;  
    long b;  
    float c;  
    char d[3];  
    int *e;  
    short *f;  
} foo;
```

1. Show how foo would be allocated in memory on an x86-64 Linux system. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an X to denote space that is allocated in the struct as padding




# Example Struct Exam Question (Cont.)

## Problem 5. (8 points):

*Struct alignment.* Consider the following C struct declaration:

```
typedef struct{  
    char a;  
    long b;  
    float c;  
    char d[3];  
    int *e;  
    short *f;  
} foo;
```

2. Rearrange the elements of foo to conserve the most space in memory. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an X to denote space that is allocated in the struct as padding


# Lecture Agenda

---

- Arrays
  - One-Dimensional
  - Multi-Dimensional (Nested)
  - Multi-Level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point


# Background

---

- History
  - x87 FP
    - Legacy, very ugly
  - SSE (Streaming SIMD Extensions) FP
    - SIMD: Single Instruction Multiple Data
    - Supported by old machines
    - Special case use of vector instructions
- AVX (Advanced Vector Extensions) FP
  - Newest version
  - Similar to SSE
  - Documented in book

# Programming with SSE3

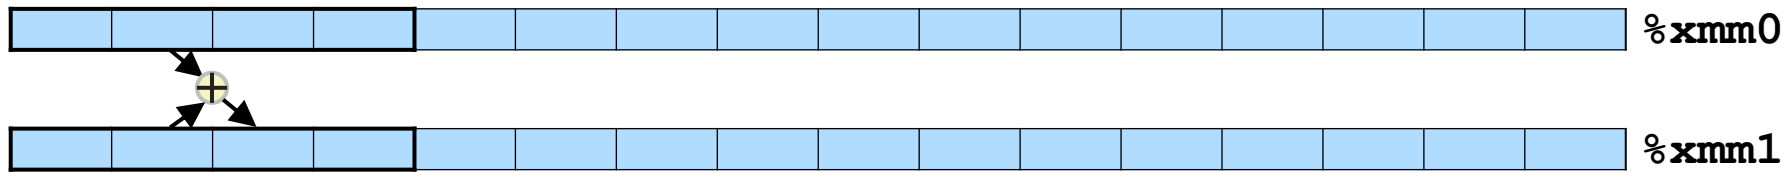
- **XMM Registers:** 16 total, each 16 bytes

- 16 single-byte integers 
- 8 16-bit integers 
- 4 32-bit integers 
- 4 single-precision floats 
- 2 double-precision floats 
- 1 single-precision float 
- 1 double-precision float 

# Scalar & SIMD Operations

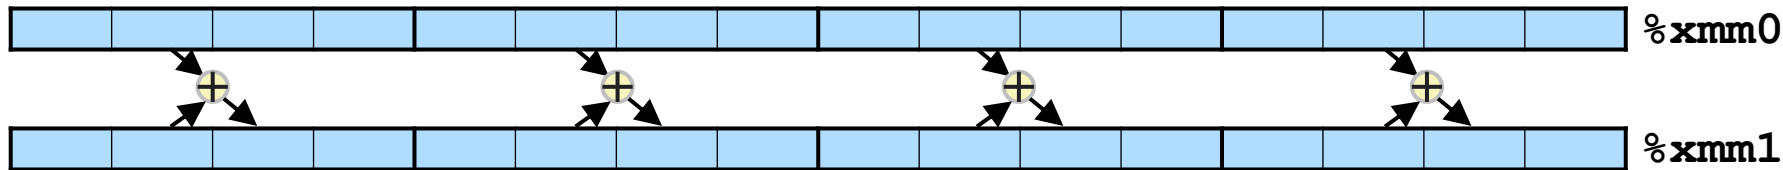
- Scalar operations: single precision

`addss %xmm0, %xmm1`



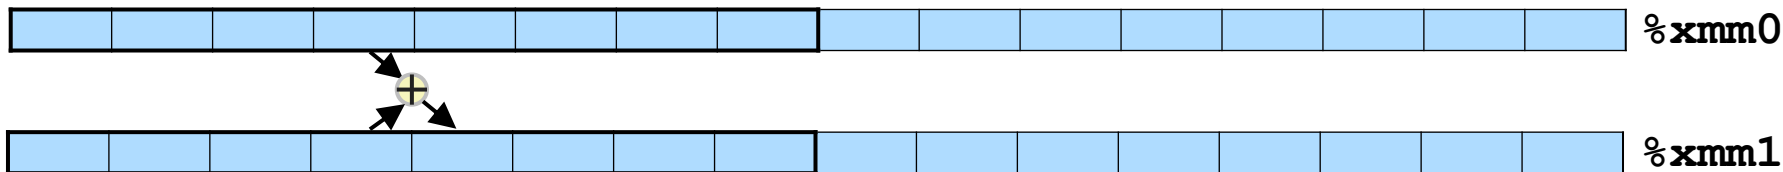
- SIMD operations: single precision

`addps %xmm0, %xmm1`



- Scalar operations: double precision

`addsd %xmm0, %xmm1`



# FP Basics

- Arguments passed in `%xmm0`, `%xmm1`, ...
- Result returned in `%xmm0`
- All XMM registers **caller-saved**

```
float fadd(float x, float y){  
    return x + y;  
}
```

```
# x in %xmm0, y in %xmm1  
addss    %xmm1, %xmm0  
ret
```

```
double dadd(double x, double y){  
    return x + y;  
}
```

```
# x in %xmm0, y in %xmm1  
addsd    %xmm1, %xmm0  
ret
```

# FP Memory Referencing

- Integer (and pointer) arguments passed in regular registers
- FP values passed in XMM registers
- Different `mov` instructions to move between XMM registers from the ones to move between memory and XMM registers

```
double dincr(double *p, double v){  
    double x = *p;  
    *p = x + v;  
    return x;  
}
```

```
# p in %rdi, v in %xmm0  
movapd    %xmm0, %xmm1    # Copy v  
movsd     (%rdi), %xmm0    # x = *p  
addsd     %xmm0, %xmm1    # t = x + v  
movsd     %xmm1, (%rdi)    # *p = t  
ret
```

# Other Aspects of FP Code

- Lots of instructions
  - Different operations, different formats, ...
- Floating-point comparisons
  - Instructions `ucomiss` and `ucomisd`
  - Set condition codes `ZF`, `PF` and `CF`
  - Zeros `OF` and `SF`
- Using constant values
  - Set `%xmm0` register to 0 with instruction `xorpd %xmm0, %xmm0`
  - Others loaded from memory

Parity Flag

UNORDERED:	{ZF,PF,CF} ← 111
GREATER_THAN:	{ZF,PF,CF} ← 000
LESS_THAN:	{ZF,PF,CF} ← 001
EQUAL:	{ZF,PF,CF} ← 100



# Summary

---

- Arrays
  - Elements packed into contiguous region of memory
  - Use index arithmetic to locate individual elements
- Structures
  - Elements packed into single region of memory
  - Access using offsets determined by compiler
  - Possible require internal and external padding to ensure alignment
- Combinations
  - Can nest structure and array code arbitrarily
- Floating point
  - Data held and operated on in XMM registers

# [CSED211] Introduction to Computer Software Systems

## Lecture 7: Data

Prof. Jisung Park



**CAOS**  
COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2023.10.11