

[CSED211] Introduction to Computer Software Systems

Lecture 8: Advanced Topics

Prof. Jisung Park



CAOS

COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.10.16


Lecture Agenda

- Unions
- Memory Layout
- Buffer Overflow
 - Vulnerability
 - Protection

Union Allocation


- Allocate according to largest element
- Can only use one field at a time

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



Memory

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```




Memory

Union Allocation


- Allocate according to largest element
- Can only use one field at a time

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



Memory

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```




Memory

Union Allocation


- Allocate according to largest element
- Can only use one field at a time

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



Memory

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```




Memory

Union Allocation


- Allocate according to largest element
- Can only use one field at a time

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



Memory


```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```



Memory

Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u) {  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as `(float) u`?

```
unsigned float2bit(float f) {  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as `(unsigned) f`?

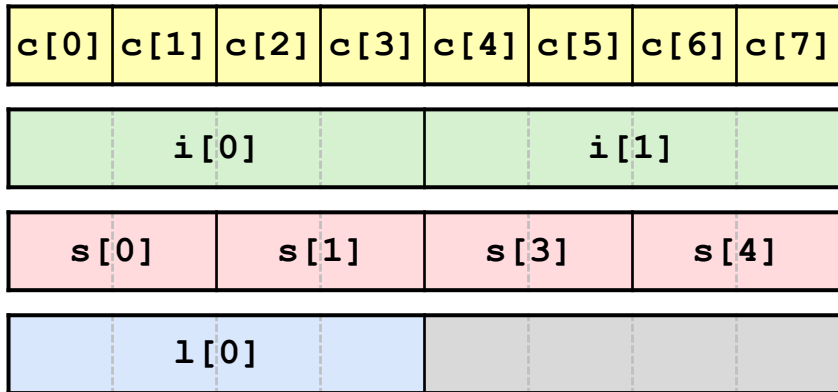
Byte Ordering Revisited

- Problem
 - `short/int/long` stored in memory as 2/4/8 consecutive bytes
 - Which is most (least) significant? – Which is the first byte: MSB or LSB?
 - Can cause problems when exchanging binary data between machines
- Big endian: the most significant byte has the lowest address
 - e.g., SPARC
- Little endian: the least significant byte has the lowest address
 - e.g., Intel x86
- Bi-endian: can be configured either way
 - e.g., ARM

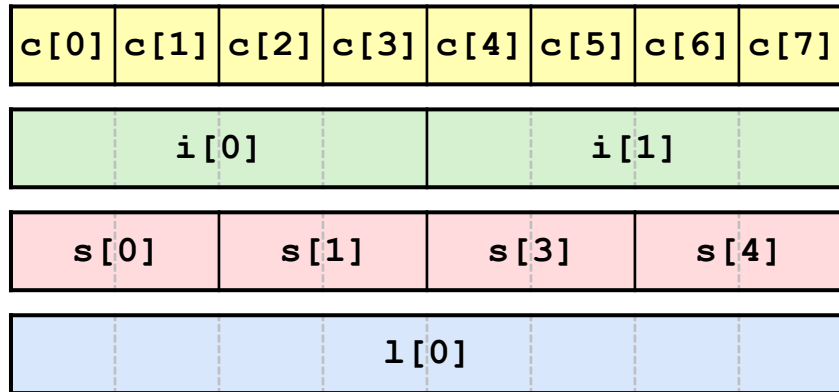
Byte Ordering Example

```
union {  
    unsigned char    c[8];  
    unsigned short   s[4];  
    unsigned int      i[2];  
    unsigned long     l[1];  
}dw;
```

How are the bytes inside **short/int/long** stored?



32-Bit Machines



64-Bit Machines

Byte Ordering Example (64-Bit Machines)

```
int j;  
for (j = 0; j < 8; j++)  
    dw.c[j] = 0xf0 + j;
```

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

```
printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",  
       dw.c[0], dw.c[1], dw.c[2], dw.c[3], dw.c[4], dw.c[5], dw.c[6], dw.c[7]);
```

```
printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",  
       dw.s[0], dw.s[1], dw.s[2], dw.s[3]);
```

s[0]		s[1]		s[3]		s[4]	
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

```
printf("Ints 0-1 == [0x%x,0x%x]\n",  
       dw.i[0], dw.i[1]);
```

i[1]				i[1]			
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

```
printf("Long 0 == [0x%lx]\n",  
       dw.l[0]);
```

l[0]							
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

Byte Ordering Example (64-Bit, Little Endian)

```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",
       dw.c[0], dw.c[1], dw.c[2], dw.c[3], dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",
       dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x]\n",
       dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx]\n",
       dw.l[0]);
```

Diagram illustrating the memory layout for the provided code, showing the byte ordering (LSB to MSB) for each data type:

Character Array (c):

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

Short Array (s):

s[0]		s[1]		s[3]		s[4]	
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7
LSB	MSB	LSB	MSB	LSB	MSB	LSB	MSB

Integer Array (i):

i[1]				i[1]			
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7
LSB		MSB		LSB		MSB	

Long (l):

l[0]							
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7
LSB				MSB			

Byte Ordering Example (64-Bit, Little Endian)

```
int j;  
for (j = 0; j < 8; j++)  
    dw.c[j] = 0xf0 + j;
```

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

```
printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",  
      dw.c[0], dw.c[1], dw.c[2], dw.c[3], dw.c[4], dw.c[5], dw.c[6], dw.c[7]);
```

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]

```
printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",  
      sh.c[0], sh.c[1], sh.c[2], sh.c[3]);
```

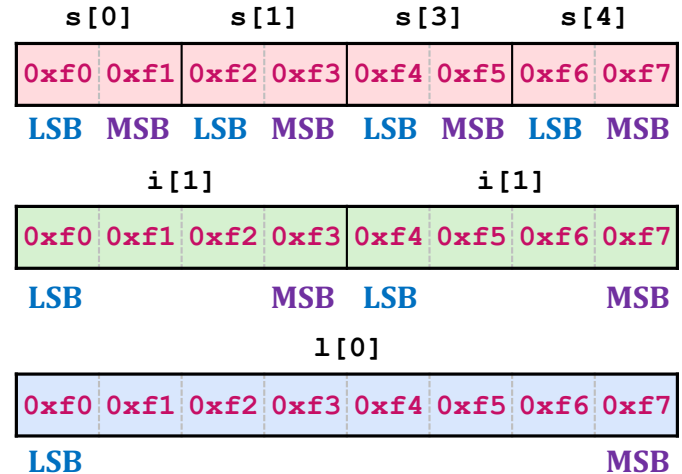
Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]

```
printf("Ints 0-1 == [0x%x,0x%x]\n",  
      i.c[0], i.c[1]);
```

Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]

```
printf("Long 0 == [0x%x]\n",  
      l.c[0]);
```

Long 0 == [0xf7f6f5f4f3f2f1f0]



Byte Ordering Example (64-Bit, Big Endian)

```
int j;  
for (j = 0; j < 8; j++)  
    dw.c[j] = 0xf0 + j;
```

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
0xf0	0xf1	0xf2	0xf3	0xf4	0xf5	0xf6	0xf7

```
printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",  
      dw.c[0], dw.c[1], dw.c[2], dw.c[3], dw.c[4], dw.c[5], dw.c[6], dw.c[7]);
```

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]

```
printf("Shorts 0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]\n",  
      sh.c[0], sh.c[1], sh.c[2], sh.c[3]);
```

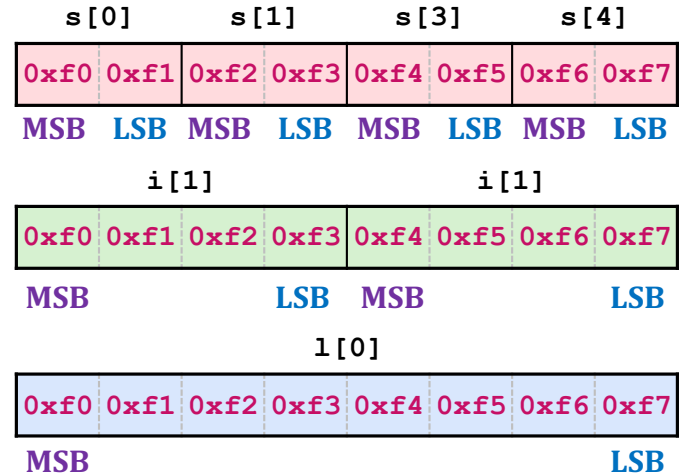
Shorts 0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]

```
printf("Ints 0-1 == [0xf0f1f2f3,0xf4f5f6f7]\n",  
      in.c[0], in.c[1]);
```

Ints 0-1 == [0xf0f1f2f3,0xf4f5f6f7]

```
printf("Long 0 == [0xf0f1f2f3f4f5f6f7]\n",  
      lg.c[0]);
```

Long 0 == [0xf0f1f2f3f4f5f6f7]



Lecture Agenda

- Unions
- Memory Layout
- Buffer Overflow
 - Vulnerability
 - Protection

x86-64 Linux Memory Layout

- **Stack**

- Runtime stack (8-MiB limit by default)
- e.g., local variables

- **Heap**


- Dynamically allocated as needed
- e.g., `malloc()`, `calloc()`, `new()`, etc.

- **Data**

- Statically allocated data
- e.g., global variables, static variables, string constants

- **Text and shared libraries**

- Executable machine instructions
- Read-only



x86-64 Linux Memory Layout

```
char big_array[1L<<24]; // 16 MB
char huge_array[1L<<31]; // 2 GB

int global = 0;

int useless(){ return 0; }


int main (){
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1L<<28); // 256 MB
    p2 = malloc(1L<<8); // 256 B
    p3 = malloc(1L<<32); // 4 GB
    p4 = malloc(1L<<8); // 256 B
    // Some print statements below
    :
}
```

local	0x00007fffffffbefc
p1	0x00007f7262a1e010
p2	0x000000008359d010
p3	0x00007f7162a1d010
p4	0x000000008359d120
big_array	0x0000000080601060
huge_array	0x0000000000601060
main()	0x000000000040060c
useless()	0x0000000000400590

0x0000000000000000

0x0000000000040000

0x00007fffffffbefc




Memory

Runaway Stack Example

```
int recurse(int x) {  
    int a[1<<15]; // 4*2^15 = 128 KiB  
    printf("x = %d.  a at %p\n", x, a);  
    a[0] = (1<<14)-1;  
    a[a[0]] = x-1;  
    if (a[a[0]] == 0)  
        return -1;  
    return recurse(a[a[0]]) - 1;  
}
```

```
$ ./runaway 67  
x = 67  a at 0x7ffd18aba930  
x = 66  a at 0x7ffd18a9a920  
x = 65  a at 0x7ffd18a7a910  
...  
x = 4   a at 0x7ffd182da540  
x = 3   a at 0x7ffd182ba530  
x = 2   a at 0x7ffd1829a520  
Segmentation fault (core dumped)
```

- Functions store local data in their stack frame
- Recursive functions cause deep nesting of frames



Lecture Agenda

- Unions
- Memory Layout
- Buffer Overflow
 - Vulnerability
 - Protection

Recall: Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

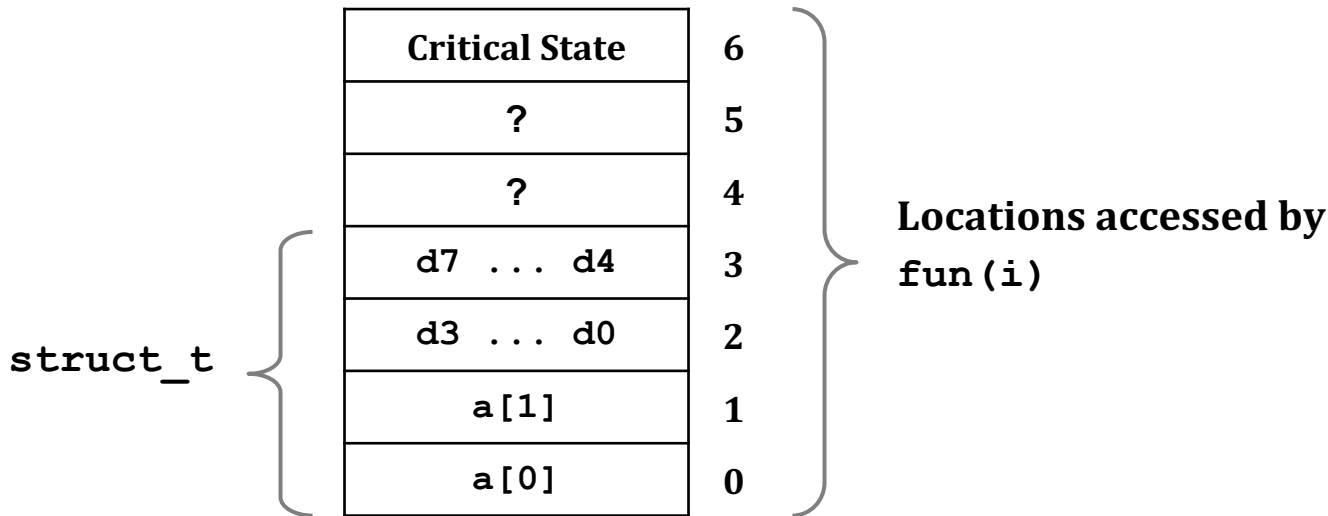
Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

Explanation:



Such Problems Are a BIG Deal

- Generally called a **buffer overflow**
 - When exceeding the memory size allocated for an array
- Why a big deal?
 - It's the **#1 technical cause** of security vulnerabilities
 - #1 overall cause is social engineering/user ignorance
- Most common form
 - Unchecked lengths on string inputs
 - Particularly for bounded character arrays on the stack
 - Sometimes referred to as **stack smashing**

Vulnerable Buffer Code

- Implementation of Unix function gets()

```
/* Get string from stdin */
char *gets(char *dest){
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- **Problem:** no way to specify on the number of characters to read
- Similar problems with other library functions
 - **strcpy()** and **strcat()**: copy strings of arbitrary length
 - **scanf()**, **fscanf()**, and **sscanf()**: when given %s conversion specification

Vulnerable Buffer Code

```
/* Echo Line */
void echo() {
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}

void call_echo() {
    echo();
}
```

How big is big enough?

```
$ ./bufdemo
Type a string: 01234567890123456789012
01234567890123456789012
```

```
$ ./bufdemo
Type a string: 012345678901234567890123
01234568901234567890123
Segmentation Fault
```

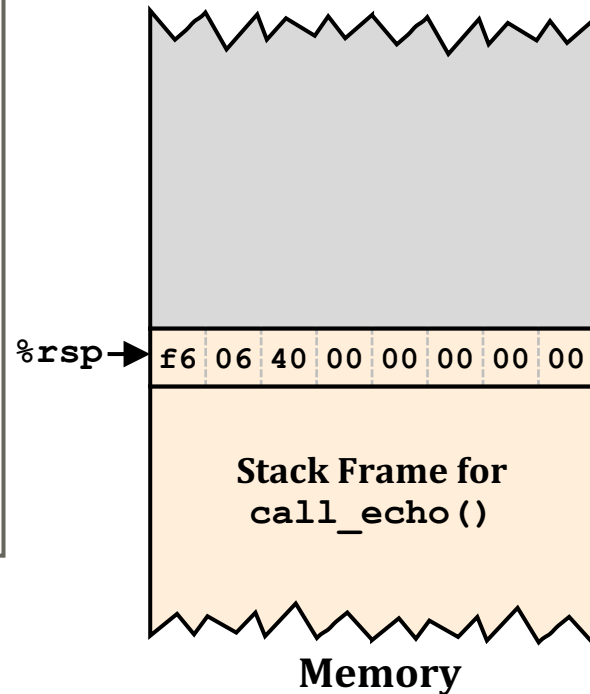
Vulnerable Buffer Code

00000000004006cf <echo>:

4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$8,%rsp
4006fa:	c3	retq	



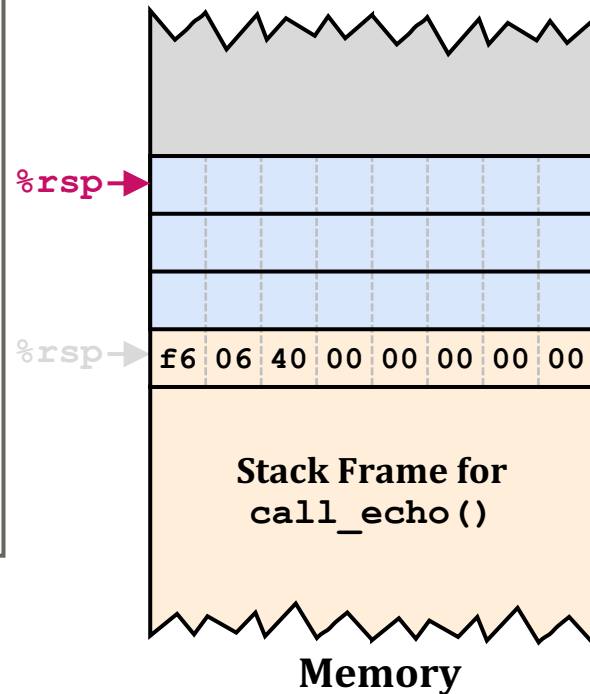
Vulnerable Buffer Code

```
00000000004006cf <echo>:
```

4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

```
00000000004006e8 <call_echo>:
```

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$8,%rsp
4006fa:	c3	retq	



Vulnerable Buffer Code

```
00000000004006cf <echo>:
```

```
4006cf: 48 83 ec 18
```

```
4006d3: 48 89 e7
```

```
4006d6: e8 a5 ff ff ff
```

```
4006db: 48 89 e7
```

```
4006de: e8 3d fe ff ff
```

```
4006e3: 48 83 c4 18
```

```
4006e7: c3
```

```
sub $24,%rsp
```

```
mov %rsp,%rdi
```

```
callq 400680 <gets>
```

```
mov %rsp,%rdi
```

```
callq 400520 <puts@plt>
```

```
add $24,%rsp
```

```
retq
```

```
00000000004006e8 <call_echo>:
```

```
4006e8: 48 83 ec 08
```

```
4006ec: b8 00 00 00 00
```

```
4006f1: e8 d9 ff ff ff
```

```
4006f6: 48 83 c4 08
```

```
4006fa: c3
```

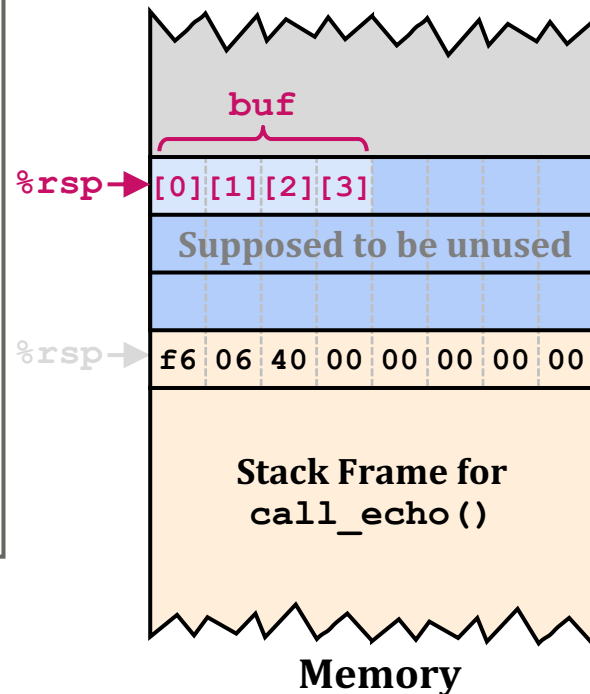
```
sub $8,%rsp
```

```
mov $0,%eax
```

```
callq 4006cf <echo>
```

```
add $8,%rsp
```

```
retq
```



Vulnerable Buffer Code

00000000004006cf <echo>:

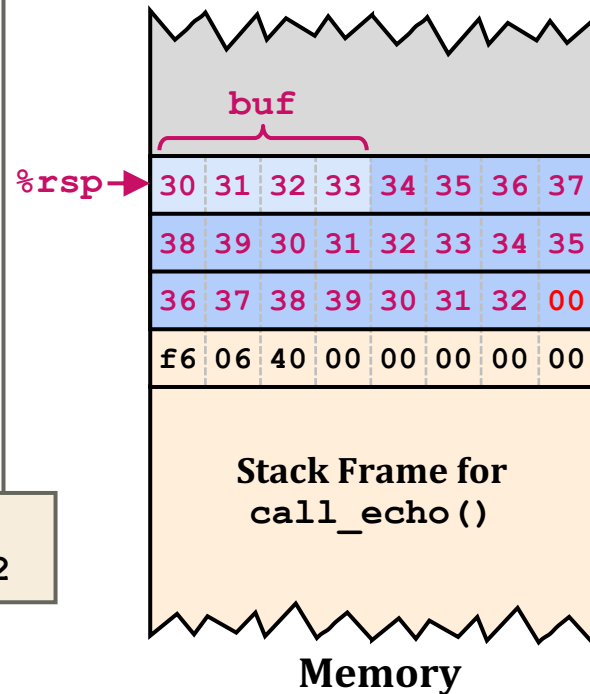
4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$8,%rsp
4006fa:	c3	retq	

\$./bufdemo

Type a string: 01234567890123456789012



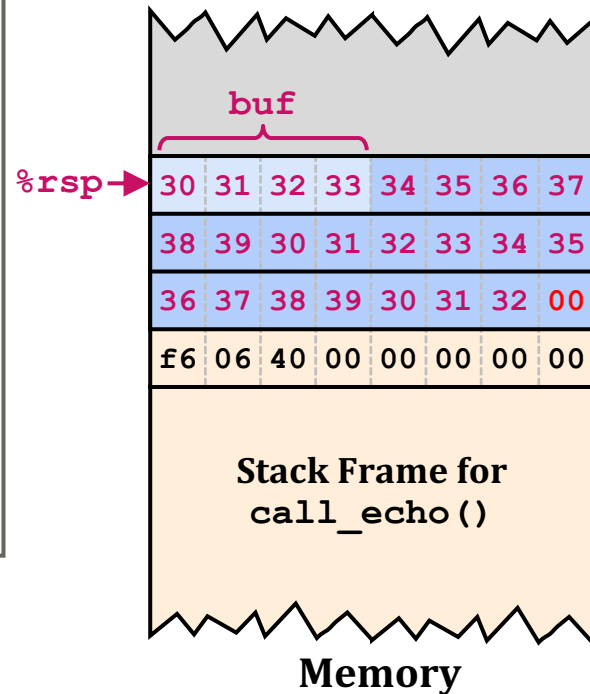
Vulnerable Buffer Code

00000000004006cf <echo>:

4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$8,%rsp
4006fa:	c3	retq	



Vulnerable Buffer Code

00000000004006cf <echo>:

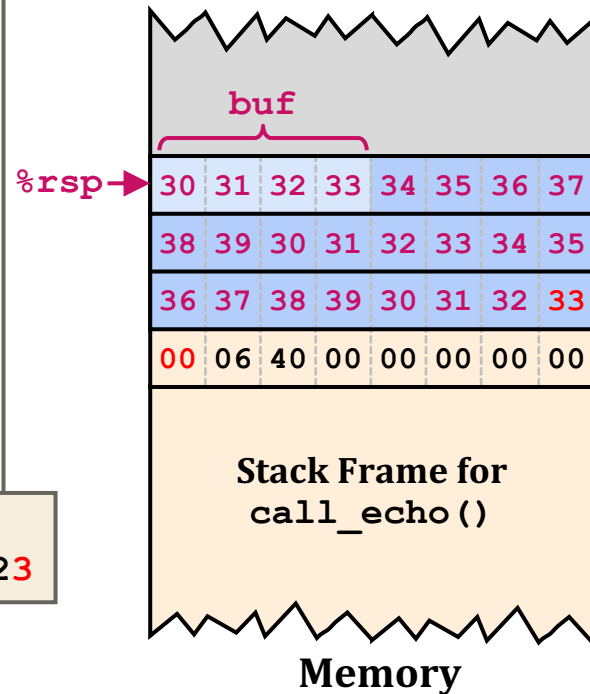
4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$8,%rsp
4006fa:	c3	retq	

\$./bufdemo

Type a string: 012345678901234567890123



Vulnerable Buffer Code

00000000004006cf <echo>:

4006cf: 48 83 c4 08

4006d3: 48 83 c4 08

4006d6: e8 a1 00 00 00

4006db: 48 83 c4 08

4006de: e8 3d 00 00 00

4006e3: 48 83 c4 08

4006e7: c3

register_tm_clones:

⋮

400600: mov %rsp,%rbp

400603: mov %rax,%rdx

400606: shr \$0x3f,%rdx

40060a: add %rdx,%rax

40060d: sar %rax

400610: jne 400614

400612: pop %rbp

400613: retq

00000000004006e8 <@plt>:

4006e8: 48 83 c4 08

4006ec: b8 00 00 00 00

4006f1: e8 d9 ff ff ff

4006f6: 48 83 c4 08

4006fa: c3

callq 4006cf <echo>

add \$8,%rsp

retq

\$./bufdemo


Type a string: 012345678901234567890123

012345678901234567890123

Segmentation Fault

Return to

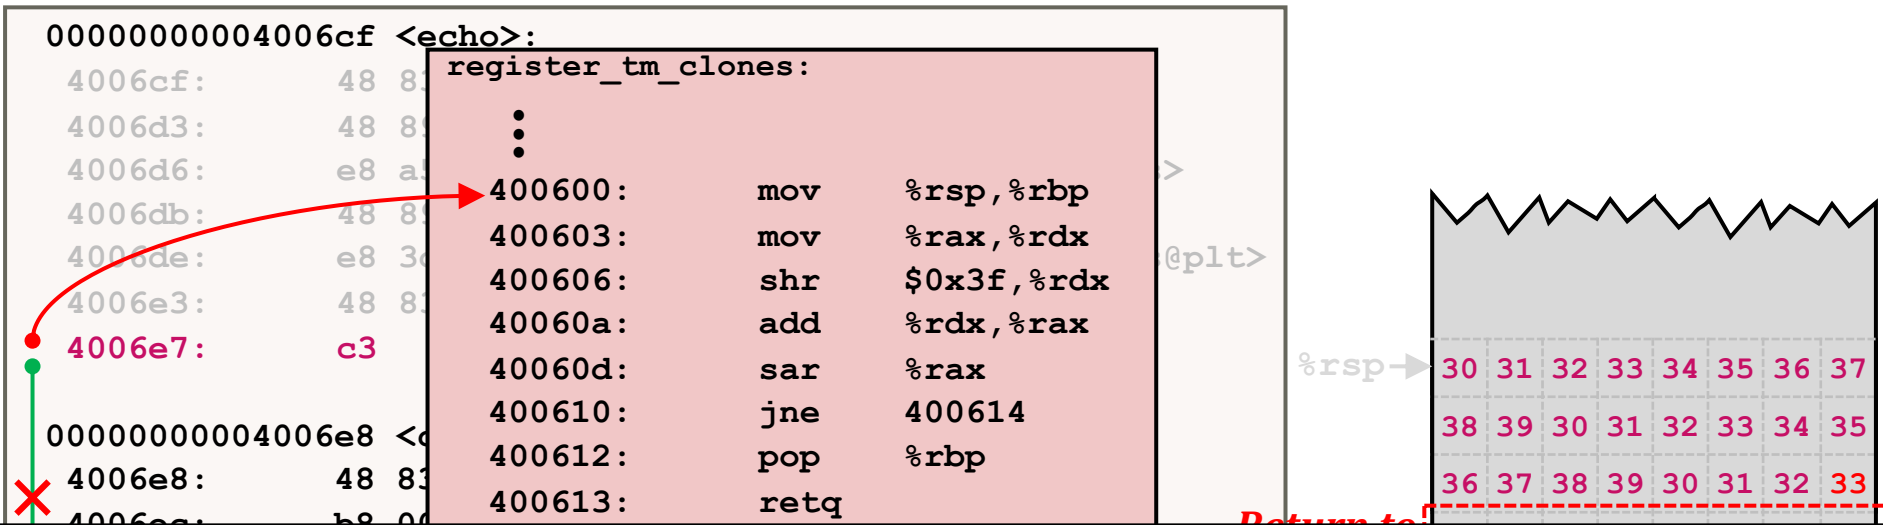
%rsp→



Stack Frame for
call_echo()

Memory

Vulnerable Buffer Code



Lots of things happen, without modifying critical state

Eventually executes `retq` back to `main()`

\$./bufdemo

Type a string: 012345678901234567890123

012345678901234567890123

Segmentation Fault

```
call echo()
```

Memory

Stack Smashing Attacks


```
void P() {  
    Q();  
    ⋮  
}
```

Return Addr. **A**

```
int Q() {  
    char buf[64];  
    gets(buf);  
    ⋮  
    return val;  
}
```

```
int S() {  
    /* Something  
    unexpected */  
    ⋮  
}
```

gets() :
*Overwrite the return address A
with an address of some other code S*



Stack after Call to `gets()`

Stack Smashing Attacks

```
void P() {  
    Q();  
    ⋮  
}
```

Return Addr. **A**


```
int Q() {  
    char buf[64];  
    gets(buf);  
    ⋮  
    return val;  
}
```

```
int S() {  
    /* Something  
    unexpected */  
    ⋮  
}
```

gets() :

*Overwrite the return address **A**
with an address of some other code **S***

*When **Q** executes ret,
will jump to the other code **S***



Stack after Call to `gets()`

Crafting Smashing String

00000000004006cf <echo>:

4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

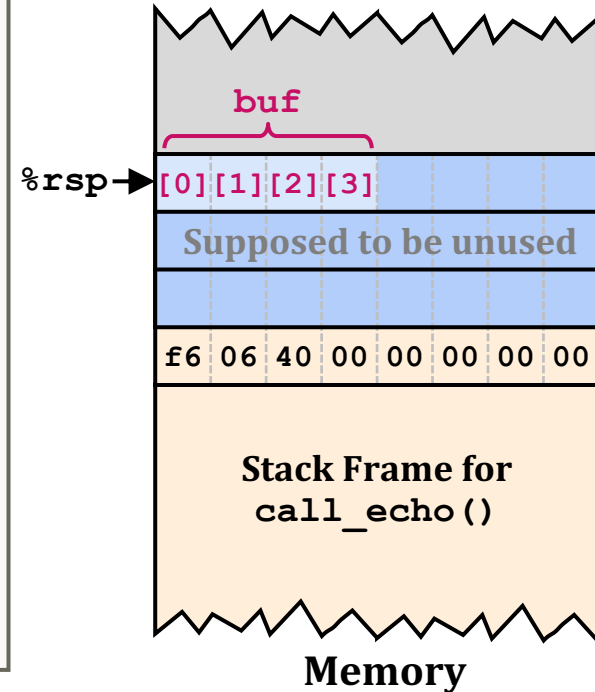
4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>

⋮

00000000004006fb <smash>:

4006fb:	48 83 ec 08
---------	-------------

⋮



Crafting Smashing String

00000000004006cf <echo>:

4006cf:	48 83 ec 18	sub	\$24,%rsp
4006d3:	48 89 e7	mov	%rsp,%rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp,%rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$24,%rsp
4006e7:	c3	retq	

00000000004006e8 <call_echo>:

4006e8:	48 83 ec 08	sub	\$8,%rsp
4006ec:	b8 00 00 00 00	mov	\$0,%eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>

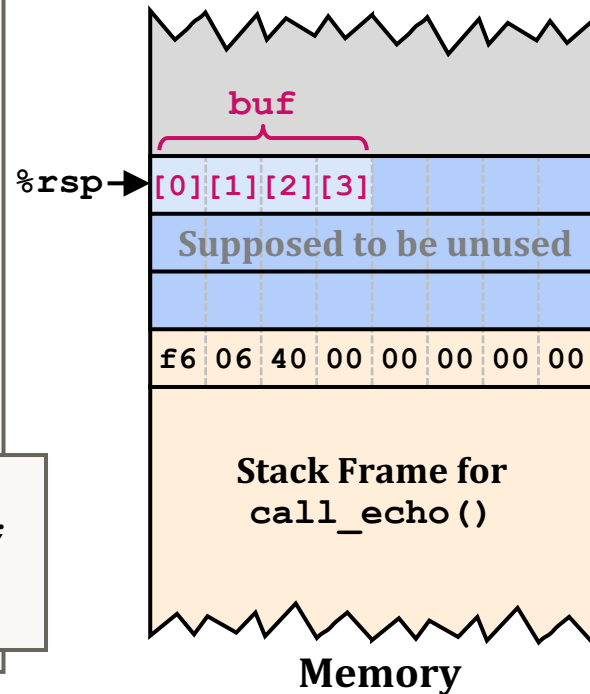
⋮

00000000004006fb <smash>:

4006fb: 48 83 ec 08

⋮

```
void smash() {  
    printf("I've been smashed!\n");  
    exit(0);  
}
```



Crafting Smashing String

00000000004006cf <echo>:

4006cf: 48 83 ec 18

4006d3: 48 89 e7

4006d6: e8 a5 ff ff ff

4006db: 48 89 e7

4006de: e8 3d fe ff ff

4006e3: 48 83 c4 18

4006e7: c3

00000000004006e8 <call_echo>:

4006e8: 48 83 ec 08

4006ec: b8 00 00 00 00

4006f1: e8 d9 ff ff ff

⋮

00000000004006fb <smash>:

4006fb: 48 83 ec 08

⋮

Attack String:

ffffffff ffffffff ffffffff ffffffff
ffffffff ffffffff fb064000 00000000

callq 400680 <gets>

mov %rsp,%rdi

callq 400520 <puts@plt>

add \$24,%rsp

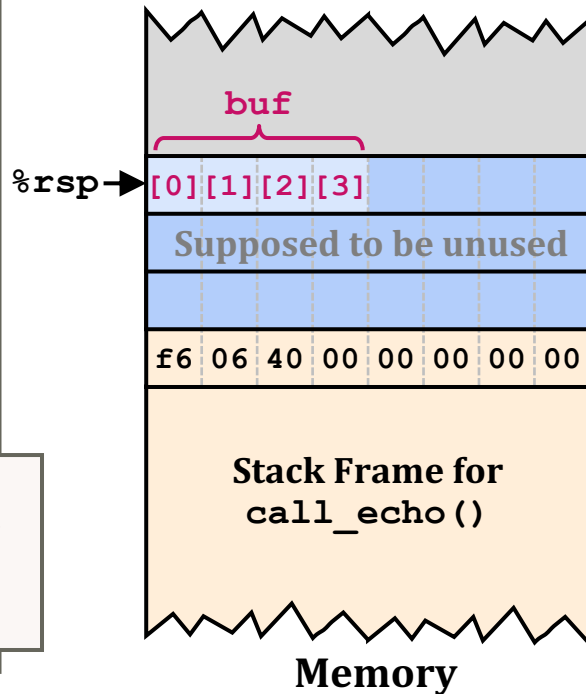
retq

sub \$8,%rsp

mov \$0,%eax

callq 4006cf <echo>

```
void smash() {  
    printf("I've been smashed!\n");  
    exit(0);  
}
```



Crafting Smashing String

00000000004006cf <echo>:

4006cf: 48 83 ec 18

4006d3: 48 89 e7

4006d6: e8 a5 ff ff ff

4006db: 48 89 e7

4006de: e8 3d fe ff ff

4006e3: 48 83 c4 18

4006e7: c3

00000000004006e8 <call_echo>:

4006e8: 48 83 ec 08

4006ec: b8 00 00 00 00

4006f1: e8 d9 ff ff ff

⋮

00000000004006fb <smash>:

4006fb: 48 83 ec 08

⋮

Attack String:

ffffffff ffffffff ffffffff ffffffff
ffffffff ffffffff fb064000 00000000

callq 400680 <gets>

mov %rsp,%rdi

callq 400520 <puts@plt>

add \$24,%rsp

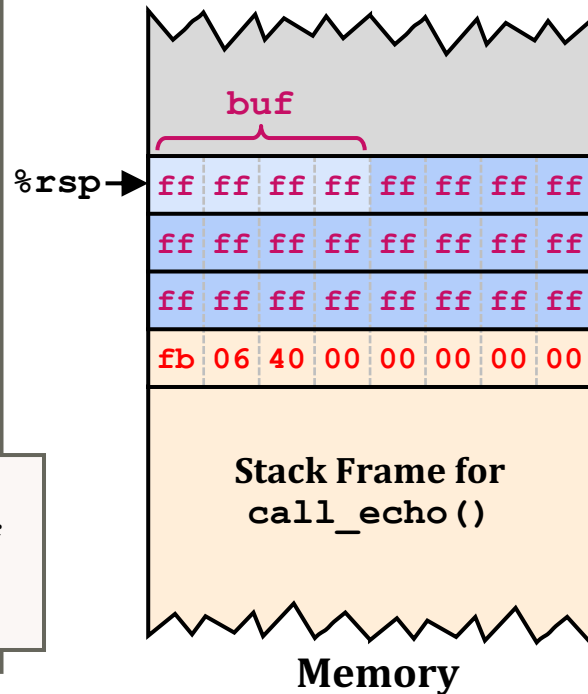
retq

sub \$8,%rsp

mov \$0,%eax

callq 4006cf <echo>

```
void smash() {  
    printf("I've been smashed!\n");  
    exit(0);  
}
```



Crafting Smashing String

00000000004006cf <echo>:

4006cf: 48 83 ec 18

4006d3: 48 89 e7

4006d6: e8 a5 ff ff ff

4006db: 48 89 e7

4006de: e8 3d fe ff ff

4006e3: 48 83 c4 18

4006e7: c3

00000000004006e8 <call_echo>:

4006e8: 48 83 ec 08

4006ec: b8 00 00 00 00

4006f1: e8 d9 ff ff ff

⋮

00000000004006fb <smash>:

4006fb: 48 83 ec 08

⋮

Attack String:

ffffffff ffffffff ffffffff ffffffff
ffffffff ffffffff fb064000 00000000

callq 400680 <gets>

mov %rsp,%rdi

callq 400520 <puts@plt>

add \$24,%rsp

retq

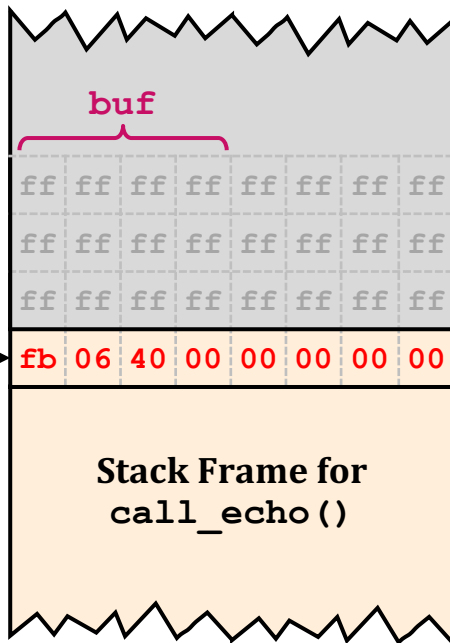
sub \$8,%rsp

mov \$0,%eax

callq 4006cf <echo>

```
void smash() {  
    printf("I've been smashed!\n");  
    exit(0);  
}
```

%rsp →



Memory


Code Injection Attacks

```
void P() {  
    Q();  
    ⋮  
}
```

Return Addr. **A**

```
int Q() {  
    char buf[64];  
    gets(buf);  
    ⋮  
    return val;  
}
```

gets() :
*Overwrite the return address A
with the address of buffer (B)
+ fill the buffer with byte representation
of executable code*



Code Injection Attacks


```
void P() {  
    Q();  
    ⋮  
}
```

Return Addr. **A**

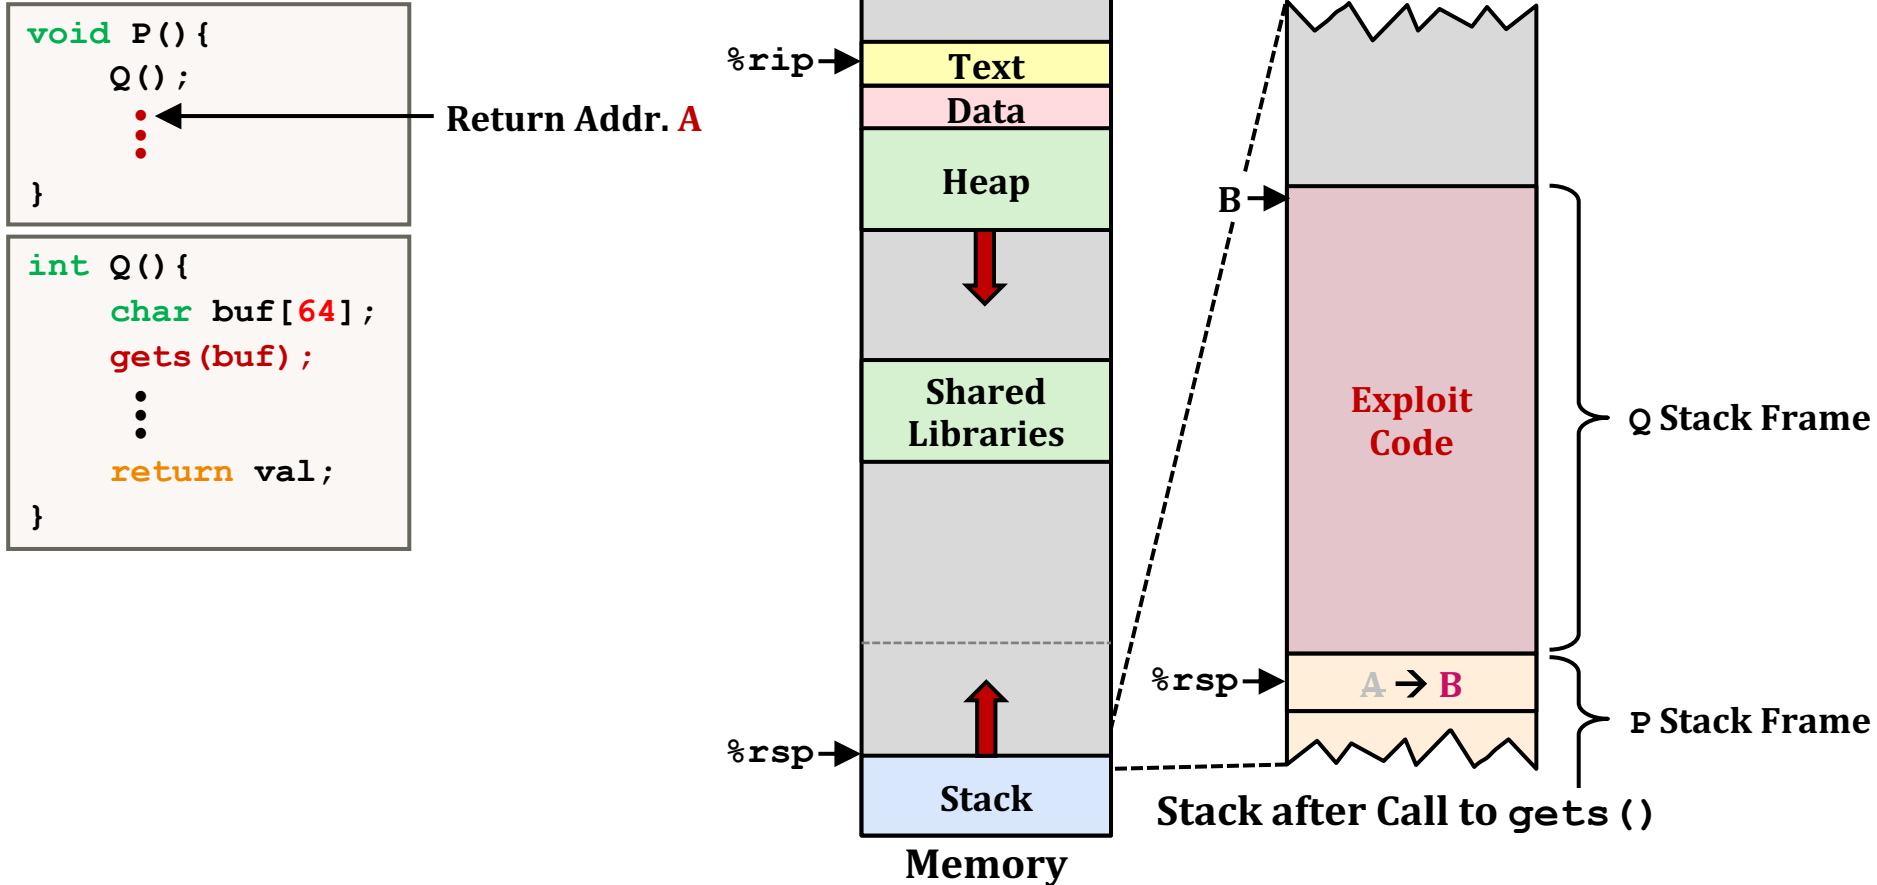
```
int Q() {  
    char buf[64];  
    gets(buf);  
    ⋮  
    return val;  
}
```

gets() :
*Overwrite the return address A
with the address of buffer (B)
+ fill the buffer with byte representation
of executable code*

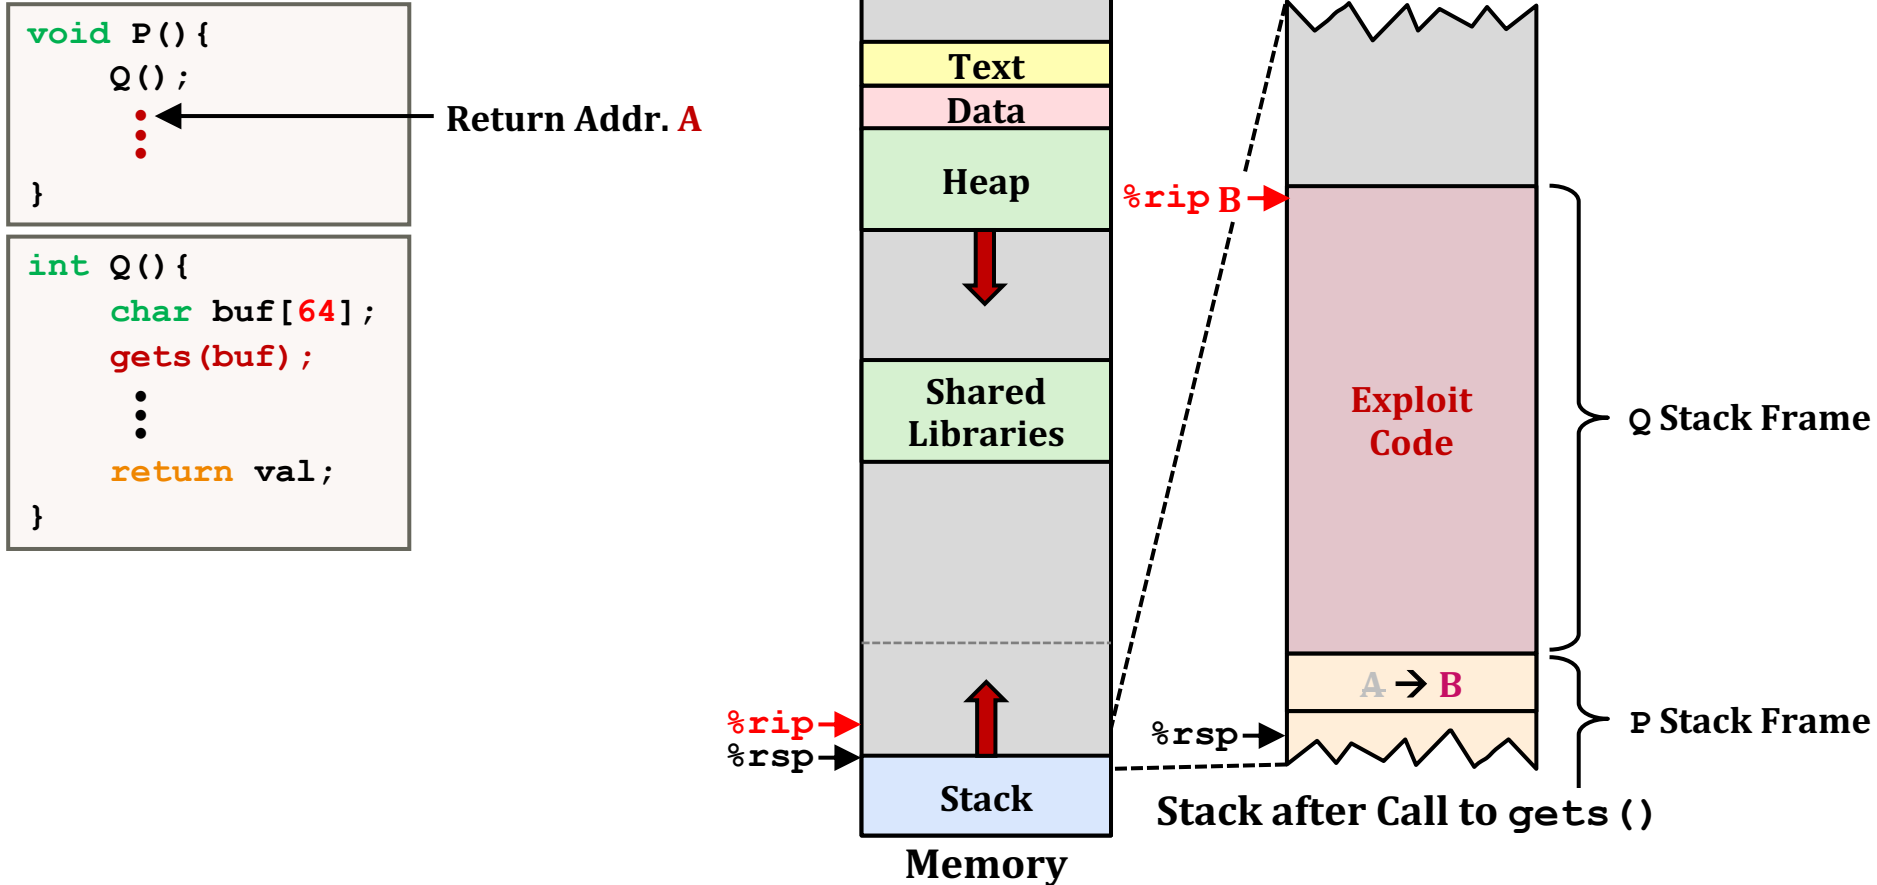
*When Q executes ret,
will jump to the exploit code*



Code Injection Attacks



Code Injection Attacks



Exploits Based on Buffer Overflows

- Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines
- Distressingly common in real programs
 - Programmers keep making the same mistakes
 - Recent measures make these attacks much more difficult
- Examples across the decades: Internet Worm (1988), IM Wars (1999), Twilight hack on Wii (2000s), and many, many more
- You will learn some of the tricks in Lab 4. Attack Lab
 - Hopefully to convince you to never leave such security holes in your program

Counter Measures Against Buffer Overflow Attacks

- Avoid overflow vulnerabilities
- Employ system-level protections
- Have compiler use stack canaries

1. Avoid Overflow Vulnerabilities in Code

```
/* Echo Line */  
void echo() {  
    char buf[4]; /* Way too small! */  
    // gets(buf); <- Vulnerable!!  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- Example: use library routines that limit string lengths
 - **fgets** instead of **gets**
 - Do **not** use **scanf with %s** conversion specification;
Use **fgets** to read the string or **%ns** where **n** is a suitable integer
 - **strncpy** instead of **strcpy**

2. System-Level Protections


- Randomized stack offsets: **ASLR (Address Space Layout Radomization)**
 - At start of program, allocate a random amount of space on stack
 - Shift stack addresses for the entire program
 - Make it difficult for adversaries to predict the beginning of the inserted code
 - e.g., 5 executions of memory allocation code: stack repositioned each time program executes

```
Local
0x7ffe4d3be87c
0x7fff75a4f9fc
0x7ffeadb7c80c
0x7ffeaea2fdac
0x7ffcd452017c
```

2. System-Level Protections (Cont.)

- Nonexecutable code segments
 - In traditional x86, a region of memory can be marked as either **read-only** or **writable**
 - Can execute anything readable
 - x86-64 added explicit **executable** permission
 - Stack marked as non-executable

*OS ensures (cooperating CPU):
any attempt to execute an instruction in non-executable region will fail*



3. Compiler-Level Protection: Stack Canaries

- Key idea
 - Place a special value, called **carnary**, on stack **just beyond buffer**
 - Check the value before exiting the function
 - The value is change → Something wrong must have happened
- GCC implementation
 - **-fstack-protector**
 - Now the default (used to be disabled earlier)

```
$ ./bufdemo-protected
Type a string: 0123456
0123456
```

```
$ ./bufdemo-protected
Type a string: 01234567
01234567
*** stack smashing detected ***
```


3. Stack Canaries: Protected Buffer Disassembly

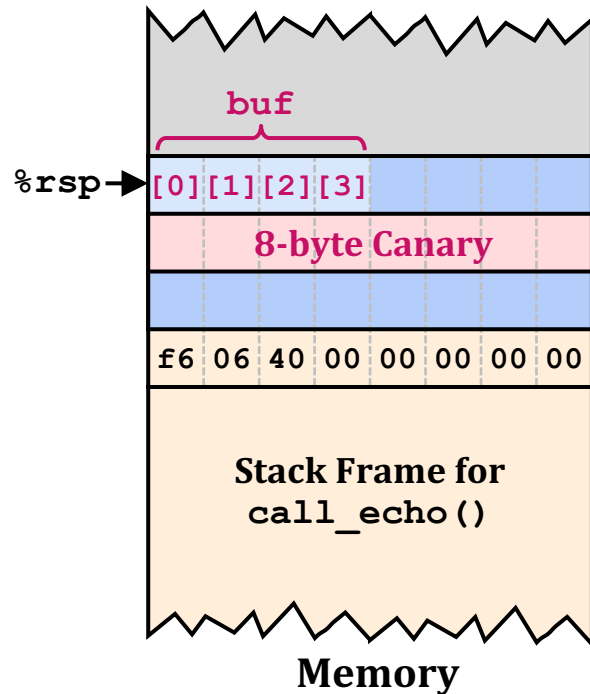
```
00000000004006cf <echo>:
```

```
40072f:      sub     $0x18,%rsp
400733:      mov     %fs:0x28,%rax
40073c:      mov     %rax,0x8(%rsp)
400741:      xor     %eax,%eax
400743:      mov     %rsp,%rdi
400746:      callq   4006e0 <gets>
40074b:      mov     %rsp,%rdi
40074e:      callq   400570 <puts@plt>
400753:      mov     0x8(%rsp),%rax
400758:      xor     %fs:0x28,%rax
400761:      je      400768 <echo+0x39>
400763:      callq   400580 <__stack_chk_fail@plt>
400768:      add     $0x18,%rsp
40076c:      retq
```

3. Stack Canaries: Setting Up Canary

```
00000000004006cf <echo>:
```

```
40072f:      sub    $0x18,%rsp
400733:      mov     %fs:0x28,%rax    # Get canary
40073c:      mov     %rax,0x8(%rsp)   # Place on Stack
400741:      xor     %eax,%eax        # Delete canary
400743:      mov     %rsp,%rdi
400746:      callq   4006e0 <gets>
40074b:      mov     %rsp,%rdi
40074e:      callq   400570 <puts@plt>
400753:      mov     0x8(%rsp),%rax
400758:      xor     %fs:0x28,%rax
400761:      je      400768 <echo+0x39>
400763:      callq   400580 <__stack_chk_fail@plt>
400768:      add     $0x18,%rsp
40076c:      retq
```




3. Stack Canaries: Checking Canary

```
00000000004006cf <echo>:
```

```
40072f:      sub    $0x18,%rsp
400733:      mov     %fs:0x28,%rax  # Get canary
40073c:      mov     %rax,0x8(%rsp) # Place on Stack
400741:      xor     %eax,%eax      # Delete canary
400743:      mov     %rsp,%rdi
400746:      callq   4006e0 <gets>
40074b:      mov     %rsp,%rdi
40074e:      callq   400570 <puts@plt>
400753:      mov     0x8(%rsp),%rax # Retrieve from stack
400758:      xor     %fs:0x28,%rax  # Compare to canary
400761:      je      400768 <echo+0x39> # If same, OK
400763:      callq   400580 <__stack_chk_fail@plt> # Fail
400768:      add     $0x18,%rsp
40076c:      retq
```

```
$ ./bufdemo-protected
Type a string: 0123456
```

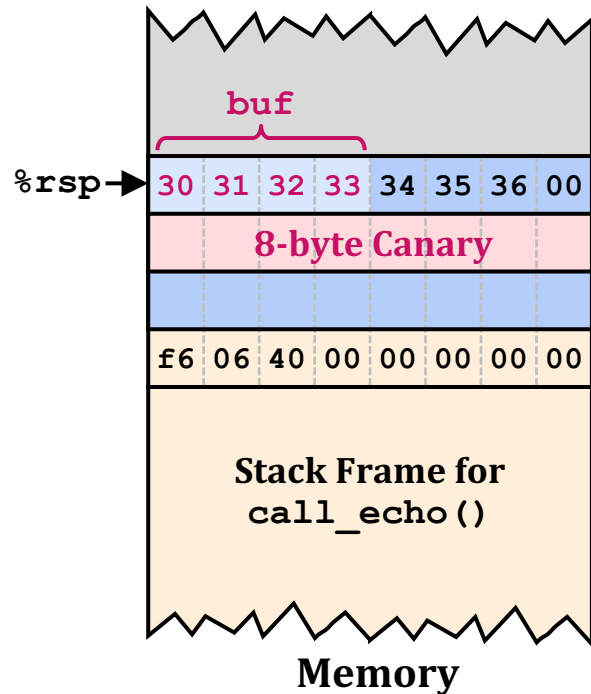


3. Stack Canaries: Checking Canary

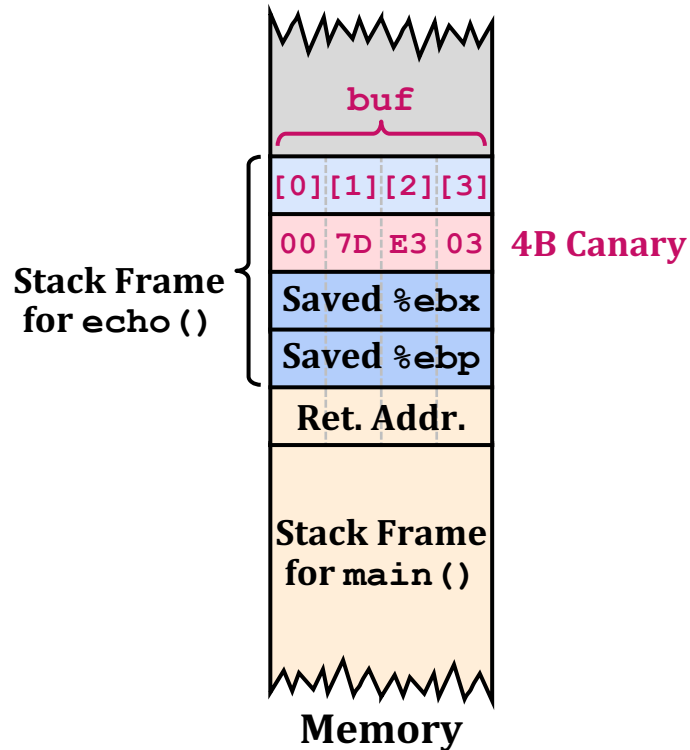
```
00000000004006cf <echo>:
```

```
40072f:      sub    $0x18,%rsp
400733:      mov     %fs:0x28,%rax  # Get canary
40073c:      mov     %rax,0x8(%rsp) # Place on Stack
400741:      xor     %eax,%eax      # Delete canary
400743:      mov     %rsp,%rdi
400746:      callq   4006e0 <gets>
40074b:      mov     %rsp,%rdi
40074e:      callq   400570 <puts@plt>
400753:      mov     0x8(%rsp),%rax # Retrieve from stack
400758:      xor     %fs:0x28,%rax  # Compare to canary
400761:      je      400768 <echo+0x39> # If same, OK
400763:      callq   400580 <__stack_chk_fail@plt> # Fail
400768:      add     $0x18,%rsp
40076c:      retq
```

```
$ ./bufdemo-protected
Type a string: 0123456
0123456
```

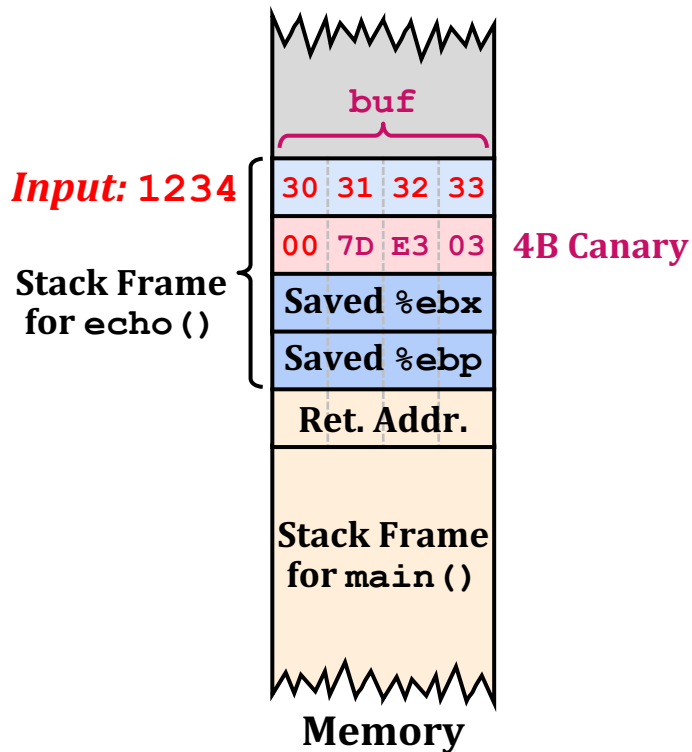


3. Stack Canaries: Example



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

3. Stack Canaries: Example



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

Benign corruption (tolerated)!:

Allows a programmer to make silent **off-by-one errors**

Return-Oriented Programming Attacks

- Challenges (for hackers)
 - Stack randomization makes it hard to predict buffer location
 - Marking stack nonexecutable makes it hard to insert binary code
- Alternative Strategy
 - Use existing code
 - e.g., library code from `stdlib`
 - String together fragments to achieve overall desired outcome
 - Does **not** overcome stack canaries
- Construct program from **gadgets**
 - Sequence of instructions ending in `ret` that is **encoded by single byte 0xc3**
 - **Code positions fixed** from run to run
 - Code is **executable**

Gadget Example#1

```
long ab_plus_c(long a, long b, long c){  
    return a * b + c;  
}
```

```
00000000004004d0 <ab_plus_c>:  
4004d0: 48 0f af fe  imul %rsi,%rdi  
4004d4: 48 8d 04 17  lea (%rdi,%rdx,1),%rax  
4004d8: c3           retq
```

$\%rax \leftarrow \%rdi + \%rdx$

Gadget Address = 0x4004d4

- Use the tail end of existing functions

Gadget Example#2

```
void setval(unsigned *p) {  
    *p = 3347663060u;  
}
```

```
<setval>:           Encodes movq %rax, %rdi  
4004d9:  c7 07 d4 48 89 c7  movl  $0xc78948d4, (%rdi)  
4004df:  c3                retq
```

$\%rdi \leftarrow \%rax$

Gadget Address = 0x4004dc

- Use the tail end of existing functions

Gadget Example#2

```
void setval(unsigned *p){  
    *p = 3347663060u;  
}
```

```
<setval>:                               Encodes movq %rax, %rdi  
4004d9:  c7 07 d4 48 89 c7  movl  $0xc78948d4, (%rdi)  
4004df:  c3                retq
```

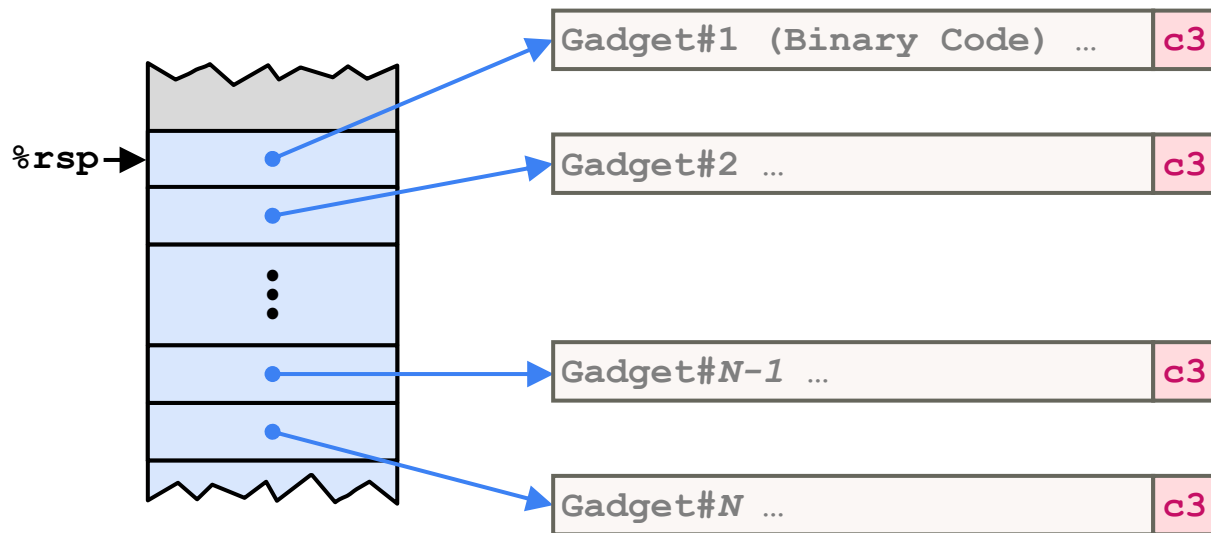
$\%rdi \leftarrow \%rax$

Gadget Address = 0x4004dc


- Use the tail end of existing functions

ROP Execution

- Trigger with ret instruction
 - Will start executing Gadget 1
 - `ret` → `pop %rip`
- Final ret in each gadget will start next one



Crafting an ROP Attack String



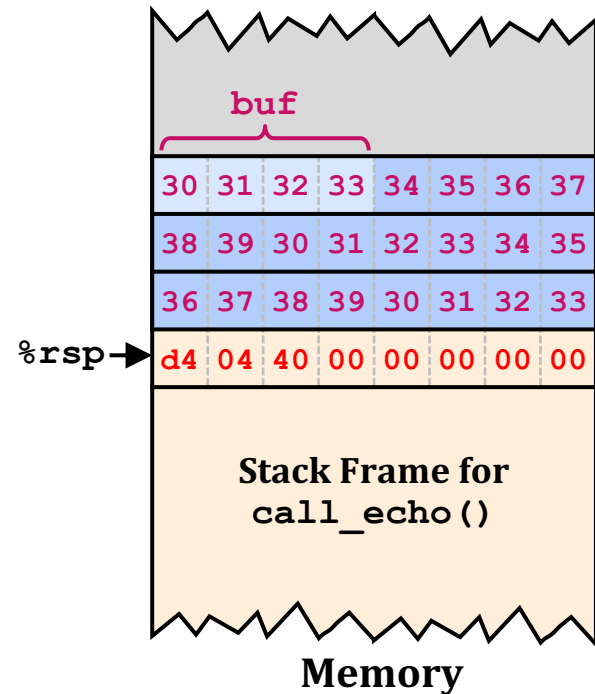
```
int echo() {  
    char buf[4];  
    gets(buf);  
    ⋮  
    return ret  
}
```

Attack: Makes echo() return %rdi+%rdx

Gadget: %rax ← %rdi + %rdx (+ ret)

```
00000000004004d0 <ab_plus_c>:  
4004d0: 48 0f af fe  imul %rsi,%rdi  
4004d4: 48 8d 04 17  lea (%rdi,%rdx,1),%rax  
4004d8: c3           retq
```

Crafting an ROP Attack String



```
int echo() {
    char buf[4];
    gets(buf);
    :
    return ret
}
```

Attack: Makes echo () return %rdi+%rdx

Gadget: `%rax ← %rdi + %rdx (+ ret)`

```
00000000004004d0 <ab_plus_c>:
 4004d0: 48 0f af fe    imul %rsi,%rdi
4004d4: 48 8d 04 17    lea (%rdi,%rdx,1),%rax
4004d8: c3             retq
```

Attack String (HEX)

```
30 31 ... 38 39 30 ... 39 30 ... 33 d4 04 40 00 00 00 00 00
```

Multiple gadgets can corrupt stack upwards

[CSED211] Introduction to Computer Software Systems

Lecture 8: Advanced Topics

Prof. Jisung Park



CAOS

COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2023.10.16